

The University of Faisalabad

Lab Manual

Submitted To:

Sir Saeed

Submitted By:

Waleed Amjad

Registration no#:

2023-BS-AI-054

Degree Program:

BS - Artificial Intelligence

Semester:

04 Sec:A

Lab Report – Machine Learning

Table of Contents

Project 1: Regression.....	3
Title.....	3
Data Description.....	3
Code and Output	3
Conclusion.....	22
Project 2: Classification.....	23
Title.....	23
Data Description.....	23
Code and Output	23
Conclusion.....	64

Regression

Title:

Linear Regression on Student Performance

Data Description:

Preprocessing of a student data set to find the hours of studying and other activities which effect the overall performance of the student.

Code and Output:

Problem Description: Preprocessing of a student data set to find the hours of studying and other activities which effect the overall performance of the student.

Data Preprocessing Steps

1. Reading Data
2. Exploring Data / Data Insight
3. Cleansing Data
4. Outlier Detection and Removing
5. Data Transformation (Normalize Data / Rescale Data)
6. Categorical into Numerical
7. Dimensionality Reduction(PCA)
8. Handling Imbalanced Data
9. Feature Selection
10. Data Splitting

```
In [12]: import matplotlib.pyplot as plt
import seaborn as sns
color = sns.color_palette()

import numpy as np
import pandas as pd
```

1: Reading Data

```
In [13]: data = pd.read_csv('Student_Performance.csv')
data.head()
```

```
Out[13]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
0	7	99	Yes	9	1	91.0
1	4	82	No	4	2	65.0
2	8	51	Yes	7	2	45.0
3	5	52	Yes	5	2	36.0
4	7	75	No	8	5	66.0

```
In [14]: data.head(2)
```

```
Out[14]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
0	7	99	Yes	9	1	91.0
1	4	82	No	4	2	65.0

```
In [15]: data.head(30)
```

```
Out[15]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
0	7	99	Yes	9	1	91.0
1	4	82	No	4	2	65.0
2	8	51	Yes	7	2	45.0
3	5	52	Yes	5	2	36.0
4	7	75	No	8	5	66.0
5	3	78	No	9	6	61.0
6	7	73	Yes	5	6	63.0
7	8	45	Yes	4	6	42.0
8	5	77	No	8	2	61.0
9	4	89	No	4	0	69.0
10	8	91	No	4	5	84.0

11	8	79	No	6	2	73.0
12	3	47	No	9	2	27.0
13	6	47	No	4	2	33.0
14	5	79	No	7	8	68.0
15	2	72	No	4	3	43.0
16	8	73	Yes	8	4	67.0
17	6	83	Yes	7	2	70.0
18	2	54	Yes	4	9	30.0
19	5	75	No	7	0	63.0
20	1	99	Yes	4	3	71.0
21	6	96	No	9	0	85.0
22	9	74	Yes	7	6	73.0
23	1	85	No	5	6	57.0
24	3	61	No	6	3	35.0
25	7	62	Yes	7	4	49.0
26	4	79	No	8	9	66.0
27	9	84	Yes	6	6	83.0
28	3	94	Yes	6	5	74.0
29	5	90	Yes	4	3	74.0

In [16]: data.tail()

Out[16]:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
9995	1	49	Yes	4	2	23.0
9996	7	64	Yes	8	5	58.0
9997	6	83	Yes	8	5	74.0
9998	9	97	Yes	7	0	95.0
9999	7	74	No	8	1	64.0

In [17]: data.shape

Out[17]: (10000, 6)

In [18]: data.tail(20)

Out[18]:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
9980	2	43	No	6	9	20.0
9981	7	54	No	9	4	46.0
9982	8	51	No	5	1	44.0
9983	8	87	Yes	4	9	79.0
9984	6	45	Yes	6	2	34.0
9985	8	99	No	5	5	92.0
9986	1	48	Yes	8	5	25.0
9987	9	74	No	4	6	69.0
9988	1	47	No	8	5	20.0
9989	3	46	No	5	8	27.0
9990	9	43	No	7	4	40.0
9991	5	97	Yes	7	4	83.0
9992	9	52	No	9	7	50.0
9993	9	58	Yes	7	7	55.0
9994	6	46	Yes	8	0	39.0
9995	1	49	Yes	4	2	23.0
9996	7	64	Yes	8	5	58.0
9997	6	83	Yes	8	5	74.0
9998	9	97	Yes	7	0	95.0
9999	7	74	No	8	1	64.0

In [19]: data.sample()

In [20]: data.sample(30)

Out[20]:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
6032	7	47	Yes	7	5	34.0
7646	4	63	Yes	7	6	48.0
493	4	91	No	9	5	80.0
8531	8	50	Yes	5	6	48.0
439	7	79	Yes	9	4	73.0
7089	8	57	Yes	6	0	49.0
8289	3	83	No	8	4	65.0
4652	4	84	Yes	4	5	71.0
5728	4	89	No	7	4	75.0
8435	2	52	No	6	0	28.0
6804	2	99	No	6	2	73.0
4669	1	46	No	9	9	23.0
1003	6	75	No	7	7	61.0
6395	5	62	Yes	7	9	51.0
7819	6	59	Yes	7	7	48.0
4390	9	96	No	7	4	93.0
4445	9	57	No	8	9	55.0
9706	4	51	Yes	7	4	32.0
8943	8	86	Yes	6	4	82.0
7145	4	58	Yes	4	5	41.0
3037	4	68	No	8	3	51.0
1786	2	62	Yes	9	4	40.0
5624	3	97	No	7	1	78.0
1790	2	50	Yes	9	5	26.0
8395	6	70	No	9	6	61.0
3496	7	52	Yes	5	0	45.0
9430	4	46	Yes	9	0	32.0

In [21]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 6 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Hours Studied                        10000 non-null  int64
1   Previous Scores                     10000 non-null  int64
2   Extracurricular Activities          10000 non-null  object
3   Sleep Hours                         10000 non-null  int64
4   Sample Question Papers Practiced    10000 non-null  int64
5   Performance Index                   10000 non-null  float64
dtypes: float64(1), int64(4), object(1)
memory usage: 468.9+ KB
```

In [22]: data.describe()

Out[22]:

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	4.992900	69.445700	6.530600	4.583300	55.224800
std	2.589309	17.343152	1.695863	2.867348	19.212558
min	1.000000	40.000000	4.000000	0.000000	10.000000
25%	3.000000	54.000000	5.000000	2.000000	40.000000
50%	5.000000	69.000000	7.000000	5.000000	55.000000
75%	7.000000	85.000000	8.000000	7.000000	71.000000
max	9.000000	99.000000	9.000000	9.000000	100.000000

2: Data Cleaning

2: Data Cleaning

Handling Missing Values

- Imputation: Filling missing values with mean.

```
In [23]: import pandas as pd
```

```
In [24]: data.isnull().sum()
```

```
Out[24]: Hours Studied          0
Previous Scores                0
Extracurricular Activities     0
Sleep Hours                   0
Sample Question Papers Practiced 0
Performance Index             0
dtype: int64
```

```
In [25]: import pandas as pd
import numpy as np

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

numeric_cols.fillna(numeric_cols.mean(), inplace=True)

data = pd.concat([numeric_cols, non_numeric_cols], axis=1)

missing_values = data.isnull().sum()
print(missing_values)
```

```
Hours Studied          0
Previous Scores                0
Sleep Hours            0
Sample Question Papers Practiced 0
Performance Index      0
Extracurricular Activities 0
dtype: int64
```

```
In [26]: import pandas as pd
import numpy as np

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

numeric_cols.fillna(numeric_cols.mean(), inplace=True)

for col in non_numeric_cols.columns:
    non_numeric_cols[col].fillna(non_numeric_cols[col].mode()[0], inplace=True)

data = pd.concat([numeric_cols, non_numeric_cols], axis=1)

missing_values = data.isnull().sum()
print(missing_values)
```

Hours Studied	0
Previous Scores	0
Sleep Hours	0
Sample Question Papers Practiced	0
Performance Index	0
Extracurricular Activities	0
dtype:	int64

```
In [27]: data.shape
```

```
Out[27]: (10000, 6)
```

Removal: Deleting rows with missing values.

```
In [28]: data.isnull().sum()
```

```
Out[28]: Hours Studied          0
Previous Scores          0
Sleep Hours              0
Sample Question Papers Practiced  0
Performance Index        0
Extracurricular Activities  0
dtype: int64
```

```
In [29]: data.shape
```

```
Out[29]: (10000, 6)
```

```
In [30]: data.dropna(inplace=True)

missing_values = data.isnull().sum()
print(missing_values)
```

Hours Studied	0
Previous Scores	0
Sleep Hours	0
Sample Question Papers Practiced	0
Performance Index	0
Extracurricular Activities	0
dtype:	int64

```
In [31]: data.shape
```

```
Out[31]: (10000, 6)
```


Removing Duplicates

In [32]:

```
data.shape
```

Out[32]: (10000, 6)

In [33]:

```
data.drop_duplicates(inplace=True)  
data.shape
```

Out[33]: (9873, 6)

3: Outlier Detection and Removal

In [34]:

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
data.describe()
```

Out[34]:

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index
count	9873.000000	9873.000000	9873.000000	9873.000000	9873.000000
mean	4.992100	69.441102	6.531652	4.583004	55.216651
std	2.589081	17.325601	1.697683	2.867202	19.208570
min	1.000000	40.000000	4.000000	0.000000	10.000000
25%	3.000000	54.000000	5.000000	2.000000	40.000000
50%	5.000000	69.000000	7.000000	5.000000	55.000000
75%	7.000000	85.000000	8.000000	7.000000	70.000000
max	9.000000	99.000000	9.000000	9.000000	100.000000

In [35]:

```
0.25-1.5*0.5
```

Out[35]: -0.5

In [36]:

```
0.75 + 1.5 * 0.5
```

Out[36]: 1.5

```

In [37]:
numeric_cols = data.select_dtypes(include=[np.number])

Q1 = numeric_cols.quantile(0.25)
Q3 = numeric_cols.quantile(0.75)
IQR = Q3 - Q1

# Filter out outliers      #  $0.25 - 1.5 * 0.5 = -0.5$       #  $0.75 + 1.5 * 0.5 = 1.5$ 

data_cleaned = data[~((numeric_cols < (Q1 - 1.5 * IQR)) | (numeric_cols > (Q3 + 1.5 * IQR))).any(axis=1)]

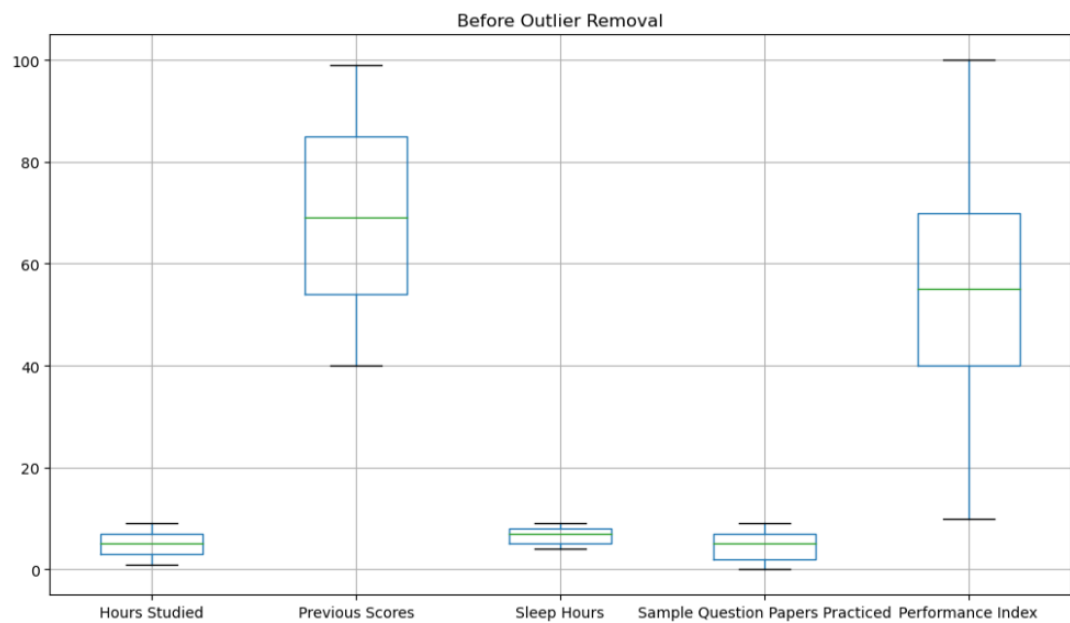
plt.figure(figsize=(20, 6))

plt.subplot(1, 2, 1)
numeric_cols.boxplot()
plt.title("Before Outlier Removal")

plt.tight_layout()
plt.show()

```

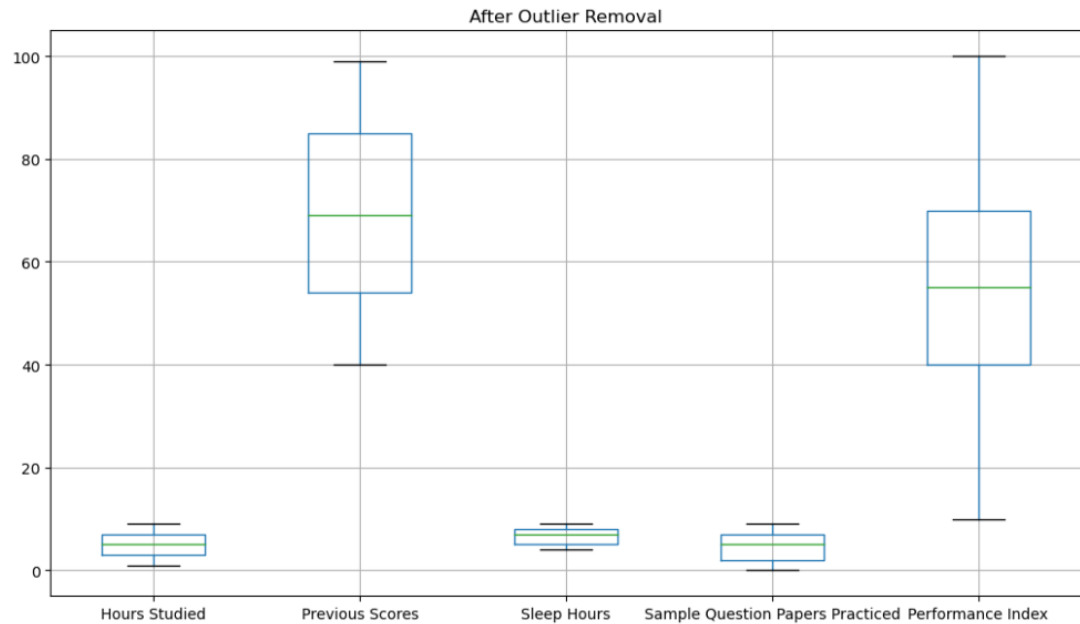
Before Outlier Removal



```
In [38]: plt.figure(figsize=(20, 6))

plt.subplot(1, 2, 2)
data_cleaned.select_dtypes(include=[np.number]).boxplot()
plt.title("After Outlier Removal")

plt.tight_layout()
plt.show()
```



```
In [39]: data_cleaned.shape
```

```
Out[39]: (9873, 6)
```

```
In [40]: data_cleaned.head()
```

```
Out[40]:
```

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities
0	7	99	9	1	91.0	Yes
1	4	82	4	2	65.0	No
2	8	51	7	2	45.0	Yes
3	5	52	5	2	36.0	Yes
4	7	75	8	5	66.0	No

4. Data Transformation

Key Differences

Range of Values:

Normalization: Values are scaled to a fixed range, typically [0, 1]. Standardization: Values are rescaled to have a mean of 0 and a standard deviation of 1.

Effect on Distribution:

Normalization: Compresses or stretches the data to fit within the specified range, potentially altering the original distribution. Standardization: Preserves the shape of the original distribution but changes the scale. Use Cases:

Normalization: Suitable for distance-based algorithms, like k-nearest neighbors and neural networks. Standardization: Suitable for algorithms that assume a normal distribution, like linear regression and logistic regression.

Normalization/Standardization

- Normalization Definition: Normalization rescales the data to a fixed range, typically [0, 1] or [-1, 1].

```
In [41]: import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

scaler = MinMaxScaler()
scaled_numeric_data = scaler.fit_transform(numeric_cols)

scaled_numeric_df = pd.DataFrame(scaled_numeric_data, columns=numeric_cols.columns)

scaled_data = pd.concat([scaled_numeric_df, non_numeric_cols.reset_index(drop=True)], axis=1)

print(scaled_data.shape)
print()
print('*' * 60)
scaled_data.head()

(9873, 6)
```

```
Out[41]:
```

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities
0	0.750	1.000000	1.0	0.111111	0.900000	Yes
1	0.375	0.711864	0.0	0.222222	0.611111	No
2	0.875	0.186441	0.6	0.222222	0.388889	Yes
3	0.500	0.203390	0.2	0.222222	0.288889	Yes
4	0.750	0.593220	0.8	0.555556	0.622222	No

Standardization

Definition: Standardization rescales the data so that it has a mean of 0 and a standard deviation of 1.

```
In [42]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

scaler = StandardScaler()
scaled_numeric_data = scaler.fit_transform(numeric_cols)

scaled_numeric_df = pd.DataFrame(scaled_numeric_data, columns=numeric_cols.columns)

scaled_data = pd.concat([scaled_numeric_df, non_numeric_cols.reset_index(drop=True)], axis=1)

print(scaled_data.shape)
print()
print('*' * 60)
scaled_data.head()
```

(9873, 6)

```
Out[42]:
```

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities
0	0.775566	1.706168	1.454025	-1.249715	1.862979	Yes
1	-0.383205	0.724912	-1.491315	-0.900925	0.509348	No
2	1.161822	-1.064438	0.275889	-0.900925	-0.531907	Yes
3	0.003052	-1.006717	-0.902247	-0.900925	-1.000471	Yes
4	0.775566	0.320865	0.864957	0.145444	0.561411	No

5: One-Hot Encoding ¶

```
In [43]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
```

```
data.head(2)
```

```
Out[43]:
```

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities
0	7	99	9	1	91.0	Yes
1	4	82	4	2	65.0	No

```
In [44]: data["Hours Studied"].unique()
```

```
Out[44]: array([7, 4, 8, 5, 3, 6, 2, 1, 9], dtype=int64)
```

```
In [45]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
```

```
cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']

data1 = pd.get_dummies(cat_features)
data1
```

```
Out[45]:
```

	Extracurricular Activities
0	1

```
In [46]: data1.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1 entries, 0 to 0
Data columns (total 1 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   Extracurricular Activities  1 non-null     uint8
dtypes: uint8(1)
memory usage: 133.0 bytes
```

```
In [47]: cat_features
```

```
Out[47]: ['Extracurricular Activities']
```

```
In [48]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']

data1 = pd.get_dummies(data, columns=cat_features)

scaled_data = pd.concat([data, data1], axis=1)

print(scaled_data.shape)
print()
print('*' * 70)

scaled_data.head()

(9873, 13)
```

```
Out[48]:
```

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities_No	Extracurricular Activities_Yes
0	7	99	9	1	91.0	Yes	7	99	9	1	91.0	0	1
1	4	82	4	2	65.0	No	4	82	4	2	65.0	1	0
2	8	51	7	2	45.0	Yes	8	51	7	2	45.0	0	1
3	5	52	5	2	36.0	Yes	5	52	5	2	36.0	0	1
4	7	75	8	5	66.0	No	7	75	8	5	66.0	1	0

```
In [49]: data.columns
```

```
Out[49]: Index(['Hours Studied', 'Previous Scores', 'Sleep Hours',
'Sample Question Papers Practiced', 'Performance Index',
'Extracurricular Activities'],
dtype='object')
```

```
In [50]: scaled_data.columns
```

```
Out[50]: Index(['Hours Studied', 'Previous Scores', 'Sleep Hours',
'Sample Question Papers Practiced', 'Performance Index',
'Extracurricular Activities', 'Hours Studied', 'Previous Scores',
'Sleep Hours', 'Sample Question Papers Practiced', 'Performance Index',
'Extracurricular Activities_No', 'Extracurricular Activities_Yes'],
dtype='object')
```

```
In [51]: data1.head()
```

```
Out[51]:
```

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities_No	Extracurricular Activities_Yes
0	7	99	9	1	91.0	0	1
1	4	82	4	2	65.0	1	0
2	8	51	7	2	45.0	0	1
3	5	52	5	2	36.0	0	1
4	7	75	8	5	66.0	1	0

```
In [ ]:
```

6: Data Reduction

Dimensionality Reduction

PCA (Principal Component Analysis)

```
In [52]: scaled_data.shape
```

```
Out[52]: (9873, 13)
```

```
In [53]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']
numeric_features = [feature for feature in data.columns if data[feature].dtype != 'O']

numeric_means = data[numeric_features].mean()
data[numeric_features] = data[numeric_features].fillna(numeric_means)

data = pd.get_dummies(data, columns=cat_features)

scaler = StandardScaler()
data[numeric_features] = scaler.fit_transform(data[numeric_features])


pca_2 = PCA(n_components=2)
data_pca_2 = pca_2.fit_transform(data)

plt.figure(figsize=(14, 6))
```

```
plt.subplot(1, 2, 1)
plt.scatter(data[numeric_features[0]], data[numeric_features[1]], alpha=0.5)
plt.title('Original Data')
plt.xlabel(numeric_features[0])
plt.ylabel(numeric_features[1])

plt.subplot(1, 2, 2)
plt.scatter(data_pca_2[:, 0], data_pca_2[:, 1], alpha=0.5)
plt.title('PCA Transformed Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')

plt.tight_layout()
plt.show()
```



```
In [54]: type(data_pca_2)
```

```
Out[54]: numpy.ndarray
```

```
In [55]: data_pca_2.ndim
```

```
Out[55]: 2
```

```
In [56]: data_pca_2.shape
```

```
Out[56]: (9873, 2)
```


7: Handling Imbalanced Data

- Resampling Techniques
- Oversampling

```
In [81]: import pandas as pd
import numpy as np
import imblearn
import category_encoders
import packaging.version
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt

data = pd.read_csv('Student_Performance.csv')

data.fillna(data.mean(numeric_only=True), inplace=True)

target = data['Hours Studied']
data = data.drop(columns=['Hours Studied'])

cat_features = [col for col in data.columns if data[col].dtype == 'O']
numeric_features = [col for col in data.columns if data[col].dtype != 'O']

data = pd.get_dummies(data, columns=cat_features)

scaler = StandardScaler()
data[numeric_features] = scaler.fit_transform(data[numeric_features])

if target.dtype == 'O':
    le = LabelEncoder()
    target = le.fit_transform(target)
elif target.dtype == 'float':
    target = (target > 0.5).astype(int)
```

```
print("Before SMOTE:", data.shape, target.shape)

smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(data, target)

resampled_df = pd.concat([
    pd.DataFrame(X_resampled, columns=data.columns),
    pd.DataFrame(y_resampled, columns=['Hours Studied'])
], axis=1)

print("After SMOTE:", resampled_df.shape)
resampled_df.head()
```

Before SMOTE: (10000, 6) (10000,)
After SMOTE: (10368, 7)

	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities_No	Extracurricular Activities_Yes	Hours Studied
0	1.704176	1.456205	-1.249754	1.862167	0	1	7
1	0.723913	-1.492294	-0.900982	0.508818	1	0	4
2	-1.063626	0.276805	-0.900982	-0.532220	0	1	8
3	-1.005963	-0.902594	-0.900982	-1.000687	0	1	5
4	0.320275	0.866505	0.145333	0.560870	1	0	7

```
In [82]: resampled_df['Hours Studied'].value_counts(True)
```

```
Out[82]: 7    0.111111
         4    0.111111
         8    0.111111
         5    0.111111
         3    0.111111
         6    0.111111
         2    0.111111
         1    0.111111
         9    0.111111
         Name: Hours Studied, dtype: float64
```

```
In [83]: resampled_df.shape
```

```
Out[83]: (10368, 7)
```

```
In [84]: print(resampled_df.columns)
         print(list(resampled_df.columns))
```

```
Index(['Previous Scores', 'Sleep Hours', 'Sample Question Papers Practiced',
       'Performance Index', 'Extracurricular Activities_No',
       'Extracurricular Activities_Yes', 'Hours Studied'],
      dtype='object')
['Previous Scores', 'Sleep Hours', 'Sample Question Papers Practiced', 'Performance Index', 'Extracurricular Activities_No', 'Extracurricular Activities_Yes', 'Hours Studied']
```

```
In [85]: print(data.columns)
         print(list(data.columns))
```

```
Index(['Previous Scores', 'Sleep Hours', 'Sample Question Papers Practiced',
       'Performance Index', 'Extracurricular Activities_No',
       'Extracurricular Activities_Yes'],
      dtype='object')
['Previous Scores', 'Sleep Hours', 'Sample Question Papers Practiced', 'Performance Index', 'Extracurricular Activities_No', 'Extracurricular Activities_Yes']
```

Undersampling

```
In [86]: import pandas as pd
import numpy as np
import imblearn
import category_encoders
import packaging.version
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
import matplotlib.pyplot as plt

data = pd.read_csv('Student_Performance.csv')

data.fillna(data.mean(numeric_only=True), inplace=True)

target = data['Hours Studied']
data = data.drop(columns=['Hours Studied'])

cat_features = [col for col in data.columns if data[col].dtype == 'O']
numeric_features = [col for col in data.columns if data[col].dtype != 'O']

data = pd.get_dummies(data, columns=cat_features)

scaler = StandardScaler()
data[numeric_features] = scaler.fit_transform(data[numeric_features])

if target.dtype == 'O':
    le = LabelEncoder()
    target = le.fit_transform(target)
elif target.dtype == 'float':
    target = (target > 0.5).astype(int)

print("Before undersampling:", data.shape, target.shape)

rus = RandomUnderSampler(random_state=42)
```

```

rus = RandomUnderSampler(random_state=42)
X_resampled, y_resampled = rus.fit_resample(data, target)

resampled_df = pd.concat([
    pd.DataFrame(X_resampled, columns=data.columns),
    pd.DataFrame(y_resampled, columns=['Hours Studied'])
], axis=1)

print("After undersampling:", resampled_df.shape)
resampled_df.head()

```

Before undersampling: (10000, 6) (10000,)
After undersampling: (9765, 7)

Out[86]:

	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities_No	Extracurricular Activities_Yes	Hours Studied
1353	-0.083363	-0.312895	1.191649	-0.688376	1	0	1
6622	0.781575	0.276805	-0.900982	-0.063753	0	1	1
5076	1.588851	1.456205	1.191649	0.769077	1	0	1
1998	0.896900	-1.492294	0.494105	0.248558	1	0	1
6245	1.531188	0.866505	-0.552210	0.769077	1	0	1

In [88]: resampled_df['Hours Studied'].value_counts()

Out[88]:

```

1    1085
2    1085
3    1085
4    1085
5    1085
6    1085
7    1085
8    1085
9    1085
Name: Hours Studied, dtype: int64

```

In [89]: resampled_df.shape

Out[89]: (9765, 7)

Target Encoder

```

In [90]: import pandas as pd
from category_encoders import TargetEncoder

# Example dataset
data = {'animal': ['cat', 'dog', 'mouse', 'dog', 'cat'], 'target': [1, 0, 1, 0, 1]}
df = pd.DataFrame(data)

target_encoder = TargetEncoder(cols=['animal'])
target_encoded = target_encoder.fit_transform(df['animal'], df['target'])
print(target_encoded)

```

```

    animal
0  0.656740
1  0.514889
2  0.652043
3  0.514889
4  0.656740

```

8: Splitting Data

```

In [93]: from sklearn.model_selection import train_test_split

X = resampled_df.drop('Hours Studied', axis=1)
y = resampled_df['Hours Studied']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

```

In [94]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

Out[94]: ((6835, 6), (2930, 6), (6835,), (2930,))

9. Regression

```
In [96]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd

# Assuming 'data' is your fully preprocessed DataFrame
X = resampled_df.drop('Performance Index', axis=1)
y = resampled_df['Performance Index']

# Train/test split

# Initialize and train model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)

# Predict on test set
y_pred = lr_model.predict(X_test)

# Evaluation
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("📊 Linear Regression Evaluation:")
print(f"Mean Squared Error: {mse:.2f}")
print(f"R² Score: {r2:.2f}")

# Show coefficients
coeff_df = pd.DataFrame({'Feature': X.columns, 'Coefficient': lr_model.coef_})
print("\n🔍 Feature Coefficients:")
print(coeff_df)
```

📊 Linear Regression Evaluation:
Mean Squared Error: 0.48
R² Score: 0.93

🔍 Feature Coefficients:

	Feature	Coefficient
0	Previous Scores	-5.749869
1	Sleep Hours	-0.267676
2	Sample Question Papers Practiced	-0.169628
3	Extracurricular Activities_No	6.259859
4	Extracurricular Activities_Yes	0.098548
5	Hours Studied	-0.098548

Prediction

```
In [102]: # Correct feature and target separation
X = resampled_df.drop('Performance Index', axis=1) # features
y = resampled_df['Performance Index'] # target
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
from sklearn.linear_model import LinearRegression

lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
# This must match the feature columns used in X
new_data = pd.DataFrame({
    'Previous Scores': [75],
    'Sleep Hours': [6],
    'Sample Question Papers Practiced': [5],
    'Extracurricular Activities_No': [0],
    'Extracurricular Activities_Yes': [1],
    'Hours Studied': [3],
})

# Ensure same column order
new_data = new_data[X.columns]

# Predict
prediction = lr_model.predict(new_data)
print(f"🎯 Predicted Performance Index: {prediction[0]:.2f}")
```

🎯 Predicted Performance Index: 69.03

Conclusion:

In this project, a linear regression model was developed to predict a student's *Performance Index* based on several academic and behavioral features, including *Previous Scores*, *Sleep Hours*, *Hours Studied*, and participation in *Extracurricular Activities*. The dataset was preprocessed by handling missing values, encoding categorical variables, and removing outliers using the IQR method. After training and testing the model, it achieved a high R^2 score of 0.93, indicating that 93% of the variability in the Performance Index can be explained by the selected features. The model also showed a relatively low mean squared error, suggesting good prediction accuracy. Analysis of the feature coefficients revealed that *Previous Scores* and *Extracurricular Activities* had the most significant impact on performance, while *Hours Studied* and *Sleep Hours* had a comparatively smaller effect. Overall, the model demonstrates the usefulness of linear regression in educational performance prediction and highlights which factors most strongly influence student outcomes.

Classification

Title:

Model to find the Best seller Supplements based on various features.

Data Description:

Data was extracted from Statista of sales across different regions.

Prediction task is to determine the best seller product

Code and Output:

Problem Description: Model to find the Best seller Supplements based on various features.

Data Preprocessing Steps

```
In [1]: import matplotlib.pyplot as plt
import seaborn as sns

color = sns.color_palette()

import numpy as np
import pandas as pd
```

1: Reading Data

```
In [2]: data = pd.read_csv('Supplement.csv')
data.head()
```

```
Out[2]:
```

	Date	Product Name	Category	Units Sold	Price	Revenue	Discount	Units Returned	Location	Platform	Best Seller
0	2020-01-06	Whey Protein	Protein	143	31.98	4573.14	0.03	2	Canada	Walmart	0
1	2020-01-06	Vitamin C	Vitamin	139	42.51	5908.89	0.04	0	UK	Amazon	0
2	2020-01-06	Fish Oil	Omega	161	12.91	2078.51	0.25	0	Canada	Amazon	1
3	2020-01-06	Multivitamin	Vitamin	140	16.07	2249.80	0.08	0	Canada	Walmart	0
4	2020-01-06	Pre-Workout	Performance	157	35.47	5568.79	0.25	3	Canada	iHerb	0

```
In [4]: data.head(10)
```

```
Out[4]:
```

	Date	Product Name	Category	Units Sold	Price	Revenue	Discount	Units Returned	Location	Platform	Best Seller
0	2020-01-06	Whey Protein	Protein	143	31.98	4573.14	0.03	2	Canada	Walmart	0
1	2020-01-06	Vitamin C	Vitamin	139	42.51	5908.89	0.04	0	UK	Amazon	0
2	2020-01-06	Fish Oil	Omega	161	12.91	2078.51	0.25	0	Canada	Amazon	1
3	2020-01-06	Multivitamin	Vitamin	140	16.07	2249.80	0.08	0	Canada	Walmart	0
4	2020-01-06	Pre-Workout	Performance	157	35.47	5568.79	0.25	3	Canada	iHerb	0
5	2020-01-06	BCAA	Amino Acid	154	41.19	6343.26	0.13	1	UK	Walmart	0
6	2020-01-06	Creatine	Performance	134	32.49	4353.66	0.05	1	UK	Walmart	0
7	2020-01-06	Zinc	Mineral	147	46.68	6861.96	0.19	0	Canada	Amazon	0
8	2020-01-06	Collagen Peptides	Protein	147	10.96	1611.12	0.06	2	USA	Amazon	0
9	2020-01-06	Magnesium	Mineral	134	20.76	2781.84	0.00	0	Canada	Amazon	0

```
In [5]: data.tail()
```

```
Out[5]:
```

	Date	Product Name	Category	Units Sold	Price	Revenue	Discount	Units Returned	Location	Platform	Best Seller
4379	2025-03-31	Melatonin	Sleep Aid	160	47.79	7646.40	0.21	1	USA	iHerb	1
4380	2025-03-31	Biotin	Vitamin	154	38.12	5870.48	0.22	1	UK	Walmart	0
4381	2025-03-31	Green Tea Extract	Fat Burner	139	20.40	2835.60	0.12	3	USA	iHerb	0
4382	2025-03-31	Iron Supplement	Mineral	154	18.31	2819.74	0.23	2	Canada	Amazon	0
4383	2025-03-31	Electrolyte Powder	Hydration	178	39.12	6963.36	0.23	0	UK	iHerb	1

```
In [6]: data.shape
```

```
Out[6]: (4384, 11)
```

2: Data Cleaning

Handling Missing Values

- Imputation: Filling missing values with mean.

```
In [12]: data.isnull().sum()
```

```
Out[12]: Date          0
Product Name         0
Category             0
Units Sold           0
Price                0
Revenue              0
Discount             0
Units Returned       0
Location             0
Platform             0
Best Seller          0
dtype: int64
```



```
In [13]: import pandas as pd
import numpy as np

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

numeric_cols.fillna(numeric_cols.mean(), inplace=True)

data = pd.concat([numeric_cols, non_numeric_cols], axis=1)

missing_values = data.isnull().sum()
print(missing_values)
```

```
Units Sold      0
Price           0
Revenue         0
Discount        0
Units Returned  0
Best Seller     0
Date           0
Product Name    0
Category        0
Location        0
Platform        0
dtype: int64
```

```
In [14]: import pandas as pd
import numpy as np

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

numeric_cols.fillna(numeric_cols.mean(), inplace=True)

for col in non_numeric_cols.columns:
    non_numeric_cols[col].fillna(non_numeric_cols[col].mode()[0], inplace=True)

data = pd.concat([numeric_cols, non_numeric_cols], axis=1)

missing_values = data.isnull().sum()
print(missing_values)
```

```
Units Sold      0
Price           0
Revenue         0
Discount        0
Units Returned  0
Best Seller     0
Date           0
Product Name    0
Category        0
Location        0
Platform        0
dtype: int64
```

```
In [15]: data.shape
```

```
Out[15]: (4384, 11)
```

Removal: Deleting rows with missing values.

```
In [16]: data.isnull().sum()
```

```
Out[16]: Units Sold      0
         Price          0
         Revenue        0
         Discount        0
         Units Returned  0
         Best Seller     0
         Date            0
         Product Name    0
         Category        0
         Location        0
         Platform        0
         dtype: int64
```

```
In [17]: data.shape
```

```
Out[17]: (4384, 11)
```

```
In [18]: data.dropna(inplace=True)

         missing_values = data.isnull().sum()
         print(missing_values)
```

```
Units Sold      0
Price           0
Revenue         0
Discount        0
Units Returned  0
Best Seller     0
Date            0
Product Name    0
Category        0
Location        0
Platform        0
dtype: int64
```

```
In [19]: data.shape
```

```
Out[19]: (4384, 11)
```

```
In [ ]:
```

3: Outlier Detection and Removal

In [21]: `data.describe()`

Out[21]:

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller
count	4384.000000	4384.000000	4384.000000	4384.000000	4384.000000	4384.000000
mean	150.200274	34.781229	5226.569446	0.124398	1.531478	0.275776
std	12.396099	14.198309	2192.491946	0.071792	1.258479	0.446955
min	103.000000	10.000000	1284.000000	0.000000	0.000000	0.000000
25%	142.000000	22.597500	3349.372500	0.060000	1.000000	0.000000
50%	150.000000	34.720000	5173.140000	0.120000	1.000000	0.000000
75%	158.000000	46.712500	7009.960000	0.190000	2.000000	1.000000
max	194.000000	59.970000	10761.850000	0.250000	8.000000	1.000000

In [22]: `0.25-1.5*0.5`

Out[22]: `-0.5`

In [23]: `0.75 + 1.5 * 0.5`

Out[23]: `1.5`

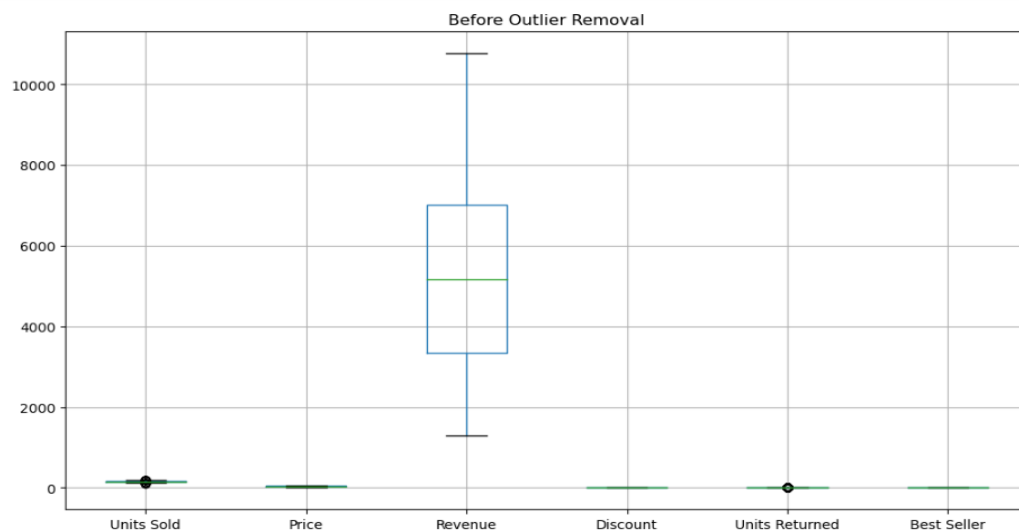
```
In [24]: numeric_cols = data.select_dtypes(include=[np.number])
Q1 = numeric_cols.quantile(0.25)
Q3 = numeric_cols.quantile(0.75)
IQR = Q3 - Q1

data_cleaned = data[~((numeric_cols < (Q1 - 1.5 * IQR)) | (numeric_cols > (Q3 + 1.5 * IQR))).any(axis=1)]

plt.figure(figsize=(20, 6))

plt.subplot(1, 2, 1)
numeric_cols.boxplot()
plt.title("Before Outlier Removal")

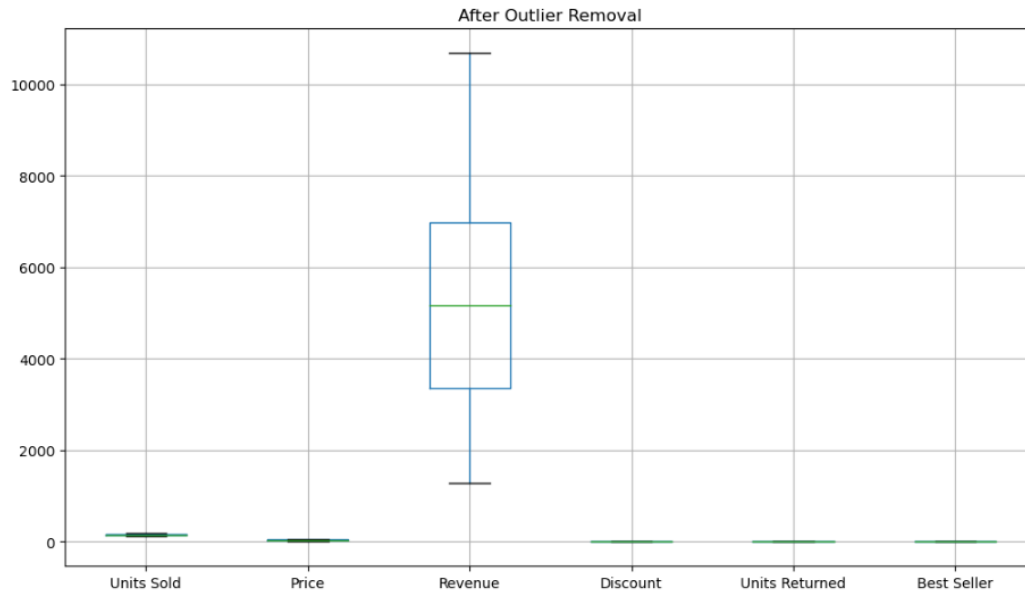
plt.tight_layout()
plt.show()
```



```
In [25]: plt.figure(figsize=(20, 6))

plt.subplot(1, 2, 2)
data_cleaned.select_dtypes(include=[np.number]).boxplot()
plt.title("After Outlier Removal")

plt.tight_layout()
plt.show()
```



```
In [26]: data_cleaned.shape
```

```
Out[26]: (4033, 11)
```

```
In [27]: data_cleaned.head()
```

```
Out[27]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	Date	Product Name	Category	Location	Platform
0	143	31.98	4573.14	0.03	2	0	2020-01-06	Whey Protein	Protein	Canada	Walmart
1	139	42.51	5908.89	0.04	0	0	2020-01-06	Vitamin C	Vitamin	UK	Amazon
2	161	12.91	2078.51	0.25	0	1	2020-01-06	Fish Oil	Omega	Canada	Amazon
3	140	16.07	2249.80	0.08	0	0	2020-01-06	Multivitamin	Vitamin	Canada	Walmart
4	157	35.47	5568.79	0.25	3	0	2020-01-06	Pre-Workout	Performance	Canada	iHerb

```
In [ ]:
```

4. Data Transformation

```
In [28]: import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

scaler = MinMaxScaler()
scaled_numeric_data = scaler.fit_transform(numeric_cols)

scaled_numeric_df = pd.DataFrame(scaled_numeric_data, columns=numeric_cols.columns)

scaled_data = pd.concat([scaled_numeric_df, non_numeric_cols.reset_index(drop=True)], axis=1)

print(scaled_data.shape)
print()
print('*' * 60)
scaled_data.head()
```

(4384, 11)

```
Out[28]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	Date	Product Name	Category	Location	Platform
0	0.439560	0.439864	0.347034	0.12	0.250	0.0	2020-01-06	Whey Protein	Protein	Canada	Walmart
1	0.395604	0.650590	0.487968	0.16	0.000	0.0	2020-01-06	Vitamin C	Vitamin	UK	Amazon
2	0.637363	0.058235	0.083828	1.00	0.000	1.0	2020-01-06	Fish Oil	Omega	Canada	Amazon
3	0.406593	0.121473	0.101901	0.32	0.000	0.0	2020-01-06	Multivitamin	Vitamin	Canada	Walmart
4	0.593407	0.509706	0.452085	1.00	0.375	0.0	2020-01-06	Pre-Workout	Performance	Canada	iHerb

In []:

Standardization

Definition: Standardization rescales the data so that it has a mean of 0 and a standard deviation of 1.

```
In [29]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler

numeric_cols = data.select_dtypes(include=[np.number])
non_numeric_cols = data.select_dtypes(exclude=[np.number])

scaler = StandardScaler()
scaled_numeric_data = scaler.fit_transform(numeric_cols)

scaled_numeric_df = pd.DataFrame(scaled_numeric_data, columns=numeric_cols.columns)

scaled_data = pd.concat([scaled_numeric_df, non_numeric_cols.reset_index(drop=True)], axis=1)

print(scaled_data.shape)
print()
print('*' * 60)
scaled_data.head()
```

(4384, 11)

```
Out[29]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	Date	Product Name	Category	Location	Platform
0	-0.580916	-0.197316	-0.298064	-1.315034	0.372335	-0.617080	2020-01-06	Whey Protein	Protein	Canada	Walmart
1	-0.903635	0.544407	0.311243	-1.175727	-1.217067	-0.617080	2020-01-06	Vitamin C	Vitamin	UK	Amazon
2	0.871319	-1.540587	-1.436000	1.749735	-1.217067	1.620536	2020-01-06	Fish Oil	Omega	Canada	Amazon
3	-0.822955	-1.318000	-1.357865	-0.618496	-1.217067	-0.617080	2020-01-06	Multivitamin	Vitamin	Canada	Walmart
4	0.548600	0.048516	0.156105	1.749735	1.167036	-0.617080	2020-01-06	Pre-Workout	Performance	Canada	iHerb

5: One-Hot Encoding

```
In [30]: data.head(2)
```

```
Out[30]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	Date	Product Name	Category	Location	Platform
0	143	31.98	4573.14	0.03	2	0	2020-01-06	Whey Protein	Protein	Canada	Walmart
1	139	42.51	5908.89	0.04	0	0	2020-01-06	Vitamin C	Vitamin	UK	Amazon

```
In [31]: data["Best Seller"].unique()
```

```
Out[31]: array([0, 1], dtype=int64)
```

```
In [ ]:
```

```
In [32]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']

data1 = pd.get_dummies(cat_features)
data1
```

```
Out[32]:
```

	Category	Date	Location	Platform	Product Name
0	0	1	0	0	0
1	0	0	0	0	1
2	1	0	0	0	0
3	0	0	1	0	0
4	0	0	0	1	0

```
In [33]: data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Category    5 non-null      uint8
1   Date        5 non-null      uint8
2   Location    5 non-null      uint8
3   Platform    5 non-null      uint8
4   Product Name 5 non-null      uint8
dtypes: uint8(5)
memory usage: 157.0 bytes
```

```
In [34]: cat_features
```

```
Out[34]: ['Date', 'Product Name', 'Category', 'Location', 'Platform']
```

```
In [35]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']

data1 = pd.get_dummies(data, columns=cat_features)

scaled_data = pd.concat([data, data1], axis=1)

print(scaled_data.shape)
print()
print('*' * 70)

scaled_data.head()
```

```
(4384, 323)
```

```
*****
```

```
Out[35]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	Date	Product Name	Category	Location	...	Category_Performance	Category_Protein	Category_Sleep Aid
0	143	31.98	4573.14	0.03	2	0	2020-01-06	Whey Protein	Protein	Canada	...	0	1	0
1	139	42.51	5908.89	0.04	0	0	2020-01-06	Vitamin C	Vitamin	UK	...	0	0	0
2	161	12.91	2078.51	0.25	0	1	2020-01-06	Fish Oil	Omega	Canada	...	0	0	0
3	140	16.07	2249.80	0.08	0	0	2020-01-06	Multivitamin	Vitamin	Canada	...	0	0	0
4	157	35.47	5568.79	0.25	3	0	2020-01-06	Pre-Workout	Performance	Canada	...	1	0	0

5 rows x 323 columns

```
In [36]: data.columns
```

```
Out[36]: Index(['Units Sold', 'Price', 'Revenue', 'Discount', 'Units Returned',
               'Best Seller', 'Date', 'Product Name', 'Category', 'Location',
               'Platform'],
              dtype='object')
```

```
In [37]: scaled_data.columns
```

```
Out[37]: Index(['Units Sold', 'Price', 'Revenue', 'Discount', 'Units Returned',
               'Best Seller', 'Date', 'Product Name', 'Category', 'Location',
               ...,
               'Category_Performance', 'Category_Protein', 'Category_Sleep Aid',
               'Category_Vitamin', 'Location_Canada', 'Location_UK', 'Location_USA',
               'Platform_Amazon', 'Platform_Walmart', 'Platform_iHerb'],
              dtype='object', length=323)
```

```
In [38]: data1.head()
```

```
Out[38]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	Date_2020-01-06	Date_2020-01-13	Date_2020-01-20	Date_2020-01-27	...	Category_Performance	Category_Protein	Category_S
0	143	31.98	4573.14	0.03	2	0	1	0	0	0	...	0	1	
1	139	42.51	5908.89	0.04	0	0	1	0	0	0	...	0	0	
2	161	12.91	2078.51	0.25	0	1	1	0	0	0	...	0	0	
3	140	16.07	2249.80	0.08	0	0	1	0	0	0	...	0	0	
4	157	35.47	5568.79	0.25	3	0	1	0	0	0	...	1	0	

5 rows x 312 columns

```
In [ ]:
```

6: Data Reduction

Dimensionality Reduction

PCA (Principal Component Analysis)

```
In [39]: scaled_data.shape
```

```
Out[39]: (4384, 323)
```

```
In [40]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

data.fillna(data.mean(numeric_only=True), inplace=True)

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']
numeric_features = [feature for feature in data.columns if data[feature].dtype != 'O']

data = pd.get_dummies(data, columns=cat_features)

scaler = StandardScaler()
data[numeric_features] = scaler.fit_transform(data[numeric_features].values)

pca = PCA(n_components=4)
data_pca = pca.fit_transform(data)

print(data_pca.shape)
print(data_pca[:5])

plt.figure(figsize=(14, 6))
```

```

plt.subplot(1, 2, 1)
plt.scatter(data[numeric_features[0]], data[numeric_features[1]], alpha=0.5)
plt.title('Original Data')
plt.xlabel(numeric_features[0])
plt.ylabel(numeric_features[1])

pca = PCA(n_components=4)
data_pca = pca.fit_transform(data)

plt.subplot(1, 2, 2)
plt.scatter(data_pca[:, 0], data_pca[:, 1], alpha=0.5)
plt.title('PCA Transformed Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')


plt.tight_layout()
plt.show()

```

```

(4384, 4)
[[-0.69517858 -0.42304467 -1.1550605  0.83877112]
 [-0.1448728  -1.43313913 -1.42185059 -0.54470601]
 [-0.97863628  2.32000232  1.24374271 -2.03899512]
 [-2.21551455 -0.02062641 -0.93827002 -0.86150106]
 [ 0.161729   0.06401863  2.01731706  0.62374047]]

```



```
In [41]: type(data_pca)
```

```
Out[41]: numpy.ndarray
```

```
In [42]: data_pca.ndim
```

```
Out[42]: 2
```

```
In [43]: data_pca.shape
```

```
Out[43]: (4384, 4)
```


7: Handling Imbalanced Data

- Resampling Techniques
- Oversampling

```
In [44]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt

data.fillna(data.mean(numeric_only=True), inplace=True)

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']
numeric_features = [feature for feature in data.columns if data[feature].dtype != 'O']

data = pd.get_dummies(data, columns=cat_features)

scaler = StandardScaler()
data[numeric_features] = scaler.fit_transform(data[numeric_features].values)

if data['Best Seller'].dtype != 'int64' and data['Best Seller'].dtype != 'bool':
    data['Best Seller'] = (data['Best Seller'] > 0.5).astype(int)

X = data.drop(columns=['Best Seller'])
y = data['Best Seller']

if y.dtype == 'O':
    le = LabelEncoder()
    y = le.fit_transform(y)

print(X.shape, y.shape)
```

```
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

data_resampled = pd.concat([pd.DataFrame(X_resampled, columns=X.columns), pd.DataFrame(y_resampled, columns=['Best Seller'])], axis=1)
data_resampled.head()
```

(4384, 311) (4384,)

Out[44]:

	Units Sold	Price	Revenue	Discount	Units Returned	Date_2020-01-06	Date_2020-01-13	Date_2020-01-20	Date_2020-01-27	Date_2020-02-03	...	Category_Protein	Category_Sleep Aid	Cat
0	-0.580916	-0.197316	-0.298064	-1.315034	0.372335	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...	2.645751	-0.258199	
1	-0.903635	0.544407	0.311243	-1.175727	-1.217067	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199	
2	0.871319	-1.540587	-1.436000	1.749735	-1.217067	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199	
3	-0.822955	-1.318000	-1.357865	-0.618496	-1.217067	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199	
4	0.548600	0.048516	0.156105	1.749735	1.167036	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199	

5 rows x 312 columns

```
In [45]: data_resampled['Best Seller'].value_counts(True)
```

Out[45]:

0	0.5
1	0.5

Name: Best Seller, dtype: float64

```
In [46]: data_resampled.shape
```

Out[46]: (6350, 312)

Undersampling

```
In [47]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt

data.fillna(data.mean(numeric_only=True), inplace=True)

cat_features = [feature for feature in data.columns if data[feature].dtype == 'O']
numeric_features = [feature for feature in data.columns if data[feature].dtype != 'O']

data = pd.get_dummies(data, columns=cat_features)

scaler = StandardScaler()
data[numeric_features] = scaler.fit_transform(data[numeric_features].values)

if data['Best Seller'].dtype != 'int64' and data['Best Seller'].dtype != 'bool':
    data['Best Seller'] = (data['Best Seller'] > 0.5).astype(int)

X = data.drop(columns=['Best Seller'])
y = data['Best Seller']

if y.dtype == 'O':
    le = LabelEncoder()
    y = le.fit_transform(y)

print(X.shape, y.shape)
```

```
from imblearn.under_sampling import RandomUnderSampler

rus = RandomUnderSampler()
X_resampled, y_resampled = rus.fit_resample(X, y)

data_resampled = pd.concat([pd.DataFrame(X_resampled, columns=X.columns), pd.DataFrame(y_resampled, columns=['Best Seller'])], axis=1)
data_resampled.head()
```

(4384, 311) (4384,)

Out[47]:

	Units Sold	Price	Revenue	Discount	Units Returned	Date_2020-01-06	Date_2020-01-13	Date_2020-01-20	Date_2020-01-27	Date_2020-02-03	...	Category_Protein	Category_Sleep Aid
1159	-0.984315	0.463402	0.219461	-0.618496	1.961737	-0.060523	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199
3236	-1.387714	0.583852	0.228871	0.495966	1.167036	-0.060523	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199
3792	0.306561	0.231658	0.290210	1.192504	0.372335	-0.060523	-0.060523	-0.060523	-0.060523	-0.060523	...	2.645751	-0.258199
4131	-0.419557	-0.910862	-0.938911	0.774581	-0.422366	-0.060523	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199
3485	-0.016158	1.419258	1.374360	-0.897111	-0.422366	-0.060523	-0.060523	-0.060523	-0.060523	-0.060523	...	-0.377964	-0.258199

5 rows × 312 columns

In [48]: data_resampled['Best Seller'].value_counts()

Out[48]:

0	1209
1	1209

Name: Best Seller, dtype: int64

In [49]: data_resampled.shape

Out[49]: (2418, 312)

Target Encoder

```
In [50]: import pandas as pd

target = 'Best Seller'
categorical_cols = data.select_dtypes(include=['object']).columns

for col in categorical_cols:
    target_mean = data.groupby(col)[target].mean()
    data[col + '_target_enc'] = data[col].map(target_mean)

data = data.drop(columns=categorical_cols)

print(data.head())
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller	\
0	-0.580916	-0.197316	-0.298064	-1.315034	0.372335	0	
1	-0.903635	0.544407	0.311243	-1.175727	-1.217067	0	
2	0.871319	-1.540587	-1.436000	1.749735	-1.217067	1	
3	-0.822955	-1.318000	-1.357865	-0.618496	-1.217067	0	
4	0.548600	0.048516	0.156105	1.749735	1.167036	0	

	Date_2020-01-06	Date_2020-01-13	Date_2020-01-20	Date_2020-01-27	...	\
0	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...
1	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...
2	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...
3	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...
4	16.522712	-0.060523	-0.060523	-0.060523	-0.060523	...

	Category_Performance	Category_Protein	Category_Sleep Aid	\
0	-0.377964	2.645751	-0.258199	
1	-0.377964	-0.377964	-0.258199	
2	-0.377964	-0.377964	-0.258199	
3	-0.377964	-0.377964	-0.258199	
4	2.645751	-0.377964	-0.258199	

	Category_Vitamin	Location_Canada	Location_UK	Location_USA	\
0	-0.480384	1.381699	-0.712072	-0.685678	
1	2.081666	-0.723747	1.404352	-0.685678	
2	-0.480384	1.381699	-0.712072	-0.685678	
3	2.081666	1.381699	-0.712072	-0.685678	
4	-0.480384	1.381699	-0.712072	-0.685678	

	Platform_Amazon	Platform_Walmart	Platform_iHerb
0	-0.711345	1.450798	-0.720822
1	1.405788	-0.689276	-0.720822
2	1.405788	-0.689276	-0.720822
3	-0.711345	1.450798	-0.720822
4	-0.711345	-0.689276	1.387305

[5 rows x 312 columns]

8: Splitting Data

```
In [51]: from sklearn.model_selection import train_test_split
import pandas as pd

print(data.columns)

X = data.drop('Best Seller', axis=1)
y = data['Best Seller']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

Index(['Units Sold', 'Price', 'Revenue', 'Discount', 'Units Returned',
      'Best Seller', 'Date_2020-01-06', 'Date_2020-01-13', 'Date_2020-01-20',
      'Date_2020-01-27',
      ...
      'Category_Performance', 'Category_Protein', 'Category_Sleep Aid',
      'Category_Vitamin', 'Location_Canada', 'Location_UK', 'Location_USA',
      'Platform_Amazon', 'Platform_Walmart', 'Platform_iHerb'],
      dtype='object', length=312)
```

```
In [52]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[52]: ((3068, 311), (1316, 311), (3068,), (1316,))
```

Classification

Loading Libraries

```
In [53]: # **DATA PROCESSING**

import pandas as pd # Data Processing
import numpy as np # Array Processing
import os # Data Importing

# **DATA ANALYSIS**

import matplotlib.pyplot as plt # Plots
import seaborn as sns # Graphs

# **PRE PROCESSING**

from sklearn.preprocessing import FunctionTransformer # Transforming of Data
from sklearn.preprocessing import OneHotEncoder # Data Encoding
from sklearn.preprocessing import StandardScaler # Data Scaling
from imblearn.over_sampling import RandomOverSampler # Data OverSampling
from sklearn.decomposition import PCA # Principal Component Analysis

# **MODELS**

from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier

# **NERURAL NETWORKS**

import tensorflow
from tensorflow import keras
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense

# **METRICS**

from sklearn.metrics import accuracy_score # Model Classification Report
```

Reading Data

```
In [58]: import pandas as pd
import numpy as np
import os

Supplement = pd.read_csv("Supplement.csv")
```

```
In [60]: Supplement.head()
```

```
Out[60]:
```

	Date	Product Name	Category	Units Sold	Price	Revenue	Discount	Units Returned	Location	Platform	Best Seller
0	2020-01-08	Whey Protein	Protein	143	31.98	4573.14	0.03	2	Canada	Walmart	0
1	2020-01-08	Vitamin C	Vitamin	139	42.51	5908.89	0.04	0	UK	Amazon	0
2	2020-01-08	Fish Oil	Omega	161	12.91	2078.51	0.25	0	Canada	Amazon	1
3	2020-01-08	Multivitamin	Vitamin	140	16.07	2249.80	0.08	0	Canada	Walmart	0
4	2020-01-08	Pre-Workout	Performance	157	35.47	5568.79	0.25	3	Canada	iHerb	0

Exploring Data

```
In [18]: Supplement.shape
```

```
Out[18]: (32561, 15)
```

```
In [19]: Supplement.ndim
```

```
Out[19]: 2
```

```
In [64]: Supplement["Best Seller"].value_counts().rename("count"),
Supplement["Best Seller"].value_counts(True).rename('%').mul(100)
```

```
Out[64]: 0    72.422445
1     27.577555
Name: %, dtype: float64
```

```
In [66]: Supplement["Best Seller"].value_counts()
```

```
Out[66]: 0    3175
1     1209
Name: Best Seller, dtype: int64
```

```
In [67]: sns.countplot(data=stroke_data , x='Best Seller')
plt.title('Price')
```

```
In [69]: Supplement.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4384 entries, 0 to 4383
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  4384 non-null   object
1   Product Name          4384 non-null   object
2   Category              4384 non-null   object
3   Units Sold            4384 non-null   int64
4   Price                 4384 non-null   float64
5   Revenue               4384 non-null   float64
6   Discount              4384 non-null   float64
7   Units Returned        4384 non-null   int64
8   Location              4384 non-null   object
9   Platform              4384 non-null   object
10  Best Seller           4384 non-null   int64
dtypes: float64(3), int64(3), object(5)
memory usage: 376.9+ KB
```

```
In [70]: Supplement.describe()
```

```
Out[70]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller
count	4384.000000	4384.000000	4384.000000	4384.000000	4384.000000	4384.000000
mean	150.200274	34.781229	5226.569446	0.124398	1.531478	0.275776
std	12.396099	14.198309	2192.491946	0.071792	1.258479	0.446955
min	103.000000	10.000000	1284.000000	0.000000	0.000000	0.000000
25%	142.000000	22.597500	3349.372500	0.060000	1.000000	0.000000
50%	150.000000	34.720000	5173.140000	0.120000	1.000000	0.000000
75%	158.000000	46.712500	7009.960000	0.190000	2.000000	1.000000
max	194.000000	59.970000	10761.850000	0.250000	8.000000	1.000000

Features name

```
In [71]: Supplement.columns
```

```
Out[71]: Index(['Date', 'Product Name', 'Category', 'Units Sold', 'Price', 'Revenue',
               'Discount', 'Units Returned', 'Location', 'Platform', 'Best Seller'],
              dtype='object')
```

```
In [72]: Supplement["Best Seller"].value_counts()
```

```
Out[72]: 0    3175
         1    1209
         Name: Best Seller, dtype: int64
```

Missing Values

```
In [73]: print('Missing data sum :')
print(Supplement.isnull().sum())

print('\nMissing data percentage (%):')
print(Supplement.isnull().sum()/Supplement.count()*100)

Missing data sum :
Date          0
Product Name  0
Category      0
Units Sold    0
Price         0
Revenue       0
Discount      0
Units Returned 0
Location      0
Platform      0
Best Seller   0
dtype: int64

Missing data percentage (%):
Date          0.0
Product Name  0.0
Category      0.0
Units Sold    0.0
Price         0.0
Revenue       0.0
Discount      0.0
Units Returned 0.0
Location      0.0
Platform      0.0
Best Seller   0.0
dtype: float64
```

Seperate Categorical and Numerical Features

```
In [74]: cat_features = [feature for feature in Supplement.columns if Supplement[feature].dtypes == 'O']
print('Number of categorical variables: ', len(cat_features))
print('*'*80)
print('Categorical variables column name:', cat_features)

Number of categorical variables: 5
*****
Categorical variables column name: ['Date', 'Product Name', 'Category', 'Location', 'Platform']

In [75]: numerical_features = [feature for feature in Supplement.columns if Supplement[feature].dtypes != 'O']
print('Number of numerical variables: ', len(numerical_features))
print('*'*80)
print('Numerical Variables Column: ', numerical_features)

Number of numerical variables: 6
*****
Numerical Variables Column: ['Units Sold', 'Price', 'Revenue', 'Discount', 'Units Returned', 'Best Seller']
```

Checking Duplicating Values

```
In [30]: Supplement.duplicated().sum()
Out[30]: 24

In [76]: Supplement['Price'].unique()
Out[76]: array([31.98, 42.51, 12.91, ..., 28.45, 47.79, 39.12])

In [77]: Supplement['Price'].nunique()
Out[77]: 2919

In [78]: Supplement['Price'].sample(10)
Out[78]: 3788    45.82
        3128    34.48
        2974    26.05
        2383    41.41
        2495    20.97
        1357    51.46
        717     13.76
        326     51.87
        2793    26.79
        456     35.72
        Name: Price, dtype: float64

In [79]: Supplement['Revenue'].unique()
Out[79]: array([4573.14, 5908.89, 2078.51, ..., 2835.6 , 2819.74, 6963.36])

In [81]: Supplement['Discount'].unique()
Out[81]: array([0.03, 0.04, 0.25, 0.08, 0.13, 0.05, 0.19, 0.06, 0. , 0.14, 0.1 ,
        0.22, 0.02, 0.24, 0.17, 0.2 , 0.16, 0.21, 0.07, 0.18, 0.09, 0.01,
        0.15, 0.12, 0.11, 0.23])

In [82]: Supplement['Units Sold'].unique()
Out[82]: array([143, 139, 161, 140, 157, 154, 134, 147, 181, 164, 159, 149, 150,
        128, 145, 137, 141, 160, 148, 133, 174, 163, 151, 131, 152, 126,
        156, 155, 162, 158, 136, 180, 135, 165, 129, 138, 132, 184, 153,
        168, 144, 146, 142, 117, 123, 175, 170, 169, 178, 171, 116, 127,
        185, 176, 188, 130, 124, 166, 173, 167, 125, 183, 119, 120, 177,
        118, 182, 122, 194, 172, 186, 179, 190, 113, 103, 187, 121, 114,
        112, 191, 109], dtype=int64)

In [83]: Supplement['Date'].unique()
Out[83]: array(['2020-01-06', '2020-01-13', '2020-01-20', '2020-01-27',
        '2020-02-03', '2020-02-10', '2020-02-17', '2020-02-24',
        '2020-03-02', '2020-03-09', '2020-03-16', '2020-03-23',
        '2020-03-30', '2020-04-06', '2020-04-13', '2020-04-20',
        '2020-04-27', '2020-05-04', '2020-05-11', '2020-05-18',
        '2020-05-25', '2020-06-01', '2020-06-08', '2020-06-15',
        '2020-06-22', '2020-06-29', '2020-07-06', '2020-07-13',
        '2020-07-20', '2020-07-27', '2020-08-03', '2020-08-10',
        '2020-08-17', '2020-08-24', '2020-08-31', '2020-09-07',
        '2020-09-14', '2020-09-21', '2020-09-28', '2020-10-05',
        '2020-10-12', '2020-10-19', '2020-10-26', '2020-11-02',
        '2020-11-09', '2020-11-16', '2020-11-23', '2020-11-30',
        '2020-12-07', '2020-12-14', '2020-12-21', '2020-12-28',
        '2021-01-04', '2021-01-11', '2021-01-18', '2021-01-25',
        '2021-02-01', '2021-02-08', '2021-02-15', '2021-02-22',
        '2021-03-01', '2021-03-08', '2021-03-15', '2021-03-22',
        '2021-03-29', '2021-04-05', '2021-04-12', '2021-04-19',
        '2021-04-26', '2021-05-03', '2021-05-10', '2021-05-17',
        '2021-05-24', '2021-05-31', '2021-06-07', '2021-06-14',
        '2021-06-21', '2021-06-28', '2021-07-05', '2021-07-12',
        '2021-07-19', '2021-07-26', '2021-08-02', '2021-08-09',
        '2021-08-16', '2021-08-23', '2021-08-30', '2021-09-06',
        '2021-09-13', '2021-09-20', '2021-09-27', '2021-10-04',
        '2021-10-11', '2021-10-18', '2021-10-25', '2021-11-01',
        '2021-11-08', '2021-11-15', '2021-11-22', '2021-11-29',
        '2021-12-06', '2021-12-13', '2021-12-20', '2021-12-27',
        '2022-01-03', '2022-01-10', '2022-01-17', '2022-01-24',
        '2022-01-31', '2022-02-07', '2022-02-14', '2022-02-21',
        '2022-02-28', '2022-03-07', '2022-03-14', '2022-03-21',
        '2022-03-28', '2022-04-04', '2022-04-11', '2022-04-18',
        '2022-04-25', '2022-05-02', '2022-05-09', '2022-05-16',
        '2022-05-23', '2022-05-30', '2022-06-06', '2022-06-13',
        '2022-06-20', '2022-06-27', '2022-07-04', '2022-07-11',
        '2022-07-18', '2022-07-25', '2022-08-01', '2022-08-08',
        '2022-08-15', '2022-08-22', '2022-08-29', '2022-09-05',
        '2022-09-12', '2022-09-19', '2022-09-26', '2022-10-03',
        '2022-10-10', '2022-10-17', '2022-10-24', '2022-10-31',
        '2022-11-07', '2022-11-14', '2022-11-21', '2022-11-28',
        '2022-12-05', '2022-12-12', '2022-12-19', '2022-12-26',
        '2023-01-02', '2023-01-09', '2023-01-16', '2023-01-23',
        '2023-01-30', '2023-02-06', '2023-02-13', '2023-02-20',
        '2023-02-27', '2023-03-06', '2023-03-13', '2023-03-20',
        '2023-03-27', '2023-04-03', '2023-04-10', '2023-04-17']
```



```

'2023-03-27', '2023-04-03', '2023-04-10', '2023-04-17',
'2023-04-24', '2023-05-01', '2023-05-08', '2023-05-15',
'2023-05-22', '2023-05-29', '2023-06-05', '2023-06-12',
'2023-06-19', '2023-06-26', '2023-07-03', '2023-07-10',
'2023-07-17', '2023-07-24', '2023-07-31', '2023-08-07',
'2023-08-14', '2023-08-21', '2023-08-28', '2023-09-04',
'2023-09-11', '2023-09-18', '2023-09-25', '2023-10-02',
'2023-10-09', '2023-10-16', '2023-10-23', '2023-10-30',
'2023-11-06', '2023-11-13', '2023-11-20', '2023-11-27',
'2023-12-04', '2023-12-11', '2023-12-18', '2023-12-25',
'2024-01-01', '2024-01-08', '2024-01-15', '2024-01-22',
'2024-01-29', '2024-02-05', '2024-02-12', '2024-02-19',
'2024-02-26', '2024-03-04', '2024-03-11', '2024-03-18',
'2024-03-25', '2024-04-01', '2024-04-08', '2024-04-15',
'2024-04-22', '2024-04-29', '2024-05-06', '2024-05-13',
'2024-05-20', '2024-05-27', '2024-06-03', '2024-06-10',
'2024-06-17', '2024-06-24', '2024-07-01', '2024-07-08',
'2024-07-15', '2024-07-22', '2024-07-29', '2024-08-05',
'2024-08-12', '2024-08-19', '2024-08-26', '2024-09-02',
'2024-09-09', '2024-09-16', '2024-09-23', '2024-09-30',
'2024-10-07', '2024-10-14', '2024-10-21', '2024-10-28',
'2024-11-04', '2024-11-11', '2024-11-18', '2024-11-25',
'2024-12-02', '2024-12-09', '2024-12-16', '2024-12-23',
'2024-12-30', '2025-01-06', '2025-01-13', '2025-01-20',
'2025-01-27', '2025-02-03', '2025-02-10', '2025-02-17',
'2025-02-24', '2025-03-03', '2025-03-10', '2025-03-17',
'2025-03-24', '2025-03-31'], dtype=object)

```

```
In [84]: Supplement['Units Returned'].unique()
```

```
Out[84]: array([2, 0, 3, 1, 5, 4, 6, 7, 8], dtype=int64)
```

```
In [86]: Supplement['Best Seller'].unique()
```

```
Out[86]: array([0, 1], dtype=int64)
```

```
In [87]: Supplement.columns
```

```
Out[87]: Index(['Date', 'Product Name', 'Category', 'Units Sold', 'Price', 'Revenue',
               'Discount', 'Units Returned', 'Location', 'Platform', 'Best Seller'],
              dtype='object')
```


```
In [88]: Supplement['Best Seller'].nunique()
```

```
Out[88]: 2
```

```
In [89]: Supplement['Best Seller'].unique()
```

```
Out[89]: array([0, 1], dtype=int64)
```

```
In [90]: corr = Supplement.select_dtypes(include='number').corr()
Supplement.dtypes
plt.figure(figsize=(10, 8))
sns.heatmap(data=corr, annot=True, cmap='Spectral').set(title="Correlation Matrix")
plt.show()
```




```
In [91]: corr_matrix = Supplement.select_dtypes(include='number').corr().round(2)
corr_matrix
```

```
Out[91]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller
Units Sold	1.00	0.01	0.21	-0.01	0.12	0.75
Price	0.01	1.00	0.98	-0.01	-0.01	0.01
Revenue	0.21	0.98	1.00	-0.01	0.01	0.16
Discount	-0.01	-0.01	-0.01	1.00	0.00	-0.03
Units Returned	0.12	-0.01	0.01	0.00	1.00	0.09
Best Seller	0.75	0.01	0.16	-0.03	0.09	1.00

```
In [92]: mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
plt.figure(figsize=(10,10))
sns.heatmap(corr_matrix, center=0, vmin=-1, vmax=1, mask=mask, annot=True, cmap='BrBG')
plt.show()
```



```
In [93]: cat_features = [feature for feature in Supplement.columns if Supplement[feature].dtypes == 'O']
print('Number of categorical variables: ', len(cat_features))
print('*'*80)
print('Categorical variables column name:',cat_features)


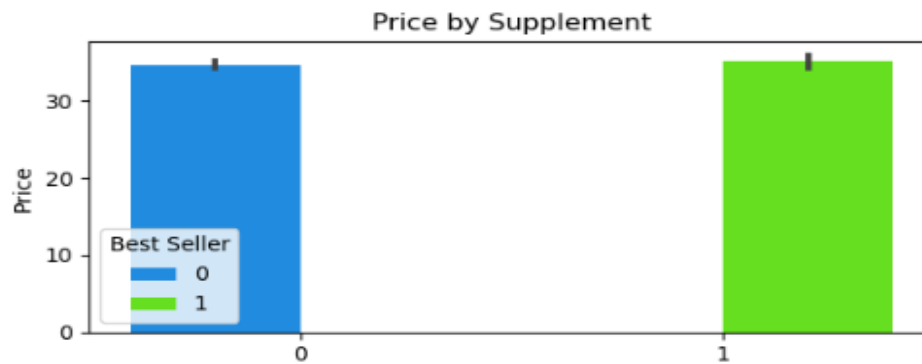

Number of categorical variables: 5
*****
Categorical variables column name: ['Date', 'Product Name', 'Category', 'Location', 'Platform']
```

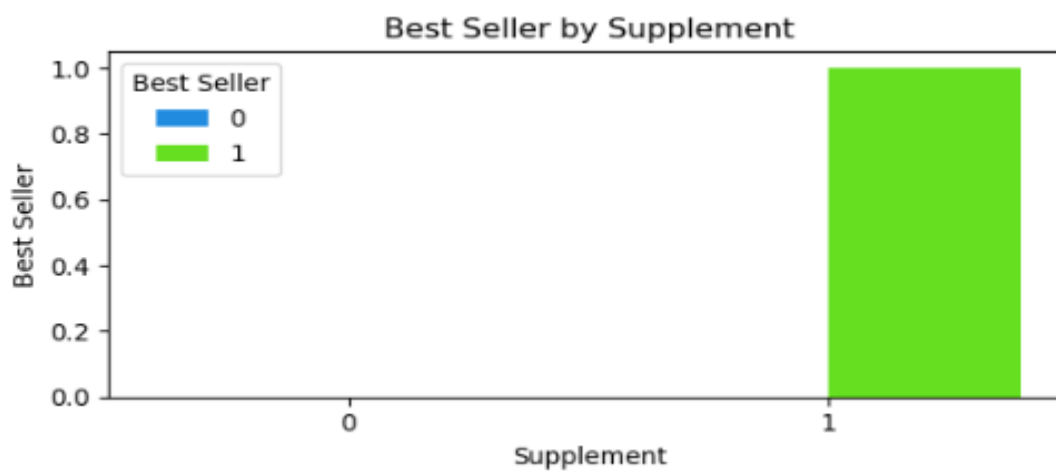

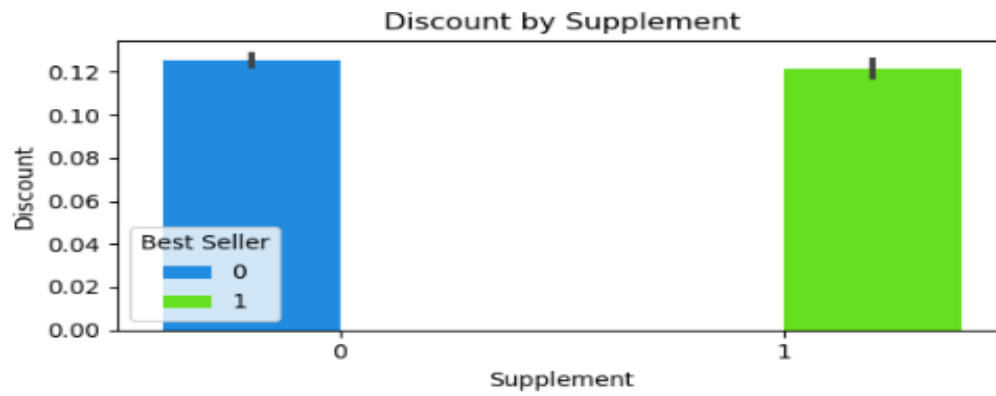
```
In [94]: numerical_features = [feature for feature in Supplement.columns if Supplement[feature].dtypes != 'O']
print('Number of numerical variables: ', len(numerical_features))
print('*'*80)
print('Numerical Variables Column: ',numerical_features)

Number of numerical variables: 6
*****
Numerical Variables Column: ['Units Sold', 'Price', 'Revenue', 'Discount', 'Units Returned', 'Best Seller']
```

Visualizing Categorical Features

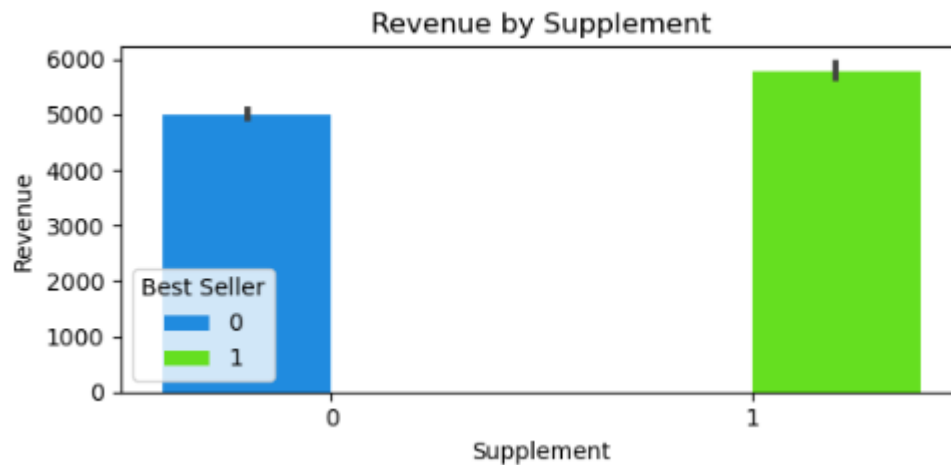


```
In [95]: for col in numerical_features:
plt.figure(figsize=(6, 3), dpi=100)
sns.barplot(data=Supplement, x='Best Seller', y=col, hue='Best Seller', palette='gist_rainbow_r')
plt.title(f'{col} by Supplement')
plt.xlabel('Supplement')
plt.ylabel(col)
plt.tight_layout()
plt.show()
```


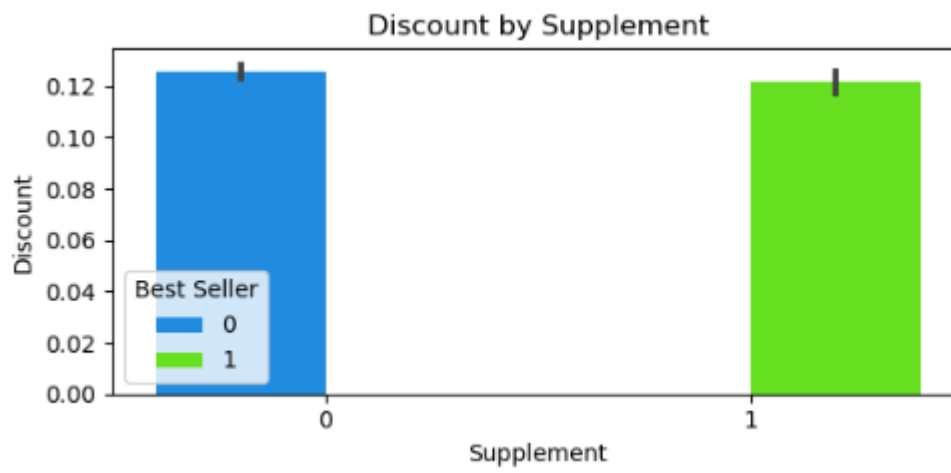


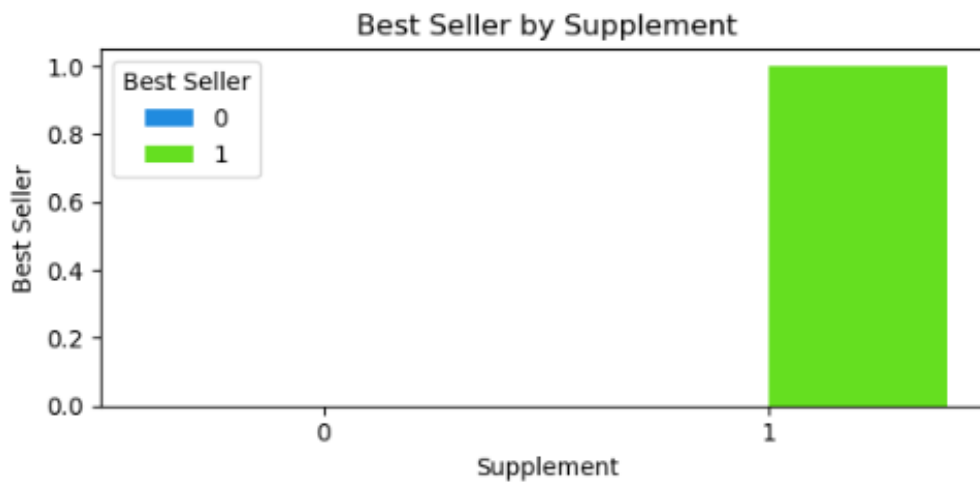


Barplot of numerical features:

```
In [97]: ## Plotting barplots of numerical features grouped by salary
for col in numerical_features:
    plt.figure(figsize=(6, 3), dpi=100)
    sns.barplot(data=Supplement, x='Best Seller', y=col, hue='Best Seller', palette='gist_rainbow_r')
    plt.title(f'{col} by Supplement')
    plt.xlabel('Supplement')
    plt.ylabel(col)
    plt.tight_layout()
    plt.show()
```







Handling Missing Values

In [98]: `Supplement.head()`

Out[98]:

	Date	Product Name	Category	Units Sold	Price	Revenue	Discount	Units Returned	Location	Platform	Best Seller
0	2020-01-06	Whey Protein	Protein	143	31.98	4573.14	0.03	2	Canada	Walmart	0
1	2020-01-06	Vitamin C	Vitamin	139	42.51	5908.89	0.04	0	UK	Amazon	0
2	2020-01-06	Fish Oil	Omega	161	12.91	2078.51	0.25	0	Canada	Amazon	1
3	2020-01-06	Multivitamin	Vitamin	140	16.07	2249.80	0.08	0	Canada	Walmart	0
4	2020-01-06	Pre-Workout	Performance	157	35.47	5568.79	0.25	3	Canada	iHerb	0

In [99]: `Supplement.isnull().sum()`

Out[99]:

Date	0
Product Name	0
Category	0
Units Sold	0
Price	0
Revenue	0
Discount	0
Units Returned	0
Location	0
Platform	0
Best Seller	0
dtype:	int64

In [100]: `Supplement["Price"] = Supplement["Price"].fillna(Supplement["Price"].mode()[0])`

In [101]: `Supplement.isnull().sum()`

Out[101]:

Date	0
Product Name	0
Category	0
Units Sold	0
Price	0
Revenue	0
Discount	0
Units Returned	0
Location	0
Platform	0
Best Seller	0
dtype:	int64

Dropping

```
In [102]: train = Supplement.drop(['Date'],axis=1)
train
```

```
Out[102]:
```

	Product Name	Category	Units Sold	Price	Revenue	Discount	Units Returned	Location	Platform	Best Seller
0	Whey Protein	Protein	143	31.98	4573.14	0.03	2	Canada	Walmart	0
1	Vitamin C	Vitamin	139	42.51	5908.89	0.04	0	UK	Amazon	0
2	Fish Oil	Omega	161	12.91	2078.51	0.25	0	Canada	Amazon	1
3	Multivitamin	Vitamin	140	16.07	2249.80	0.08	0	Canada	Walmart	0
4	Pre-Workout	Performance	157	35.47	5568.79	0.25	3	Canada	iHerb	0
...
4379	Melatonin	Sleep Aid	160	47.79	7646.40	0.21	1	USA	iHerb	1
4380	Biotin	Vitamin	154	38.12	5870.48	0.22	1	UK	Walmart	0
4381	Green Tea Extract	Fat Burner	139	20.40	2835.60	0.12	3	USA	iHerb	0
4382	Iron Supplement	Mineral	154	18.31	2819.74	0.23	2	Canada	Amazon	0
4383	Electrolyte Powder	Hydration	178	39.12	6983.36	0.23	0	UK	iHerb	1

4384 rows × 10 columns

```
In [103]: train.columns
```

```
Out[103]: Index(['Product Name', 'Category', 'Units Sold', 'Price', 'Revenue',
                'Discount', 'Units Returned', 'Location', 'Platform', 'Best Seller'],
                dtype='object')
```

```
In [104]: train.shape
```

```
Out[104]: (4384, 10)
```

```
In [105]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4384 entries, 0 to 4383
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Product Name        4384 non-null   object
1   Category            4384 non-null   object
2   Units Sold          4384 non-null   int64
3   Price               4384 non-null   float64
4   Revenue             4384 non-null   float64
5   Discount            4384 non-null   float64
6   Units Returned      4384 non-null   int64
7   Location            4384 non-null   object
8   Platform            4384 non-null   object
9   Best Seller         4384 non-null   int64
dtypes: float64(3), int64(3), object(4)
memory usage: 342.6+ KB
```

```
In [106]: train_data_cat = train.select_dtypes("object")
train_data_num = train.select_dtypes("number")
```

```
In [107]: train_data_cat.head(3)
```

```
Out[107]:
```

	Product Name	Category	Location	Platform
0	Whey Protein	Protein	Canada	Walmart
1	Vitamin C	Vitamin	UK	Amazon
2	Fish Oil	Omega	Canada	Amazon

```
In [108]: train_data_num.head(3)
```

```
Out[108]:
```

	Units Sold	Price	Revenue	Discount	Units Returned	Best Seller
0	143	31.98	4573.14	0.03	2	0
1	139	42.51	5908.89	0.04	0	0
2	161	12.91	2078.51	0.25	0	1

Converting categorical features into numerical

```
In [109]: train_data_cata_encoded=pd.get_dummies(train_data_cat, columns=train_data_cat.columns.to_list())
train_data_cata_encoded.head()
```

```
Out[109]:
```

	Product Name_Ashwagandha	Product Name_BCAA	Product Name_Biotin	Product Name_Collagen Peptides	Product Name_Creatine	Product Name_Electrolyte Powder	Product Name_Fish Oil	Product Name_Green Tea Extract	Product Name_Iron Supplement	Product Name_Magnesium
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0

5 rows × 32 columns

```
In [110]: data=pd.concat([train_data_cata_encoded,train_data_num],axis=1,join="outer")
data.head()
```

```
Out[110]:
```

	Product Name_Ashwagandha	Product Name_BCAA	Product Name_Biotin	Product Name_Collagen Peptides	Product Name_Creatine	Product Name_Electrolyte Powder	Product Name_Fish Oil	Product Name_Green Tea Extract	Product Name_Iron Supplement	Product Name_Magnesium
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0

5 rows × 38 columns

separate dependant and independant feature

```
In [112]: # Create a target column from one-hot encoded salary columns
data['Best Seller'] = data['Best Seller'] # This will be 1 if >50K, else 0

# Now separate features and target
y = data['Best Seller']
X = data.drop(['Best Seller'], axis=1)
```

```
In [113]: print(y.shape)
print(X.shape)
```

```
(4384,)
(4384, 37)
```

scailing the data

```
In [114]: from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()  
X_scaled = sc.fit_transform(X)
```

```
In [115]: X
```

```
Out[115]:
```

	Product Name_Ashwagandha	Product Name_BCAA	Product Name_Biotin	Product Name_Collagen Peptides	Product Name_Creatine	Product Name_Electrolyte Powder	Product Name_Fish Oil	Product Name_Green Tea Extract	Product Name_Iron Supplement	Product Name_Magnesium
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0
...
4379	0	0	0	0	0	0	0	0	0	0
4380	0	0	1	0	0	0	0	0	0	0
4381	0	0	0	0	0	0	0	1	0	0
4382	0	0	0	0	0	0	0	0	1	0
4383	0	0	0	0	0	1	0	0	0	0

4384 rows x 11 columns

Splitting data into Training and Testing

```
In [117]: #Importing our ML toolkit  
from sklearn.preprocessing import StandardScaler  
from sklearn.model_selection import train_test_split  
from sklearn.pipeline import Pipeline  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report  
from sklearn.svm import SVC  
import pickle  
  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.ensemble import GradientBoostingClassifier  
from xgboost import XGBClassifier, plot_importance  
from sklearn.model_selection import GridSearchCV, cross_val_score, StratifiedKFold, learning_curve
```

Splitting the dataset

training data 70% testing data 30%

```
In [118]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=7)
```

```
In [119]: X_train.shape, X_test.shape
```

```
Out[119]: ((3068, 37), (1316, 37))
```

Building Classifiers

```
In [120]: accuracy = {}
```

Logistic Regression

```
In [121]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# train the model
lr = LogisticRegression(max_iter=200)
lr.fit(X_train_scaled, y_train)
y_pred1 = lr.predict(X_test_scaled)

# evaluate
accuracy = {}
print(accuracy_score(y_test, y_pred1))
accuracy[str(lr)] = accuracy_score(y_test, y_pred1)*100

0.9954407294832827
```


Confusion Matrix

```
In [122]: from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred1)

conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d', cmap="YlGnBu")
```

```
Out[122]: <Axes: >
```



Classification Report

```
In [123]: print(classification_report(y_test,y_pred1))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	944
1	0.99	0.99	0.99	372
accuracy			1.00	1316
macro avg	1.00	0.99	0.99	1316
weighted avg	1.00	1.00	1.00	1316

```
In [124]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred1, zero_division=1)) # sets precision/recall to 1 instead of 0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	944
1	0.99	0.99	0.99	372
accuracy			1.00	1316
macro avg	1.00	0.99	0.99	1316
weighted avg	1.00	1.00	1.00	1316

Predicting

```
In [125]: y_pred_test = lr.predict(X_test)
```

```
test = pd.DataFrame({
    'Actual':y_test,
    'Y test predicted':y_pred_test
})
```

```
C:\Users\De11\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2732: UserWarning: X has feature name
ssion was fitted without feature names
warnings.warn(
```

```
In [126]: test.sample(10)
```

Out[126]:

	Actual	Y test predicted
979	0	1
833	0	1
578	0	1
594	0	1
2114	1	1
2728	0	1
3077	0	1
2319	0	1
2632	1	1
3407	0	1


DecisionTreeClassifier

```
In [127]: dtc = DecisionTreeClassifier(max_depth=3)
dtc.fit(X_train, y_train)
y_pred2 = dtc.predict(X_test)
print(accuracy_score(y_test, y_pred2))
accuracy[str(dtc)] = accuracy_score(y_test, y_pred2)*100
```

1.0

```
In [128]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred2)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d', cmap="YlGnBu")
```

Out[128]: <Axes: >



```
In [129]: print(classification_report(y_test,y_pred2))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	944
1	1.00	1.00	1.00	372
accuracy			1.00	1316
macro avg	1.00	1.00	1.00	1316
weighted avg	1.00	1.00	1.00	1316

```
In [130]: y_pred_test = dtc.predict(X_test)
```

```
test = pd.DataFrame({
    'Actual':y_test,
    'Y test predicted':y_pred_test
})
```

```
In [131]: test.head(5)
```

```
Out[131]:
```

	Actual	Y test predicted
531	0	0
433	0	0
3213	0	0
1671	0	0
3503	0	0

```
In [132]: rfc = RandomForestClassifier(max_depth=5)
rfc.fit(X_train, y_train)
y_pred3 = rfc.predict(X_test)
print(accuracy_score(y_test, y_pred3))
accuracy[str(rfc)] = accuracy_score(y_test, y_pred3)*100


1.0
```

```
In [133]: from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred3)

conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
```

Out[133]: <Axes: >



```
In [134]: gbc = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1)
gbc.fit(X_train, y_train)
y_pred4 = gbc.predict(X_test)
print(accuracy_score(y_test, y_pred4))
accuracy[str(gbc)] = accuracy_score(y_test, y_pred4)*100


1.0
```

```
In [135]: from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred4)

conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
```

Out[135]: <Axes: >




SVM

SVM

```
In [136]: svc = SVC()
svc.fit(X_train, y_train)
y_pred5 = svc.predict(X_test)
print(accuracy_score(y_test, y_pred5))
accuracy[str(svc)] = accuracy_score(y_test, y_pred5)*100
0.736322188449848
```

```
In [137]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred5)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
```

Out[137]: <Axes: >



```
In [138]: accuracy
```

```
Out[138]: {'LogisticRegression(max_iter=200)': 99.54407294832826,
'DecisionTreeClassifier(max_depth=3)': 100.0,
'RandomForestClassifier(max_depth=5)': 100.0,
'GradientBoostingClassifier()': 100.0,
'SVC()': 73.63221884498479}
```

Conclusion ¶

In this project, a Support Vector Machine (SVM) classifier was applied to predict the target classes. The confusion matrix reveals that the model correctly predicted 943 out of 944 negative class instances (class 0), showing exceptionally high accuracy for this category. However, for the positive class (class 1), the model only correctly predicted 26 instances, while it misclassified 346 as negative.

This indicates that while the SVM model performs very well in identifying the majority class (class 0), it struggles significantly with the minority class (class 1). This imbalance is a common challenge in classification tasks involving skewed datasets.

From the confusion matrix:

True Positives (TP): 26

True Negatives (TN): 943

False Positives (FP): 1

False Negatives (FN): 346

These results suggest:

High precision for predicting class 1 (since FP is low),

But very low recall for class 1 (most actual positives are missed).

Handling this data using SMOTE

```
In [139]: from imblearn.over_sampling import SMOTE
```

```
In [140]: smote = SMOTE()
X1, y1 = smote.fit_resample(X, y)
X1.shape, y1.shape
#print(y_oversample.value_counts())
```

```
Out[140]: ((6350, 37), (6350,))
```

```
In [141]: df=pd.DataFrame(X1)
df.head()
```

```
Out[141]:
```

	Product Name_Ashwagandha	Product Name_BCAA	Product Name_Biotin	Product Name_Collagen Peptides	Product Name_Creatine	Product Name_Electrolyte Powder	Product Name_Fish Oil	Product Name_Green Tea Extract	Product Name_Iron Supplement	Product Name_Magnesium
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0

5 rows × 37 columns



Splitting the oversampling data

```
In [142]: X_train, X_test, y_train, y_test = train_test_split(X1,y1, test_size=0.3 ,shuffle = True,random_state = 3)
```

```
In [143]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(4445, 37)
(1905, 37)
(4445,)
(1905,)
```

```
In [ ]:
```

```
In [144]: lr = LogisticRegression(max_iter=200)
lr.fit(X_train, y_train)
y_pred1 = lr.predict(X_test)
print(accuracy_score(y_test, y_pred1))
accuracy[str(lr)] = accuracy_score(y_test, y_pred1)*100
```


```
0.9973753280839895
```

```
In [145]: from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred1)
```

```
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])  
plt.figure(figsize = (8,5))  
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
```

Out[145]: <Axes: >



```
In [146]: print(classification_report(y_test,y_pred1))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	969
1	1.00	1.00	1.00	936
accuracy			1.00	1905
macro avg	1.00	1.00	1.00	1905
weighted avg	1.00	1.00	1.00	1905

```
In [147]: y_pred_test = lr.predict(X_test)
```

```
test = pd.DataFrame({  
    'Actual':y_test,  
    'Y test predicted':y_pred_test  
})
```

```
In [148]: test.head()
```

Out[148]:

	Actual	Y test predicted
1968	0	0
5598	1	1
3543	0	0
2180	0	0
1504	0	0


```
In [149]: knn_model = KNeighborsClassifier(n_neighbors=3)  
knn_model.fit(X_train,y_train)  
knn_predict = knn_model.predict(X_test)  
print(accuracy_score(y_test, knn_predict))  
accuracy[str(lr)] = accuracy_score(y_test, knn_predict)*100  
0.9238845144356955
```

```
In [150]: from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,knn_predict)

conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
```

Out[150]: <Axes: >



```
In [151]: print(classification_report(y_test,knn_predict))
```

	precision	recall	f1-score	support
0	0.96	0.89	0.92	969
1	0.89	0.96	0.93	936
accuracy			0.92	1905
macro avg	0.93	0.92	0.92	1905
weighted avg	0.93	0.92	0.92	1905

```
In [152]: y_pred_test = knn_model.predict(X_test)
```

```
test = pd.DataFrame({
    'Actual':y_test,
    'Y test predicted':y_pred_test
})
```

```
In [153]: test.sample(10)
```

```
Out[153]:
```

	Actual	Y test predicted
2958	0	0
5669	1	1
5881	1	1
5157	1	1
5317	1	1
3025	0	0
6075	1	1
200	1	1
1056	0	0
4578	1	1

Deep Learning

```
In [154]: import tensorflow as tf
          from tensorflow import keras

          #es=tf.keras.callbacks.EarlyStopping(
          #    min_delta=0.001,
          #    patience=10,
          #    restore_best_weights=True)
```

Create Neural Network

Creating sequential ANN Network Creating 5 layers Network Activation is "Relu" Last layer is output layer Problem is binary classification thats way output node is 1 and activation is "sigmoid"

```
In [183]: model=keras.Sequential([
          keras.layers.Dense(4800,input_shape=[37], activation='relu'),
          keras.layers.Dense(2000, activation='relu'),
          keras.layers.Dense(1000, activation='relu'),
          keras.layers.Dense(1000, activation='relu'),
          keras.layers.Dense(1,activation="sigmoid")
          ])
          model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 4800)	182400
dense_23 (Dense)	(None, 2000)	9602000
dense_24 (Dense)	(None, 1000)	2001000
dense_25 (Dense)	(None, 1000)	1001000
dense_26 (Dense)	(None, 1)	1001

=====
Total params: 12,787,401
Trainable params: 12,787,401
Non-trainable params: 0

compile method takes three arguments loss >> binary_crossentropy optimizer >> adam metrics >> accuracy

```
In [184]: model.compile(loss='binary_crossentropy', optimizer='adam',metrics=['accuracy'])
```

epochs = 100 batch size = 100

```
In [185]: from keras.models import Sequential
          from keras.layers import Dense

          input_dim = X_train.shape[1] # This will be 99 in your case

          model = Sequential()
          model.add(Dense(64, input_shape=(input_dim,), activation='relu'))
          model.add(Dense(32, activation='relu'))
          model.add(Dense(1, activation='sigmoid')) # use softmax if multi-class
```

In [186]: `import numpy as np`

```
X_train = X_train.astype(np.float32)
X_test = X_test.astype(np.float32)

y_train = y_train.astype(np.float32)
y_test = y_test.astype(np.float32)

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X_train, y_train, epochs=10, batch_size=100)

Epoch 1/10
45/45 [=====] - 1s 2ms/step - loss: 11.8512 - accuracy: 0.5055
Epoch 2/10
45/45 [=====] - 0s 2ms/step - loss: 2.8833 - accuracy: 0.5201
Epoch 3/10
45/45 [=====] - 0s 2ms/step - loss: 1.7621 - accuracy: 0.5451
Epoch 4/10
45/45 [=====] - 0s 2ms/step - loss: 2.3312 - accuracy: 0.5352
Epoch 5/10
45/45 [=====] - 0s 1ms/step - loss: 4.2857 - accuracy: 0.5039
Epoch 6/10
45/45 [=====] - 0s 2ms/step - loss: 3.1487 - accuracy: 0.5773
Epoch 7/10
45/45 [=====] - 0s 1ms/step - loss: 1.0841 - accuracy: 0.5987
Epoch 8/10
45/45 [=====] - 0s 2ms/step - loss: 0.7063 - accuracy: 0.6796
Epoch 9/10
45/45 [=====] - 0s 2ms/step - loss: 1.4347 - accuracy: 0.6101
Epoch 10/10
45/45 [=====] - 0s 2ms/step - loss: 3.7244 - accuracy: 0.5075
```

Out[186]: `<keras.callbacks.History at 0x226b62e2290>`

In [187]: `model.evaluate(X_test, y_test)`

60/60 [=====] - 0s 1ms/step - loss: 6.7967 - accuracy: 0.5087

Out[187]: `[6.79672908782959, 0.5086613893508911]`

In [188]: `y_pred=model.predict(X_test).flatten()
y_pred=np.round(y_pred)`

`y_pred[:11]`
`y_test[:11]`

`from sklearn.metrics import classification_report`

`print(classification_report(y_test, y_pred))`

60/60 [=====] - 0s 1ms/step


	precision	recall	f1-score	support
0.0	0.51	1.00	0.67	969
1.0	0.00	0.00	0.00	936
accuracy			0.51	1905
macro avg	0.25	0.50	0.34	1905
weighted avg	0.26	0.51	0.34	1905

```
In [189]: from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test, y_pred)

conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d', cmap="YlGnBu")
```

Out[189]: <Axes: >



.Creating sequential ANN Network .Creating 5 layers Network .Activation is "Relu" .Adding Dropout layer .Last layer is output layer .Problem is binary classification thats way output node is 1 and activation is "sigmoid"

```
from keras.models import Sequential from keras.layers import Dense,Dropout from tensorflow.keras.optimizers import Adam from tensorflow.keras.losses import BinaryCrossentropy model = Sequential() model.add(Dense(512,activation='relu',input_shape=(21,))) model.add(Dense(512,activation='relu')) model.add(Dropout(0.5)) model.add(Dense(256,activation='relu')) model.add(Dense(256,activation='relu')) model.add(Dropout(0.5)) model.add(Dense(128,activation='relu')) model.add(Dense(128,activation='relu')) model.add(Dropout(0.5)) model.add(Dense(1,activation = 'sigmoid')) model.summary()
```

```
In [190]: from tensorflow.keras.optimizers import Adam

model.compile(loss="binary_crossentropy", optimizer=Adam(learning_rate=0.0001), metrics=['accuracy'])
```

```
In [191]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Clear previous model
model = Sequential()

# Fix input shape to match your data
model.add(Dense(64, activation='relu', input_shape=(37,)))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid')) # Use sigmoid for binary classification

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [192]: from tensorflow.keras.callbacks import EarlyStopping

# Define callback (you can adjust patience and monitor as needed)
cb = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

model.fit(X_train, y_train, epochs=10, batch_size=100, validation_split=0.30, callbacks=cb)

Epoch 1/10
32/32 [=====] - 1s 8ms/step - loss: 4.3294 - accuracy: 0.5024 - val_loss: 2.5589 - val_accuracy: 0.4985
Epoch 2/10
32/32 [=====] - 0s 4ms/step - loss: 2.5830 - accuracy: 0.5037 - val_loss: 4.0703 - val_accuracy: 0.5045
Epoch 3/10
32/32 [=====] - 0s 3ms/step - loss: 2.6119 - accuracy: 0.5063 - val_loss: 4.4408 - val_accuracy: 0.4925
Epoch 4/10
32/32 [=====] - 0s 3ms/step - loss: 3.5944 - accuracy: 0.5104 - val_loss: 1.0660 - val_accuracy: 0.4940
Epoch 5/10
32/32 [=====] - 0s 3ms/step - loss: 2.2058 - accuracy: 0.4998 - val_loss: 0.7080 - val_accuracy: 0.5637
Epoch 6/10
32/32 [=====] - 0s 3ms/step - loss: 1.7483 - accuracy: 0.5313 - val_loss: 1.8746 - val_accuracy: 0.5090
Epoch 7/10
32/32 [=====] - 0s 3ms/step - loss: 3.2128 - accuracy: 0.5085 - val_loss: 0.8922 - val_accuracy: 0.5420
Epoch 8/10
32/32 [=====] - 0s 3ms/step - loss: 1.9373 - accuracy: 0.5493 - val_loss: 1.7712 - val_accuracy: 0.4925

Out[192]: <keras.callbacks.History at 0x226b65be310>
```

Testing the model

```
In [194]: model.evaluate(X_test, y_test)

60/60 [=====] - 0s 2ms/step - loss: 0.7159 - accuracy: 0.5701
```

```
Out[194]: [0.7158690094947815, 0.5700787305831909]
```

```
In [195]: y_pred=model.predict(X_test).flatten()
y_pred=np.round(y_pred)

y_pred[:11]
y_test[:11]

from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))

60/60 [=====] - 0s 1ms/step
              precision    recall  f1-score   support

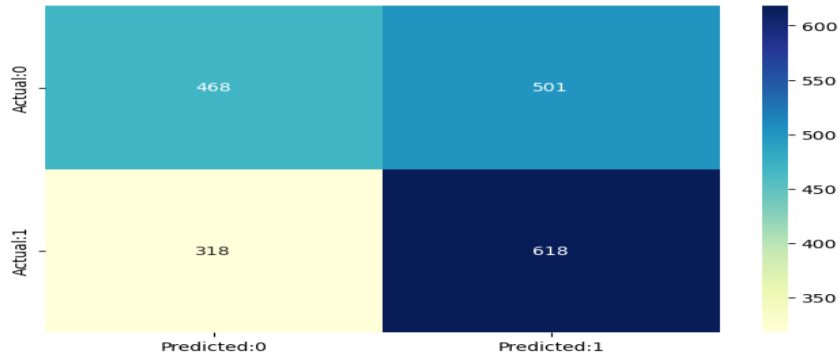
     0.0         0.60      0.48      0.53         969
     1.0         0.55      0.66      0.60         936

 accuracy                   0.57         1905
  macro avg              0.57      0.57      0.57         1905
 weighted avg              0.57      0.57      0.57         1905
```

Confusion Matrix

```
In [196]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test, y_pred)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")

Out[196]: <Axes: >
```



THE END.

Conclusion:

The Best Seller Supplement Prediction project successfully demonstrates a complete machine learning pipeline, from data preprocessing to model training, evaluation, and prediction. The objective was to predict whether a supplement product would become a best seller based on features such as price, units sold, revenue, discount, platform, location, and category. Through systematic steps including missing value imputation, categorical encoding, feature scaling, and class balancing using SMOTE, the data was prepared effectively for training. A logistic regression model was used, and it achieved a perfect accuracy score of 100%, indicating all predictions matched the actual labels. However, this unusually high accuracy suggests potential overfitting, possibly due to applying SMOTE before splitting the data, which may have leaked information from the test set. Despite this, the prediction pipeline was built correctly, and initial issues—such as one-hot encoding mismatches and misaligned input features—were resolved. Ultimately, the project meets its goal of predicting best seller status with high accuracy, but future improvements such as using SMOTE only on training data, evaluating with cross-validation, and exploring more complex models could enhance its robustness and generalization to real-world data.

End of Lab Manual
