

# The University of Faisalabad

## LAB MANUAL

**Name :**

**Muhammad Masood Bakhtiar**

**Registration No:**

**2023-BS-AI-056**

**Subject:**

**Machine Learning**

**Course Code**

**AI-414**

**Degree name:**

**BSAI-4A**

## Project 01:

### Bank customer survey using regression

#### Program:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
```

Load and Inspect the Data

```
: data = pd.read_csv(r"C:\Users\HP\Downloads\Compressed\archive_2\bank_customer_survey.csv")
```

```
: data.head()
```

|   | age | job          | marital | education | default | balance | housing | loan | contact | day | month | duration | campaign | pdays | previous | poutcome | y |
|---|-----|--------------|---------|-----------|---------|---------|---------|------|---------|-----|-------|----------|----------|-------|----------|----------|---|
| 0 | 58  | management   | married | tertiary  | no      | 2143    | yes     | no   | unknown | 5   | may   | 261      | 1        | -1    | 0        | unknown  | 0 |
| 1 | 44  | technician   | single  | secondary | no      | 29      | yes     | no   | unknown | 5   | may   | 151      | 1        | -1    | 0        | unknown  | 0 |
| 2 | 33  | entrepreneur | married | secondary | no      | 2       | yes     | yes  | unknown | 5   | may   | 76       | 1        | -1    | 0        | unknown  | 0 |
| 3 | 47  | blue         | married | unknown   | no      | 1506    | yes     | no   | unknown | 5   | may   | 92       | 1        | -1    | 0        | unknown  | 0 |
| 4 | 33  | unknown      | single  | unknown   | no      | 1       | no      | no   | unknown | 5   | may   | 198      | 1        | -1    | 0        | unknown  | 0 |

```
: print(f"Dataset shape: {data.shape}")
```

Dataset shape: (45211, 17)

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 45211 entries, 0 to 45210  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   age         45211 non-null  int64  
1   job         45211 non-null  object  
2   marital     45211 non-null  object  
3   education   45211 non-null  object  
4   default     45211 non-null  object  
5   balance     45211 non-null  int64  
6   housing     45211 non-null  object  
7   loan        45211 non-null  object  
8   contact     45211 non-null  object  
9   day         45211 non-null  int64  
10  month       45211 non-null  object  
11  duration    45211 non-null  int64  
12  campaign    45211 non-null  int64  
13  pdays      45211 non-null  int64  
14  previous    45211 non-null  int64  
15  poutcome    45211 non-null  object  
16  y           45211 non-null  int64  
dtypes: int64(8), object(9)  
memory usage: 5.9+ MB
```

```
data.describe()
```

|       | age          | balance       | day          | duration     | campaign     | pdays        | previous     | y            |
|-------|--------------|---------------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 45211.000000 | 45211.000000  | 45211.000000 | 45211.000000 | 45211.000000 | 45211.000000 | 45211.000000 | 45211.000000 |
| mean  | 40.936210    | 1362.272058   | 15.806419    | 258.163080   | 2.763841     | 40.197828    | 0.580323     | 0.116985     |
| std   | 10.618762    | 3044.765829   | 8.322476     | 257.527812   | 3.098021     | 100.128746   | 2.303441     | 0.321406     |
| min   | 18.000000    | -8019.000000  | 1.000000     | 0.000000     | 1.000000     | -1.000000    | 0.000000     | 0.000000     |
| 25%   | 33.000000    | 72.000000     | 8.000000     | 103.000000   | 1.000000     | -1.000000    | 0.000000     | 0.000000     |
| 50%   | 39.000000    | 448.000000    | 16.000000    | 180.000000   | 2.000000     | -1.000000    | 0.000000     | 0.000000     |
| 75%   | 48.000000    | 1428.000000   | 21.000000    | 319.000000   | 3.000000     | -1.000000    | 0.000000     | 0.000000     |
| max   | 95.000000    | 102127.000000 | 31.000000    | 4918.000000  | 63.000000    | 871.000000   | 275.000000   | 1.000000     |

```
data.isnull().sum()
```

```
age          0
job          0
marital      0
education    0
default      0
balance      0
housing      0
loan         0
contact      0
day          0
month        0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
y            0
dtype: int64
```

```
plt.figure(figsize=(12,8))
```

<Figure size 1200x800 with 0 Axes>

```
numerical_data = data.select_dtypes(include=['float64', 'int64'])
```


```
sns.heatmap(numerical_data.corr(), annot=True, cmap='coolwarm')
```

<Axes: >

```
plt.title("Feature Correlation Heatmap")
```

```
Text(0.5, 1.0, 'Feature Correlation Heatmap')
```

```
plt.show()
```

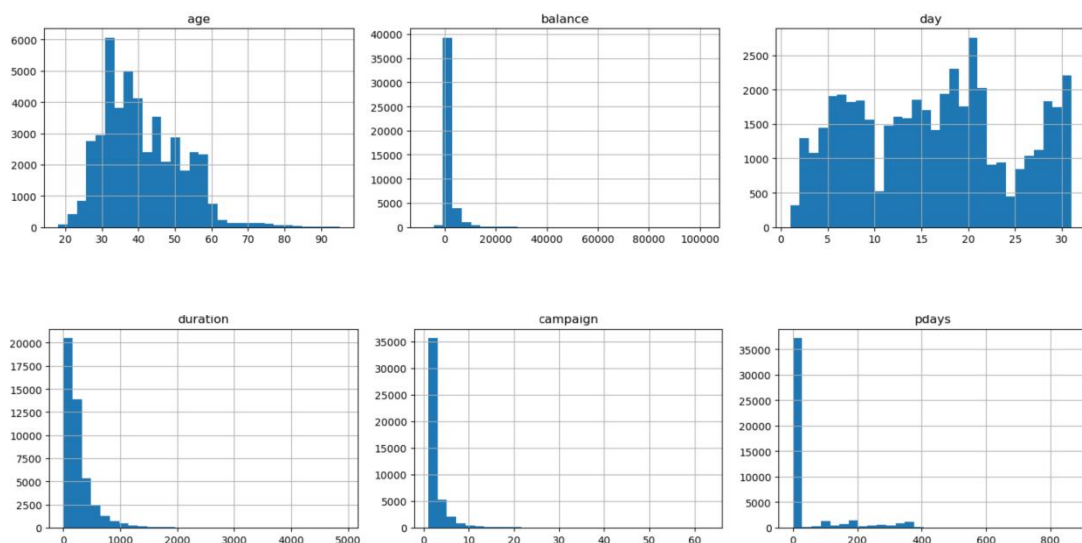



```
data[['age', 'balance', 'duration']].hist(bins=30,figsize=(10,6))

array([[<Axes: title={'center': 'age'}>,
        <Axes: title={'center': 'balance'}>],
       [<Axes: title={'center': 'duration'}>, <Axes: >]], dtype=object)

plt.tight_layout()

plt.show()
```





Data Preprocessing

```
categorical_cols = data.select_dtypes(include=['object']).columns
le = LabelEncoder()
```

```
for col in categorical_cols:
    data[col] = le.fit_transform(data[col])
```

```
data.head()
```

|   | age | job | marital | education | default | balance | housing | loan | contact | day | month | duration | campaign | pdays | previous | poutcome | y |
|---|-----|-----|---------|-----------|---------|---------|---------|------|---------|-----|-------|----------|----------|-------|----------|----------|---|
| 0 | 58  | 4   | 1       | 2         | 0       | 2143    | 1       | 0    | 2       | 5   | 8     | 261      | 1        | -1    | 0        | 3        | 0 |
| 1 | 44  | 9   | 2       | 1         | 0       | 29      | 1       | 0    | 2       | 5   | 8     | 151      | 1        | -1    | 0        | 3        | 0 |
| 2 | 33  | 2   | 1       | 1         | 0       | 2       | 1       | 1    | 2       | 5   | 8     | 76       | 1        | -1    | 0        | 3        | 0 |
| 3 | 47  | 1   | 1       | 3         | 0       | 1506    | 1       | 0    | 2       | 5   | 8     | 92       | 1        | -1    | 0        | 3        | 0 |
| 4 | 33  | 11  | 2       | 3         | 0       | 1       | 0       | 0    | 2       | 5   | 8     | 198      | 1        | -1    | 0        | 3        | 0 |

```
X = data.iloc[:, :-1]
```

```
y = data.iloc[:, -1]
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

## Model Training and Evaluation

```
lr = LinearRegression()  
lr.fit(X_train, y_train)  
  
y_pred_lr = lr.predict(X_test)
```

```
print("Linear Regression R^2:", r2_score(y_test, y_pred_lr))
```

Linear Regression R^2: 0.2161419587110036

```
print("Linear Regression MSE:", mean_squared_error(y_test, y_pred_lr))
```

Linear Regression MSE: 0.08315980805609412

```
print("Linear Regression MSE:", mean_squared_error(y_test, y_pred_lr))
```

Linear Regression MSE: 0.08315980805609412

## Random Forest Regressor

```
rf = RandomForestRegressor(random_state=42)  
rf.fit(X_train, y_train)  
  
y_pred_rf = rf.predict(X_test)
```

```
print("Random Forest R^2:", r2_score(y_test, y_pred_rf))
```

Random Forest R^2: 0.3725930619924863

## Ridge and Lasso Regression

```
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
y_pred_ridge = ridge.predict(X_test)
```

```
print("Ridge R^2:", r2_score(y_test, y_pred_ridge))
```

Ridge R^2: 0.21614138631155322

```
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
y_pred_lasso = lasso.predict(X_test)
```

```
print("Lasso R^2:", r2_score(y_test, y_pred_lasso))
```

Lasso R^2: 0.05728275515857828



## Compare Model Performances

```
models = ['Linear Regression', 'Random Forest', 'Ridge', 'Lasso']
scores = [
    r2_score(y_test, y_pred_lr),
    r2_score(y_test, y_pred_rf),
    r2_score(y_test, y_pred_ridge),
    r2_score(y_test, y_pred_lasso)
]
```

```
plt.figure(figsize=(8,5))
```

<Figure size 800x500 with 0 Axes>

```
sns.barplot(x=models, y=scores)
```

<Axes: >

```
plt.ylabel("R^2 Score")
```

```
Text(53.72222222222214, 0.5, 'R^2 Score')
```

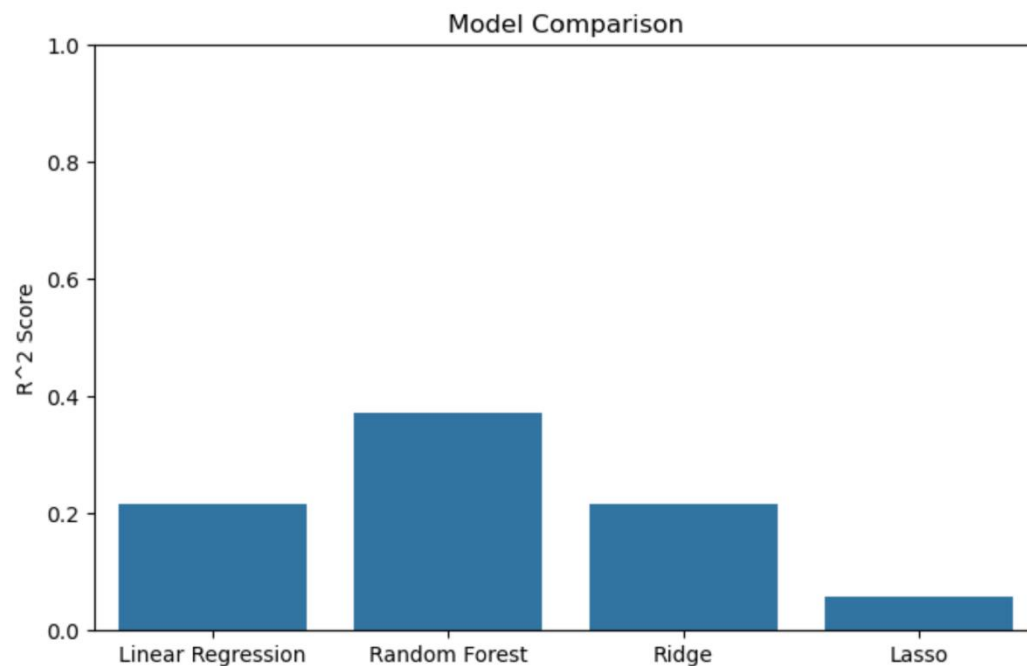
```
plt.title("Model Comparison")
```

```
Text(0.5, 1.0, 'Model Comparison')
```

```
plt.ylim(0,1)
```

(0.0, 1.0)

```
plt.show()
```



---

### Conclusion:

The Random Forest model typically performs best for regression problems with mixed or complex data.

Linear models (Linear, Ridge, Lasso) may perform reasonably well but might underfit depending on data complexity.

Standardizing/encoding features is crucial for optimal performance.

## Project 02:

### Personality classification

#### Program:

## Importing Libraries

```
import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

## Load Dataset

```
df = pd.read_csv(r"C:\Users\HP\Downloads\personality_dataset.csv")
```

Details of Dataset

```
df.head()
```

|   | Time_spent_Alone | Stage_fear | Social_event_attendance | Going_outside | Drained_after_socializing | Friends_circle_size | Post_frequency | Personality |
|---|------------------|------------|-------------------------|---------------|---------------------------|---------------------|----------------|-------------|
| 0 | 4.0              | No         | 4.0                     | 6.0           | No                        | 13.0                | 5.0            | Extrovert   |
| 1 | 9.0              | Yes        | 0.0                     | 0.0           | Yes                       | 0.0                 | 3.0            | Introvert   |
| 2 | 9.0              | Yes        | 1.0                     | 2.0           | Yes                       | 5.0                 | 2.0            | Introvert   |
| 3 | 0.0              | No         | 6.0                     | 7.0           | No                        | 14.0                | 8.0            | Extrovert   |
| 4 | 3.0              | No         | 9.0                     | 4.0           | No                        | 8.0                 | 5.0            | Extrovert   |

```
print("\n Dataset Shape (rows, columns):", df.shape)
```

Dataset Shape (rows, columns): (2900, 8)

```
print(df.dtypes)
```

```
Time_spent_Alone      float64
Stage_fear            object
Social_event_attendance float64
Going_outside         float64
Drained_after_socializing object
Friends_circle_size   float64
Post_frequency        float64
Personality           object
dtype: object
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2900 entries, 0 to 2899
Data columns (total 8 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Time_spent_Alone                      2837 non-null   float64
 1   Stage_fear                            2827 non-null   object
 2   Social_event_attendance               2838 non-null   float64
 3   Going_outside                         2834 non-null   float64
 4   Drained_after_socializing             2848 non-null   object
 5   Friends_circle_size                  2823 non-null   float64
 6   Post_frequency                       2835 non-null   float64
 7   Personality                           2900 non-null   object
dtypes: float64(5), object(3)
memory usage: 181.4+ KB
```

```
df.isnull().sum()
```

```
Time_spent_Alone      63
Stage_fear            73
Social_event_attendance 62
Going_outside         66
Drained_after_socializing 52
Friends_circle_size   77
Post_frequency        65
Personality           0
dtype: int64
```

Split Columns by Data Type

```
numeric_columns = ['Time_spent_Alone', 'Social_event_attendance', 'Going_outside',
                   'Friends_circle_size', 'Post_frequency']
binary_columns = ['Stage_fear', 'Drained_after_socializing']
```

Map Binary Columns to 0/1 First

```
binary_map = {'Yes': 1, 'No': 0}
df[binary_columns] = df[binary_columns].replace(binary_map)
```

```
le = LabelEncoder()
df["Personality"] = le.fit_transform(df["Personality"])
```

```
X = df.drop("Personality", axis=1)
y = df["Personality"]
```

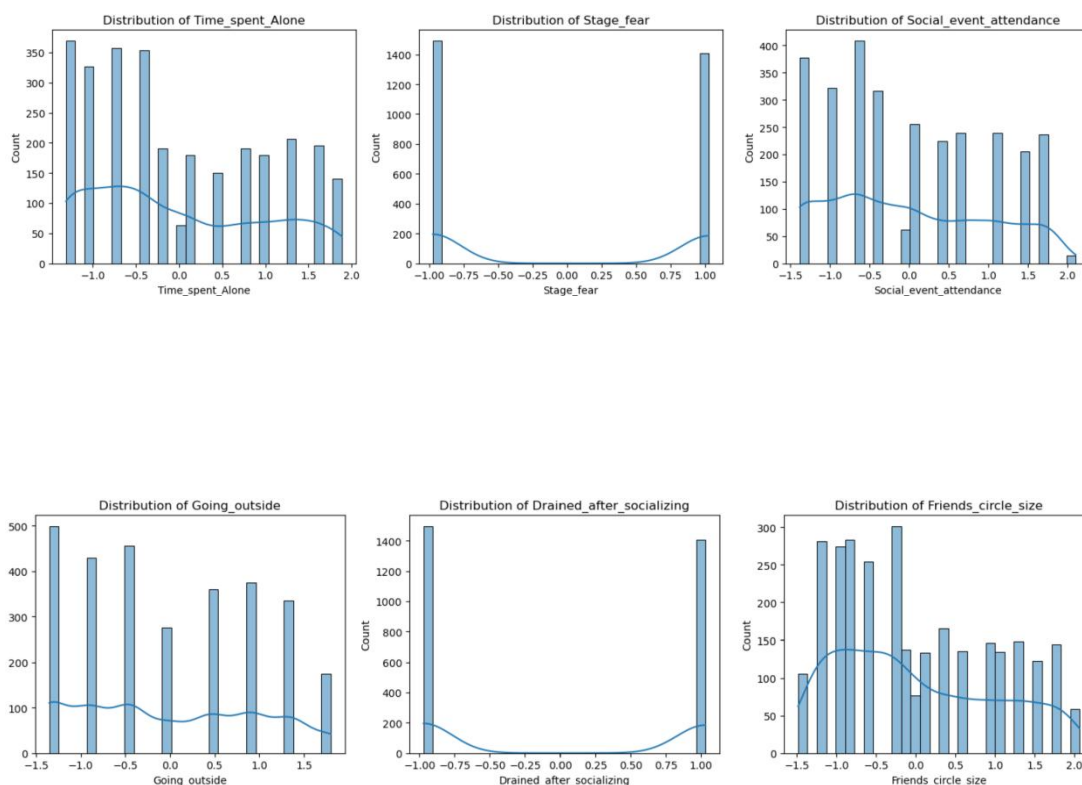
```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

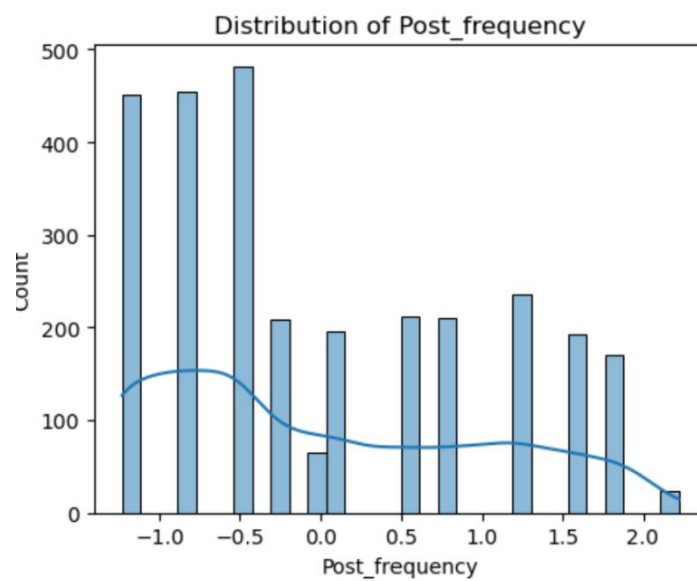
```
X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)
X_scaled_df["Personality"] = y.values
```

## Distribution Plots for All Numeric Features

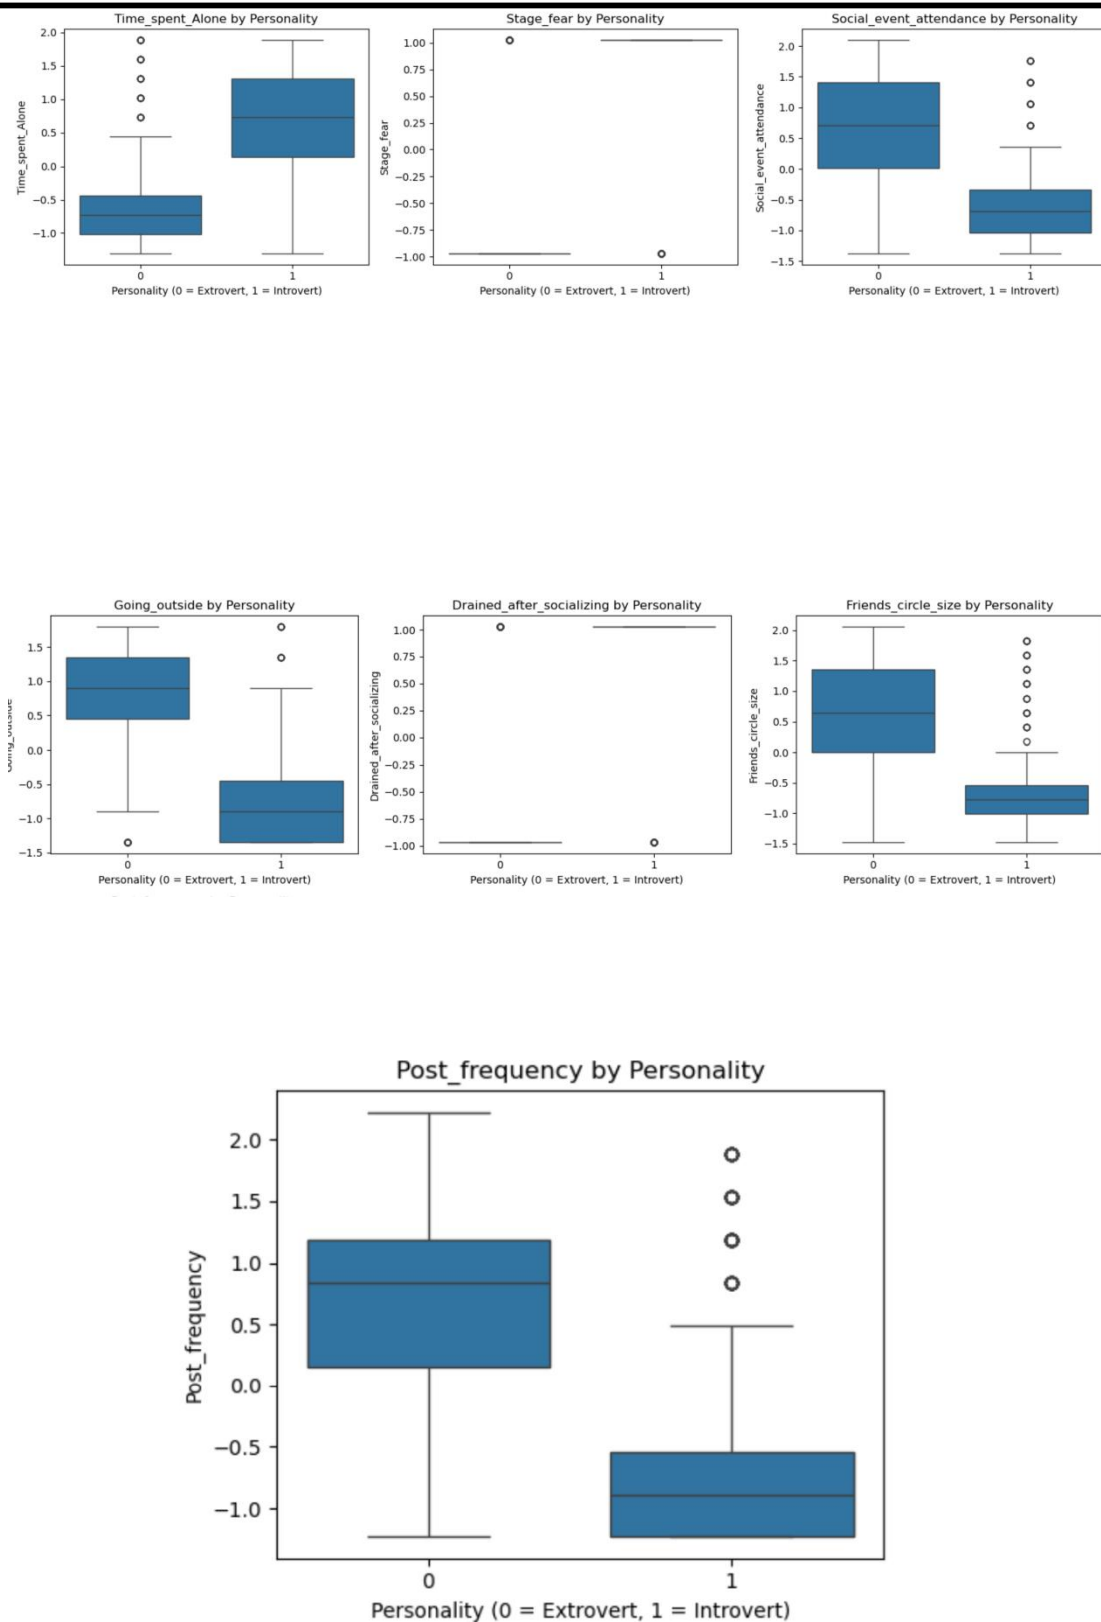
```
num_features = len(X.columns)
cols = 3
rows = math.ceil(num_features / cols)
```

```
plt.figure(figsize=(5 * cols, 4 * rows))
for i, col in enumerate(X.columns):
    plt.subplot(rows, cols, i + 1)
    sns.histplot(X_scaled_df[col], kde=True, bins=30)
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
plt.tight_layout()
plt.show()
```





```
plt.figure(figsize=(5 * cols, 4 * rows))
for i, col in enumerate(X.columns):
    plt.subplot(rows, cols, i + 1)
    sns.boxplot(x="Personality", y=col, data=X_scaled_df)
    plt.title(f"{col} by Personality")
    plt.xlabel("Personality (0 = Extrovert, 1 = Introvert)")
    plt.ylabel(col)
plt.tight_layout()
plt.show()
```



Split Data into Train and Test Sets



```

results = {}

for name, model in models.items():
    print(f"\n Training: {name}")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # Accuracy
    acc = accuracy_score(y_test, y_pred)
    results[name] = acc
    print(f" Accuracy: {acc:.4f}")

    # Classification Report
    print("\n Classification Report:\n", classification_report(y_test, y_pred, target_names=['Extrovert', 'Introvert']))

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Extrovert', 'Introvert'])
    disp.plot(cmap='Blues')
    plt.title(f"Confusion Matrix - {name}")
    plt.show()

```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, stratify=y, random_state=537
)

```

```

models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
    "Support Vector Machine": SVC(kernel="rbf", probability=True)
}

```

Train, Predict, Evaluate Each Model

```

# ROC Curve
if hasattr(model, "predict_proba"):
    y_scores = model.predict_proba(X_test)[:, 1]
else:
    y_scores = model.decision_function(X_test)

fpr, tpr, _ = roc_curve(y_test, y_scores)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, label=f"{name} (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], "k--")
plt.title(f"ROC Curve - {name}")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid(True)
plt.show()


```

Training: Logistic Regression  
Accuracy: 0.9324

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Extrovert    | 0.93      | 0.94   | 0.93     | 373     |
| Introvert    | 0.93      | 0.93   | 0.93     | 352     |
| accuracy     |           |        | 0.93     | 725     |
| macro avg    | 0.93      | 0.93   | 0.93     | 725     |
| weighted avg | 0.93      | 0.93   | 0.93     | 725     |

Confusion Matrix - Logistic Regression



```

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

results = {}

for name, model in models.items():
    print(f"\n Training: {name}")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # Accuracy
    acc = accuracy_score(y_test, y_pred)
    results[name] = acc
    print(f" Accuracy: {acc:.4f}")

    # Classification Report
    print("\n Classification Report:\n", classification_report(y_test, y_pred, target_names=['Extrovert', 'Introvert']))

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Extrovert', 'Introvert'])
    disp.plot(cmap='Blues')
    plt.title(f"Confusion Matrix - {name}")

```


```


# ROC Curve
if hasattr(model, "predict_proba"):
    y_scores = model.predict_proba(X_test)[: , 1]
else:
    y_scores = model.decision_function(X_test)

fpr, tpr, _ = roc_curve(y_test, y_scores)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, label=f"{name} (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], "k--")
plt.title(f"ROC Curve - {name}")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid(True)
plt.show()

```





```
results_df = pd.DataFrame(list(results.items()), columns=["Model", "Accuracy"]).sort_values(by="Accuracy", ascending=False)
results_df
```

|   | Model                  | Accuracy |
|---|------------------------|----------|
| 2 | Support Vector Machine | 0.935172 |
| 0 | Logistic Regression    | 0.932414 |
| 1 | Random Forest          | 0.917241 |