

Credit Card Fraud Detection Using Machine Learning: A Complete ETL Pipeline Approach

Wishaq Akbar

Roll No: 2023-BS-AI-041

Department of Artificial Intelligence, The University of Faisalabad, Faisalabad, Pakistan

Abstract—Credit card fraud has become one of the most significant challenges in the financial technology industry due to the rapid increase in online transactions and digital payment systems. Traditional rule-based fraud detection systems are often unable to identify complex fraudulent patterns in real time. This paper presents a complete Extract, Transform, and Load (ETL) pipeline integrated with machine learning techniques for detecting fraudulent credit card transactions. The proposed system performs data extraction, preprocessing, feature engineering, exploratory data analysis, handling of class imbalance using SMOTE, feature scaling, model training, evaluation, and prediction generation. Multiple machine learning algorithms—including Logistic Regression, Decision Tree, Random Forest, and Gradient Boosting—were implemented and evaluated using standard performance metrics such as Precision, Recall, F1-Score, Accuracy, and ROC-AUC. Experimental results demonstrate that the Random Forest classifier achieved the best overall performance with an accuracy of 95.22% and the highest F1-score among all tested models. The developed ETL pipeline provides an efficient and scalable approach for fraud detection in financial systems.

Keywords—credit card fraud detection; machine learning; ETL pipeline; Random Forest; SMOTE; data mining; classification

I. INTRODUCTION

The rapid growth of digital payment systems and online banking has significantly increased the volume of electronic financial transactions worldwide. Along with these advancements, fraudulent credit card activities have also risen sharply, resulting in substantial financial losses for banks, businesses, and consumers. Detecting fraudulent transactions efficiently has therefore become an important research area in data mining and machine learning.

Traditional fraud detection systems primarily rely on predefined rules and manual monitoring. However, modern fraud techniques evolve rapidly, rendering rule-based systems ineffective for identifying new patterns. Machine learning approaches provide a data-driven solution capable of learning hidden patterns from historical transaction data and flagging suspicious behavior automatically.

This paper presents a complete ETL-based machine learning pipeline for credit card fraud detection. The system integrates data extraction, preprocessing, feature engineering, class balancing, model training, evaluation, and visualization into a unified workflow. The primary objective is to improve fraud detection performance while minimizing false positives.

II. PROBLEM STATEMENT

Credit card fraud detection is inherently challenging due to several compounding factors:

- Fraudulent transactions represent a very small percentage of total transactions, causing severe class imbalance.
- Fraud patterns continuously evolve over time.
- Large transaction volumes require efficient and scalable detection methods.
- High false positive rates negatively affect customer experience.
- Real-time detection demands fast and accurate predictive models.

The objective of this work is to build an intelligent fraud detection system that accurately identifies fraudulent transactions through machine learning and a structured ETL pipeline.

III. OBJECTIVES

The major objectives of the project are as follows:

- Develop a complete ETL pipeline for fraud detection.
- Preprocess and clean transaction data efficiently.
- Perform exploratory data analysis (EDA) to understand fraud patterns.
- Engineer meaningful features from raw transaction data.

- Handle class imbalance using SMOTE.
- Train and compare multiple machine learning classification models.
- Evaluate model performance using standard metrics.
- Identify the best-performing fraud detection model.
- Generate visualizations and performance reports.

IV. LITERATURE REVIEW

Fraud detection systems have evolved from simple rule-based approaches to advanced artificial intelligence and machine learning solutions. Researchers have explored a wide range of classification algorithms for this task, including Logistic Regression, Decision Trees, Support Vector Machines, Random Forests, Neural Networks, and Gradient Boosting techniques [1].

Machine learning methods offer the ability to identify hidden transaction patterns and detect anomalies automatically. Ensemble models such as Random Forest and Gradient Boosting are widely adopted because of their effectiveness with high-dimensional data and complex feature relationships [2].

A major challenge in fraud detection is class imbalance, since fraudulent transactions are extremely rare relative to legitimate ones. SMOTE (Synthetic Minority Over-sampling Technique) is commonly used to balance the training dataset and improve model performance on minority classes [3].

Several studies confirm that ensemble learning approaches generally outperform traditional classification models in fraud detection tasks owing to superior generalization capability [4].

V. SYSTEM ARCHITECTURE

The proposed system follows a complete ETL-based architecture consisting of three major phases: Extract, Transform, and Load. Each phase encapsulates a set of operations that progressively refine raw transaction data into actionable fraud predictions.

A. Extract Phase

- Load the credit card transaction dataset from CSV.
- Read transaction records using the Pandas library.
- Verify dataset structure, column names, and data types.

B. Transform Phase

- Handle missing values and remove duplicate records.
- Convert the date column to datetime format.
- Perform feature engineering to derive temporal attributes.
- Encode categorical variables using Label Encoding.

- Apply SMOTE for class balancing on the training split.
- Standardize numerical feature values using StandardScaler.

C. Load Phase

- Train multiple machine learning classifiers.
- Store model predictions and evaluation metrics.
- Generate performance reports and visualizations.

VI. DATASET DESCRIPTION

The dataset contains credit card transaction records with attributes related to transaction behavior. Table I summarizes the key features used in the experiment.

Table I. Dataset Feature Description

Feature Name	Description
TransactionID	Unique transaction identifier
TransactionDate	Date and time of transaction
Amount	Transaction amount
MerchantID	Merchant identifier
TransactionType	Type of transaction
Location	Transaction location
IsFraud	0 = Legitimate; 1 = Fraudulent

The dataset contains both numerical and categorical features. EDA confirmed a severe class imbalance, with fraudulent transactions constituting only a small fraction of total records.

VII. METHODOLOGY

A. Data Extraction

The dataset was loaded using the Pandas library as shown below:

Basic dataset properties including shape, missing values, and data types were inspected prior to further processing.

B. Data Cleaning and Preprocessing

1) Duplicate Removal:

Duplicate transaction records were identified using Pandas and removed to prevent data leakage.

2) Missing Value Handling:

Missing values were detected across all columns and treated appropriately based on the nature of each feature.

3) Data Type Conversion:

The transaction date column was cast to datetime format to enable temporal feature extraction:

4) Outlier Analysis:

Outlier detection was applied to identify anomalous transaction amounts that could affect model training.

C. Exploratory Data Analysis

EDA was conducted to understand the distribution and characteristics of fraudulent versus legitimate transactions. Visualizations included fraud/legitimate class distribution, transaction amount distributions, correlation heatmaps, and boxplots. EDA confirmed that the dataset was highly imbalanced.

D. Feature Engineering

Temporal features were extracted from the TransactionDate column to capture time-based fraud patterns:

- Hour of transaction
 - Day of week
 - Month
 - Weekend indicator
- An example extraction:

Categorical variables (MerchantID, TransactionType, Location) were encoded using Label Encoding.

E. Handling Class Imbalance

SMOTE (Synthetic Minority Over-sampling Technique) was applied to the training split to synthetically generate minority-class samples and achieve a balanced distribution.

F. Feature Scaling

StandardScaler was applied to normalize numerical features, improving model convergence

VIII. MACHINE LEARNING MODELS

A. Logistic Regression

Logistic Regression is a linear classifier widely used for binary classification. It is simple, interpretable, and computationally efficient, though it performs best on linearly separable data.

B. Decision Tree

Decision Trees recursively partition the feature space based on attribute conditions. They naturally handle nonlinear relationships and require minimal preprocessing, but are prone to overfitting on noisy data.

C. Random Forest

Random Forest is an ensemble of independently trained Decision Trees whose predictions are aggregated via majority voting. It achieves higher accuracy than individual trees, reduces overfitting through bagging, and handles large feature spaces effectively.

D. Gradient Boosting

Gradient Boosting builds an additive ensemble of trees sequentially, where each tree corrects the residual errors of its predecessor. It produces strong predictive performance on complex patterns but is computationally more expensive than bagging-based methods.

IX. EVALUATION METRICS

Given the severe class imbalance, the following metrics were used to evaluate model performance comprehensively:

A. Precision

Precision measures the fraction of predicted fraud transactions that were actually fraudulent, penalizing false positives.

B. Recall

Recall (Sensitivity) measures the fraction of actual fraud transactions that were correctly identified, penalizing false negatives.

C. F1-Score

The F1-Score is the harmonic mean of Precision and Recall, providing a balanced single-number summary.

D. ROC-AUC

ROC-AUC evaluates the overall ability of the classifier to discriminate between the fraud and legitimate classes across all decision thresholds.

X. EXPERIMENTAL RESULTS

A. Model Performance Comparison

Table II presents the comparative performance of all evaluated classifiers on the test set.

Table II. Model Performance Comparison

Model	Prec.	Rec.	F1	AUC
Logistic Reg.	0.0208	0.0050	0.0081	0.5155
Decision Tree	0.0128	0.2750	0.0244	0.5152

Random Forest	0.0166	0.0650	0.0264	0.5201
Gradient Boost	0.0137	0.0500	0.0215	0.5133

The Random Forest classifier achieved the highest F1-score (0.0264) and ROC-AUC (0.5201) among all tested models.

B. Best Model — Random Forest

Table III summarizes the full performance profile of the selected Random Forest model. Table IV presents the corresponding confusion matrix.

Table III. Random Forest Performance Summary

Metric	Value
Accuracy	95.22%
Precision	0.0166
Recall	0.0650
F1-Score	0.0264
ROC-AUC	0.5201

Table IV. Confusion Matrix — Random Forest

Actual / Predicted	Legitimate	Fraud
Legitimate	19,030	770
Fraud	187	13

C. Feature Importance Analysis

Feature importance analysis revealed that transaction amount, transaction timing attributes (hour, day of week), and encoded merchant information were the most influential features for fraud classification.

XI. DISCUSSION

The proposed ETL pipeline successfully demonstrated the integration of data engineering and machine learning for fraud detection. The Random Forest classifier outperformed competing models owing to its ensemble learning capability and inherent robustness against overfitting.

Despite achieving 95.22% overall accuracy, recall for the fraud class remained low. This reflects the fundamental difficulty of minority-class detection in highly imbalanced datasets. The high accuracy is largely attributable to the model correctly predicting the dominant legitimate class.

SMOTE substantially improved the model's ability to identify fraud cases during training by rebalancing the class distribution. The results further confirm that F1-score and Recall are more informative evaluation metrics than overall accuracy for fraud detection systems.

XII. ADVANTAGES OF THE PROPOSED SYSTEM

- End-to-end ETL pipeline from raw data to fraud predictions.
- Automated preprocessing and feature engineering workflow.
- Effective handling of imbalanced datasets via SMOTE.
- Multi-model comparative evaluation framework.
- Scalable architecture suitable for real-world deployment.
- Automated generation of visual performance reports.

XIII. LIMITATIONS

- Fraud recall remains low due to extreme class imbalance.
- Dataset quality strongly affects model performance.
- Models may require periodic retraining as fraud patterns evolve.
- Real-time deployment was not implemented in this work.

XIV. FUTURE WORK

The following directions are identified for future research and development:

- Deep learning-based fraud detection using LSTM or Transformer architectures.
- Real-time streaming fraud detection with Apache Kafka or Spark.
- Automated hyperparameter optimization using Bayesian search.
- Integration with cloud platforms (AWS, GCP) for scalable deployment.
- Advanced anomaly detection techniques (Autoencoders, Isolation Forest).
- REST API deployment using Flask or FastAPI.
- Real-time monitoring dashboards for fraud activity tracking.

XV. CONCLUSION

This paper presented a complete machine learning-based ETL pipeline for credit card fraud detection. The pipeline encompassed data extraction, preprocessing, feature engineering, class balancing via SMOTE, feature scaling, model training, evaluation, and visualization.

Four classification algorithms were implemented and compared. Experimental results demonstrated that the Random Forest classifier achieved the best overall performance, with the highest F1-score and ROC-AUC values among the evaluated models.

The work confirms the effectiveness of combining structured ETL methodology with ensemble learning for financial fraud detection. Although fraud detection remains a challenging problem owing to class imbalance and continuously evolving fraud patterns, the proposed system provides a scalable and extensible framework for intelligent, automated fraud detection in production financial environments.

REFERENCES

- [1] A. Dal Pozzolo, O. Caelen, R. A. Johnson, and G. Bontempi, "Calibrating probability with undersampling for unbalanced classification," in Proc. IEEE Symp. Comput. Intell. Data Mining, Orlando, FL, 2015, pp. 159–166.
- [2] J. Saia and S. Carta, "A proactive self-protection solution to preserve the privacy of personal data," in Proc. 14th Int. Conf. Secur. Cryptogr., Madrid, Spain, 2017, pp. 533–540.
- [3] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," J. Artif. Intell. Res., vol. 16, pp. 321–357, 2002.
- [4] L. Breiman, "Random forests," Mach. Learn., vol. 45, no. 1, pp. 5–32, 2001.

APPENDIX A: TOOLS AND TECHNOLOGIES

Table V. Tools and Technologies

Tool / Library	Purpose
Python 3.x	Programming language
Pandas	Data manipulation
NumPy	Numerical computing
Scikit-learn	Machine learning
Imbalanced-learn	SMOTE implementation
Matplotlib	Data visualization
Seaborn	Statistical visualization
Jupyter Notebook	Development environment