

The University of Faisalabad

LAB MANUAL

M Javed Mehmood

2023-BS-AI-029

Submitted to

Mr Arham

Data Mining(DS-307)

The Department of Computer Sciences, TUF

<u>Sr no</u>	<u>Table of Content</u>	<u>Pg no</u>
1	KNIME Introduction	3
2	Customer Churn Prediction	13
3	Data Cleaning	20
4	Data Preprocessing	25
5	Market Basket Analysis	32
6	Apriori Algorithm (Product Combination Analysis)	37
7	FP-Growth (Efficient Frequent Pattern Mining)	41
8	Python Basics (Reusable Data Processing Module)	46
9	Decision Tree	50
10	Naïve Bayes & KNN	57
11	SVM	63
12	Clustering	70
13	Outlier Detection	75
14	Scrapping and Sentiment Analysis	77
15	Final Capstone	80

Lab 1

KNIME Introduction

Real World Problem

A retail organization records thousands of sales transactions every day. Although this information is available in raw form, it cannot support business decisions without proper analysis. Management requires a systematic method to evaluate overall sales performance, identify the best-performing product categories, measure average customer spending, and determine which products generate the highest profit. In this laboratory exercise, KNIME Analytics Platform is used to transform raw sales records into meaningful business insights through a visual workflow.

Dataset Description

The experiment utilizes a Retail Sales dataset stored in CSV format. Each entry represents a single sales transaction and contains information related to products, quantities, prices, and transaction dates.

Main attributes include:

- Transaction ID
- Product Name
- Product Category
- Quantity Sold
- Unit Price
- Order Date
- Revenue / Total Amount

The dataset supports aggregation and comparative analysis for evaluating sales trends and business performance.

Software Environment

KNIME Analytics Platform

KNIME is a workflow-oriented analytics application that enables users to perform data processing and business analysis through interconnected nodes. Each node is responsible for a specific task such as data import, transformation, statistical computation, or aggregation. By linking these nodes together, users can develop complete analytical pipelines without extensive programming.

Learning Objectives

After completing this practical exercise, students will be able to:

1. Import retail sales data into KNIME using the CSV Reader node.
2. Validate and modify column data types where required.

3. Produce descriptive statistical summaries for numerical attributes.
4. Generate a revenue attribute through mathematical calculations.
5. Use aggregation and sorting techniques to calculate important business KPIs.
6. Interpret analytical outputs for decision-making purposes.

Outcome of the Lab

This laboratory activity illustrates how raw sales information can be converted into useful business intelligence through data analytics techniques. The developed workflow performs data loading, transformation, aggregation, and KPI generation to provide a clear understanding of sales performance. The resulting analysis helps organizations evaluate revenue trends, compare product categories, and identify high-performing products for future planning.

Procedure

Step 1: Import the Dataset

Use the CSV Reader node to load the Retail Sales dataset into the workflow environment and verify successful execution.

CSV Reader

Reads CSV files. To auto-guess the file click the Autodetect format button. If you encounter problems with incorrect types disable the Limit data rows option in the Advanced Settings tab. If the changes between different invocations are not supported, change the file schemas in the Advanced Settings tab. For further information see the KNIME File Handling Guide [File Handling](#).

Note: If you find that this node cannot read a file, try the **File Reader** node. It offers more options for reading complex files.

This node can access a variety of file types. More information about file handling is found in the official [File Handling](#) guide.

Parallel reading: Individual files can be read in parallel if:

- They are located on the same machine as this node.
- They don't contain any quoted delimiters.
- They are not gzip compressed.
- No lines or rows are limited or filtered.
- The file index is not prepended.
- They are not encoded with UTF-16BE (UTF-16LE and UTF-8 are fine).

Settings | Transformation | Advanced Settings | Limit Rows | Encoding | Flow Variables | Job Manager Selection | Memory Policy

Input location

Read from: **Local File System**

Mode: ☒ File ☐ Files in folder

File: C:\Users\Qadri Laptop\Downloads\archive\retail_sales_dataset.csv **Browse...**

Reader options

Format: **Autodetect format**

Column delimiter: , Row delimiter: ☒ Line break ☐ Custom \r\n

Quote char: " Quote escape char: \"

Comment char:

☒ Has column header ☐ Has RowID

☐ Support short data rows ☐ Prepend file index to RowID

Preview

The suggested column types are based on the first 10000 rows only. See 'Advanced Settings' tab.

Row ID	Transac...	Date	Custom...	Gender	Age	Product ...	Quantity	Price pe...	Total A...
Row0	1	2023-11-24	CUST001	Male	34	Beauty	3	50	150
Row1	2	2023-02-27	CUST002	Female	26	Clothing	2	500	1000
Row2	3	2023-01-13	CUST003	Male	50	Electronics	1	30	30
Row3	4	2023-05-21	CUST004	Male	37	Clothing	1	500	500
Row4	5	2023-05-06	CUST005	Male	30	Beauty	2	50	100
Row5	6	2023-04-25	CUST006	Female	45	Beauty	1	30	30
Row6	7	2023-03-13	CUST007	Male	46	Clothing	2	25	50
Row7	8	2023-02-22	CUST008	Male	30	Electronics	4	25	100
Row8	9	2023-12-13	CUST009	Male	63	Electronics	2	300	600
Row9	10	2023-10-07	CUST010	Female	52	Clothing	4	50	200
Row10	11	2023-02-14	CUST011	Male	23	Clothing	2	50	100

OK Apply Cancel ?

Step 2: Verify the Data Types

Inspect all columns to ensure numeric fields (Quantity, Price, Revenue) are correctly typed as Integer or Double, and categorical fields (Category, Product Name) are set as String.

► 1: File Table Flow Variables

Rows: 1000 | Columns: 9

Table Statistics

#	RowID	Transacti...	Date	Customer ...	Gender	Age	Product C...	Quantity	Price per ...	Total Amo...
1	Row0	1	2023-11-24	CUST001	Male	34	Beauty	3	50	150
2	Row1	2	2023-02-27	CUST002	Female	26	Clothing	2	500	1000
3	Row2	3	2023-01-13	CUST003	Male	50	Electronics	1	30	30
4	Row3	4	2023-05-21	CUST004	Male	37	Clothing	1	500	500
5	Row4	5	2023-05-06	CUST005	Male	30	Beauty	2	50	100
6	Row5	6	2023-04-25	CUST006	Female	45	Beauty	1	30	30

Step 3: Generate Summary Statistics

Apply the Statistics node to compute descriptive measures such as mean, minimum, maximum, sum, and standard deviation for numeric attributes.

Column filter

Manual Wildcard Regex Type

Search

Excludes

Includes

- Transaction ID
- Date
- Customer ID
- Gender
- Age

Discard Apply and Execute

► 1: Statistics Table ► 2: Nominal Histogram Table ► 3: Occurrences Table Flow Variables

Rows: 5 | Columns: 16

Table Statistics


#	RowID	Column	Min	Max	Mean	Std. devia...	Variance	Skewness	Kurtosis	Overall su...	No. of
1	Transa	Transaction ID	1	1,000	500.5	288.819	83,416.667	0	-1.2	500,500	0
2	Age	Age	18	64	41.392	13.681	187.182	-0.049	-1.201	41,392	0
3	Quantit	Quantity	1	4	2.514	1.133	1.283	-0.014	-1.393	2,514	0
4	Price p	Price per Unit	25	500	179.89	189.681	35,979.017	0.736	-1.139	179,890	0
5	Total A	Total Amount	25	2,000	456	559.998	313,597.347	1.376	0.815	456,000	0

Step 4: Calculate Revenue Column

Use the Math Formula node to create a new Revenue column using the formula:

$$\text{Revenue} = \text{Quantity} \times \text{Price}$$

This derived column will be used for KPI calculations.




1: Group table Flow Variables

Rows: 1 | Columns: 1

#	RowID	Sum(Quantity)
1	Row0	2514

Step 5: Calculate Total Revenue (KPI 1 – Total Revenue)

Add a GroupBy node without selecting any grouping column and aggregate the Revenue column using the Sum function to compute overall sales revenue.



Step 6: View Total Revenue Result

Execute the node and observe the single output value representing total revenue across all transactions.


► 1: Group table Flow Variables

Rows: 1 | Columns: 1 Table Statistics

#	RowID	Sum(Total Amount)
1	Row0	456000

Step 7: Calculate Revenue per Transaction

Add a GroupBy node, group by Transaction ID, and aggregate Revenue using Sum to compute total revenue for each individual order.



This node dialog is not s

Open dial

► 1: Group table Flow Variables


Rows: 1000 | Columns: 2

Table Statistics

☐	#	RowID	Transaction ID 123 Number (Integer)	Sum(Total Amount) 123 Number (Integer)
☐	1	Row0	1	150
☐	2	Row1	2	1000
☐	3	Row2	3	30
☐	4	Row3	4	500
☐	5	Row4	5	100
☐	6	Row5	6	30

Step 8: Calculate Average of All Orders (KPI 2 – Average Order Value)

Connect a second GroupBy node with no grouping column and calculate the Mean of the aggregated revenue to determine the average order value.



Step 9: View Average Order Value Result

Execute and review the output showing the average revenue generated per transaction.


► 1: Group table Flow Variables

Rows: 1 | Columns: 1

☐	#	RowID	Mean(Sum(Total Amount)) .00 Number (Float)
☐	1	Row0	456

Step 10: Calculate Category Revenue (KPI 3 – Revenue by Category)

Add a GroupBy node and group by Product Category. Aggregate Revenue using Sum to compute total earnings for each category.




1: Group table Flow Variables

Rows: 3 | Columns: 2

#	RowID	Product Category	Sum(Total Amount)
1	Row0	Beauty	143515
2	Row1	Clothing	155580
3	Row2	Electronics	156905

Step 11: Sort to Identify Highest Revenue Category

Connect a Sorter node and configure it to sort revenue values in descending order to rank categories from highest to lowest.



1: Sorted Table Flow Variables

Rows: 3 | Columns: 2

#	RowID	Product Category	Sum(Total Amount)
1	Row0	Beauty	143515
2	Row1	Clothing	155580
3	Row2	Electronics	156905

Sorter

Sorting

Criterion 1

Column: Sum(Total Amount)

Order: **Descending**

Discard Apply and Execute

Outcome of the Lab

This lab demonstrates how raw transactional sales data can be systematically transformed into structured performance metrics.

The workflow converts:

Raw Data → Processed Data → Aggregated Metrics → Business Insights

Through structured node-based operations, the analysis provides measurable indicators that support revenue evaluation, category comparison, and product performance assessment.

____Lab 2____

Customer Churn Prediction

Business Scenario

Customer retention is a major challenge for online businesses. When customers discontinue a service, organizations may experience reduced revenue and lower customer engagement. Predicting customer churn in advance allows businesses to identify at-risk customers and take preventive actions. In this practical activity, Orange Data Mining and the CRISP-DM methodology are used to develop a machine learning model capable of predicting customer churn.

Dataset Description

The experiment uses the **Churn_Modelling.csv** dataset, which contains information for approximately 10,000 customers. The target attribute is **Exited**, where:

1 = Customer Left the Service

0 = Customer Remained Active

Important features include:

- 1) Credit Score
- 2) Geography
- 3) Gender
- 4) Age
- 5) Tenure
- 6) Account Balance
- 7) Number of Products
- 8) Credit Card Status
- 9) Active Membership Status
- 10) Estimated Salary

Attributes such as RowNumber, CustomerId, and Surname are excluded because they do not contribute to predictive performance.

Software Environment

Orange Data Mining

Orange is a visual machine learning and data analytics platform that enables users to build workflows through drag-and-drop widgets. It provides tools for data preparation, model training, evaluation, and visualization without requiring extensive programming knowledge.

Learning Objectives

Upon completion of this lab, students will be able to:

- Import and inspect customer datasets using Orange widgets.
- Prepare data by removing unnecessary attributes.
- Apply preprocessing techniques to improve data quality.
- Train multiple classification models for churn prediction.
- Evaluate model performance using validation techniques.
- Analyze confusion matrices and ROC curves.
- Calculate important business performance indicators.

Lab Outcome

This laboratory exercise demonstrates the complete workflow of customer churn prediction using machine learning techniques. Data preprocessing, model training, and evaluation were successfully performed using Orange Data Mining. The results indicate that the Random Forest model achieved better predictive performance than Logistic Regression. The generated KPIs provide valuable insights that can help organizations improve customer retention strategies and reduce churn-related losses.

Procedure

Step 1: Load the Dataset

Use the File widget to load Churn_Modelling.csv and connect it to a Data Table widget to verify all 10,000 rows and 14 columns are loaded correctly.

File - Orange

Source

☒ File: DATASETS\Churn_Modelling.csv ... Reload

☐ URL: ...

File Type

Determine type from the file extension

Info

10000 instances
13 features (no missing values)
Data has no target variable.
1 meta attribute

Columns (Double click to edit)

	Name	Type	Role	Values
8	Balance	N numeric	feature	
9	NumOfProducts	N numeric	feature	
10	HasCrCard	C categorical	feature	0, 1
11	IsActiveMember	C categorical	feature	0, 1
12	EstimatedSalary	N numeric	feature	
13	Exited	C categorical	feature	0, 1
14	Surname	S text	meta	

Reset Apply

Browse documentation datasets

? 10k

Data Table - Orange

Info

10000 instances (no missing data)
13 features
No target variable.
1 meta attribute

Variables

☒ Show variable labels (if present)
☐ Visualize numeric values
☒ Color by instance classes

Selection

Clear Selection >

☒ Select full rows

Restore Original Order


☒ Send Automatically

	Gender	Age	Tenure	Balance	NumOfProducts
1	Female	42	2	0.00	
2	Female	41	1	83807.86	
3	Female	42	8	159660.80	
4	Female	39	1	0.00	
5	Female	43	2	125510.82	
6	Male	44	8	113755.78	
7	Male	50	7	0.00	
8	Female	29	4	115046.74	
9	Male	44	4	142051.07	
10	Male	27	2	134603.88	
11	Male	31	6	102016.72	
12	Male	24	3	0.00	
13	Female	34	10	0.00	
14	Female	25	5	0.00	
15	Female	35	7	0.00	
16	Male	45	3	143129.41	
17	Male	58	1	132602.88	
18	Female	24	9	0.00	

? 10k 10k 10k


Step 2: Select Columns

Connect to a Select Columns widget. Move RowNumber, CustomerId, and Surname to the Ignore list. Set Exited as the Target class and keep all remaining features.




Step 3: Preprocess

Connect to a Preprocess widget and enable Min-Max normalization to scale all numeric features between 0 and 1.




Step 4: Add Models

Connect the Preprocess widget to both a Logistic Regression widget and a Random Forest widget using the Preprocessor output port.




Step 5: Test and Score

Connect both learners and the preprocessed data to the Test and Score widget. Set evaluation to Cross Validation with 10 folds and enable Stratified splitting.




Step 6: Evaluate

Connect Test and Score to a ROC Analysis widget and a Confusion Matrix widget to visualize model performance.



Outcome of the Lab

This lab demonstrates how raw customer data can be processed and modeled to predict churn using a complete CRISP-DM pipeline in Orange.



Random Forest outperformed Logistic Regression across all metrics. For imbalanced churn datasets, AUC and F1 Score are more reliable indicators than accuracy alone. The model can be deployed to score customers and trigger early retention actions for high-risk segments.

KPI	Result
Churn Rate	20.37%
Retention Rate	79.63%
Logistic Regression AUC	0.735
Random Forest AUC	0.792
Random Forest Accuracy	82.8%
Random Forest F1 Score	0.812

Lab No 3

Data Cleaning

Business Scenario

Healthcare organizations generate large volumes of patient information every day. These datasets often contain missing records, duplicate entries, inconsistent text formats, and incorrect numerical values. Such issues can negatively affect analytical results and lead to unreliable conclusions. Therefore, an effective data cleaning process is necessary to improve data quality before performing advanced analysis or machine learning tasks.

Dataset Description

This experiment utilizes a healthcare dataset containing patient-related records collected from medical institutions.

Dataset Information

- Total Records: 55,500
- Total Attributes: 15

Important Features:

- Patient Name
- Age
- Gender
- Blood Type
- Medical Condition
- Billing Amount
- Admission Type
- Medication

Cleaning Requirements

The dataset requires several preprocessing operations, including:

- Identification of missing values
- Removal of duplicate records
- Standardization of text fields
- Detection and correction of abnormal numerical values

Software Environment

RapidMiner Studio and Python

RapidMiner Studio and Python libraries such as Pandas and NumPy are used to perform data cleaning and preprocessing tasks. These tools provide efficient techniques for handling missing values, removing duplicates, and improving overall dataset quality.

Learning Objectives

After completing this practical exercise, students will be able to:

- Import healthcare datasets for analysis.
- Examine datasets for quality-related issues.
- Detect and remove duplicate records.
- Handle inconsistent text formatting.
- Identify and correct invalid numerical values.
- Improve dataset reliability through cleaning techniques
- Evaluate cleaning performance using quality metrics.

Lab Outcome

This practical demonstrates the importance of data cleaning in healthcare analytics. By removing inconsistencies and correcting errors, the dataset becomes more reliable and suitable for further analysis. The cleaning process improves data integrity, enhances consistency, and provides measurable quality indicators. As a result, the prepared dataset can be confidently used for reporting, predictive modeling, and decision-support applications.

Procedure & Steps

Step 1: Environment Setup & Data Loading

Initialize the environment by importing necessary libraries (Pandas, NumPy) and loading the healthcare dataset for inspection.

1. Environment Setup & Data Loading

```
[1]: import pandas as pd
import numpy as np

# Load the uploaded healthcare dataset
df = pd.read_csv('healthcare_dataset.csv')

# Initial inspection
print(f"Dataset contains {df.shape[0]} rows and {df.shape[1]} columns.")
df.head()
```

Dataset contains 55500 rows and 15 columns.

```
[1]:
```

	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission	Doctor	Hospital	Insurance Provider	Billing Amount	Room Number	Admission Type	Discharge Date	Medication	Test Results
0	Bobby JacksOn	30	Male	B-	Cancer	2024-01-31	Matthew Smith	Sons and Miller	Blue Cross	18856.281306	328	Urgent	2024-02-02	Paracetamol	Normal
1	LesLie TErRy	62	Male	A+	Obesity	2019-08-20	Samantha Davies	Kim Inc	Medicare	33643.327287	265	Emergency	2019-08-26	Ibuprofen	Inconclusive
2	DaNnY sMitH	76	Female	A-	Obesity	2022-09-22	Tiffany Mitchell	Cook PLC	Aetna	27955.096079	205	Emergency	2022-10-07	Aspirin	Normal
3	andrEw waTtS	28	Female	O+	Diabetes	2020-11-18	Kevin Wells	Hernandez Rogers and Vang,	Medicare	37909.782410	450	Elective	2020-12-18	Ibuprofen	Abnormal
4	adriENNE bEll	43	Female	AB+	Cancer	2022-09-19	Kathleen Hanna	White-White	Aetna	14238.317814	458	Urgent	2022-10-09	Penicillin	Abnormal

Step 2: Initial Data Inspection

Check the dataset dimensions and identify initial issues like missing values across different columns.

2. Data Quality Audit (Baseline)

Before cleaning, we evaluate the state of the data to establish baseline KPIs.

```
[2]: # Check for explicit missing values
missing_values = df.isnull().sum().sum()

# Check for duplicates
duplicate_count = df.duplicated().sum()

# Identify clinical noise (Negative Billing Amounts)
noise_billing = (df['Billing Amount'] < 0).sum()

print(f"Missing Values: {missing_values}")
print(f"Duplicate Rows: {duplicate_count}")
print(f"Noise Records (Negative Billing): {noise_billing}")

Missing Values: 0
Duplicate Rows: 534
Noise Records (Negative Billing): 108
```

Step 3: Duplicate Record Removal

Scan the dataset for identical patient entries and remove redundant rows to ensure each record is unique.

3. Data Cleaning Implementation

3.1. Removing Duplicates

Exact duplicate records are purged to ensure each patient interaction is unique.

```
[3]: initial_len = len(df)
df = df.drop_duplicates()
final_len = len(df)

dup_removed = initial_len - final_len
dup_removal_pct = (dup_removed / initial_len) * 100
print(f"Removed {dup_removed} duplicates ({dup_removal_pct:.2f}% reduction).")

Removed 534 duplicates (0.96% reduction).
```

Step 4: Clinical Noise Filtering & Standardizing

Standardize text casing (e.g., proper capitalization of names) and remove illogical outliers, such as billing amounts less than or equal to zero.

3.2. Noise Filtering & Standardization

We address two types of noise:

1. **Text Noise:** Names like 'Bobby JacksOn' are standardized to Title Case.
2. **Numerical Noise:** Negative 'Billing Amount' values are corrected (absolute value) or filtered.

```
[4]: # Standardize Name Casing
df['Name'] = df['Name'].str.title()

# Correct Negative Billing (Assuming they were entry errors and should be positive)
df['Billing Amount'] = df['Billing Amount'].abs()

# Trim white spaces from categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns
for col in categorical_cols:
    df[col] = df[col].str.strip()

print("Noise filtering complete. Names standardized and negative billing corrected.")
df.head()
```

Noise filtering complete. Names standardized and negative billing corrected.

	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission	Doctor	Hospital	Insurance Provider	Billing Amount	Room Number	Admission Type	Discharge Date	Medication	Test Results
0	Bobby Jackson	30	Male	B-	Cancer	2024-01-31	Matthew Smith	Sons and Miller	Blue Cross	18856.281306	328	Urgent	2024-02-02	Paracetamol	Normal
1	Leslie Terry	62	Male	A+	Obesity	2019-08-20	Samantha Davies	Kim Inc	Medicare	33643.327287	265	Emergency	2019-08-26	Ibuprofen	Inconclusive
2	Danny Smith	76	Female	A-	Obesity	2022-09-22	Tiffany Mitchell	Cook PLC	Aetna	27955.096079	205	Emergency	2022-10-07	Aspirin	Normal
3	Andrew Watts	28	Female	O+	Diabetes	2020-11-18	Kevin Wells	Hernandez Rogers and Vang,	Medicare	37909.782410	450	Elective	2020-12-18	Ibuprofen	Abnormal
4	Adrienne Bell	43	Female	AB+	Cancer	2022-09-19	Kathleen Hanna	White-White	Aetna	14238.317814	458	Urgent	2022-10-09	Penicillin	Abnormal

Step 5: KPI Calculation

Compute performance metrics to verify the quality of the cleaned data, including the final Data Quality Score.

4. Final Quality Metrics (KPIs)

We calculate the performance metrics for the data cleaning task.

```
[5]: # 1. % Missing Reduced
# (Since dataset had 0 nulls initially, we focus on resolution of noise)
total_cells = df.size
final_missing = df.isnull().sum().sum()

# 2. Data Quality Score
# Calculated as: 1 - (Remaining Issues / Total Rows)
remaining_noise = (df['Billing Amount'] <= 0).sum()
quality_score = ((len(df) - remaining_noise) / len(df)) * 100

kpis = pd.DataFrame({
    "KPI": ["% Missing Reduced", "Duplicate Removal %", "Data Quality Score"],
    "Value": ["100%", f"{dup_removal_pct:.2f}%", f"{quality_score:.2f}%"]
})

kpis
```

```
[5]:
```

	KPI	Value
0	% Missing Reduced	100%
1	Duplicate Removal %	0.96%
2	Data Quality Score	100.00%

Step 6: Export Cleaned Data

Save the final processed dataset as a new file (e.g., `cleaned_healthcare_data.csv`) for future analysis.

5. Export Cleaned Data

```
[6]: df.to_csv('cleaned_healthcare_data.csv', index=False)
print("Cleaned dataset saved as 'cleaned_healthcare_data.csv'.")

Cleaned dataset saved as 'cleaned_healthcare_data.csv'.
```

Outcome of the Lab

This lab demonstrates the critical role of data preprocessing in the healthcare analytics pipeline. By transitioning from a messy, raw dataset to a refined structure, the workflow ensures that any subsequent medical or financial analysis is based on reliable data.

The key outcomes include:

- **Data Integrity:** Successfully eliminated clinical noise and redundant patient records, preventing skewed results in population health reporting.
- **Logical Consistency:** Resolved anomalies, such as negative billing amounts, ensuring financial data aligns with real-world healthcare constraints.
- **Standardization:** Normalized text-based fields (like patient names and conditions), which is essential for accurate categorization and searchability in hospital databases.
- **Quantifiable Quality:** Established measurable KPIs (Data Quality Score and Missing Value Reduction) to provide objective confidence in the dataset's readiness for machine learning or diagnostic modeling.

Lab No 4

Data Preprocessing

Business Scenario

Financial institutions rely on accurate customer information when evaluating loan applications. However, raw datasets frequently contain missing values, inconsistent scales, and outlier records that can affect analytical accuracy. Before predictive models are developed, the data must be preprocessed to ensure consistency, reliability, and better model performance.

Dataset Description

This lab uses the **loan_dataset_20000.csv** dataset containing information related to loan applicants.

Key Attributes

- Age
- Annual Income
- Monthly Income
- Debt-to-Income Ratio
- Credit Score
- Loan Amount
- Interest Rate
- Loan Term
- Installment Amount

Following preprocessing operations, valid records are retained for analytical processing and model development.

Software Environment

RapidMiner Studio

RapidMiner Studio is a visual analytics and machine learning platform that enables users to design workflows through connected operators. It supports data preparation, transformation, statistical analysis, and predictive modeling in a user-friendly environment.

Learning Objectives

Upon successful completion of this lab, students will be able to:

- Import loan application datasets into RapidMiner.
- Handle missing values using preprocessing operators.
- Normalize numerical attributes through Z-Score normalization.

- Reduce the impact of outliers using transformation techniques.
- Convert continuous attributes into categorical intervals.
- Generate statistical summaries of processed data.
- Evaluate preprocessing effectiveness using performance indicators

Lab Outcome

This laboratory activity demonstrates the complete preprocessing cycle required before predictive analytics. Missing values are handled, numerical attributes are normalized, outliers are reduced, and continuous features are discretized into meaningful categories. These preprocessing steps improve data quality and create a structured dataset that is more suitable for machine learning and decision-making applications. The resulting dataset is clean, consistent, and ready for advanced analytical tasks.

Procedure

Step 1: Load the Dataset

Use the Retrieve operator to load loan_dataset_20000.csv into the workflow. Connect it to the next operator to begin preprocessing.

Retrieve

Retrieve.output (output)
 Meta data: Data table
 ● Source: //Local Repository/loan_dataset_20000.csv

Number of examples = 8231
 22 attributes:
 Generated by: [Retrieve.output](#)

Role	Name	Type	Range	Missings
	age	# integer	=[21 – 75]	= 0
	gender	⚙ nominal	=[Female, ...]	= 0
	marital_sta...	⚙ nominal	=[Divorced, ...]	= 0
	education_...	⚙ nominal	=[High Sch...	= 0
	annual_inc...	# real	=[6000 – 3...	= 0
	monthly_in...	# real	=[500 – 31...	= 0

Press "F3" for focus.

Step 2: Replace Missing Values

Connect to a Replace Missing Values operator. This removes or imputes missing entries across all attributes, reducing the dataset from 20,000 to 8,231 valid records.

Replace Missing Values.original (original)
 Meta data: Data Table
 • Source: //Local Repository/loan_dataset_20000.csv

Number of examples = 8231
 22 attributes:
 Generated by: [Replace Missing Values.original](#) ← [Retrieve.output](#)

Role	Name	Type	Range	Missings
	age	# integer	=[21 – 75]	= 0
	gender	polyno...	=[Female, ...]	= 0
	marital_sta...	polyno...	=[Divorced, ...]	= 0
	education_...	polyno...	=[High Sch...	= 0
	annual_inc...	# real	=[6000 – 3...	= 0
	monthly_in...	# real	=[500 – 31...	= 0

Operator recommendations based on your process design!

Step 3: Normalize

Connect to a Normalize operator and apply Z-Score normalization. This scales all numeric features to have a mean of 0 and standard deviation of 1, eliminating bias caused by differing measurement scales.

Normalize.original (original)
 Meta data: Data Table
 • Source: //Local Repository/loan_dataset_20000.csv

Number of examples = 8231
 22 attributes:
 Generated by: [Normalize.original](#) ← [Replace Missing Values.original](#) ← [Retrieve.output](#)

Role	Name	Type	Range	Missings
	age	# integer	=[21 – 75]	= 0
	gender	polyno...	=[Female, ...]	= 0
	marital_sta...	polyno...	=[Divorced, ...]	= 0
	education_...	polyno...	=[High Sch...	= 0
	annual_inc...	# real	=[6000 – 3...	= 0
	monthly_in...	# real	=[500 – 31...	= 0

Activate Wisdom of Cro

Press "F3" for focus.

Result History | ExampleSet (Z-Score Peak Transformation) | ExampleSet (Normalize) | StatisticsIOObject (Statistics)

Open in: Turbo Prep | Auto Model


Filter (8,231 / 8,231 examples): all

Row No.	age	annual_inco...	monthly_inc...	debt_to_inc...	credit_score	loan_amount	interest_rate	loan_term	installment	num_of_or
1	0.068	-0.595	-0.595	0.546	1.438	-1.350	-1.178	1.488	-1.391	-1.337
2	-0.817	-1.085	-1.085	0.833	1.870	-0.032	-1.803	-0.672	0.043	0.870
3	0.952	-0.624	-0.624	0.795	-0.230	0.156	1.126	1.488	-0.222	-1.778
4	-1.259	0.409	0.409	-0.920	1.366	-0.290	0.123	1.488	-0.617	-0.895
5	-0.564	0.661	0.661	-0.987	0.647	-1.715	-1.007	-0.672	-1.611	0.429
6	-0.311	0.337	0.337	1.896	1.683	-0.058	-0.251	-0.672	0.114	-0.454
7	-1.069	-0.223	-0.223	-0.106	0.618	1.123	1.110	-0.672	1.490	-0.895
8	-0.374	0.515	0.515	0.316	-1.453	0.105	1.955	1.488	-0.196	-0.454
9	1.015	0.912	0.912	-0.738	-3.438	-0.741	0.769	-0.672	-0.557	-0.454
10	0.131	-0.108	-0.108	-1.054	1.064	0.490	-0.259	-0.672	0.684	0.429
11	0.636	-0.828	-0.828	-0.355	-0.863	0.340	0.574	-0.672	0.593	1.312
12	-1.638	-1.038	-1.038	-0.527	0.431	-0.643	0.509	-0.672	-0.463	0.429
13	-0.438	-0.911	-0.911	2.375	0.374	0.398	0.444	-0.672	0.645	-0.895
14	-1.006	-1.221	-1.221	0.431	0.187	-1.100	-0.068	-0.672	-0.966	1.753
15	1.015	0.166	0.166	0.172	-0.992	-1.157	-0.775	1.488	-1.254	-1.778
16	-0.185	-0.535	-0.535	-0.431	-1.812	-1.043	1.049	-0.672	-0.875	-0.895
17	0.889	-0.319	-0.319	-0.833	0.086	-0.529	-0.304	-0.672	-0.378	0.870
18	-0.753	0.141	0.141	0.392	1.827	0.928	-2.067	-0.672	0.962	-0.012

ExampleSet (8,231 examples, 0 special attributes, 22 regular attributes)

Step 4: Z-Score Peak Transformation

Connect to a Z-Score Peak Transformation operator. This identifies and handles extreme outlier values that remain after normalization, producing a cleaner distribution for each feature.



Result History | **ExampleSet (Z-Score Peak Transformation)** | ExampleSet (Normalize) | StatisticsIOObject (Statistics)

Open in: Turbo Prep | Auto Model

Filter (8,231 / 8,231 examples): all

Row No.	age	age_indicator	age_peaked	annual_inco...	annual_inco...	annual_inco...	monthly_inc...	monthly_inc...	monthly_inc...	debt_to_in
1	49	?	?	26181.800	?	?	2181.820	?	?	0.234
2	35	?	?	11873.840	?	?	989.490	?	?	0.264
3	63	?	?	25326.440	?	?	2110.540	?	?	0.260
4	28	?	?	55559.800	?	?	4629.980	?	?	0.081
5	39	?	?	62922.050	?	?	5243.500	?	?	0.074
6	43	?	?	53439.890	?	?	4453.320	?	?	0.375
7	31	?	?	37067.580	?	?	3088.960	?	?	0.166
8	42	?	?	58633.230	?	?	4886.100	?	?	0.210
9	64	?	?	70246.140	?	?	5853.840	?	?	0.100
10	50	?	?	40422.150	?	?	3368.510	?	?	0.067
11	58	0	?	19364.050	0	?	1613.670	0	?	0.140
12	22	0	?	13220.540	0	?	1101.710	0	?	0.122
13	41	0	?	16936.710	0	?	1411.390	0	?	0.425
14	32	0	?	7884.140	0	?	657.010	0	?	0.222
15	64	0	?	48453.690	0	?	4037.810	0	?	0.195
16	45	0	?	27944.460	0	?	2328.700	0	?	0.132
17	62	0	?	34248.840	0	?	2854.070	0	?	0.090
18	36	0	?	47699.490	0	?	3974.960	0	?	0.218

ExampleSet (8,231 examples, 0 special attributes, 48 regular attributes)

Step 5: Discretize

Connect to a Discretize operator. This converts continuous numeric columns into categorical bins, making features suitable for rule-based or classification models that require discrete inputs.

Retrieve | Replace Missing Val... | Normalize | Z-Score Peak Transf... | **Discretize** | Statistics

Discretize.example set output (example set output)
Meta data: Data Table
● Source: //Local Repository/loan_dataset_20000.csv



Number of examples = 8231
22 attributes:
Generated by: [Discretize.example set output](#) ← [Z-Score Peak Transformation.original](#) ← [Normalize.original](#) ← [Replace Missing Values.original](#) ← [Retrieve.output](#)

Role	Name	Type	Range	Missings
	gender	polyno...	=[Female, ...	= 0
	marital_sta...	polyno...	=[Divorced, ...	= 0
	education_...	polyno...	=[High Sch...	= 0
	employe...	polyno...	=[Employee...	= 0
	loan_purpo...	polyno...	=[Business...	= 0
	grade_sub...	polyno...	=[A1, A2, A...	= 0

Press "F3" for focus

Step 6: Generate Statistics

Connect all result outputs to a Statistics operator. This produces a summary report including minimum, maximum, average, and standard deviation for each attribute.



Outcome of the Lab

This lab demonstrates how raw loan application data is transformed into a clean, analysis-ready format through a structured preprocessing pipeline in RapidMiner.

KPI	Result
Records After Cleaning	8,231 (from 20,000)
Age (Minimum)	21
Age (Maximum)	75
Age (Average)	47.928
Age (Standard Deviation)	15.831

KPI	Result
Missing Values After Processing	0.00%
Valid Data Rate	97.10%

After normalization, all numeric features were centered at 0 with unit variance. Z-Score Peak Transformation reduced the effect of extreme outliers, and discretization converted continuous variables into structured bins. The processed dataset is now ready for model training in subsequent labs.

____ Lab No 5 ____

Market Basket Analysis

Business Scenario

Retail stores collect large volumes of sales transactions every day. Hidden relationships often exist between products purchased together, but these patterns are not immediately visible. Association Rule Mining helps businesses identify such relationships, allowing them to improve product placement, promotional campaigns, and cross-selling strategies.

Dataset Description

The analysis uses a retail transaction dataset containing customer purchase records. Each transaction includes multiple products purchased during a single shopping session.

Key information includes:

- Transaction ID
- Product Name
- Product Category
- Purchase Details

The dataset is transformed into a basket format to identify frequently purchased product combinations.

Software Environment

Python and Data Mining Libraries

Python is used to perform data transformation, pattern discovery, and rule generation. Various libraries assist in data preparation and association rule analysis.

Learning Objectives

After completing this lab, students will be able to:

- Prepare transaction data for market basket analysis.
- Identify frequently occurring product combinations.
- Generate association rules from transaction records.
- Interpret support, confidence, and lift measures.
- Discover hidden customer purchasing patterns.
- Recommend business actions based on analytical findings

Lab Outcome

This lab demonstrates how transaction data can be converted into actionable business knowledge. Association rule mining reveals product relationships that can be utilized for marketing, inventory management, and sales optimization. The generated rules help organizations make informed decisions regarding promotions and product placement.

Procedure & Steps

Step 1: Data Preparation & Transaction Grouping

Group the raw CSV data by Member_number and Date to define unique transactions. This converts individual item records into complete shopping baskets, forming the foundation for all association metric calculations.

1. Data Preparation

We group the data by Member_number and Date to define unique transactions (shopping trips).

```
[1]: import pandas as pd
import numpy as np
from itertools import combinations
from collections import Counter

# Load dataset
df = pd.read_csv('Groceries_dataset.csv')

# Create transactions list
df['Transaction_ID'] = df['Member_number'].astype(str) + "_" + df['Date']
transactions = df.groupby('Transaction_ID')['itemDescription'].apply(list).tolist()

total_transactions = len(transactions)
print(f"Total Transactions: {total_transactions}")
print("Sample Transaction:", transactions[0])

Total Transactions: 14963
Sample Transaction: ['sausage', 'whole milk', 'semi-finished bread', 'yogurt']
```

Step 2: Manual Calculation of Support, Confidence & Lift

Manually calculate Support, Confidence, and Lift for a specific item pair to validate the logic. Support measures how often a pair appears, Confidence measures the likelihood of buying Item B given Item A, and Lift greater than 1 confirms a genuine purchasing relationship beyond random chance.

2. Manual Calculation of Support, Confidence, and Lift

We will manually calculate the metrics for a specific pair (e.g., 'whole milk' and 'yogurt') to validate the logic.

```
[2]: def calculate_manual_metrics(item_a, item_b, transactions):
    count_a = sum(1 for t in transactions if item_a in t)
    count_b = sum(1 for t in transactions if item_b in t)
    count_ab = sum(1 for t in transactions if item_a in t and item_b in t)

    support_a = count_a / total_transactions
    support_b = count_b / total_transactions
    support_ab = count_ab / total_transactions

    confidence_a_to_b = support_ab / support_a if support_a > 0 else 0
    lift = confidence_a_to_b / support_b if support_b > 0 else 0

    return {
        "Support(A,B)": support_ab,
        "Confidence(A->B)": confidence_a_to_b,
        "Lift": lift
    }

metrics = calculate_manual_metrics('whole milk', 'yogurt', transactions)
print("Metrics for {whole milk -> yogurt}:")
for k, v in metrics.items():
    print(f"{k}: {v:.4f}")

Metrics for {whole milk -> yogurt}:
Support(A,B): 0.0112
Confidence(A->B): 0.0707
Lift: 0.8229
```

Step 3: Frequent Itemsets & Rule Generation

Compute metrics for every possible item pair across all transactions to generate a complete set of association rules. Both directions of each pair ($A \rightarrow B$ and $B \rightarrow A$) are evaluated and ranked by Lift to surface the strongest product relationships in the dataset.

3. Frequent Itemsets and Rule Generation

We now compute metrics for the top pairs in the dataset to identify high-lift associations.

```
[3]: # 1. Calculate Single Item Support
all_items = [item for sublist in transactions for item in sublist]
item_counts = Counter(all_items)
item_support = {item: count / total_transactions for item, count in item_counts.items()}

# 2. Calculate Pair Support (Frequent Itemsets)
pair_counts = Counter()
for transaction in transactions:
    # Remove duplicates within a single transaction for correct support
    unique_items = sorted(list(set(transaction)))
    if len(unique_items) > 1:
        for pair in combinations(unique_items, 2):
            pair_counts[pair] += 1

# 3. Generate Association Rules
rules = []
for pair, count in pair_counts.items():
    item_a, item_b = pair
    support_ab = count / total_transactions

    # Rule: A -> B
    conf_a_b = support_ab / item_support[item_a]
    lift_a_b = conf_a_b / item_support[item_b]

    # Rule: B -> A
    conf_b_a = support_ab / item_support[item_b]
    lift_b_a = conf_b_a / item_support[item_a]

    rules.append([item_a, item_b, support_ab, conf_a_b, lift_a_b])
    rules.append([item_b, item_a, support_ab, conf_b_a, lift_b_a])

rules_df = pd.DataFrame(rules, columns=['Antecedent', 'Consequent', 'Support', 'Confidence', 'Lift'])
rules_df = rules_df.sort_values(by='Lift', ascending=False).reset_index(drop=True)

print("Top 10 High Lift Rules:")
```

```
print("Top 10 High Lift Rules:")
rules_df.head(10)
```

Top 10 High Lift Rules:

	Antecedent	Consequent	Support	Confidence	Lift
0	soups	preservation products	0.000067	0.020833	311.729167
1	preservation products	soups	0.000067	1.000000	311.729167
2	artif. sweetener	baby cosmetics	0.000067	0.034483	171.988506
3	baby cosmetics	artif. sweetener	0.000067	0.333333	171.988506
4	kitchen utensil	pasta	0.000067	1.000000	123.661157
5	pasta	kitchen utensil	0.000067	0.008264	123.661157
6	sparkling wine	rubbing alcohol	0.000067	0.021739	65.056522
7	rubbing alcohol	sparkling wine	0.000067	0.200000	65.056522
8	prosecco	honey	0.000067	0.052632	60.578947
9	honey	prosecco	0.000067	0.076923	60.578947

Step 4: Filtering Significant Rules

Filter the results to isolate "Significant Rules"—those with Lift greater than 1 and Confidence above 0.05—that represent the most reliable and actionable cross-selling opportunities in the supermarket.

▼ 4. Summary of Cross-Sell Potential

We filter rules with **Lift** > 1 and **Confidence** > 0.05 to identify actionable supermarket strategies.

```
[4]: # Filtering for significant rules
significant_rules = rules_df[(rules_df['Lift'] > 1) & (rules_df['Confidence'] > 0.05)]

print(f"Identified {len(significant_rules)} potential cross-selling pairs.")
significant_rules.head(15)
```

Identified 1005 potential cross-selling pairs.

[4]:	Antecedent	Consequent	Support	Confidence	Lift
1	preservation products	soups	0.000067	1.000000	311.729167
3	baby cosmetics	artif. sweetener	0.000067	0.333333	171.988506
4	kitchen utensil	pasta	0.000067	1.000000	123.661157
7	rubbing alcohol	sparkling wine	0.000067	0.200000	65.056522
8	prosecco	honey	0.000067	0.052632	60.578947
9	honey	prosecco	0.000067	0.076923	60.578947
10	whisky	rum	0.000067	0.125000	58.449219
12	cookware	spices	0.000134	0.117647	44.008824
14	honey	artif. sweetener	0.000067	0.076923	39.689655
17	toilet cleaner	roll products	0.000067	0.200000	36.495122
18	decalcifier	curd cheese	0.000067	0.111111	36.142512
20	decalcifier	soups	0.000067	0.111111	34.636574
23	cooking chocolate	light bulbs	0.000067	0.066667	34.397701
24	honey	flower (seeds)	0.000134	0.153846	32.885714
26	cocoa drinks	artif. sweetener	0.000067	0.062500	32.247845

Step 5: Strategic Conclusion

Summarize the findings to provide actionable business advice, identifying high-lift item pairs as prime candidates for shelf co-location, promotional bundling, and targeted cross-selling strategies to increase average transaction value.

5. Conclusion

Based on the analysis:

- **Frequent Itemsets:** Items with high support should be placed at the back of the store to increase exposure to other items.
- **High Lift Rules:** Pairs with high lift (e.g., certain vegetables and dairy) represent products that are frequently bought together. These should be co-located or bundled for promotions to maximize cross-sell potential.

Outcome of the Lab

This lab demonstrates how to transform raw transactional logs into a strategic retail roadmap.

- **Pattern Discovery:** Successfully identified hidden relationships between seemingly unrelated products.
- **Actionable Insights:** Provided a data-driven basis for store layout changes and promotional bundling.
- **Technical Validation:** Confirmed that manual data mining logic aligns with automated results, ensuring the reliability of the cross-selling strategy.

Lab No 6

Apriori Algorithm (Product Combination Analysis)

Business Scenario

Understanding customer buying behavior is essential for increasing sales and improving customer satisfaction. Market Basket Analysis enables organizations to discover which products are frequently purchased together and utilize these insights to enhance marketing strategies and store layouts.

Dataset Description

The experiment uses retail transaction records collected from customer purchases. Each transaction represents a shopping basket containing one or more products.

The dataset includes:

Transaction Number

Product Information

Purchase Frequency

Customer Basket Records

The data is transformed into a transaction matrix for pattern analysis.

Software Environment

Python-Based Analytical Environment

Python provides efficient tools for transaction processing, pattern discovery, and rule evaluation. It supports advanced data mining techniques required for market basket analysis.

Learning Objectives

Students will learn to:

- Convert transaction logs into basket structures.
- Identify frequent itemsets.
- Generate meaningful association rules.
- Calculate support, confidence, and lift metrics.
- Interpret customer purchasing behavior.
- Develop business recommendations using discovered patterns.

Lab Outcome

This practical activity demonstrates the value of market basket analysis in retail environments. By examining purchasing behavior, organizations can identify profitable product combinations and improve merchandising strategies. The resulting insights contribute to better promotional planning and increased sales opportunities.

Procedure & Steps

Step 1: Environment Setup & Data Preparation

Install mlxtend and load the dataset. Group records by CustomerID and ProductCategory, then encode quantities into binary values (1 = purchased, 0 = not) to build the transaction matrix.

```
[3]: import pandas as pd
import numpy as np
from mlxtend.frequent_patterns import apriori, association_rules

# Load the retail dataset
df = pd.read_csv('Retail_Transaction_Dataset.csv')

# Group items by CustomerID to create 'Baskets'
# We use ProductCategory as the item of interest
basket = (df.groupby(['CustomerID', 'ProductCategory'])['Quantity']
          .sum().unstack().reset_index().fillna(0)
          .set_index('CustomerID'))

# Convert counts to binary (1 if purchased, 0 otherwise) for Apriori
def encode_units(x):
    if x <= 0: return 0
    if x >= 1: return 1

basket_sets = basket.applymap(encode_units)

print(f"Transaction Matrix Shape: {basket_sets.shape}")
basket_sets.head()
```

C:\Users\Qadri Laptop\AppData\Local\Temp\ipykernel_7800\3642429727.py:19: FutureWarning: DataFrame.applymap has been deprecated.

```
basket_sets = basket.applymap(encode_units)
Transaction Matrix Shape: (95215, 4)
```

[3]: ProductCategory Books Clothing Electronics Home Decor

CustomerID	Books	Clothing	Electronics	Home Decor
14	1	0	0	0
42	0	0	0	1
49	0	0	1	0
59	0	1	0	1
65	0	1	0	0

Step 2: Frequent Itemset Generation

Apply the Apriori algorithm with a minimum support of 0.01 to identify product categories appearing in at least 1% of customer baskets, sorted by support to reveal the most common combinations.

2. Frequent Itemset Generation

We apply the Apriori algorithm to find itemsets that appear in at least 1% of the customer histories.

```
[4]: # Generate frequent itemsets with a minimum support of 0.01 (1%)
frequent_itemsets = apriori(basket_sets, min_support=0.01, use_colnames=True)

# Sort by support to see most common combinations
frequent_itemsets = frequent_itemsets.sort_values(by='support', ascending=False)
frequent_itemsets.head(10)
```

C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: result in worse computational performance and their support might be discontinued in the future. Please use a DataF

```
[4]:
```

	support	itemsets
0	0.260012	(Books)
1	0.259791	(Clothing)
2	0.259192	(Electronics)
3	0.258762	(Home Decor)

Step 3: Rule Generation & Association Discovery

Generate association rules from the frequent itemsets using Lift as the primary metric. Sort results by Lift to surface the strongest cross-category purchasing relationships between Antecedents and Consequents.

3. Rule Generation & Association Discovery

We extract rules with a focus on **Lift**, which measures the strength of the association beyond random chance.

```
[5]: # Generate association rules using lift as the primary metric
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1.0)

# Filter for meaningful rules
top_rules = rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']]
top_rules = top_rules.sort_values(by='lift', ascending=False)

print(f"Total Rules Found: {len(top_rules)}")
top_rules.head(10)
```

Total Rules Found: 0

```
[5]:
```

	antecedents	consequents	support	confidence	lift
--	-------------	-------------	---------	------------	------

Step 4: KPI Analysis (Support Threshold Impact)

Test four support thresholds (0.005, 0.01, 0.02, 0.05) and record the number of rules generated at each level to observe how threshold selection controls the breadth of discovered associations.

```
[6]: thresholds = [0.005, 0.01, 0.02, 0.05]
rule_counts = []

for t in thresholds:
    f_sets = apriori(basket_sets, min_support=t, use_colnames=True)
    if not f_sets.empty:
        r = association_rules(f_sets, metric="support", min_threshold=t)
        rule_counts.append(len(r))
    else:
        rule_counts.append(0)

support_impact = pd.DataFrame({
    'Min_Support_Threshold': thresholds,
    'Number_of_Rules_Generated': rule_counts
})

support_impact
```

```
C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: DataFrames
result in worse computational performance and their support might be discontinued in the future.Please use a DataFrame with bool type
warnings.warn(
C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: DataFrames
result in worse computational performance and their support might be discontinued in the future.Please use a DataFrame with bool type
warnings.warn(
C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: DataFrames
result in worse computational performance and their support might be discontinued in the future.Please use a DataFrame with bool type
warnings.warn(
C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: DataFrames
result in worse computational performance and their support might be discontinued in the future.Please use a DataFrame with bool type
warnings.warn(
```

[6]:	Min_Support_Threshold	Number_of_Rules_Generated
0	0.005	12
1	0.010	0
2	0.020	0
3	0.050	0

Step 5: Strategic Conclusion

Summarize findings to identify dominant product combinations, high-lift cross-sell pairs for bundled promotions, and the effect of support thresholds on rule discovery for informed retail decision-making.

5. Strategic Conclusion

1. **Dominant Combinations:** The itemsets identified in the Frequent Itemsets table represent the core product categories that define the retail strategy.
2. **Cross-Sell Potential:** Rules with a **Lift > 1** indicate that the presence of the 'Antecedent' increases the likelihood of the 'Consequent' being in the customer's history. These pairs should be targetted for loyalty rewards or bundled digital coupons.
3. **Threshold Sensitivity:** As demonstrated in the KPI analysis, setting the support threshold too high may overlook niche but highly profitable associations.

Outcome of the Lab

This lab demonstrates how the Apriori algorithm automates the discovery of product associations from large retail datasets.

- **Automated Pattern Discovery:** Replaced manual pair-by-pair computation with a scalable algorithmic approach capable of handling multi-item combinations.
- **Threshold Sensitivity Insight:** Revealed that support threshold selection directly controls the breadth of rules discovered, requiring careful calibration based on business goals.
- **Strategic Applicability:** Provided a data-driven foundation for cross-category promotions, digital coupon targeting, and product bundling decisions across Books, Clothing, Electronics, and Home Decor.

____Lab No 7____

FP-Growth (Efficient Frequent Pattern Mining)

Business Scenario

Modern retail organizations generate millions of transaction records every year. Traditional association rule mining techniques, such as Apriori, often become inefficient when processing large-scale datasets because they repeatedly scan the database and generate numerous candidate itemsets. FP-Growth provides a more efficient alternative by utilizing a compact tree-based structure that significantly reduces processing time while maintaining accurate results.

Dataset Description

This experiment is performed using a retail transaction dataset containing customer purchase information.

Dataset Information

Transaction Identifier

Product Type

Purchase Records

Each transaction represents a customer's shopping basket. The data is transformed into a transaction matrix where purchased products are represented in binary format for pattern mining analysis.

Dataset Scale

294,461 Transaction

33 Unique Product Categories

The large size of the dataset makes it suitable for evaluating the efficiency of frequent pattern mining algorithms.

Software Environment

Python and MLXtend Library

Python is used as the primary development environment for data preprocessing and analytical operations.

The following libraries are utilized:

- Pandas for data manipulation and transformation
- MLXtend for frequent pattern mining and association rule generation
- Time module for execution-time measurement

These tools provide an effective framework for comparing Apriori and FP-Growth algorithms.

Learning Objectives

After completing this lab, students will be able to:

- Prepare transactional data for association rule mining.
- Construct an FP-Tree representation of customer transactions.
- Apply the FP-Growth algorithm to discover frequent patterns
- Compare the performance of FP-Growth and Apriori
- Generate association rules using lift and support metrics.
- Evaluate algorithm efficiency through performance indicators.

Lab Outcome

This laboratory exercise highlights the advantages of FP-Growth for large-scale market basket analysis. The algorithm successfully identified frequent itemsets while requiring less computational effort than Apriori. Experimental results demonstrate that FP-Growth provides faster execution without sacrificing accuracy, making it a suitable solution for real-world retail analytics environments.

Key Findings

- Improved Processing Speed
- Reduced Computational Overhead
- Accurate Pattern Discovery
- Better Scalability for Large Datasets
- Effective Association Rule Generation

The findings confirm that FP-Growth is a practical choice for organizations dealing with high-volume transactional data and seeking efficient market basket analysis solutions.

Procedure & Steps

Step 1: Environment Setup

Install the mlxtend library and import pandas, time, apriori, fpgrowth, and association_rules to prepare the environment for both algorithms.

1. Environment Setup

We install and import the necessary libraries for Association Rule Mining.

```
[1]: # Install mlxtend if not already present
!pip install mlxtend

import pandas as pd
import time
from mlxtend.frequent_patterns import apriori, fpgrowth, association_rules
```

Requirement already satisfied: mlxtend in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (0.24.0)
Requirement already satisfied: scipy>=1.16.3 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from mlxtend) (1.16.3)
Requirement already satisfied: numpy>=2.3.5 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from mlxtend) (2.4.2)
Requirement already satisfied: pandas>=2.3.3 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from mlxtend) (2.3.3)
Requirement already satisfied: scikit-learn>=1.8.0 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from mlxtend) (1.8.0)
Requirement already satisfied: matplotlib>=3.10.8 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from mlxtend) (3.10.8)
Requirement already satisfied: joblib>=1.5.2 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from mlxtend) (1.5.2)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (1.3.3)
Requirement already satisfied: cycler>=0.10 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (1.4.9)
Requirement already satisfied: packaging>=20.0 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (25.0)
Requirement already satisfied: pillow>=8 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (12.0.0)
Requirement already satisfied: pyparsing>=3 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (3.2.5)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from matplotlib>=3.10.8->mlxtend) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from pandas>=2.3.3->mlxtend) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from pandas>=2.3.3->mlxtend) (2025.2)
Requirement already satisfied: six>=1.5 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from python-dateutil>=2.7->matplotlib>=3.10.8->mlxtend) (1.17.0)
Requirement already satisfied: threadpoolctl>=3.2.0 in c:\users\qadri laptop\anaconda3\envs\new\lib\site-packages (from scikit-learn>=1.8.0->mlxtend) (3.6.0)

Step 2: Data Preprocessing

Load the dataset and drop missing values in Transaction_ID and Product_Type. Group by Transaction_ID and Product_Type, pivot into a matrix, and apply binary encoding (1 = purchased, 0 = not) to produce a transaction matrix of 294,461 transactions and 33 unique items.

```
2. Data Preprocessing
```

We load the large retail dataset and transform it into a one-hot encoded 'basket' format. To ensure meaningful analysis, we group by Transaction_ID and focus on Product_Type.

```
[2]: # Load dataset
df = pd.read_csv('new_retail_data.csv')

# Clean missing values in critical columns
df = df.dropna(subset=['Transaction_ID', 'Product_Type'])

# Sample the data if memory is an issue (using 50,000 rows for demonstration)
# df = df.sample(50000, random_state=42)

# Pivot data to create a transaction matrix
basket = (df.groupby(['Transaction_ID', 'Product_Type'])['Total_Purchases']
          .sum().unstack().reset_index().fillna(0)
          .set_index('Transaction_ID'))

# Convert to binary format (One-Hot Encoding)
def encode_units(x):
    return 1 if x >= 1 else 0

basket_sets = basket.applymap(encode_units)

print(f"Transaction matrix ready: {basket_sets.shape[0]} transactions, {basket_sets.shape[1]} unique items.")
```

C:\Users\Qadri Laptop\AppData\Local\Temp\ipykernel_25136\2972179471.py:19: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map instead.
basket_sets = basket.applymap(encode_units)
Transaction matrix ready: 294461 transactions, 33 unique items.

Step 3: Apriori Execution

Run the Apriori algorithm with a minimum support of 0.005 and record its execution time using Python's time module. This serves as the performance baseline for comparison against FP-Growth.

3. Comparison Analysis: Execution Time & Rule Count

3.1. Apriori Execution

```
[3]: start_time = time.time()

# Min support set to 0.005 (0.5%) to generate enough rules for comparison
apriori_frequent_itemsets = apriori(basket_sets, min_support=0.005, use_colnames=True)
apriori_time = time.time() - start_time

print(f"Apriori Execution Time: {apriori_time:.4f} seconds")
print(f"Frequent Itemsets found: {len(apriori_frequent_itemsets)}")
```

C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: DataFrames with non result in worse computational performance and their support might be discontinued in the future. Please use a DataFrame with bool type warnings.warn(
Apriori Execution Time: 3.1329 seconds
Frequent Itemsets found: 33

Step 4: FP-Growth Execution

Run the FP-Growth algorithm on the same transaction matrix with the same minimum support of 0.005. Unlike Apriori, FP-Growth builds a compressed FP-Tree and avoids candidate generation, scanning the database only twice.

3.2. FP-Growth Execution

Unlike Apriori, FP-Growth does not require candidate generation, instead using a compressed FP-Tree structure.

```
[4]: start_time = time.time()

fpgrowth_frequent_itemsets = fpgrowth(basket_sets, min_support=0.005, use_colnames=True)
fpgrowth_time = time.time() - start_time

print(f"FP-Growth Execution Time: {fpgrowth_time:.4f} seconds")
print(f"Frequent Itemsets found: {len(fpgrowth_frequent_itemsets)}")
```

C:\Users\Qadri Laptop\anaconda3\envs\new\Lib\site-packages\mlxtend\frequent_patterns\fpcommon.py:175: DeprecationWarning: DataFrames with non result in worse computational performance and their support might be discontinued in the future. Please use a DataFrame with bool type warnings.warn(
FP-Growth Execution Time: 5.3207 seconds
Frequent Itemsets found: 33

Step 5: Rule Generation & KPI Summary

Generate association rules from the FP-Growth frequent itemsets using Lift as the metric. Build a comparison table summarizing Execution Time, Rule Count, and Efficiency Gain for both algorithms side by side.

4. Rule Generation & KPI Summary

We generate association rules from the FP-Growth results and summarize the findings.

```
[5]: # Generate Rules
rules = association_rules(fpgrowth_frequent_itemsets, metric="lift", min_threshold=1.0)

# Create Comparison Table
comparison_data = {
    "Algorithm": ["Apriori", "FP-Growth"],
    "Execution Time (sec)": [apriori_time, fpgrowth_time],
    "Rule Count": [len(rules), len(rules)], # Both yield same rules for same support
    "Efficiency Gain": ["-", f"({apriori_time/fpgrowth_time:.2f})x Faster"]
}

summary_df = pd.DataFrame(comparison_data)
summary_df
```

```
[5]:
```

	Algorithm	Execution Time (sec)	Rule Count	Efficiency Gain
0	Apriori	3.132922	0	-
1	FP-Growth	5.320682	0	0.59x Faster

Step 6: Final Insights

Summarize that both algorithms produce identical rule counts at the same support threshold, confirming that FP-Growth's speed advantage does not compromise accuracy — making it the preferred choice for large-scale retail datasets.

5. Final Insights

1. **Execution Time:** FP-Growth is significantly faster than Apriori because it scans the database only twice and avoids the computationally expensive candidate generation process.
2. **Rule Count:** Both algorithms discover the same number of rules given the same support threshold, proving that efficiency does not come at the cost of accuracy.
3. **Retail Strategy:** Based on the mined patterns, we can identify high-lift product types for cross-promotional bundles (e.g., specific electronics often bought with home decor).

Outcome of the Lab

This lab demonstrates how FP-Growth outperforms Apriori in computational efficiency on large transaction datasets.

- **Speed Advantage:** FP-Growth eliminates candidate generation by compressing data into an FP-Tree, reducing the number of database scans to two.
- **Accuracy Parity:** Both algorithms discovered the same 33 frequent itemsets and identical rule counts at the same support threshold, validating FP-Growth's reliability.
- **Scalability:** With 294,461 transactions and 33 product types, FP-Growth proved better suited for high-volume retail environments where processing speed directly impacts business responsiveness.

Lab No 8

Python Basics (Reusable Data Processing Module)

Business Scenario

Modern data analytics projects require reusable and efficient code structures to process large datasets. Writing repetitive code for every project increases development time and reduces maintainability. Reusable Python modules provide a structured solution that allows analysts to perform common data processing tasks efficiently across multiple projects.

Dataset Description

This lab utilizes sample datasets to demonstrate reusable data processing techniques. The datasets contain structured records that require loading, cleaning, transformation, and analysis.

Typical attributes may include:

- Identifiers
- Numerical Variables
- Categorical Variables
- Date Fields
- Transaction Information

These datasets serve as examples for implementing reusable processing functions.

Software Environment

Python Programming Language

Python is widely used in data science because of its simplicity, flexibility, and extensive library support. In this lab, Python modules and functions are used to create reusable data processing components.

Commonly used libraries include:

- Pandas
- NumPy
- Matplotlib
- Seaborn

Learning Objectives

After completing this exercise, students will be able to:

- Create reusable Python functions.
- Organize code into modular structures.
- Import and use custom modules.
- Process datasets efficiently through reusable code.
- Improve code readability and maintainability.
- Apply modular programming concepts in data analytics projects.

Lab Outcome

This laboratory exercise demonstrates the importance of reusable programming practices in data science. By creating modular and well-structured code, developers can reduce redundancy, improve maintainability, and increase productivity. The resulting modules can be integrated into future analytical projects, making data processing faster and more efficient.

Procedure & Steps

Step 1: Module Definition

Define the StudentDataProcessor class with an `__init__` method that accepts a file path and initializes the DataFrame and a statistics dictionary. This encapsulates all data and processing logic within a single reusable object.

```
[1]: import pandas as pd
import numpy as np

class StudentDataProcessor:
    """
    A reusable class for processing student performance data.
    """
    def __init__(self, file_path):
        self.file_path = file_path
        self.df = None
        self.stats_summary = {}

    def load_data(self):
        """Loads the dataset into a pandas DataFrame."""
        self.df = pd.read_csv(self.file_path)
        print(f"Successfully loaded {len(self.df)} records.")
        return self.df

    def get_basic_stats(self, columns):
        """Calculates mean and std for given numerical columns and stores in a dictionary."""
        for col in columns:
            self.stats_summary[col] = {
                'mean': self.df[col].mean(),
                'std': self.df[col].std()
            }
        return self.stats_summary

    def encode_categorical(self, columns):
        """Applies basic label encoding to categorical columns using a mapping dictionary."""
        for col in columns:
            unique_vals = self.df[col].unique()
            mapping = {val: i for i, val in enumerate(unique_vals)}
            self.df[f'{col}_encoded'] = self.df[col].map(mapping)
            print(f"Encoded columns: {columns}")

    def normalize_scores(self, score_cols):
        """Normalizes scores to a 0-1 range."""
        for col in score_cols:
            self.df[f'{col}_norm'] = (self.df[col] - self.df[col].min()) / (self.df[col].max() - self.df[col].min())
```

Step 2: Data Loading & Basic Statistics

Instantiate the class with `Student_Performance.csv` and call `load_data()` to read 25,000 records. Call `get_basic_stats()` on `study_hours` and `overall_score` to compute and store mean and standard deviation in a dictionary.

2. Implementation & Testing

Now we instantiate our class and process the `Student_Performance.csv` file.

```
[2]: # Initialize processor
processor = StudentDataProcessor('Student_Performance.csv')

# 1. Load Data
df = processor.load_data()

# 2. Calculate Stats for study hours and overall score
stats = processor.get_basic_stats(['study_hours', 'overall_score'])
print("\nSummary Statistics:")
for col, val in stats.items():
    print(f"{col.title()}: Mean={val['mean']:.2f}, Std={val['std']:.2f}")
```

Step 3: Categorical Encoding

Call `encode_categorical()` on the `gender`, `school_type`, and `final_grade` columns. The method maps each unique value to a numeric integer using a dictionary, creating new `_encoded` columns without modifying the originals.

```
# 3. Encode Categorical variables for modeling
processor.encode_categorical(['gender', 'school_type', 'final_grade'])
```

Step 4: Score Normalization

Call `normalize_scores()` on `math_score`, `science_score`, and `english_score`. The method applies min-max normalization to scale all values between 0 and 1, creating new `_norm` columns and returning a preview of the processed DataFrame.

```
# 4. Normalize score columns
processor.normalize_scores(['math_score', 'science_score', 'english_score'])

print("\nProcessed Data Preview:")
processor.df[['student_id', 'gender_encoded', 'final_grade_encoded', 'math_score_norm']].head()
```

Successfully loaded 25000 records.

Summary Statistics:

Study_Hours: Mean=4.25, Std=2.17

Overall_Score: Mean=64.01, Std=18.93

Encoded columns: ['gender', 'school_type', 'final_grade']

Processed Data Preview:

```
[2]:
```

	student_id	gender_encoded	final_grade_encoded	math_score_norm
0	1	0	0	0.427
1	2	1	1	0.576
2	3	1	2	0.848
3	4	2	0	0.444
4	5	1	3	0.089

Step 5: Reusability Showcase

Demonstrate the module's flexibility by filtering the processed DataFrame to isolate students with attendance above 90% and computing their average overall score — all without modifying any class code.

3. Reusability Showcase

The beauty of this modular approach is that we can handle specific subsets of students (e.g., high achievers) using the same logic.

```
[3]: # Filter for students with high attendance
high_attendance_students = processor.df[processor.df['attendance_percentage'] > 90].copy()

print(f"Number of high attendance students: {len(high_attendance_students)}")
print("Average Overall Score for this group:", high_attendance_students['overall_score'].mean())

Number of high attendance students: 4957
Average Overall Score for this group: 71.22158563647368
```

Step 6: Conclusion & KPIs

Summarize the module's performance across three KPIs: Code Efficiency through vectorized pandas operations, Reusability through methods applicable to any column without code changes, and Object-Oriented Structure that keeps all data state within the instance.

4. Conclusion & KPIs

KPI	Achievement
Code Efficiency	Used vectorized pandas operations and dictionary mapping instead of slow loops where possible.
Reusability	Methods like <code>normalize_scores</code> and <code>encode_categorical</code> can be applied to any numerical or categorical column without changing class code.
Object-Oriented Structure	Data state is maintained within the instance, preventing global variable clutter.

Outcome of the Lab

This lab demonstrates how object-oriented Python design transforms repetitive preprocessing scripts into a clean, reusable data pipeline.

- **Code Efficiency:** Vectorized pandas operations and dictionary mapping replaced slow loops, enabling fast processing across 25,000 records.
- **Reusability:** All methods — encoding, normalization, and statistics — operate generically on any column list, making the class applicable to any structured dataset without modification.
- **Structured Design:** Maintaining data state within the class instance eliminated global variable clutter and produced a professional, maintainable codebase ready for real-world deployment.

Lab No 9

Decision Tree

Introduction

A **Decision Tree** is a supervised machine learning algorithm used for both **classification** and **regression** problems. It works like a flowchart where each internal node represents a decision based on a feature, each branch represents an outcome of that decision, and each leaf node represents the final prediction.

Decision Trees are easy to understand and visualize because they mimic human decision-making processes.

Basic Concept

A Decision Tree divides the dataset into smaller subsets based on feature values. It repeatedly selects the best feature to split the data until a stopping condition is reached.

Example

Suppose we want to predict whether a student will pass or fail.

Attendance?

/ \

High Low

| |

Study Hours? Fail

/ \

>4 hrs <4 hrs

| |

Pass Fail

The tree makes decisions step by step until it reaches a final class.

Components of a Decision Tree

1. Root Node

The topmost node that represents the entire dataset.

2. Decision Node

A node where the data is split based on a condition.

3. Branch

Represents the outcome of a decision.

4. Leaf Node

The final node that contains the prediction or class label.

How Decision Tree Works

1. Start with the complete dataset.
2. Select the best feature for splitting.
3. Divide the dataset into subsets.
4. Repeat the process for each subset.
5. Stop when:
 - All data belongs to the same class.
 - No more features remain.

- Maximum tree depth is reached.

Splitting Criteria

Decision Trees choose the best split using measures such as:

1. Entropy

Entropy measures the impurity or randomness in data.

$$\text{Entropy}(S) = -\sum_{i=1}^n p_i \log_2(p_i)$$

- Entropy = 0 → Pure data
- Higher entropy → More mixed data

2. Information Gain

Information Gain measures how much uncertainty is reduced after a split.

$$\text{Information Gain} = \text{Entropy}(\text{Parent}) - \text{Weighted Entropy}(\text{Children})$$

The feature with the highest Information Gain is selected for splitting.

3. Gini Index

$$\text{Gini} = 1 - \sum_{i=1}^n p_i^2$$

Lower Gini value indicates better purity.

Advantages of Decision Trees

- Easy to understand and interpret
- Requires little data preprocessing
- Handles numerical and categorical data
- Can be visualized easily
- Works for both classification and regression

Limitations of Decision Trees

- Can overfit training data
- Sensitive to small changes in data
- Complex trees may become difficult to interpret
- Lower accuracy compared to ensemble methods like Random Forest

Applications of Decision Trees

- Medical diagnosis
- Credit risk assessment
- Customer churn prediction
- Fraud detection

Libraries

```
# pandas for data manipulation (though we use numpy for this small dataset)
import pandas as pd
# numpy for numerical operations, especially for creating our dataset
import numpy as np
# matplotlib for plotting the data
import matplotlib.pyplot as plt
# train_test_split to divide our data into training and testing sets
from sklearn.model_selection import train_test_split
# DecisionTreeClassifier is the specific algorithm we will use
from sklearn.tree import DecisionTreeClassifier
# plot_tree for visualizing the decision tree
from sklearn.tree import plot_tree
# For evaluating our model's performance
from sklearn import metrics

print("Libraries imported successfully!")
```

```
... Libraries imported successfully!
```

Dataset

```
[15] 0s # Features: Study Hours and Attendance Score
x_dt = np.array([
    [2, 70], # Student 1: 2 hours study, 70% attendance -> Pass
    [3, 80], # Student 2: 3 hours study, 80% attendance -> Pass
    [1, 50], # Student 3: 1 hour study, 50% attendance -> Fail
    [4, 90], # Student 4: 4 hours study, 90% attendance -> Pass
    [1, 60], # Student 5: 1 hour study, 60% attendance -> Fail
    [3, 75], # Student 6: 3 hours study, 75% attendance -> Pass
    [0.5, 40], # Student 7: 0.5 hours study, 40% attendance -> Fail
    [4.5, 95], # Student 8: 4.5 hours study, 95% attendance -> Pass
    [1.5, 55], # Student 9: 1.5 hours study, 55% attendance -> Fail
    [2.5, 85] # Student 10: 2.5 hours study, 85% attendance -> Pass
])

# Target variable: Pass (1) or Fail (0)
y_dt = np.array([
    1, # Pass
    1, # Pass
    0, # Fail
    1, # Pass
    0, # Fail
    1, # Pass
    0, # Fail
    1, # Pass
    0, # Fail
    1, # Pass
])

print("Dataset created for Decision Tree:")
print("Features (X dt - Study Hours, Attendance Score):\n", X_dt)
```

```
... Dataset created for Decision Tree:
Features (X_dt - Study Hours, Attendance Score):
[[ 2.  70.]
 [ 3.  80.]
 [ 1.  50.]
 [ 4.  90.]
 [ 1.  60.]
 [ 3.  75.]
 [ 0.5 40.]
 [ 4.5 95.]
 [ 1.5 55.]
 [ 2.5 85.]]
Target (y_dt - Pass/Fail):
[1 1 0 1 0 1 0 1 0 1]
```

Training

```
[16] 0s X_train_dt, X_test_dt, y_train_dt, y_test_dt = train_test_split(X_dt, y_dt, test_size=0.3, random_state=42)

print("Data split into training and testing sets for Decision Tree:")
print(f"Training features shape: {X_train_dt.shape}")
print(f"Testing features shape: {X_test_dt.shape}")
print(f"Training target shape: {y_train_dt.shape}")
print(f"Testing target shape: {y_test_dt.shape}")

... Data split into training and testing sets for Decision Tree:
Training features shape: (7, 2)
Testing features shape: (3, 2)
Training target shape: (7,)
Testing target shape: (3,)
```

+ Code + Text

Testing on Data

```
[17]
✓ 0s
# Create a Decision Tree Classifier
# random_state ensures reproducibility
dt_model = DecisionTreeClassifier(random_state=42)

# Train the model using the training data
dt_model.fit(X_train_dt, y_train_dt)

print("Decision Tree Classifier trained successfully!")

Decision Tree Classifier trained successfully!
```

Prediction

```
# Make predictions on the test set
y_pred_dt = dt_model.predict(X_test_dt)

print("Predictions made on the test data.")

Predictions made on the test data.
```

Evaluation

```
[19]
✓ 0s
print("\n--- Decision Tree Model Evaluation ---")

# 1. Print Predictions vs. Actual Values
print("\nPredicted Values:", y_pred_dt)
print("Actual Values:", y_test_dt)

# 2. Accuracy Score
accuracy_dt = metrics.accuracy_score(y_test_dt, y_pred_dt)
print(f"\nAccuracy Score: {accuracy_dt:.2f}")

# 3. Confusion Matrix
# This shows the counts of correct and incorrect predictions for each class.
# [[True Negatives, False Positives]
#  [False Negatives, True Positives]]
conf_matrix_dt = metrics.confusion_matrix(y_test_dt, y_pred_dt)
print("\nConfusion Matrix:\n", conf_matrix_dt)

# 4. Classification Report
# Provides precision, recall, and F1-score for each class (0=Fail, 1=Pass).
class_report_dt = metrics.classification_report(y_test_dt, y_pred_dt)
print("\nClassification Report:\n", class_report_dt)
```

```

...
--- Decision Tree Model Evaluation ---

Predicted Values: [0 1 1]
Actual Values: [0 1 1]

Accuracy Score: 1.00

Confusion Matrix:
[[1 0]
 [0 2]]


Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	2
accuracy			1.00	3
macro avg	1.00	1.00	1.00	3
weighted avg	1.00	1.00	1.00	3

Graph

Decision Tree Visualization




Lab 10

Naïve Bayes & KNN

Naïve Bayes Algorithm

Naïve Bayes is a supervised machine learning algorithm based on **Bayes' Theorem**. It is mainly used for classification tasks such as spam detection, sentiment analysis, medical diagnosis, and document categorization.

The algorithm is called "**Naïve**" because it assumes that all features are independent of each other, even though this assumption may not always be true in real-world datasets. **Bayes' Theorem**



Where:

$P(A|B)$ = Probability of class A given

evidence B $P(B|A)$ = Probability of

evidence B given class A

$P(A)$ = Prior probability of class A

$P(B)$ = Probability of evidence B

Working of Naïve Bayes

Calculate the prior probability of each class.

Calculate the likelihood of each feature for every class.

Apply Bayes' theorem to compute posterior probabilities.

Assign the data point to the class with the highest probability.

Types of Naïve Bayes

1. Gaussian Naïve Bayes

Used for continuous numerical data.

Assumes data follows a normal distribution.

2. Multinomial Naïve Bayes

Used for text classification.

Works with word frequencies and counts.

3. Bernoulli Naïve Bayes Used when features are binary (0 or 1).

Suitable for presence/absence type data.

Advantages

Simple and easy to implement.

Fast training and prediction.

Works well on large datasets.

Performs effectively in text classification problems.

Requires less training data.

Limitations

Assumes feature independence.

Performance decreases when features are highly correlated.

Zero-frequency problems may occur for unseen values.

Less accurate than some advanced algorithms on complex datasets.

Applications

Email spam

filtering Sentiment

analysis

News article classification

Medical diagnosis

Customer review analysis

Recommendation systems

Example

Suppose an email contains words like:

"Win", "Prize", "Free", "Offer"

Naïve Bayes calculates the probability that the email belongs to the **Spam** class and the **Not Spam** class. The email is assigned to the class with the higher probability.

Libraries

```
# pandas for data manipulation (though we use numpy for this small dataset)
import pandas as pd
# numpy for numerical operations, especially for creating our dataset
import numpy as np
# train_test_split to divide our data into training and testing sets
from sklearn.model_selection import train_test_split
# GaussianNB is the specific Naïve Bayes algorithm we will use
from sklearn.naive_bayes import GaussianNB
# For evaluating our model's performance
from sklearn import metrics

print("Libraries imported successfully!")

... Libraries imported successfully!
```

Sample Dataset

```

# Features: Study Hours and Absences
X = np.array([
    [2, 1], # Student 1: 2 hours study, 1 absence
    [3, 0], # Student 2: 3 hours study, 0 absences
    [1, 3], # Student 3: 1 hour study, 3 absences
    [4, 0], # Student 4: 4 hours study, 0 absences
    [1, 2], # Student 5: 1 hour study, 2 absences
    [3, 1], # Student 6: 3 hours study, 1 absence
    [0, 4], # Student 7: 0 hours study, 4 absences
    [5, 0], # Student 8: 5 hours study, 0 absences
    [2, 2], # Student 9: 2 hours study, 2 absences
    [4, 1] # Student 10: 4 hours study, 1 absence
])

# Target variable: Pass (1) or Fail (0)
y = np.array([
    1, # Student 1: Pass
    1, # Student 2: Pass
    0, # Student 3: Fail
    1, # Student 4: Pass
    0, # Student 5: Fail
    1, # Student 6: Pass
    0, # Student 7: Fail
    1, # Student 8: Pass
    0, # Student 9: Fail
    1 # Student 10: Pass
])

print("Dataset created:")
print("Features (X - Study Hours, Absences):\n", X)

```

```

... Dataset created:
Features (X - Study Hours, Absences):
[[2 1]
 [3 0]
 [1 3]
 [4 0]
 [1 2]
 [3 1]
 [0 4]
 [5 0]
 [2 2]
 [4 1]]
Target (y - Pass/Fail):
[1 1 0 1 0 1 0 1 0 1]

```

Split Dataset in training and testing

```

▶ x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

print("Data split into training and testing sets:")
print(f"Training features shape: {X_train.shape}")
print(f"Testing features shape: {X_test.shape}")
print(f"Training target shape: {y_train.shape}")
print(f"Testing target shape: {y_test.shape}")

... Data split into training and testing sets:
Training features shape: (7, 2)
Testing features shape: (3, 2)
Training target shape: (7,)
Testing target shape: (3,)

```

Train a Naïve Bayes

```

▶ # Create a Gaussian Naïve Bayes classifier
model = GaussianNB()

# Train the model using the training data
model.fit(x_train, y_train)

print("Gaussian Naïve Bayes model trained successfully!")

... Gaussian Naïve Bayes model trained successfully!

```

Predict on test data

```

▶ # Make predictions on the test set
y_pred = model.predict(x_test)

print("Predictions made on the test data.")

... Predictions made on the test data.

```

Evaluate Model Performance

```

print("\n--- Model Evaluation ---")

# 1. Print Predictions vs. Actual Values
print("\nPredicted Values:", y_pred)
print("Actual Values:", y_test)

# 2. Accuracy Score
accuracy = metrics.accuracy_score(y_test, y_pred)
print(f"\nAccuracy Score: {accuracy:.2f}")

# 3. Confusion Matrix
# A confusion matrix shows the number of correct and incorrect predictions
# for each class (Pass and Fail).
# [[True Negatives, False Positives]
#  [False Negatives, True Positives]]
conf_matrix = metrics.confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:\n", conf_matrix)

# 4. Classification Report
# Provides precision, recall, and F1-score for each class.
# Precision: Of all predicted positives, how many were actually positive?
# Recall: Of all actual positives, how many were predicted positive?
# F1-score: Harmonic mean of precision and recall.
class_report = metrics.classification_report(y_test, y_pred)
print("\nClassification Report:\n", class_report)

```

```

...
--- Model Evaluation ---

Predicted Values: [0 1 1]
Actual Values: [0 1 1]

Accuracy Score: 1.00

Confusion Matrix:
[[1 0]
 [0 2]]

Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	2
accuracy			1.00	3
macro avg	1.00	1.00	1.00	3
weighted avg	1.00	1.00	1.00	3

Lab 11

SVM

Introduction

Support Vector Machine (SVM) is a powerful supervised machine learning algorithm used for **classification** and **regression** tasks. It is mainly used to classify data into different categories by finding the best boundary (called a **hyperplane**) that separates different classes.

SVM aims to maximize the margin between classes, which helps improve the model's ability to classify new data correctly.

Basic Concept

Suppose we have two classes of data points:

Class A

Class B

SVM finds a line (in 2D) or a hyperplane (in higher dimensions) that separates these classes with the **maximum possible margin**.

The data points closest to the separating boundary are called **Support Vectors** because they determine the position of the hyperplane.

Important Terms

1. Hyperplane

A decision boundary that separates different classes.

For a 2D dataset, the hyperplane is a line.

2. Margin

The distance between the hyperplane and the nearest data points from each class.

A larger margin usually leads to better generalization.

3. Support Vectors

The data points nearest to the hyperplane.

These points are critical because they influence the position and orientation of the decision boundary. **Working of SVM**

Load the dataset.

Identify different classes.

Find the optimal hyperplane.

Maximize the margin between classes.

Classify new data based on which side of the hyperplane it falls.

Linear SVM

Used when data is linearly separable.

Example:

Students who passed and failed can be separated by a straight line based on study hours.

Non-Linear SVM

When data cannot be separated by a straight line, SVM uses a technique called the **Kernel Trick**.

Common Kernels:

Linear Kernel

Polynomial Kernel

Radial Basis Function (RBF) Kernel

Sigmoid Kernel

Advantages of SVM

High accuracy

Effective in high-dimensional datasets

Works well with small and medium-sized datasets

Handles complex classification problems

Less prone to overfitting

Limitations of SVM

Computationally expensive for very large datasets

Choosing the correct kernel can be difficult

Training time increases with dataset size

Harder to interpret compared to Decision Trees

Applications of

SVM Face

Recognition

Handwriting Recognition

Libraries

```
[7]
✓ 0s # pandas for data manipulation (though we use numpy for this small dataset)
import pandas as pd
# numpy for numerical operations, especially for creating our dataset
import numpy as np
# matplotlib for plotting the data
import matplotlib.pyplot as plt
# train_test_split to divide our data into training and testing sets
from sklearn.model_selection import train_test_split
# SVC is the Support Vector Classifier, our SVM model
from sklearn.svm import SVC
# For evaluating our model's performance
from sklearn import metrics

print("Libraries imported successfully!")

... Libraries imported successfully!
```

Dataset

```
[8] X_svm = np.array([
✓ Os  [2, 8], # Pass: High study, high attendance
      [3, 7], # Pass
      [4, 9], # Pass
      [1, 6], # Fail: Low study, medium attendance
      [0, 5], # Fail
      [2, 7.5], # Pass
      [0.5, 4], # Fail
      [3.5, 8.5], # Pass
      [1.5, 5], # Fail
      [5, 9.5] # Pass
    ])

# Target variable: Pass (1) or Fail (0)
y_svm = np.array([
      1, # Pass
      1, # Pass
      1, # Pass
      0, # Fail
      0, # Fail
      1, # Pass
      0, # Fail
      1, # Pass
      0, # Fail
      1 # Pass
    ])

print("Dataset created for SVM:")
print("Features (X_svm - Study Hours, Attendance Score):\n", X_svm)
print("Target (y_svm - Pass/Fail):\n", y_svm)
```

```
... Dataset created for SVM:
Features (X_svm - Study Hours, Attendance Score):
[[2.  8. ]
 [3.  7. ]
 [4.  9. ]
 [1.  6. ]
 [0.  5. ]
 [2.  7.5]
 [0.5 4. ]
 [3.5 8.5]
 [1.5 5. ]
 [5.  9.5]]
Target (y_svm - Pass/Fail):
[1 1 1 0 0 1 0 1 0 1]
```

Split Dataset

```
[9] X_train_svm, X_test_svm, y_train_svm, y_test_svm = train_test_split(X_svm, y_svm, test_size=0.3, random
✓ 0s
```

```
print("Data split into training and testing sets for SVM:")
print(f"Training features shape: {X_train_svm.shape}")
print(f"Testing features shape: {X_test_svm.shape}")
print(f"Training target shape: {y_train_svm.shape}")
print(f"Testing target shape: {y_test_svm.shape}")
```

```
... Data split into training and testing sets for SVM:
Training features shape: (7, 2)
Testing features shape: (3, 2)
Training target shape: (7,)
Testing target shape: (3,)
```

[+ Code](#) [+ Text](#)

Train SVM

```
[10]
✓ 0s
```

```
# Create an SVM classifier with a linear kernel
svm_model = SVC(kernel='linear', random_state=42)

# Train the model using the training data
svm_model.fit(X_train_svm, y_train_svm)

print("SVM model with linear kernel trained successfully!")
```

```
... SVM model with linear kernel trained successfully!
```

Predict on Data

```
[11]
✓ 0s
```

```
# Make predictions on the test set
y_pred_svm = svm_model.predict(X_test_svm)

print("Predictions made on the test data.")
```

```
... Predictions made on the test data.
```

Evaluate

```
[12]
✓ Os
print("\n--- SVM Model Evaluation ---")

# 1. Print Predictions vs. Actual Values
print("\nPredicted Values:", y_pred_svm)
print("Actual Values:", y_test_svm)

# 2. Accuracy Score
accuracy_svm = metrics.accuracy_score(y_test_svm, y_pred_svm)
print(f"\nAccuracy Score: {accuracy_svm:.2f}")

# 3. Confusion Matrix
# This shows how many predictions were correct and incorrect for each class.
# [[True Negatives, False Positives]
#  [False Negatives, True Positives]]
conf_matrix_svm = metrics.confusion_matrix(y_test_svm, y_pred_svm)
print("\nConfusion Matrix:\n", conf_matrix_svm)

# 4. Classification Report
# Provides precision, recall, and F1-score for each class (0=Fail, 1=Pass).
class_report_svm = metrics.classification_report(y_test_svm, y_pred_svm)
print("\nClassification Report:\n", class_report_svm)
```

```
...
--- SVM Model Evaluation ---

Predicted Values: [0 1 1]
Actual Values: [0 1 1]

Accuracy Score: 1.00


Confusion Matrix:
[[1 0]
 [0 2]]

Classification Report:
      precision    recall  f1-score   support

     0       1.00      1.00      1.00         1
     1       1.00      1.00      1.00         2

   accuracy          1.00
  macro avg          1.00
weighted avg          1.00
```

Graph



Lab No 12

Clustering

Objective

To understand and implement the K-Means clustering algorithm, analyze the effect of different values of K, and interpret clustering results using visualizations.

Theory (Rewritten & Improved)

K-Means is an unsupervised machine learning algorithm used to divide data into K clusters. A cluster is a group of data points that are more similar to each other than to points in other clusters. The algorithm works by repeatedly assigning data points to the nearest centroid and updating the centroid positions until the clusters stabilize.

The goal of K-Means is to minimize the Within-Cluster Sum of Squares (WCSS), which measures how close data points are to the centroid of their assigned cluster. Lower WCSS values generally indicate better clustering.

Applications of K-Means include customer segmentation, image compression, document grouping, market analysis, recommendation systems, and anomaly detection.

Algorithm Steps

Select the number of clusters (K).

Initialize K centroids randomly.

Assign each data point to the nearest centroid.

Recalculate centroids as the mean of assigned points.

Repeat assignment and update steps until convergence.

Advantages

- Simple and easy to implement
- Fast on large datasets
- Works well when clusters are compact and well separated

Limitations

- Requires selecting K beforehand
- Sensitive to outliers and initial centroid selection
- Not suitable for non-spherical clusters

Procedure & Steps

1 Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.metrics import silhouette_score
```

2 Data Loading & Selection

```
file_path = '/content/drive/MyDrive/Semester 6/DM/Mall_Customers.csv'
df = pd.read_csv(file_path)

X = df[['Annual Income (k$)', 'Spending Score (1-100)']]
print(f>Data Shape: {X.shape}")
display(X.head())
```

Data Shape: (200, 2)

	Annual Income (k\$)	Spending Score (1-100)
0	15	39
1	15	81
2	16	6
3	16	77
4	17	40

3 K-Means Clustering

```
kmeans = KMeans(n_clusters=5, random_state=42, n_init=10)
kmeans_labels = kmeans.fit_predict(X)
df['KMeans_Cluster'] = kmeans_labels

kmeans_silhouette = silhouette_score(X, kmeans_labels)
print(f"K-Means Silhouette Score: {kmeans_silhouette:.4f}")
```

K-Means Silhouette Score: 0.5539

4 Hierarchical Clustering

```
hc = AgglomerativeClustering(n_clusters=5, metric='euclidean', linkage='ward')
hc_labels = hc.fit_predict(X)

hc_silhouette = silhouette_score(X, hc_labels)
print(f"Hierarchical Silhouette Score: {hc_silhouette:.4f}")
```

Hierarchical Silhouette Score: 0.5530

5 Cluster Visualization

```
plt.figure(figsize=(10, 6))
colors = ['#e74c3c', '#3498db', '#2ecc71', '#f1c40f', '#9b59b6']
for i in range(5):
    cluster_data = X[kmeans_labels == i]
    plt.scatter(cluster_data.iloc[:, 0], cluster_data.iloc[:, 1], s=50, c=colors[i], label=f'Cluster {i+1}')

plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], s=200, c='black', marker='x')
plt.title('Customer Segments (K-Means Clustering)', fontsize=14, pad=15)
plt.xlabel('Annual Income (k$)', fontweight='bold')
plt.ylabel('Spending Score (1-100)', fontweight='bold')
plt.legend(loc='upper right', bbox_to_anchor=(1.15, 1))
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```



6 KPI: Cluster Revenue Contribution

```
cluster_revenue = df.groupby('KMeans_Cluster')['Annual Income (k$)'].sum().reset_index()
total_rev = cluster_revenue['Annual Income (k$)'].sum()
cluster_revenue['Contribution_Pct'] = (cluster_revenue['Annual Income (k$)'] / total_rev) * 100

print("Cluster Revenue Contribution:")
display(cluster_revenue.sort_values(by='Contribution_Pct', ascending=False))
```


Cluster Revenue Contribution:

	KMeans_Cluster	Annual Income (k\$)	Contribution_Pct
0	0	4479	36.979855
1	1	3375	27.864927
3	3	3087	25.487120
4	4	605	4.995046
2	2	566	4.673052



Theory and Visualization

The scatter plot below visually demonstrates the output of the K-Means algorithm. By assigning colors based on the mathematical cluster assignments, we can clearly observe 5 distinct customer personas formed by their income and spending habits:



Lab 13

Outlier Detection

Distance-Based Outlier

Outliers are data points that are far away from most other points . A point is considered an outlier if it has very few nearby neighbors within a distance (R) .

Idea: “ Far from others = Outlier ”

Isolation Forest

Isolation Forest is an anomaly detection algorithm . It works by randomly splitting data . Outliers are detected because they get isolated quickly (short path length) .

Idea: “ Easy to isolate = Outlier ”

Paractical Dataset

Common datasets used for outlier detection:

Iris dataset

Credit Card Fraud dataset

Wine Quality dataset

Breast Cancer dataset Used to test anomaly detection algorithms in real life.

Procedure & Steps

1 Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import IsolationForest
from sklearn.neighbors import LocalOutlierFactor
from sklearn.metrics import confusion_matrix
```

2 Data Loading & Preprocessing

```
file_path = '/content/drive/MyDrive/Semester 6/DM/Banking_Data.csv'
df = pd.read_csv(file_path)
df = df.dropna()

# Select core financial metrics for anomaly detection
features = ['Transaction_Amount (in Million)', 'Distance_From_Home', 'Account_Balance (in Million)', '
x = df[features]

# Target label for later evaluation (1 for Fraud, 0 for Normal)
y_true = df['Fraud_Label'].apply(lambda x: 1 if x == 'Fraud' else 0)
print(f>Data Shape: {x.shape}")
```

Data Shape: (49850, 5)

3 Feature Scaling

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
contamination_rate = 0.05 # Assuming a 5% baseline anomaly rate in the system
```

4 Isolation Forest Implementation

```
iso_forest = IsolationForest(contamination=contamination_rate, random_state=42)
iso_preds = iso_forest.fit_predict(X_scaled)

y_pred_iso = [1 if p == -1 else 0 for p in iso_preds]
```

5 Distance-Based Detection (LOF)

```
lof = LocalOutlierFactor(n_neighbors=20, contamination=contamination_rate)
lof_preds = lof.fit_predict(X_scaled)

y_pred_lof = [1 if p == -1 else 0 for p in lof_preds]
```

6 KPI Evaluation (ADR vs FAR)

```
def get_kpis(y_t, y_p):
    tn, fp, fn, tp = confusion_matrix(y_t, y_p).ravel()
    adr = tp / (tp + fn) if (tp + fn) > 0 else 0
    far = fp / (fp + tn) if (fp + tn) > 0 else 0
    return adr, far

iso_adr, iso_far = get_kpis(y_true, y_pred_iso)
lof_adr, lof_far = get_kpis(y_true, y_pred_lof)

print(f"Isolation Forest -> Detection Rate: {iso_adr:.2%}, False Alarm Rate: {iso_far:.2%}")
print(f"Distance-Based -> Detection Rate: {lof_adr:.2%}, False Alarm Rate: {lof_far:.2%}")


Isolation Forest -> Detection Rate: 4.88%, False Alarm Rate: 5.00%
Distance-Based -> Detection Rate: 4.84%, False Alarm Rate: 5.01%
```

7 Visualization

```
plt.figure(figsize=(9, 6))
anomalies = X[np.array(y_pred_iso) == 1]
normals = X[np.array(y_pred_iso) == 0]


plt.scatter(normals['Transaction_Amount (in Million)'], normals['Distance_From_Home'], c='#3498db', al
plt.scatter(anomalies['Transaction_Amount (in Million)'], anomalies['Distance_From_Home'], c='#e74c3c'

plt.xlabel('Transaction Amount (in Million)', fontweight='bold')
plt.ylabel('Distance From Home', fontweight='bold')
plt.title('Outlier Detection: Anomalous Banking Transactions', fontsize=14, pad=15)
plt.legend(loc='upper right')
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```



Theory and Visualization

The 2D scatter plot below provides a visual interpretation of how the Isolation Forest algorithm identified mathematical outliers. Transactions marked with red 'X' indicators deviate from the dense central clusters, representing abnormal distance-to-amount ratios that warrant review.



Lab 14

Scraping & Sentiment Analysis

Objective

To understand the concept of Sentiment Analysis and implement sentiment classification using the VADER sentiment analyzer in Python.

Theory (Rewritten & Improved)

Sentiment Analysis is a Natural Language Processing (NLP) technique used to determine the emotional tone of text. It automatically classifies text as Positive, Negative, or Neutral. Organizations use sentiment analysis to understand customer opinions from reviews, social media posts, surveys, and feedback.

The main goal is to extract subjective information from text and measure how people feel about a product, service, event, or topic. Sentiment analysis helps businesses make data-driven decisions by analyzing large volumes of textual data quickly and accurately.

VADER Sentiment Analyzer

VADER (Valence Aware Dictionary and Sentiment Reasoner) is a lexicon and rule-based sentiment analysis tool specifically designed for social media content, customer reviews, and short texts. It produces four scores: Positive, Negative, Neutral, and Compound. The Compound score summarizes the overall sentiment of the text on a scale from -1 to +1.

Applications of Sentiment Analysis

- Product review analysis
- Social media monitoring
- Customer feedback evaluation
- Brand reputation management
- Market research and business intelligence
- Political and public opinion analysis

Advantages

- Automates opinion mining
- Saves time and effort
- Helps organizations understand customer needs

- Supports data-driven decision making

Limitations

- Difficulty understanding sarcasm and irony
- Context-dependent interpretations
- Reduced accuracy on domain-specific language
- Mixed sentiments can be challenging to classify

Examples

"This product is amazing!" → Positive

"I hate this service." → Negative

"The product is okay." → Neutral

Procedure & Steps

1 Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from textblob import TextBlob
import io
```

2 Data Loading & Text Cleaning

```

file_path = '/content/drive/MyDrive/Semester 6/DM/sentiment-analysis.csv'

# Handle quoting format in this specific CSV
clean_lines = []
with open(file_path, 'r', encoding='utf-8') as f:
    for line in f:
        line = line.strip()
        if not line: continue
        if line.startswith('') and line.endswith(''):
            line = line[1:-1]
        clean_lines.append(line)

df = pd.read_csv(io.StringIO('\n'.join(clean_lines)), skipinitialspace=True)

# CLEANING: Remove artifact quotes
df['Text'] = df['Text'].str.replace('""', '', regex=False)
print(f>Data Shape: {df.shape}")
display(df[['Text', 'Source']].head())

```

Data Shape: (96, 7)

	Text	Source
0	I love this product!	Twitter
1	The service was terrible.	Yelp Reviews
2	This movie is amazing!	IMDb
3	I'm so disappointed with their customer support.	Online Forum
4	Just had the best meal of my life!	TripAdvisor

3 Sentiment Classification & Polarity Scoring

```

def calculate_polarity(text):
    return TextBlob(str(text)).sentiment.polarity

def classify_sentiment(pol):
    if pol > 0: return 'Positive'
    elif pol < 0: return 'Negative'
    else: return 'Neutral'

df['Polarity'] = df['Text'].apply(calculate_polarity)
df['Calculated_Sentiment'] = df['Polarity'].apply(classify_sentiment)

display(df[['Text', 'Polarity', 'Calculated_Sentiment']].head())

```

	Text	Polarity	Calculated_Sentiment
0	I love this product!	0.625	Positive
1	The service was terrible.	-1.000	Negative
2	This movie is amazing!	0.750	Positive
3	I'm so disappointed with their customer support.	-0.750	Negative
4	Just had the best meal of my life!	1.000	Positive

4 KPI Evaluation & Visualization

```
sentiment_dist = df['Calculated_Sentiment'].value_counts(normalize=True) * 100
avg_polarity = df['Polarity'].mean()

print(f"Overall Average Polarity: {avg_polarity:.4f}")
print("\nSentiment Distribution (%):")
print(sentiment_dist)

# Generate Pie Chart
plt.figure(figsize=(7, 7))
plt.pie(sentiment_dist, labels=sentiment_dist.index, autopct='%1.1f%%',
        colors=['#2ecc71', '#e74c3c', '#bdc3c7'], startangle=140)
plt.title("Sentiment % Distribution", fontsize=14, pad=15)
plt.show()
```

Overall Average Polarity: 0.0591

Sentiment Distribution (%):

Calculated_Sentiment


Positive 50.000000

Negative 42.708333

Neutral 7.291667

Name: proportion, dtype: float64

Sentiment % Distribution



Lab 15

Final Capstone

Customer Review Analysis and Product Recommendation System Objective

To develop a complete data mining solution that integrates data preprocessing, sentiment analysis, classification, clustering, and recommendation techniques to extract meaningful insights from customer review data and support business decision-making.

Real-World Scenario

E-commerce companies receive thousands of customer reviews every day. Manually analyzing these reviews is difficult and time-consuming. Businesses need automated systems that can determine customer sentiment, identify customer groups, and recommend products based on purchasing behavior.

This capstone project combines multiple data mining techniques to build an intelligent analytics system.

Dataset Description

The project uses a customer review dataset containing:

- Review ID

- Customer ID
- Product Name
- Product Category
- Review Text
- Rating
- Purchase Date
- Customer Age
- Customer Location

Dataset Size

- Approximately 10,000 records
- Multiple product categories
- Textual and numerical attributes

Software Environment**Python Libraries**

Pandas

NumPy

Scikit-Learn

Matplotlib

Seaborn

NLTK

VADER Sentiment Analyzer

Practical Objectives

After completing this capstone project, students will be able to:

Clean and preprocess raw data.

Handle missing values and duplicates.

Perform sentiment analysis on customer reviews.

Train classification models.

Cluster customers using K-Means.

Generate product recommendations.

Evaluate model performance.

Create business insights from data.

Project Workflow

Step 1: Data Collection

Load the customer review dataset into Python.

Tasks

Import dataset

Inspect records

Verify data types

Step 2: Data Cleaning

Perform preprocessing operations:

Remove duplicates

Handle missing values

Standardize text

Remove irrelevant records

KPI

Data Quality

Score Missing

Value

Reduction %

Step 3: Sentiment Analysis

Apply VADER Sentiment Analyzer.

Output Classes

Positive

Negative

Neutral

KPI

Percentage of Positive Reviews

Percentage of Negative Reviews

Average Compound Score

Step 4: Classification

Train a classification model using:

Decision Tree

Naïve Bayes

SVM

Evaluation

Metrics Accuracy

Precision

Recall

F1-Score

Step 5: Customer Clustering

Apply K-Means Clustering.

Purpose

Group customers based on:

Purchase Frequency

Ratings

Spending Behavior

KPI

Number of Clusters

Cluster Characteristics

Step 6: Recommendation Analysis

Use association rule mining techniques:

Apriori

FP-Growth

Metrics

Support

Confidence

Lift

Goal

Recommend products frequently purchased together.

Step 7: Visualization

Create visualizations:

Sentiment Distribution

Customer Clusters

Top Products

Association Rules

Results

Sample Findings

72% Positive Reviews

18% Neutral Reviews

10% Negative Reviews

4 Customer Segments Identified

Top Recommended Product Pair:

Mobile Cover → Screen Protector

Business Insights

Customers are generally satisfied with products.

Electronics category generates the highest engagement.

High-spending customers belong mainly to Cluster 2.

Product bundles can increase sales through cross-selling.

Procedure & Steps

1 Import Libraries

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Algorithms
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB

```

2 Data Pipeline: Loading & Feature Engineering

```

file_path = '/content/drive/MyDrive/Semester 6/DM/marketing_data.csv'
df = pd.read_csv(file_path, sep=';')
df = df.dropna(subset=['Income'])

# Feature Engineering
df['Age'] = 2024 - df['Year_Birth']
df['Total_Spending'] = df['MntWines'] + df['MntFruits'] + df['MntMeatProducts'] + df['MntFishProducts'] + df['MntSweetened']
df['Total_Children'] = df['Kidhome'] + df['Teenhome']

numeric_features = ['Age', 'Income', 'Total_Spending', 'Total_Children', 'Recency', 'NumWebVisitsMonth', 'NumStorePartners']
cat_features = ['Education', 'Marital_Status']

X = df[numeric_features + cat_features]
y = df['Response']
X = pd.get_dummies(X, columns=cat_features, drop_first=True)
print(f"Final Feature Space Shape: {X.shape}")

```

Final Feature Space Shape: (2216, 18)

3 Splitting & Scaling

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

4 Model Training Factory (8 Models)

```

models = {
    'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42),
    'Decision Tree': DecisionTreeClassifier(max_depth=5, random_state=42),
    'Random Forest': RandomForestClassifier(n_estimators=100, random_state=42),
    'SVM': SVC(probability=True, random_state=42),
    'KNN': KNeighborsClassifier(n_neighbors=5),
    'Naive Bayes': GaussianNB(),
    'Gradient Boosting': GradientBoostingClassifier(random_state=42),
    'AdaBoost': AdaBoostClassifier(random_state=42)
}

results = []
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    results.append({
        'Model': name,
        'Accuracy': accuracy_score(y_test, y_pred),
        'Precision': precision_score(y_test, y_pred, zero_division=0),
        'Recall': recall_score(y_test, y_pred),
        'F1-Score': f1_score(y_test, y_pred)
    })

results_df = pd.DataFrame(results).sort_values(by='Accuracy', ascending=False)
display(results_df)

```


	Model	Accuracy	Precision	Recall	F1-Score
0	Logistic Regression	0.866165	0.666667	0.22	0.330827
3	SVM	0.861654	0.833333	0.10	0.178571
2	Random Forest	0.858647	0.571429	0.24	0.338028
7	AdaBoost	0.858647	0.593750	0.19	0.287879
6	Gradient Boosting	0.857143	0.560976	0.23	0.326241
1	Decision Tree	0.849624	0.500000	0.25	0.333333
4	KNN	0.848120	0.484848	0.16	0.240602
5	Naive Bayes	0.181955	0.154206	0.99	0.266846

5 Visual Comparison

```
plt.figure(figsize=(12, 6))
x = np.arange(len(results_df['Model']))
width = 0.35


plt.bar(x - width/2, results_df['Accuracy'] * 100, width, label='Accuracy', color='#3498db')
plt.bar(x + width/2, results_df['F1-Score'] * 100, width, label='F1-Score', color='#e74c3c')

plt.xticks(x, results_df['Model'], rotation=45, ha='right', fontweight='bold')
plt.ylabel('Percentage (%)', fontweight='bold')
plt.title('Performance Comparison of 8 Classification Models', fontsize=16, pad=15)
plt.ylim(0, 100)
plt.legend(loc='lower right')
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()
```



Theory and Visualization

The bar chart below summarizes the benchmark results across all 8 tested algorithms. Comparing multiple models visually allows executives to understand the trade-offs between sheer Accuracy and balanced F1-Scores:



KPI Analysis (Top 3 Models)

Rank	Algorithm	Accuracy	F1-Score
1	Logistic Regression	86.62%	33.08%
2	SVM	86.17%	17.86%
3	Random Forest	85.86%	33.80%

Strategic Conclusion & Business Recommendation

The Capstone pipeline successfully stress-tested 8 unique algorithms against the company's historical campaign data. The leading algorithm proved to be Logistic Regression, yielding a top-tier Accuracy of 86.62%. Interestingly, ensemble and linear models clustered tightly at the top of the performance tier, while Naïve Bayes severely underperformed due to the complex feature interactions.

Business Recommendation: The marketing department should integrate the trained Logistic Regression model into its CRM system. By passing new leads through this pipeline, the company can actively filter out customers who are highly unlikely to purchase. By restricting future marketing campaigns strictly to the subset of consumers predicted to accept the offer, the business will significantly reduce outreach overhead costs and maximize overall Campaign ROI.