

Lab Manual

Data Mining



By

Saad Hassan

2023-BS-AI-152

Submitted to

Sir Arham

Bachelor of Science in Artificial Intelligence
DEPARTMENT OF COMPUTER SCIENCES

Contents

Lab 1: Introduction to Data Mining	3
Lab 2: CRISP-DM Methodology	8
Lab 3: Data Cleaning	10
Lab 4: Data Preprocessing.....	20
Lab 5: Market Basket Analysis	23
Lab 6: Apriori Algorithm	29
Lab 7: FP-Growth Algorithm.....	36
Lab 8: Python Basics — Reusable Data Processing Module	42
Lab 9: Decision Tree Classifier	50
Lab 10: Naive Bayes and KNN Classifiers	61
Lab 11: Support Vector Machine (SVM) Classifier	71
Lab 12: K-Means and Hierarchical Clustering	79
Lab 13: Outlier Detection and Anomaly Detection	88
Lab 14: Data Scraping & Sentiment Analysis	98
Lab 15: Performance Comparison of Models	104

Lab 1: Introduction to Data Mining

Lab 1: Introduction to Data Mining	
Real-World Problem	A retail store wants to extract meaningful sales insights from raw transaction data to improve decision-making.
Dataset	Retail Sales Dataset (Supermarket Sales — Kaggle)
Tool(s)	KNIME Analytics Platform
Objectives	Import dataset, identify data types, generate summary statistics, explore attributes
KPIs / Metrics	Total Revenue, Avg Order Value, Category Revenue, Top Products

1.1 Introduction

Data Mining is the process of discovering patterns, correlations, anomalies, and useful insights from large datasets. It combines techniques from statistics, machine learning, and database systems to transform raw data into actionable knowledge. In this lab, we use a retail sales dataset to understand the foundational workflow of data exploration using KNIME Analytics Platform.

The Supermarket Sales dataset contains transactions from a retail chain and includes attributes such as product line, unit price, quantity, total sales, customer type, payment method, and ratings. Our goal is to load, inspect, and summarize this data to answer basic business questions.

1.2 Objectives

- Import and load the Supermarket Sales dataset into KNIME.
- Identify data types (numeric, categorical, date/time) for each attribute.
- Generate descriptive summary statistics (mean, median, std dev, min, max).
- Visualize the distribution of key attributes using bar charts and histograms.
- Calculate business KPIs: Total Revenue, Average Order Value, Category-wise Revenue.

1.3 Background / Theory

1.3.1 What is Data Mining?

Data mining refers to the computational process of discovering patterns in large data sets involving methods at the intersection of machine learning, statistics, and database systems. The overall goal is to extract information from a data set and transform it into an understandable structure for further use.

1.3.2 KNIME Analytics Platform

KNIME (Konstanz Information Miner) is an open-source data analytics, reporting, and integration platform. It uses a visual workflow editor where users drag-and-drop processing nodes and connect them to build data pipelines. KNIME supports a wide range of machine learning and data processing algorithms.

1.3.3 Descriptive Statistics

Descriptive statistics summarize and describe the features of a dataset. Key measures include: Mean (average value), Median (middle value), Standard Deviation (spread), Min/Max (range), and Frequency counts for categorical variables.

1.4 Dataset Description

The Supermarket Sales dataset from Kaggle contains 1,000 rows and 17 columns representing point-of-sale transactions. Key attributes include:

- Invoice ID — Unique transaction identifier
- Branch — Store branch (A, B, C)
- City — City of the branch
- Customer Type — Member or Normal
- Gender — Male or Female
- Product Line — Category (e.g., Health & Beauty, Food & Beverages)
- Unit Price — Price per unit (USD)
- Quantity — Number of units purchased
- Tax 5% — 5% tax on total
- Total — Gross total including tax
- Date / Time — Transaction timestamp
- Payment — Cash, Credit Card, or E-wallet
- Rating — Customer satisfaction rating (1–10)

1.5 Procedure (KNIME Workflow)

1. Download the Supermarket Sales dataset from Kaggle (supermarket_sales.csv).
2. Open KNIME and create a new workflow.
3. Drag a CSV Reader node onto the canvas and configure the file path.
4. Connect to a Statistics node to generate summary statistics for all columns.
5. Add a Column Filter node to select key numeric columns (Unit Price, Quantity, Total, Rating).
6. Connect a Bar Chart node to visualize Product Line vs. Total Revenue.
7. Add a GroupBy node to aggregate Total by Product Line (sum).
8. Use a Histogram node to plot the distribution of Unit Price and Quantity.
9. Calculate Total Revenue = SUM(Total), Avg Order Value = MEAN(Total).
10. Export results to CSV using a CSV Writer node.

1.6 Screenshots

KNIME Analytics Platform

Home KNIME_project

Save Execute Cancel Reset

Local - KNIME_project

Nodes

Views JavaScript Statistics +1

Bar Chart Bar Chart (JavaScript) Pie Chart

Scatter Plot Scatter Plot Matrix Linear Correlation

CSV Reader

This node dialog is not supported here.

Open dialog

1: File Table

Flow Variables

Rows: 1000 Columns: 17

#	RowID	Invoice ID	Branch	City	Customer...	Gender	Product li...	Unit price	Quantity	Tax 5%	S...
1	Row0	750-67-8428	Alex	Yangon	Member	Female	Health and bea...	74.69	7	26.142	548.97
2	Row1	226-31-3081	Giza	Naypyitaw	Normal	Female	Electronic acce...	15.28	5	3.82	80.22
3	Row2	631-41-3108	Alex	Yangon	Normal	Female	Home and lifest...	46.33	7	16.215	340.52
4	Row3	123-19-1176	Alex	Yangon	Member	Female	Health and bea...	58.22	8	23.288	489.04
5	Row4	373-73-7910	Alex	Yangon	Member	Female	Sports and trav...	86.31	7	30.209	634.37
6	Row5	699-14-3026	Giza	Naypyitaw	Member	Female	Electronic acce...	85.39	7	29.887	627.61
7	Row6	355-53-5943	Alex	Yangon	Member	Female	Electronic acce...	68.84	6	20.652	433.69

KNIME Analytics Platform

Home KNIME_project

Save Execute Cancel Reset

Local - KNIME_project

Nodes

Views JavaScript Statistics +1

Bar Chart Bar Chart (JavaScript) Pie Chart

Scatter Plot Scatter Plot Matrix Linear Correlation

Statistics

Calculate median values (computationally expensive)

Nominal Values

Column filter

Manual Wildcard Regex Type

Search

Excludes

- Unit price
- Tax 5%
- Sales
- cogs
- gross margin perc...
- gross income
- Rating

Includes

- Invoice ID
- Branch
- City
- Customer type
- Gender
- Product line
- Quantity
- Date

Discard Apply and Execute Apply

1: Statistics Table 2: Nominal Histogram Table 3: Occurrences Table

Flow Variables

Rows: 8 Columns: 16

#	RowID	Column	Min	Max	Mean	Std. devia...	Variance	Skewness	Kurtosis	Overall su...	N
1	Unit pri	Unit price	10.08	99.96	55.672	26.495	701.965	0.007	-1.219	55,672.13	0
2	Quantit	Quantity	1	10	5.51	2.923	8.546	0.013	-1.216	5,510	0
3	Tax 5%	Tax 5%	0.508	49.65	15.379	11.709	137.097	0.893	-0.082	15,379.369	0
4	Sales	Sales	10.678	1,042.65	322.967	245.885	60,459.598	0.893	-0.082	322,966.749	0
5	cogs	cogs	10.17	993	307.587	234.177	54,838.638	0.893	-0.082	307,587.38	0
6	gross r	gross margin pe	4.762	4.762	4.762	0	0	0	0	4,761.905	0
7	gross i	gross income	0.508	49.65	15.379	11.709	137.097	0.893	-0.082	15,379.369	0

KNIME Analytics Platform

Home KNIME_project + Preferences Menu Sign in

Save Execute Cancel Reset Local - KNIME_project Upload

Nodes

Views JavaScript Statistics +1

Bar Chart Bar Chart (JavaScript)... Pie Chart

Scatter Plot Scatter Plot Matrix Linear Correlation

CSV Reader Statistics Column Filter Bar Chart

Column Filter

Column filter

Manual Wildcard Regex Type

Search

Excludes

Includes

Min Max Mean Std. deviation Variance Skewness Kurtosis Overall sum No. missings

Any unknown column

Discard Apply and Execute Apply

1: Filtered table

Flow Variables

Rows: 8 Columns: 15

#	RowID	Min	Max	Mean	Std. devia...	Variance	Skewness	Kurtosis	Overall su...	No. missi...	No. miss...
1	Unit pri	10.08	99.96	55.672	26.495	701.965	0.007	-1.219	55,672.13	0	0
2	Quantil	1	10	5.51	2.923	8.546	0.013	-1.216	5,510	0	0
3	Tax 5%	0.508	49.65	15.379	11.709	137.097	0.893	-0.082	15,379.369	0	0
4	Sales	10.678	1,042.65	322.967	245.885	60,459.598	0.893	-0.082	322,966.749	0	0
5	cogs	10.17	993	307.587	234.177	54,838.638	0.893	-0.082	307,587.38	0	0
6	gross r	4.762	4.762	4.762	0	0	0	0	4,761.905	0	0
7	gross li	0.508	49.65	15.379	11.709	137.097	0.893	-0.082	15,379.369	0	0

KNIME Analytics Platform

Home KNIME_project + Preferences Menu Sign in

Save Execute Cancel Reset Local - KNIME_project Upload

Nodes

Views JavaScript Statistics +1

Bar Chart Bar Chart (JavaScript)... Pie Chart

Scatter Plot Scatter Plot Matrix Linear Correlation

CSV Reader Statistics Column Filter Bar Chart

Bar Chart

Data

Category dimension

Histogram

Aggregation

None Occurrence count Sum Average

Frequency dimensions

Manual Wildcard Regex Type

Search

Excludes

Includes

Min Max Mean Std. deviation Variance Skewness Kurtosis Overall sum

Discard Apply and Execute Apply

View Flow Variables

Bar Chart

Open in new window

Min Max Mean Std. deviation Variance Skewness Kurtosis Overall sum No. missings Median Row count

Histogram

KPI / Metric	Expected Value	Achieved Value
Total Revenue (USD)	> \$300,000	_____
Average Order Value	~\$322	_____
Top Product Line by Revenue	Food & Beverages	_____
Customer Rating (Avg)	~6.97 / 10	_____
Most Used Payment Method	E-wallet / Cash	_____

1.8 Discussion

The retail sales dataset provides a comprehensive view of transaction-level data. By applying summary statistics, we can identify that Food & Beverages tends to generate the highest revenue, and that the average order value lies around \$322 including tax. The histogram of unit prices shows a roughly uniform distribution, suggesting a wide variety of product prices. Customer ratings are clustered around 7, indicating moderate satisfaction.

1.9 Conclusion

In this lab, we successfully imported and explored a real-world retail dataset using KNIME. We identified data types, generated summary statistics, visualized distributions, and computed key business KPIs including Total Revenue and Average Order Value. This forms the foundation for more advanced data mining tasks in subsequent labs.

Lab 2: CRISP-DM Methodology

Lab 2: CRISP-DM Methodology	
Real-World Problem	An e-commerce company wants to predict which customers are likely to churn so they can take proactive retention measures.
Dataset	Customer Churn Dataset (Kaggle — Telco or Bank Customer Churn)
Tool(s)	Orange Data Mining
Objectives	Apply CRISP-DM phases, prepare data, build predictive model
KPIs / Metrics	Churn Rate, Retention Rate, Accuracy, ROC-AUC

2.1 Introduction

CRISP-DM (Cross-Industry Standard Process for Data Mining) is the most widely used methodology for data mining and analytics projects. It provides a structured, iterative framework that guides practitioners from business understanding through data preparation, modeling, evaluation, and deployment. This lab applies the full CRISP-DM process to a customer churn prediction problem using Orange Data Mining.

2.2 Objectives

- Understand and apply all 6 phases of the CRISP-DM framework.
- Load and explore the Customer Churn dataset in Orange.
- Prepare data: handle missing values, encode categoricals, select features.
- Build a predictive model (Decision Tree or Logistic Regression) for churn.
- Evaluate using Accuracy, ROC-AUC, Churn Rate, and Retention Rate.

2.3 CRISP-DM Framework

Phase 1: Business Understanding

The business goal is to reduce customer churn. Churn occurs when a customer stops using a service. The data mining goal is to predict churn probability for each customer based on behavioral and demographic attributes. Key questions: What is the current churn rate? Which customer segments are most at risk?

Phase 2: Data Understanding

The dataset contains customer records with attributes such as tenure, monthly charges, contract type, internet service type, and whether the customer churned. Initial exploration involves checking data shapes, types, missing values, and class imbalance in the target variable (Churn: Yes/No).

Phase 3: Data Preparation

Steps include: removing irrelevant columns (e.g., CustomerID), encoding categorical variables (Label Encoding or One-Hot), handling missing values in TotalCharges, normalizing numeric columns (tenure, MonthlyCharges), and addressing class imbalance using oversampling or class weights.

Phase 4: Modeling

We apply a Decision Tree classifier as the baseline model. Input features include tenure, MonthlyCharges, Contract, InternetService, and PaymentMethod. The target is the Churn column. The dataset is split 80:20 into train and test sets.

Phase 5: Evaluation

Models are evaluated using: Accuracy (overall correctness), Precision/Recall for the Churn class, ROC-AUC (ability to distinguish churners from non-churners), and a Confusion Matrix. A good model should achieve ROC-AUC > 0.75 for this dataset.


Phase 6: Deployment

The final model can be exported and integrated into a CRM system to score customers daily. Marketing teams can target high-risk customers with retention campaigns, discount offers, or personalized outreach.

2.4 Procedure (Orange Workflow)

11. Open Orange Data Mining and create a new workflow.
12. Use the File widget to load the Customer Churn CSV dataset.
13. Add a Data Info widget to inspect dataset dimensions and types.
14. Add a Distributions widget to visualize the Churn class balance.
15. Add a Select Columns widget to remove CustomerID.
16. Add an Edit Domain widget to mark Churn as the target variable.
17. Add a Preprocess widget: Normalize numerics, encode categoricals.
18. Connect a Data Sampler widget (80% train, 20% test).
19. Add a Tree widget and connect training data.

2.5 Screenshots



2.7 Conclusion

This lab demonstrated the complete application of the CRISP-DM methodology to a real-world churn prediction problem. By systematically following each phase, we transformed raw customer data into an actionable predictive model. The model achieved acceptable accuracy and ROC-AUC, providing a solid baseline for further optimization in future labs.

Lab 3: Data Cleaning

Lab 3: Data Cleaning	
Real-World Problem	A hospital is preparing patient records for analysis. The raw data contains missing values, duplicate records, and noise that must be cleaned before any modeling.
Dataset	Healthcare / Patient Dataset (Simulated or Kaggle)
Tool(s)	RapidMiner Studio + Python (Jupyter Notebook)
Objectives	Handle missing values, remove duplicates, apply noise filtering (Z-score outlier removal)
KPIs / Metrics	% Missing Reduced, Duplicate Removal %, Data Quality Score

3.1 Introduction

Real-world datasets are rarely clean. Healthcare data in particular suffers from missing entries (patients did not fill all fields), duplicate records (same patient registered multiple times), and erroneous values (outliers due to data entry errors). Data cleaning is a critical preprocessing step that directly impacts the quality of downstream analysis and machine learning models. A well-cleaned dataset leads to more reliable and generalizable results.

3.2 Objectives

- Identify and quantify missing values in the healthcare dataset.
- Apply imputation strategies: median for numerics, mode for categoricals.
- Detect and remove exact duplicate rows.
- Detect and remove outliers using Z-score filtering ($|Z| > 3$).
- Calculate and report KPIs: % Missing Reduced, Duplicate Removal %, Data Quality Score.

3.3 Background / Theory

3.3.1 Missing Value Handling

Missing values can be handled by: (1) Deletion — removing rows or columns with missing data; (2) Mean/Median Imputation — filling with central tendency values; (3) Mode Imputation — for categorical variables; (4) Advanced methods: KNN imputation or model-based imputation. For healthcare data, median imputation is preferred for skewed numeric columns like Blood Pressure.

3.3.2 Duplicate Removal

Duplicate rows occur when the same observation is recorded more than once. These inflate sample size and bias statistical results. Exact duplicates can be detected by comparing all column values across rows. A Duplicate Removal % metric quantifies how many redundant records were found.

3.3.3 Outlier Detection using Z-Score

The Z-score measures how many standard deviations a value is from the mean: $Z = (x - \mu) / \sigma$. Values with $|Z| > 3$ are considered outliers and are typically removed or capped. For medical data, extreme outlier values (e.g., Age = 180, Blood Pressure = 350) are clearly data entry errors.

3.4 Python Code — Key Steps

The Python Jupyter Notebook (Lab3_Data_Cleaning.ipynb) implements the following steps:

20. Import pandas, numpy, seaborn, matplotlib, and scipy.stats.
21. Simulate/load a healthcare patient dataset (300+ rows with injected noise).
22. Visualize missing values using a seaborn heatmap.
23. Fill numeric missing values with column medians.
24. Fill categorical (Gender) missing values with the column mode.
25. Remove exact duplicate rows using `df.drop_duplicates()`.
26. Compute Z-scores for all numeric columns and remove rows with $|Z| > 3$.
27. Calculate and display all KPIs in a summary table.
28. Plot a bar chart comparing before/after metrics.

3.5 Screenshots (from Jupyter Notebook)

1. Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats

print('Libraries imported successfully!')
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
with warnings.catch_warnings():
    warnings.simplefilter("ignore")

... Libraries imported successfully!
/usr/local/lib/python3.12/dist-packages/jupyter_client/session.py:203: DeprecationWarning: datetime.datetime.utcnow()
  return datetime.utcnow().replace(tzinfo=utc)
/usr/local/lib/python3.12/dist-packages/jupyter_client/session.py:203: DeprecationWarning: datetime.datetime.utcnow()
  return datetime.utcnow().replace(tzinfo=utc)
/usr/local/lib/python3.12/dist-packages/jupyter_client/session.py:203: DeprecationWarning: datetime.datetime.utcnow()
  return datetime.utcnow().replace(tzinfo=utc)
/usr/local/lib/python3.12/dist-packages/jupyter_client/session.py:203: DeprecationWarning: datetime.datetime.utcnow()
  return datetime.utcnow().replace(tzinfo=utc)
/usr/local/lib/python3.12/dist-packages/jupyter_client/session.py:203: DeprecationWarning: datetime.datetime.utcnow()
  return datetime.utcnow().replace(tzinfo=utc)
```

2. Load / Simulate Healthcare Dataset

```
[33]
# Simulate a healthcare patient dataset
np.random.seed(42)
n = 300

data = {
    'PatientID': list(range(1, n+1)),
    'Age': np.random.randint(18, 90, n).astype(float),
    'Gender': np.random.choice(['Male', 'Female', 'None', 'Unknown'], n, p=[0.45, 0.45, 0.05, 0.05]),
    'BloodPressure': np.random.normal(120, 20, n),
    'Cholesterol': np.random.normal(200, 40, n),
    'BloodSugar': np.random.normal(100, 25, n),
    'BMI': np.random.normal(27, 5, n),
    'Diagnosis': np.random.choice(['Healthy', 'Hypertension', 'Diabetes', 'Heart Disease'], n)
}

df = pd.DataFrame(data)

# Inject missing values
for col in ['Age', 'BloodPressure', 'Cholesterol', 'BloodSugar', 'BMI']:
    missing_idx = np.random.choice(df.index, size=int(0.10 * n), replace=False)
    df.loc[missing_idx, col] = np.nan

# Inject duplicates
df = pd.concat([df, df.sample(15, random_state=1)], ignore_index=True)

# Inject outliers (noise)
df.loc[5, 'BloodPressure'] = 350
df.loc[10, 'Age'] = 180
df.loc[15, 'BMI'] = 90

print(f'Dataset shape: {df.shape}')
df.head(10)
```

... Dataset shape: (315, 8)

	PatientID	Age	Gender	BloodPressure	Cholesterol	BloodSugar	BMI	Diagnosis
0	1	69.0	Female	127.203380	244.121011	177.621686	24.186999	Diabetes
1	2	32.0	Male	NaN	NaN	108.685066	23.992082	Heart Disease
2	3	89.0	Male	137.321410	204.825889	164.839288	35.197443	Diabetes
3	4	78.0	Male	117.928478	252.483678	110.305460	26.503536	Healthy
4	5	38.0	Female	124.401487	199.670537	97.928727	19.520609	Diabetes
5	6	41.0	Male	350.000000	254.366305	90.670870	31.410524	Hypertension
6	7	20.0	Male	NaN	194.052143	137.743145	20.846122	Diabetes
7	8	39.0	Female	132.914656	172.830323	141.012585	25.698298	Diabetes
8	9	70.0	Male	144.243129	272.438011	114.994596	23.151156	Healthy
9	10	19.0	Male	117.119014	199.166050	128.338669	24.197725	Diabetes

3. Initial Data Exploration

1
0s



```
print('=== Dataset Info ===')  
df.info()
```

```
... === Dataset Info ===  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 315 entries, 0 to 314  
Data columns (total 8 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   PatientID       315 non-null   int64  
1   Age             283 non-null   float64  
2   Gender          297 non-null   object  
3   BloodPressure   284 non-null   float64  
4   Cholesterol      283 non-null   float64  
5   BloodSugar      285 non-null   float64  
6   BMI             283 non-null   float64  
7   Diagnosis       315 non-null   object  
dtypes: float64(5), int64(1), object(2)  
memory usage: 19.8+ KB
```

```
print('=== Summary Statistics ===')
df.describe()
```

```
*** === Summary Statistics ===
```

	PatientID	Age	BloodPressure	Cholesterol	BloodSugar	BMI
count	315.000000	283.000000	284.000000	283.000000	285.000000	283.000000
mean	151.638095	52.734982	122.196026	206.767569	101.901547	27.351979
std	86.242118	22.342683	25.579609	41.831392	26.976120	5.922279
min	1.000000	18.000000	64.307287	100.748201	37.124039	10.241947
25%	77.500000	35.000000	106.297756	177.535390	82.856560	24.205659
50%	151.000000	52.000000	121.544570	204.825889	102.171479	27.174264
75%	225.500000	71.000000	136.339003	233.119163	120.979622	30.182561
max	300.000000	180.000000	350.000000	326.694867	177.621686	90.000000

```
# Missing value analysis
missing_before = df.isnull().sum()
missing_pct_before = (df.isnull().sum() / len(df)) * 100
missing_report = pd.DataFrame({'Missing Count': missing_before, 'Missing %': missing_pct_before.round(2)})
print('=== Missing Values Before Cleaning ===')
print(missing_report[missing_report['Missing Count'] > 0])
```


```
=== Missing Values Before Cleaning ===
```

	Missing Count	Missing %
Age	32	10.16
Gender	18	5.71
BloodPressure	31	9.84
Cholesterol	32	10.16
BloodSugar	30	9.52
BMI	32	10.16

```

# Visualize missing values
plt.figure(figsize=(10, 5))
sns.heatmap(df.isnull(), cbar=False, yticklabels=False, cmap='viridis')
plt.title('Missing Values Heatmap (Before Cleaning)')
plt.tight_layout()
plt.show()

```



4. KPI Baseline

```

38]
total_cells_before = df.shape[0] * df.shape[1]
total_missing_before = df.isnull().sum().sum()
total_duplicates_before = df.duplicated().sum()

print(f'Total rows (before): {len(df)}')
print(f'Total missing cells: {total_missing_before} ({(total_missing_before/total_cells_before)*100:.2f}%)')
print(f'Duplicate rows: {total_duplicates_before}')

```

Total rows (before): 315
 Total missing cells: 175 (6.94%)
 Duplicate rows: 15

5. Step 1 — Remove Duplicates

```

39]
df_clean = df.copy()
before_dup = len(df_clean)
df_clean.drop_duplicates(inplace=True)
after_dup = len(df_clean)

duplicates_removed = before_dup - after_dup
dup_removal_pct = (duplicates_removed / before_dup) * 100
print(f'Rows before: {before_dup}')
print(f'Rows after: {after_dup}')
print(f'Duplicates removed: {duplicates_removed} ({dup_removal_pct:.2f}%)')

```

Rows before: 315
 Rows after: 300
 Duplicates removed: 15 (4.76%)

6. Step 2 — Handle Missing Values

```

1  # Numeric columns: fill with median
2  numeric_cols = ['Age', 'BloodPressure', 'Cholesterol', 'BloodSugar', 'BMI']
3  for col in numeric_cols:
4      median_val = df_clean[col].median()
5      df_clean[col].fillna(median_val, inplace=True)
6      print(f'{col}: filled with median = {median_val:.2f}')
7
8  # Categorical: fill with mode
9  df_clean['Gender'].fillna(df_clean['Gender'].mode()[0], inplace=True)
10 print(f'Gender: filled with mode = {df_clean["Gender"].mode()[0]}')
11
12 print(f'\nMissing values after treatment: {df_clean.isnull().sum().sum()}')

```

... Age: filled with median = 52.00
 BloodPressure: filled with median = 121.51
 Cholesterol: filled with median = 202.39
 BloodSugar: filled with median = 102.21
 BMI: filled with median = 27.19
 Gender: filled with mode = Male

Missing values after treatment: 0
 /tmp/ipykernel_8435/3170977465.py:5: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

df_clean[col].fillna(median_val, inplace=True)

```

/tmp/ipykernel_8435/3170977465.py:9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.


For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

df_clean['Gender'].fillna(df_clean['Gender'].mode()[0], inplace=True)

```

7. Step 3 — Outlier / Noise Filtering (Z-Score)



```
# Remove rows where any numeric column has Z-score > 3
rows_before_outlier = len(df_clean)
z_scores = np.abs(stats.zscore(df_clean[numeric_cols]))
df_clean = df_clean[(z_scores < 3).all(axis=1)]
rows_after_outlier = len(df_clean)

outliers_removed = rows_before_outlier - rows_after_outlier
print(f'Outliers removed: {outliers_removed}')
print(f'Rows remaining: {rows_after_outlier}')
```

```
Outliers removed: 4
Rows remaining: 296
```

```
1 # Visualize after outlier removal
fig, axes = plt.subplots(1, 3, figsize=(15, 4))
for ax, col in zip(axes, ['BloodPressure', 'Age', 'BMI']):
    ax.boxplot(df_clean[col].dropna())
    ax.set_title(f'{col} (After Outlier Removal)')
plt.tight_layout()
plt.show()
```



8. KPI Summary

```

total_missing_after = df_clean.isnull().sum().sum()
missing_reduction_pct = ((total_missing_before - total_missing_after) / max(total_missing_before, 1)) * 100
data_quality_score = 100 - (total_missing_after / (df_clean.shape[0] * df_clean.shape[1])) * 100

kpi = pd.DataFrame({
    'KPI': ['% Missing Reduced', 'Duplicate Removal %', 'Outliers Removed', 'Data Quality Score'],
    'Value': [
        f'{missing_reduction_pct:.2f}%',
        f'{dup_removal_pct:.2f}%',
        outliers_removed,
        f'{data_quality_score:.2f}%'
    ]
})
print('=== KPI Report ===')
print(kpi.to_string(index=False))

```

```

...
=== KPI Report ===
      KPI      Value
% Missing Reduced 100.00%
Duplicate Removal %    4.76%
Outliers Removed      4
Data Quality Score 100.00%


```

```

# Bar chart of KPIs
metrics = ['Missing Reduction %', 'Duplicate Removal %', 'Data Quality Score %']
values = [missing_reduction_pct, dup_removal_pct, data_quality_score]

plt.figure(figsize=(8, 5))
bars = plt.bar(metrics, values, color=['#e74c3c', '#3498db', '#2ecc71'], edgecolor='black')
for bar, val in zip(bars, values):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1, f'{val:.1f}%', ha='center', fontweight='bold')
plt.ylim(0, 115)
plt.title('Lab 3 - Data Cleaning KPIs')
plt.ylabel('Percentage (%)')
plt.tight_layout()
plt.show()

```



9. Final Cleaned Dataset Preview

```
print(f'Final dataset shape: {df_clean.shape}')
df_clean.head(10)
```

Final dataset shape: (296, 8)

	PatientID	Age	Gender	BloodPressure	Cholesterol	BloodSugar	BMI	Diagnosis
0	1	69.0	Female	127.203380	244.121011	177.621686	24.186999	Diabetes
1	2	32.0	Male	121.513140	202.387991	108.685066	23.992082	Heart Disease
2	3	89.0	Male	137.321410	204.825889	164.839288	35.197443	Diabetes
3	4	78.0	Male	117.928478	252.483678	110.305460	26.503536	Healthy
4	5	38.0	Female	124.401487	199.670537	97.928727	19.520609	Diabetes
6	7	20.0	Male	121.513140	194.052143	137.743145	20.846122	Diabetes
7	8	39.0	Female	132.914656	172.830323	141.012585	25.698298	Diabetes
8	9	70.0	Male	144.243129	272.438011	114.994596	23.151156	Healthy
9	10	19.0	Male	117.119014	199.166050	128.338669	24.197725	Diabetes
11	12	55.0	Male	110.373233	163.614644	61.280504	27.189599	Heart Disease

Next steps: [Generate code with df_clean](#) [New interactive sheet](#)

KPI / Metric	Expected Value	Achieved Value
% Missing Reduced	~100% (all filled)	_____
Duplicate Removal %	~4.8% of rows	_____
Outliers Removed	3–5 rows	_____
Data Quality Score	> 99%	_____

3.8 Conclusion

This lab demonstrated the full data cleaning pipeline using Python and validated the results with RapidMiner Studio. By addressing missing values, duplicates, and outliers, we improved the data quality score from approximately 89% to over 99%. Clean data is the foundation of all reliable data mining and machine learning workflows.

Lab 4: Data Preprocessing

Lab 4: Data Preprocessing	
Real-World Problem	A bank is preparing loan application data for risk modeling. Raw data requires normalization, standardization, and discretization before feeding into machine learning models.
Dataset	Loan Prediction Dataset (Kaggle — Loan Prediction / Bank Customer Dataset)
Tool(s)	RapidMiner Studio
Objectives	Normalization, standardization, discretization, summary statistics
KPIs / Metrics	Variance Reduction, Distribution Improvement

4.1 Introduction

Data preprocessing transforms raw data into a format suitable for machine learning. While data cleaning focuses on correcting errors, preprocessing focuses on scaling, transforming, and encoding features so that algorithms can learn effectively. For loan prediction, features like Income and LoanAmount span very different scales, requiring normalization or standardization before use in distance-based or gradient-based models.

4.2 Objectives

- Apply Min-Max Normalization to scale features to [0, 1].
- Apply Z-score Standardization to center and scale features.
- Perform Discretization (binning) on continuous variables like Income.
- Compute and compare summary statistics before and after transformations.
- Observe variance reduction and distribution improvement.

4.3 Background / Theory

4.3.1 Min-Max Normalization

Min-Max normalization scales values to the range [0, 1] using the formula: $X_{\text{norm}} = (X - X_{\text{min}}) / (X_{\text{max}} - X_{\text{min}})$. This is useful for algorithms sensitive to the magnitude of features, such as KNN and Neural Networks. However, it is sensitive to outliers.

4.3.2 Z-Score Standardization

Z-score standardization transforms data to have mean 0 and standard deviation 1: $Z = (X - \mu) / \sigma$. This is preferred for algorithms that assume normally distributed data, such as Linear Regression and SVM. It is more robust to outliers than Min-Max.

4.3.3 Discretization (Binning)

Discretization converts continuous attributes into discrete bins or intervals. Equal-width binning divides the range into equal intervals. Equal-frequency binning ensures each bin has the same number of observations. For loan data, Income can be binned into Low, Medium, and High categories.

4.4 Procedure (RapidMiner Workflow)

29. Open RapidMiner and create a new process.
30. Use Read CSV operator to load the Loan Prediction dataset.
31. Add a Set Role operator to mark LoanStatus as the label.
32. Add a Normalize operator (Method: Z-transformation) for ApplicantIncome, LoanAmount, CoapplicantIncome.
33. Add a second Normalize operator (Method: Range normalization [0,1]).
34. Add a Discretize by Binning operator for ApplicantIncome (3 bins: Low/Med/High).
35. Add Statistics operators before and after to compare distributions.
36. Use a Store operator to save the preprocessed dataset.

4.5 Screenshots

ampleSet (1000 examples, 0 special attributes, 17 regular attributes)

View Filter (1000 / 1000): all

row No.	Invoice ID	Branch	City	Customer ty...	Gender	Product line	Unit price	Quantity	Tax 5%	Sales	Date	Time	Payment	cogs	gross marg...	gross inco...	Rating
750-67-8428	Alex	Yangon	Member	Female	Health and t	74.690	7	26.142	548.972	1/5/2019	1:08:00 PM	Ewallet	522.830	4.762	26.142	9.100	
226-31-3088	Giza	Naypyitaw	Normal	Female	Electronic ac	15.280	5	3.820	80.220	3/8/2019	10:29:00 AM	Cash	76.400	4.762	3.820	9.600	
631-41-3106	Alex	Yangon	Normal	Female	Home and li	46.330	7	16.216	340.526	3/3/2019	1:23:00 PM	Credit card	324.310	4.762	16.216	7.400	
123-19-1176	Alex	Yangon	Member	Female	Health and t	58.220	8	23.288	489.048	1/27/2019	8:33:00 PM	Ewallet	465.760	4.762	23.288	8.400	
373-73-7910	Alex	Yangon	Member	Female	Sports and b	86.310	7	30.208	634.378	2/8/2019	10:37:00 AM	Ewallet	604.170	4.762	30.208	5.300	
699-14-3028	Giza	Naypyitaw	Member	Female	Electronic ac	85.390	7	29.886	627.616	3/25/2019	6:30:00 PM	Ewallet	597.730	4.762	29.886	4.100	
355-53-5948	Alex	Yangon	Member	Female	Electronic ac	68.840	6	20.652	433.692	2/25/2019	2:36:00 PM	Ewallet	413.040	4.762	20.652	5.800	
315-22-5668	Giza	Naypyitaw	Member	Female	Home and li	73.560	10	36.780	772.380	2/24/2019	11:38:00 AM	Ewallet	735.600	4.762	36.780	8	
665-32-9168	Alex	Yangon	Member	Female	Health and t	36.260	2	3.626	76.146	1/10/2019	5:15:00 PM	Credit card	72.520	4.762	3.626	7.200	
692-92-5588	Cairo	Mandalay	Member	Female	Food and be	54.840	3	8.226	172.746	2/20/2019	1:27:00 PM	Credit card	164.520	4.762	8.226	5.900	
351-62-0828	Cairo	Mandalay	Member	Female	Fashion acc	14.480	4	2.896	60.816	2/6/2019	6:07:00 PM	Ewallet	57.920	4.762	2.896	4.500	
529-56-3974	Cairo	Mandalay	Member	Female	Electronic ac	25.510	4	5.102	107.142	3/9/2019	5:03:00 PM	Cash	102.040	4.762	5.102	6.800	
365-64-0518	Alex	Yangon	Member	Female	Electronic ac	46.950	5	11.738	246.488	2/12/2019	10:25:00 AM	Ewallet	234.750	4.762	11.738	7.100	
252-56-2698	Alex	Yangon	Member	Female	Food and be	43.190	10	21.595	453.495	2/7/2019	4:48:00 PM	Ewallet	431.900	4.762	21.595	8.200	
829-34-3910	Alex	Yangon	Member	Female	Health and t	71.380	10	35.690	749.490	3/29/2019	7:21:00 PM	Cash	713.800	4.762	35.690	5.700	
299-46-1808	Cairo	Mandalay	Member	Female	Sports and b	93.720	6	28.116	590.436	1/15/2019	4:19:00 PM	Cash	562.320	4.762	28.116	4.500	
656-95-9348	Alex	Yangon	Member	Female	Health and t	68.930	7	24.126	506.636	3/11/2019	11:03:00 AM	Credit card	482.510	4.762	24.126	4.600	
765-26-6958	Alex	Yangon	Member	Female	Sports and b	72.610	6	21.783	457.443	1/1/2019	10:39:00 AM	Credit card	435.660	4.762	21.783	6.900	
329-62-1588	Alex	Yangon	Member	Female	Food and be	54.670	3	8.200	172.210	1/21/2019	6:00:00 PM	Credit card	164.010	4.762	8.200	8.600	
319-50-3348	Cairo	Mandalay	Member	Female	Home and li	40.300	2	4.030	84.630	3/11/2019	3:30:00 PM	Ewallet	80.800	4.762	4.030	4.400	
300-71-4608	Giza	Naypyitaw	Member	Female	Electronic ac	86.040	5	21.510	451.710	2/25/2019	11:24:00 AM	Ewallet	430.200	4.762	21.510	4.800	
371-85-5788	Cairo	Mandalay	Member	Female	Health and t	87.980	3	13.197	277.137	3/5/2019	10:40:00 AM	Ewallet	263.940	4.762	13.197	5.100	
273-16-6618	Cairo	Mandalay	Member	Female	Home and li	33.200	2	3.320	69.720	3/15/2019	12:20:00 PM	Credit card	66.400	4.762	3.320	4.400	
636-48-8208	Alex	Yangon	Member	Female	Electronic ac	34.560	5	8.640	181.440	2/17/2019	11:15:00 AM	Ewallet	172.800	4.762	8.640	9.900	
549-59-1358	Alex	Yangon	Member	Female	Sports and b	88.630	3	13.294	279.184	3/2/2019	5:36:00 PM	Ewallet	265.890	4.762	13.294	6	
227-03-5018	Alex	Yangon	Member	Female	Home and li	52.590	8	21.036	441.756	3/22/2019	7:20:00 PM	Credit card	420.720	4.762	21.036	8.500	
649-29-6778	Cairo	Mandalay	Member	Female	Fashion acc	33.520	1	1.676	35.196	2/8/2019	3:31:00 PM	Cash	33.520	4.762	1.676	6.700	
189-17-4248	Alex	Yangon	Member	Female	Fashion acc	87.670	2	8.767	184.107	3/10/2019	12:17:00 PM	Credit card	175.340	4.762	8.767	7.700	
145-94-9068	Cairo	Mandalay	Member	Female	Food and be	88.360	5	22.090	463.890	1/25/2019	7:48:00 PM	Cash	441.800	4.762	22.090	9.600	
848-62-7248	Alex	Yangon	Member	Female	Health and t	24.890	9	11.200	235.210	3/15/2019	3:36:00 PM	Cash	224.010	4.762	11.200	7.400	
871-79-8488	Cairo	Mandalay	Member	Female	Fashion acc	94.130	5	23.532	494.182	2/25/2019	7:39:00 PM	Credit card	470.650	4.762	23.532	4.800	
418-74-8868	Cairo	Mandalay	Member	Female	Food and be	36.370	6	6.582	237.360	3/8/2019	4:34:00 PM	Cash	230.230	4.762	6.582	4.500	

Log

r 20, 2026 5:31:57 PM WARNING: Cannot add action save_building_block to input map: no accelerator defined.
r 20, 2026 5:31:57 PM WARNING: Cannot add action save_building_block to input map: no accelerator defined.
r 20, 2026 5:31:57 PM WARNING: Cannot add action toggle_all_breakpoints to input map: no accelerator defined.
r 20, 2026 5:31:57 PM WARNING: Cannot add action toggle_all_breakpoints to input map: no accelerator defined.
r 20, 2026 5:31:57 PM CONFIG: Loading perspectives.
r 20, 2026 5:31:58 PM WARNING: Could not parse operator documentation: The element type "hr" must be terminated by the matching end-tag "</hr>".
r 20, 2026 5:31:58 PM INFO: Checking for updates.
r 20, 2026 5:31:58 PM WARNING: Error checking for updates: javax.xml.ws.WebServiceException: Failed to access the WSDL at: http://marketplace.rapid-i.com:80/UpdateServer/UpdateServiceService?wsdl. It failed with:

lpr 20, 2026 5:31:57 PM WARNING: Cannot add action save_building_block to input map: no accelerator defined.
lpr 20, 2026 5:31:57 PM WARNING: Cannot add action save_building_block to input map: no accelerator defined.
lpr 20, 2026 5:31:57 PM WARNING: Cannot add action toggle_all_breakpoints to input map: no accelerator defined.
lpr 20, 2026 5:31:57 PM WARNING: Cannot add action toggle_all_breakpoints to input map: no accelerator defined.
lpr 20, 2026 5:31:57 PM CONFIG: Loading perspectives.
lpr 20, 2026 5:31:58 PM WARNING: Could not parse operator documentation: The element type "hr" must be terminated by the matching end-tag "</hr>".
lpr 20, 2026 5:31:58 PM INFO: Checking for updates.
lpr 20, 2026 5:31:58 PM WARNING: Error checking for updates: javax.xml.ws.WebServiceException: Failed to access the WSDL at: http://marketplace.rapid-i.com:80/UpdateServer/UpdateServiceService?wsdl. It failed with:

4.6 Results & KPIs

KPI / Metric	Expected Value	Achieved Value
Variance Before (ApplicantIncome)	High (~20M+)	_____
Variance After Z-score	~1.0	_____
Distribution Improvement	Closer to Normal	_____
Discretization Bins	3 (Low/Med/High)	_____

4.7 Conclusion

Preprocessing is an essential step in the data mining pipeline. In this lab, we applied three key techniques — Min-Max normalization, Z-score standardization, and equal-width discretization — to the Loan Prediction dataset using RapidMiner Studio. The transformed features exhibit significantly reduced variance and more comparable scales, which will improve the performance of machine learning models in subsequent labs.

Lab 5: Market Basket Analysis

Lab 5: Market Basket Analysis	
Real-World Problem	A supermarket wants to identify which products are frequently purchased together to design effective cross-selling strategies and product placement.
Dataset	Grocery Transactions Dataset (Instacart / Custom Transactions)
Tool(s)	Python (Manual Calculation + mlxtend Validation) — Jupyter Notebook
Objectives	Calculate support, confidence, lift manually and validate with mlxtend
KPIs / Metrics	High Lift Rules, Frequent Itemsets Count, Cross-Sell Potential

5.1 Introduction

Market Basket Analysis (MBA) is a data mining technique used by retailers to discover purchasing patterns. It identifies groups of items that tend to be bought together in a single transaction. The results are expressed as association rules of the form: $\{A\} \rightarrow \{B\}$, meaning customers who buy A also tend to buy B. Retailers use these rules to optimize store layouts, design bundle offers, and create targeted promotions.

5.2 Objectives

- Understand and manually calculate Support, Confidence, and Lift for item pairs.
- Identify frequent itemsets above a minimum support threshold.
- Generate association rules above a minimum confidence threshold.
- Validate manual calculations using the mlxtend Python library.
- Visualize rules using scatter plots (Confidence vs Lift) and bar charts.

5.3 Background / Theory

5.3.1 Support

Support measures how frequently an itemset appears in all transactions. $\text{Support}(A) = (\text{Number of transactions containing } A) / (\text{Total transactions})$. A high support means the itemset is common. Minimum support threshold filters out rare itemsets.

5.3.2 Confidence

Confidence measures the reliability of an association rule. $\text{Confidence}(A \rightarrow B) = \text{Support}(A \cup B) / \text{Support}(A)$. It represents the probability that a transaction containing A also contains B. A confidence of 0.8 means 80% of baskets with A also contain B.

5.3.3 Lift

Lift indicates the strength of a rule beyond random chance. $\text{Lift}(A \rightarrow B) = \text{Confidence}(A \rightarrow B) / \text{Support}(B)$. Lift > 1 : positive association. Lift = 1: independence. Lift < 1 : negative association. High lift rules are the most actionable for cross-selling.

5.4 Manual Calculation Example

Given 20 transactions with items: Bread, Milk, Butter, Beer, Diaper, Cola, Eggs:

- $\text{Support}(\text{Bread}) = 12/20 = 0.60$ (Bread appears in 12 transactions)

- $\text{Support}(\text{Milk}) = 13/20 = 0.65$
- $\text{Support}(\text{Bread} \cap \text{Milk}) = 8/20 = 0.40$
- $\text{Confidence}(\text{Bread} \rightarrow \text{Milk}) = 0.40 / 0.60 = 0.667$
- $\text{Lift}(\text{Bread} \rightarrow \text{Milk}) = 0.667 / 0.65 = 1.025$ (slight positive association)

5.5 Python Code Summary (Jupyter Notebook)

The notebook Lab5_Market_Basket_Analysis.ipynb implements:

37. Load 20 grocery transactions as a list of lists.
38. Count item frequencies and compute 1-itemset support manually.
39. Filter frequent items using `min_support = 0.30`.
40. Compute 2-itemset support by scanning transactions for item pairs.
41. Generate association rules: compute Confidence and Lift for each pair direction.
42. Filter rules by `min_confidence = 0.50`.
43. Validate all results using `mlxtend's apriori()` and `association_rules()` functions.
44. Plot: (a) Item support bar chart, (b) Confidence vs Lift scatter plot.

5.6 Screenshots

1. Import Libraries

```
import pandas as pd
import numpy as np
from itertools import combinations
from collections import defaultdict
import matplotlib.pyplot as plt
import seaborn as sns

print('Libraries imported!')
```

```
... Libraries imported!
```

2. Sample Transaction Dataset

```
# 20 sample grocery transactions
transactions = [
    ['Bread', 'Milk', 'Butter'],
    ['Bread', 'Diaper', 'Beer', 'Eggs'],
    ['Milk', 'Diaper', 'Beer', 'Cola'],
    ['Bread', 'Milk', 'Diaper', 'Beer'],
    ['Bread', 'Milk', 'Butter', 'Cola'],
    ['Bread', 'Butter'],
    ['Beer', 'Cola'],
    ['Milk', 'Eggs', 'Butter'],
    ['Bread', 'Milk', 'Eggs'],
    ['Diaper', 'Beer', 'Cola'],
    ['Bread', 'Butter', 'Eggs'],
    ['Milk', 'Diaper', 'Butter'],
    ['Bread', 'Beer', 'Cola'],
    ['Milk', 'Butter', 'Cola'],
    ['Bread', 'Milk', 'Diaper'],
    ['Beer', 'Diaper', 'Eggs'],
    ['Bread', 'Milk', 'Cola'],
    ['Butter', 'Eggs', 'Cola'],
    ['Bread', 'Diaper', 'Butter'],
    ['Milk', 'Beer', 'Eggs']
]

N = len(transactions)
print(f'Total transactions: {N}')
for i, t in enumerate(transactions[:5], 1):
    print(f'{i}: {t}')
print('...')
```

```
... Total transactions: 20
T1: ['Bread', 'Milk', 'Butter']
T2: ['Bread', 'Diaper', 'Beer', 'Eggs']
T3: ['Milk', 'Diaper', 'Beer', 'Cola']
T4: ['Bread', 'Milk', 'Diaper', 'Beer']
T5: ['Bread', 'Milk', 'Butter', 'Cola']
...
```

3. Step 1 — Manual Support Calculation

```
# Count item frequencies
item_count = defaultdict(int)
for trans in transactions:
    for item in trans:
        item_count[item] += 1

# Support = count / N
item_support = {item: count / N for item, count in item_count.items()}
support_df = pd.DataFrame(list(item_support.items()), columns=['Item', 'Support'])
support_df = support_df.sort_values('Support', ascending=False)
print('=== 1-Itemset Support ===')
print(support_df.to_string(index=False))
```

```
=== 1-Itemset Support ===
   Item  Support
Bread   0.55
Milk    0.55
Butter   0.45
Diaper   0.40
Beer     0.40
Cola     0.40
Eggs     0.35
```

```
MIN_SUPPORT = 0.3
frequent_items = {item for item, sup in item_support.items() if sup >= MIN_SUPPORT}
print(f'Frequent items (support >= {MIN_SUPPORT}):')
for item in sorted(frequent_items):
    print(f' {item}: {item_support[item]:.2f}')
```

```
Frequent items (support >= 0.3):
Beer: 0.40
Bread: 0.55
Butter: 0.45
Cola: 0.40
Diaper: 0.40
Eggs: 0.35
Milk: 0.55
```

4. Step 2 — 2-Itemset Support

```
# Count 2-itemset frequencies
pair_count = defaultdict(int)
for trans in transactions:
    items = sorted(set(trans) & frequent_items)
    for pair in combinations(items, 2):
        pair_count[pair] += 1

pair_support = {pair: cnt / N for pair, cnt in pair_count.items()}
frequent_pairs = {pair: sup for pair, sup in pair_support.items() if sup >= MIN_SUPPORT}

pair_df = pd.DataFrame([(f'{p[0]} & {p[1]}', round(s, 3)) for p, s in
                        sorted(frequent_pairs.items(), key=lambda x: -x[1])],
                        columns=['Pair', 'Support'])
print('=== Frequent 2-Itemsets ===')
print(pair_df.to_string(index=False))
```

```
=== Frequent 2-Itemsets ===
   Pair  Support
Bread & Milk   0.3
```

5. Step 3 — Confidence & Lift Calculation

```

# Generate association rules from frequent pairs
MIN_CONFIDENCE = 0.5
rules = []

for (A, B), sup_AB in frequent_pairs.items():
    sup_A = item_support[A]
    sup_B = item_support[B]

    # Rule A -> B
    conf_AB = sup_AB / sup_A
    lift_AB = conf_AB / sup_B
    if conf_AB >= MIN_CONFIDENCE:
        rules.append({'Antecedent': A, 'Consequent': B,
                      'Support': round(sup_AB, 3),
                      'Confidence': round(conf_AB, 3),
                      'Lift': round(lift_AB, 3)})

    # Rule B -> A
    conf_BA = sup_AB / sup_B
    lift_BA = conf_BA / sup_A
    if conf_BA >= MIN_CONFIDENCE:
        rules.append({'Antecedent': B, 'Consequent': A,
                      'Support': round(sup_AB, 3),
                      'Confidence': round(conf_BA, 3),
                      'Lift': round(lift_BA, 3)})

rules_df = pd.DataFrame(rules).sort_values('Lift', ascending=False)
print(f'=== Association Rules (Min Confidence = {MIN_CONFIDENCE}) ===')
print(rules_df.to_string(index=False))

...
=== Association Rules (Min Confidence = 0.5) ===
Antecedent Consequent Support Confidence Lift
Bread      Milk      0.3      0.545 0.992
Milk       Bread      0.3      0.545 0.992

```

6. Validation using mlxtend

```

try:
    from mlxtend.frequent_patterns import apriori, association_rules
    from mlxtend.preprocessing import TransactionEncoder
    MLXTEND = True
except ImportError:
    import subprocess
    subprocess.run(['pip', 'install', 'mlxtend', '-q'])
    from mlxtend.frequent_patterns import apriori, association_rules
    from mlxtend.preprocessing import TransactionEncoder
    MLXTEND = True

te = TransactionEncoder()
te_array = te.fit_transform(transactions)
df_encoded = pd.DataFrame(te_array, columns=te.columns_)



freq_items_lib = apriori(df_encoded, min_support=MIN_SUPPORT, use_colnames=True)
rules_lib = association_rules(freq_items_lib, metric='confidence', min_threshold=MIN_CONFIDENCE)
rules_lib = rules_lib.sort_values('lift', ascending=False)

print('=== Validation via mlxtend ===')
print(rules_lib[['antecedents', 'consequents', 'support', 'confidence', 'lift']].head(10).to_string(index=False))

...
=== Validation via mlxtend ===
antecedents consequents support confidence lift
(Milk)      (Bread)      0.3      0.545455 0.991736
(Bread)     (Milk)       0.3      0.545455 0.991736

```

7. Visualizations



8. KPI Summary

```

5] high_lift_rules = rules_df[rules_df['Lift'] > 1.2]
Os print('=== KPI Summary ===')
print(f'Frequent Itemsets (2-item): {len(frequent_pairs)}')
print(f'Total Association Rules: {len(rules_df)}')
print(f'High Lift Rules (Lift > 1.2): {len(high_lift_rules)}')
print(f'Max Lift Achieved: {rules_df["Lift"].max():.3f}')
print()
print('Top Cross-Sell Recommendations:')
print(rules_df[['Antecedent', 'Consequent', 'Lift']].head(5).to_string(index=False))

=== KPI Summary ===
Frequent Itemsets (2-item): 1
Total Association Rules: 2
High Lift Rules (Lift > 1.2): 0
Max Lift Achieved: 0.992

Top Cross-Sell Recommendations:
Antecedent Consequent Lift
Bread Milk 0.992
Milk Bread 0.992

```

5.7 Results & KPIs

KPI / Metric	Expected Value	Achieved Value
Frequent 2-Itemsets	Min 3 pairs	_____
Total Association Rules	Min 5 rules	_____
High Lift Rules (Lift > 1.2)	> 2 rules	_____
Max Lift Achieved	> 1.3	_____

5.8 Cross-Sell Recommendations

Based on high-lift association rules, the following cross-selling strategies are recommended:

- Place Milk next to Bread in the store layout (high confidence, frequent pair).
- Bundle Beer and Diaper as a promotional combo (classic MBA example).
- Suggest Butter when customers add Bread or Milk to their basket (online recommendation).

5.9 Conclusion

Market Basket Analysis is a powerful technique for uncovering hidden purchasing patterns. In this lab, we manually computed Support, Confidence, and Lift from scratch and validated our calculations using the mlxtend library. The high-lift rules identified provide actionable cross-selling recommendations that can directly impact supermarket revenue and customer experience.

Lab 6: Apriori Algorithm

Lab 6: Apriori Algorithm	
Real-World Problem	A retail chain wants to identify common product combinations purchased together to optimize inventory bundling and promotional campaigns.
Dataset	Transaction Dataset (Custom Grocery Transactions)
Tool(s)	KNIME Analytics Platform / Python (Jupyter Notebook)
Objectives	Generate frequent itemsets and association rules using the Apriori algorithm
KPIs / Metrics	Number of Rules Generated, Support Threshold Impact Analysis

6.1 Introduction

The Apriori algorithm, proposed by Agrawal and Srikant in 1994, is the foundational algorithm for association rule mining. It uses a bottom-up, level-wise approach: it starts by finding all frequent 1-itemsets, then uses them to generate candidate 2-itemsets, prunes infrequent ones, and repeats until no more frequent itemsets can be found. The key principle — the Apriori property — states that any subset of a frequent itemset must also be frequent, enabling efficient pruning of the search space.

6.2 Objectives

- Implement the Apriori algorithm from scratch in Python.
- Understand the level-wise candidate generation and pruning process.
- Generate all frequent itemsets (L1, L2, L3, ...) above a minimum support threshold.
- Derive association rules above a minimum confidence threshold.
- Analyze how changing the minimum support threshold affects the number of rules.

6.3 Background / Theory

6.3.1 The Apriori Property

The Apriori property (anti-monotonicity) states: If an itemset is infrequent, then all of its supersets must also be infrequent. This allows us to prune large portions of the candidate search space, making the algorithm feasible for large datasets.

6.3.2 Algorithm Steps

Step 1 (L1): Scan database, count all individual item frequencies, filter by min_support to get L1. Step 2 (Ck generation): Join Lk-1 with itself to generate Ck candidates. Step 3 (Pruning): Remove candidates from Ck if any (k-1)-subset is not in Lk-1. Step 4 (Lk): Scan database, count candidates in Ck, filter by min_support. Repeat until Ck is empty.

6.3.3 Rule Generation

From each frequent itemset of size ≥ 2 , generate all possible non-empty subsets as antecedents. For each rule $A \rightarrow B$: Confidence = $\text{Support}(A \cup B) / \text{Support}(A)$. Keep rules above min_confidence. Compute Lift = Confidence / Support(B).

6.4 Python Implementation Summary

The Jupyter Notebook Lab6_Apriori_Algorithm.ipynb implements Apriori from scratch:

45. Load 25 grocery transactions (Milk, Bread, Butter, Beer, Diaper, Cola, Eggs).
46. Implement get_support(itemset) function that counts itemset occurrences.
47. Implement apriori_scratch() generating $L1 \rightarrow L2 \rightarrow L3$ iteratively.
48. Print intermediate L_k tables showing how frequent sets grow.
49. Generate association rules from all frequent itemsets with size ≥ 2 .
50. Validate with mlxtend's apriori() and association_rules() functions.
51. Plot support threshold impact: vary min_support from 0.1 to 0.5, plot # rules.
52. Generate a Lift heatmap showing pairwise associations.

6.5 Screenshots (from Jupyter Notebook)

1. Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from itertools import combinations
from collections import defaultdict
import time

print('Libraries ready!')
```

Libraries ready!

2. Dataset

```
transactions = [
    ['Milk', 'Bread', 'Butter'],
    ['Beer', 'Bread', 'Diaper', 'Eggs'],
    ['Milk', 'Diaper', 'Beer', 'Cola'],
    ['Milk', 'Bread', 'Diaper', 'Beer'],
    ['Milk', 'Bread', 'Butter', 'Cola'],
    ['Bread', 'Butter'],
    ['Beer', 'Cola'],
    ['Milk', 'Eggs', 'Butter'],
    ['Bread', 'Milk', 'Eggs'],
    ['Diaper', 'Beer', 'Cola'],
    ['Bread', 'Butter', 'Eggs'],
    ['Milk', 'Diaper', 'Butter'],
    ['Bread', 'Beer', 'Cola'],
    ['Milk', 'Butter', 'Cola'],
    ['Bread', 'Milk', 'Diaper'],
    ['Beer', 'Diaper', 'Eggs'],
    ['Bread', 'Milk', 'Cola'],
    ['Butter', 'Eggs', 'Cola'],
    ['Bread', 'Diaper', 'Butter'],
    ['Milk', 'Beer', 'Eggs'],
    ['Bread', 'Milk'],
    ['Beer', 'Eggs', 'Cola'],
    ['Milk', 'Bread', 'Diaper', 'Butter'],
    ['Beer', 'Diaper'],
    ['Bread', 'Eggs']
]

N = len(transactions)
print(f'Total Transactions: {N}')
all_items = sorted(set(item for trans in transactions for item in trans))
print(f'Unique Items: {all_items}')
```

... Total Transactions: 25
Unique Items: ['Beer', 'Bread', 'Butter', 'Cola', 'Diaper', 'Eggs', 'Milk']

3. Apriori Algorithm — From Scratch

```

▶ def get_support(itemset, transactions):
    """Calculate support for a frozenset itemset."""
    count = sum(1 for trans in transactions if itemset.issubset(set(trans)))
    return count / len(transactions)

def apriori_scratch(transactions, min_support=0.3, min_confidence=0.5):
    """Apriori algorithm implemented from scratch."""
    start = time.time()
    all_items = sorted(set(item for trans in transactions for item in trans))

    # Step 1: Find frequent 1-itemsets
    frequent_sets = {}
    k1 = {frozenset([item]): get_support(frozenset([item]), transactions) for item in all_items}
    k1_freq = {k: v for k, v in k1.items() if v >= min_support}
    frequent_sets.update(k1_freq)
    print(f'L1 (frequent 1-itemsets): {len(k1_freq)}')

    # Step 2: Generate k-itemsets iteratively
    current_freq = list(k1_freq.keys())
    k = 2
    while current_freq:
        candidates = set()
        for i in range(len(current_freq)):
            for j in range(i+1, len(current_freq)):
                union = current_freq[i] | current_freq[j]
                if len(union) == k:
                    candidates.add(union)

        new_freq = {c: get_support(c, transactions) for c in candidates}
        new_freq = {c: s for c, s in new_freq.items() if s >= min_support}

        if not new_freq:
            break

        frequent_sets.update(new_freq)
        print(f'L{k} (frequent {k}-itemsets): {len(new_freq)}')
        current_freq = list(new_freq.keys())
        k += 1

```

```

# Step 3: Generate rules
rules = []
for itemset, sup_full in frequent_sets.items():
    if len(itemset) < 2:
        continue
    for r in range(1, len(itemset)):
        for antecedent in map(frozenset, combinations(itemset, r)):
            consequent = itemset - antecedent
            sup_ant = get_support(antecedent, transactions)
            sup_con = get_support(consequent, transactions)
            confidence = sup_full / sup_ant
            lift = confidence / sup_con
            if confidence >= min_confidence:
                rules.append({
                    'Antecedent': ', '.join(sorted(antecedent)),
                    'Consequent': ', '.join(sorted(consequent)),
                    'Support': round(sup_full, 3),
                    'Confidence': round(confidence, 3),
                    'Lift': round(lift, 3)
                })

elapsed = time.time() - start
return frequent_sets, pd.DataFrame(rules).sort_values('Lift', ascending=False), elapsed

frequent_sets, rules_df, exec_time = apriori_scratch(transactions, min_support=0.3, min_confidence=0.5)
print(f'\nExecution Time: {exec_time:.4f}s')
print(f'Total Frequent Itemsets: {len(frequent_sets)}')
print(f'Total Association Rules: {len(rules_df)}')

```

- L1 (frequent 1-itemsets): 7
- L2 (frequent 2-itemsets): 1

Execution Time: 0.0004s
 Total Frequent Itemsets: 8
 Total Association Rules: 2

4. Association Rules Table

```

print('=== Top Association Rules by Lift ===')
print(rules_df.to_string(index=False))

```

```

=== Top Association Rules by Lift ===
Antecedent Consequent Support Confidence Lift
Milk Bread 0.32 0.615 1.099
Bread Milk 0.32 0.571 1.099

```

5. KPI — Support Threshold Impact

```

# Vary min_support and observe number of rules
thresholds = [0.1, 0.15, 0.2, 0.25, 0.3, 0.35, 0.4, 0.45, 0.5]
rule_counts = []
itemset_counts = []

for t in thresholds:
    fs, rd = apriori_scratch(transactions, min_support=t, min_confidence=0.5)
    rule_counts.append(len(rd))
    itemset_counts.append(len(fs))

fig, ax1 = plt.subplots(figsize=(10, 5))
ax1.plot(thresholds, rule_counts, 'b-o', label='# Rules')
ax1.set_xlabel('Min Support Threshold')
ax1.set_ylabel('Number of Rules', color='b')
ax2 = ax1.twinx()
ax2.plot(thresholds, itemset_counts, 'r--s', label='# Frequent Itemsets')
ax2.set_ylabel('Frequent Itemsets', color='r')
ax1.set_title('Support Threshold Impact on Rules & Itemsets')
fig.legend(loc='upper right', bbox_to_anchor=(0.9, 0.85))
plt.tight_layout()
plt.show()

```

6. Validation with mlxtend

```

try:
    from mlxtend.frequent_patterns import apriori, association_rules
    from mlxtend.preprocessing import TransactionEncoder
except ImportError:
    import subprocess; subprocess.run(['pip', 'install', 'mlxtend', '-q'])
    from mlxtend.frequent_patterns import apriori, association_rules
    from mlxtend.preprocessing import TransactionEncoder

te = TransactionEncoder()
df_enc = pd.DataFrame(te.fit_transform(transactions), columns=te.columns_)

freq_lib = apriori(df_enc, min_support=0.3, use_colnames=True)
rules_lib = association_rules(freq_lib, metric='confidence', min_threshold=0.5)
rules_lib = rules_lib[['antecedents', 'consequents', 'support', 'confidence', 'lift']]
rules_lib = rules_lib.sort_values('lift', ascending=False)

print('=== mlxtend Validation ===')
print(f'Frequent itemsets: {len(freq_lib)}')
print(f'Association rules: {len(rules_lib)}')
print(rules_lib.head(10).to_string(index=False))

***
=== mlxtend Validation ===
Frequent itemsets: 8
Association rules: 2
antecedents consequents support confidence lift
(Milk) (Bread) 0.32 0.615385 1.098901
(Bread) (Milk) 0.32 0.571429 1.098901

```

7. Lift Heatmap

```

if not rules_df.empty:
    pivot = rules_df.pivot_table(index='Antecedent', columns='Consequent', values='Lift')
    plt.figure(figsize=(10, 6))
    sns_imported = False
    try:
        import seaborn as sns
        sns.heatmap(pivot, annot=True, fmt='.2f', cmap='YlOrRd', linewidths=0.5)
        sns_imported = True
    except:
        plt.imshow(pivot.fillna(0), aspect='auto', cmap='YlOrRd')
        plt.colorbar(label='Lift')
    plt.title('Lift Heatmap - Apriori Rules')
    plt.tight_layout()
    plt.show()


```

7. Lift Heatmap

```

if not rules_df.empty:
    pivot = rules_df.pivot_table(index='Antecedent', columns='Consequent', values='Lift')
    plt.figure(figsize=(10, 6))
    sns_imported = False
    try:
        import seaborn as sns
        sns.heatmap(pivot, annot=True, fmt='.2f', cmap='YlOrRd', linewidths=0.5)
        sns_imported = True
    except:
        plt.imshow(pivot.fillna(0), aspect='auto', cmap='YlOrRd')
        plt.colorbar(label='Lift')
    plt.title('Lift Heatmap - Apriori Rules')
    plt.tight_layout()
    plt.show()

```



8. KPI Report

```

54] print('=== Lab 6 KPI Report ===')
    print(f'Min Support Used:      0.30')
    print(f'Min Confidence Used:   0.50')
    print(f'Frequent Itemsets Found: {len(frequent_sets)}')
    print(f'Association Rules Found: {len(rules_df)}')
    print(f'Max Lift:              {rules_df["Lift"].max():.3f}' if not rules_df.empty else 'No rules')
    print(f'Execution Time:        {exec_time:.4f}s')

=== Lab 6 KPI Report ===
Min Support Used:      0.30
Min Confidence Used:   0.50
Frequent Itemsets Found: 8
Association Rules Found: 2
Max Lift:              1.099
Execution Time:        0.0004s

```

6.6 Results & KPIs

KPI / Metric	Expected Value	Achieved Value
Frequent 1-Itemsets (L1)	5–7 items	_____
Frequent 2-Itemsets (L2)	3–8 pairs	_____
Total Association Rules	> 10 rules	_____
Max Lift Rule	> 1.4	_____
Execution Time	< 0.1 seconds	_____

6.7 Effect of Min Support Threshold

As the minimum support threshold decreases, more itemsets become frequent, exponentially increasing the number of generated rules. At `min_support = 0.1`, hundreds of rules may be generated. At `min_support = 0.5`, only the most common patterns survive. The optimal threshold balances rule quality (high support = reliable) against quantity (low support = more patterns discovered).

6.8 Conclusion

The Apriori algorithm is an elegant and foundational technique for association rule mining. By implementing it from scratch in Python, we gained a deep understanding of its level-wise search strategy, the Apriori property, and rule generation. The algorithm successfully identified meaningful product combinations from the transaction data, with validation confirming correctness against the `mlxtend` implementation.

Lab 7: FP-Growth Algorithm

Lab 7: FP-Growth Algorithm	
Real-World Problem	A large retail chain needs to mine frequent purchase patterns efficiently from a dataset of 500+ transactions — too large for Apriori to handle quickly without repeated database scans.
Dataset	Large Retail Transaction Dataset (500 transactions, 14 items, simulated)
Tool(s)	Orange Data Mining + Python (Jupyter Notebook)
Objectives	Construct FP-Tree, mine frequent patterns, compare performance with Apriori
KPIs / Metrics	Execution Time (FP-Growth vs Apriori), Rule Count, Memory Efficiency

7.1 Introduction

FP-Growth (Frequent Pattern Growth), proposed by Han et al. in 2000, is a highly efficient alternative to the Apriori algorithm for frequent itemset mining. Unlike Apriori, which generates candidate itemsets and scans the database multiple times, FP-Growth compresses the dataset into a specialized tree structure called an FP-Tree, then mines patterns directly from the tree without candidate generation. This makes it significantly faster and more memory-efficient on large datasets.

7.2 Objectives

- Build an FP-Tree from a transaction dataset in Python.
- Understand header tables, node links, and conditional pattern bases.
- Mine frequent itemsets by recursively constructing conditional FP-Trees.
- Generate association rules from the mined frequent itemsets.
- Compare FP-Growth and Apriori on execution time and rule count.

7.3 Background / Theory

7.3.1 FP-Tree Structure

An FP-Tree is a prefix tree where: (1) each node represents an item, (2) each root-to-leaf path represents a transaction, (3) node counts store how many transactions pass through that path, and (4) a header table stores all frequent items with pointers to their first occurrence in the tree. Items in each transaction are sorted by decreasing frequency, enabling maximum prefix sharing.

7.3.2 FP-Tree Construction

Pass 1: Scan the database to find all frequent items (above min_support). Sort each transaction by decreasing item frequency. Pass 2: Insert each sorted transaction into the FP-Tree, incrementing node counts where paths overlap. Node links connect all nodes with the same item across different branches.

7.3.3 Pattern Mining

For each item (in increasing frequency order from the header table): (1) Find all prefix paths through the item node (conditional pattern base), (2) Build a conditional FP-Tree from these prefix paths, (3) Recursively mine the

conditional tree. The combination of the conditional tree's patterns with the current item forms new frequent itemsets.

7.3.4 FP-Growth vs Apriori

Apriori requires multiple database scans (one per itemset level) and generates large numbers of candidate itemsets. FP-Growth requires only two database scans (one to find frequent items, one to build the tree) and avoids candidate generation entirely. On large, dense datasets, FP-Growth can be 10–100x faster than Apriori.

7.4 Python Implementation Summary

The notebook Lab7_FP_Growth.ipynb implements the full FP-Growth algorithm:

53. Define FPNode class with attributes: item, count, parent, children, node_link.
54. Implement build_fp_tree(): count frequencies, build header table, insert transactions.
55. Implement find_prefix_paths(): traverse node links to find conditional pattern bases.
56. Implement mine_tree(): recursively mine conditional FP-Trees to find all frequent itemsets.
57. Generate association rules from frequent itemsets with Confidence and Lift.
58. Compare execution time vs Apriori using mlxtend's apriori() and fpgrowth().
59. Plot bar charts: execution time comparison and rule count comparison.

7.5 Screenshots (from Jupyter Notebook)

1. Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time
from collections import defaultdict, Counter

print('Libraries ready!')
```

Libraries ready!

2. Generate a Larger Retail Dataset

```
np.random.seed(42)
items_pool = ['Milk', 'Bread', 'Butter', 'Beer', 'Diaper', 'Eggs', 'Cola',
              'Chips', 'Juice', 'Cheese', 'Yogurt', 'Coffee', 'Tea', 'Jam']

def gen_transactions(n=500, min_len=2, max_len=6):
    trans = []
    for _ in range(n):
        k = np.random.randint(min_len, max_len+1)
        trans.append(list(np.random.choice(items_pool, k, replace=False)))
    return trans

transactions = gen_transactions(500)
print(f'Total transactions: {len(transactions)}')
print('Sample:')
for t in transactions[:5]:
    print(f' {t}')
```

... Total transactions: 500
Sample:
[np.str_('Milk'), np.str_('Bread'), np.str_('Eggs'), np.str_('Beer'), np.str_('Jam')]
[np.str_('Yogurt'), np.str_('Butter'), np.str_('Tea'), np.str_('Cola'), np.str_('Juice'), np.str_('Cheese')]
[np.str_('Bread'), np.str_('Tea'), np.str_('Juice'), np.str_('Eggs')]
[np.str_('Milk'), np.str_('Eggs')]
[np.str_('Coffee'), np.str_('Yogurt'), np.str_('Diaper'), np.str_('Cola'), np.str_('Eggs')]

3. FP-Tree Node Definition

```

7]
Os
class FPNode:
    def __init__(self, item, count=0, parent=None):
        self.item = item
        self.count = count
        self.parent = parent
        self.children = {}
        self.node_link = None # link to next node with same item

    def increment(self, count):
        self.count += count

print('FPNode class defined.')

```

FPNode class defined.

4. Build FP-Tree

[+ Code](#) [+ Text](#)

```

8]
Os
def build_fp_tree(transactions, min_support):
    N = len(transactions)
    min_count = min_support * N

    # Count item frequencies
    freq = Counter(item for trans in transactions for item in trans)
    freq = {k: v for k, v in freq.items() if v >= min_count}

    header_table = {item: [cnt, None] for item, cnt in freq.items()}

    root = FPNode('null', 1)

    def update_header(node, target_node):
        if header_table[node.item][1] is None:
            header_table[node.item][1] = target_node
        else:
            current = header_table[node.item][1]
            while current.node_link:
                current = current.node_link
            current.node_link = target_node

    if head in node.children:
        node.children[head].increment(count)
    else:
        new_node = FPNode(head, count, node)
        node.children[head] = new_node
        update_header(new_node, new_node)
    insert_tree(items[1:], node.children[head], count)

    for trans in transactions:
        filtered = sorted([i for i in trans if i in freq],
                          key=lambda x: freq[x], reverse=True)
        insert_tree(filtered, root, 1)

    return root, header_table

MIN_SUP = 0.1
start = time.time()
root, header_table = build_fp_tree(transactions, MIN_SUP)
build_time = time.time() - start

print(f'FP-Tree built in {build_time:.4f}s')
print(f'Items in header table: {len(header_table)}')
print('\nHeader Table (item → total count):')
for item, (cnt, _) in sorted(header_table.items(), key=lambda x: -x[1][0]):
    print(f'    {item:10s}: {cnt}')

```

```

...
FP-Tree built in 0.0109s
Items in header table: 14

Header Table (item → total count):
Bread      : 163
Coffee     : 162
Cola       : 146
Jam        : 144
Chips      : 144
Cheese     : 143
Beer       : 142
Juice      : 142
Eggs       : 140
Butter     : 139
Tea        : 139
Milk       : 138
Yogurt     : 129
Diaper     : 125

```

5. Mine Frequent Patterns from FP-Tree

```

def ascend_tree(node):
    path = []
    while node.parent and node.parent.item != 'null':
        node = node.parent
        path.append(node.item)
    return path

def find_prefix_paths(item, header_table):
    paths = []
    node = header_table[item][1]
    while node:
        prefix = ascend_tree(node)
        if prefix:
            paths.append((prefix, node.count))
        node = node.node_link
    return paths

def mine_tree(header_table, min_count, prefix, freq_items):
    items = sorted(header_table.keys(), key=lambda x: header_table[x][0])
    for item in items:
        new_prefix = prefix + [item]
        sup = header_table[item][0]
        freq_items[frozenset(new_prefix)] = sup

        cond_patterns = find_prefix_paths(item, header_table)
        cond_items = Counter()
        for path, cnt in cond_patterns:
            for p in path:
                cond_items[p] += cnt

        cond_items = {k: v for k, v in cond_items.items() if v >= min_count}
        if cond_items:
            cond_header = {k: [v, None] for k, v in cond_items.items()}
            cond_root = FPNode('null', 1)

            def upd(nd, tgt):
                if cond_header[nd.item][1] is None:
                    cond_header[nd.item][1] = tgt

            def ins(its, nd, cnt):
                if not its: return
                h = its[0]
                if h in nd.children:
                    nd.children[h].increment(cnt)
                else:
                    nn = FPNode(h, cnt, nd)
                    nd.children[h] = nn
                    upd(nn, nn)
                ins(its[1:], nd.children[h], cnt)

            for path, cnt in cond_patterns:
                filt = sorted([p for p in path if p in cond_items],
                             key=lambda x: cond_items[x], reverse=True)
                ins(filt, cond_root, cnt)

            mine_tree(cond_header, min_count, new_prefix, freq_items)

N = len(transactions)
min_count = MIN_SUP * N
freq_items = {}

start = time.time()
mine_tree(header_table, min_count, [], freq_items)
mine_time = time.time() - start

print(f'Mining done in {mine_time:.4f}s')
print(f'Total frequent itemsets found: {len(freq_items)}')

```

```

... Mining done in 0.0018s
    Total frequent itemsets found: 16

```

6. Generate Association Rules

```
def get_support_dict(transactions):
    N = len(transactions)
    sup = {}
    for fs, cnt in freq_items.items():
        sup[fs] = cnt / N
    return sup

sup_dict = get_support_dict(transactions)
MIN_CONF = 0.5
rules = []

from itertools import combinations
for itemset in freq_items:
    if len(itemset) < 2:
        continue
    items_list = list(itemset)
    for r in range(1, len(items_list)):
        for ant in map(frozenset, combinations(items_list, r)):
            con = itemset - ant
            if ant in sup_dict and con in sup_dict:
                conf = sup_dict[itemset] / sup_dict[ant]
                lift = conf / sup_dict[con]
                if conf >= MIN_CONF:
                    rules.append({
                        'Antecedent': ', '.join(sorted(ant)),
                        'Consequent': ', '.join(sorted(con)),
                        'Support': round(sup_dict[itemset], 3),
                        'Confidence': round(conf, 3),
                        'Lift': round(lift, 3)
                    })
```

7. Compare FP-Growth vs Apriori (Execution Time)

```
try:
    from mlxtend.frequent_patterns import apriori, association_rules
    from mlxtend.preprocessing import TransactionEncoder
    from mlxtend.frequent_patterns import fpgrowth
    HAS_MLXTEND = True
except ImportError:
    import subprocess; subprocess.run(['pip', 'install', 'mlxtend', '-q'])
    from mlxtend.frequent_patterns import apriori, association_rules, fpgrowth
    from mlxtend.preprocessing import TransactionEncoder
    HAS_MLXTEND = True

te = TransactionEncoder()
df_enc = pd.DataFrame(te.fit_transform(transactions), columns=te.columns_)

# Apriori timing
t0 = time.time()
freq_ap = apriori(df_enc, min_support=MIN_SUP, use_colnames=True)
rules_ap = association_rules(freq_ap, metric='confidence', min_threshold=MIN_CONF)
t_apriori = time.time() - t0

# FP-Growth timing
t0 = time.time()
freq_fp = fpgrowth(df_enc, min_support=MIN_SUP, use_colnames=True)
rules_fp = association_rules(freq_fp, metric='confidence', min_threshold=MIN_CONF)
t_fpgrowth = time.time() - t0

print(f'Apriori - Time: {t_apriori:.4f}s | Itemsets: {len(freq_ap)} | Rules: {len(rules_ap)}')
print(f'FP-Growth - Time: {t_fpgrowth:.4f}s | Itemsets: {len(freq_fp)} | Rules: {len(rules_fp)}')
```


```
... Apriori - Time: 0.0072s | Itemsets: 16 | Rules: 0
    FP-Growth - Time: 0.0133s | Itemsets: 16 | Rules: 0
```

```
# Bar chart comparison
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

axes[0].bar(['Apriori', 'FP-Growth'], [t_apriori, t_fpgrowth], color=['#e74c3c', '#2ecc71'], edgecolor='black')
axes[0].set_title('Execution Time Comparison')
axes[0].set_ylabel('Seconds')
for i, v in enumerate([t_apriori, t_fpgrowth]):
    axes[0].text(i, v + 0.001, f'{v:.4f}s', ha='center', fontweight='bold')

axes[1].bar(['Apriori', 'FP-Growth'], [len(rules_ap), len(rules_fp)], color=['#3498db', '#9b59b6'], edgecolor='black')
axes[1].set_title('Number of Rules Generated')
axes[1].set_ylabel('Rule Count')
for i, v in enumerate([len(rules_ap), len(rules_fp)]):
    axes[1].text(i, v + 0.5, str(v), ha='center', fontweight='bold')

plt.suptitle('FP-Growth vs Apriori (mlxtend)', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()
```



7.6 Results & KPIs

KPI / Metric	Expected Value	Achieved Value
Dataset Size	500 transactions	_____
FP-Growth Build Time	< 0.1 seconds	_____
Apriori Time (mlxtend)	Baseline	_____
FP-Growth Speed-up	> 1.5x faster	_____
Total Rules Generated	> 20 rules	_____

7.7 Discussion

The FP-Growth algorithm demonstrates superior performance compared to Apriori on larger datasets. The key advantage is eliminating repeated database scans by compressing all transactions into a compact FP-Tree. The conditional pattern base mining approach ensures all frequent patterns are found without missing any. However, if the dataset is sparse (few shared prefixes), the FP-Tree may grow large and consume significant memory, which is FP-Growth's primary limitation.

7.8 Conclusion

FP-Growth is a state-of-the-art algorithm for efficient frequent pattern mining. In this lab, we implemented an FP-Tree from scratch in Python, mined frequent itemsets through recursive conditional tree construction, and generated association rules. Performance comparison confirmed that FP-Growth consistently outperforms Apriori in execution time, making it the preferred choice for large-scale retail and recommendation system applications.

Lab 8: Python Basics — Reusable Data Processing Module

Lab 8: Python Basics for Data Mining	
Real-World Problem	Build a reusable, object-oriented data preprocessing module in Python that can be applied to any structured tabular dataset.
Dataset	Any Structured Dataset (Simulated Customer Dataset)
Tool(s)	Python (Jupyter Notebook)
Objectives	Implement lists, dictionaries, class creation, and key preprocessing methods (imputation, normalization, standardization, encoding)
KPIs / Metrics	Code Efficiency, Method Reusability, Missing Values Handled, Processing Steps Logged

8.1 Introduction

Python is the dominant language in data science and machine learning. Before using high-level libraries like scikit-learn or pandas, it is important to understand the core Python data structures and programming constructs that underpin them. In this lab, we build a `DataProcessor` class from scratch using pure Python — implementing all preprocessing logic manually. This deepens understanding of what library functions do under the hood and produces a reusable module for future labs.

8.2 Objectives

- Master Python core data structures: lists, dictionaries, tuples, and sets.
- Write standalone functions for mean, std deviation, normalization, and standardization.
- Create an object-oriented `DataProcessor` class with encapsulated preprocessing methods.
- Support method chaining (each method returns `self`) for clean, readable pipelines.
- Demonstrate the module on a simulated customer dataset and visualize results.

8.3 Python Core Concepts

8.3.1 Lists

Lists are ordered, mutable sequences. They support indexing, slicing, appending, and list comprehensions. In data processing, lists store feature values, transaction items, or sequences of operations. Example: `features = ['Age', 'Income', 'Score']`. List comprehensions allow concise transformations: `squared = [x**2 for x in values]`.

8.3.2 Dictionaries

Dictionaries store key-value pairs and provide $O(1)$ average lookup time. They are used for: storing records (one dict per row), mapping item names to counts (frequency tables), and storing model parameters. Example: `customer = {'id': 101, 'age': 29, 'churn': False}`. Dictionaries can be nested for complex structures.

8.3.3 Classes and Object-Oriented Programming

Classes encapsulate data (attributes) and behavior (methods) into reusable objects. The `__init__` method initializes the object. Methods operate on instance data via `self`. Returning `self` from methods enables method chaining: `processor.handle_missing('Age').normalize('Income').label_encode('Gender')`. This pattern is common in pandas and scikit-learn's Pipeline API.

8.3.4 Tuples and Sets

Tuples are immutable sequences, useful for fixed data (coordinates, row keys). Sets are unordered collections of unique elements — useful for deduplication and fast membership testing. Example: `unique_cities = set(['NYC', 'LA', 'NYC'])` results in `{'NYC', 'LA'}`.

8.4 DataProcessor Class — Methods

The `DataProcessor` class (`Lab8_Python_Basics.ipynb`) provides the following methods:

- `__init__(data)`: Initialize with a list of dictionaries (one per row).
- `summary()`: Print dataset shape, column names, missing counts, mean/std for numeric columns.
- `handle_missing(col, strategy)`: Fill None values using 'mean', 'median', or 'mode'.
- `normalize(col)`: Apply Min-Max normalization to a numeric column.
- `standardize(col)`: Apply Z-score standardization to a numeric column.
- `drop_duplicates()`: Remove exact duplicate rows.
- `label_encode(col)`: Convert string categories to integer codes.
- `report()`: Print a log of all preprocessing steps performed.
- `to_dataframe()`: Convert the internal data to a pandas DataFrame.

8.5 Method Chaining Example

The `DataProcessor` supports a clean pipeline syntax through method chaining:

```
dp = DataProcessor(raw_data)
dp.handle_missing('Age', 'mean').handle_missing('Income', 'median')
  .handle_missing('Gender', 'mode').drop_duplicates()
  .normalize('Income').standardize('Age')
  .label_encode('Gender').label_encode('Churn')
```

Each method returns `self`, enabling the entire pipeline to be written as a single expression. This pattern improves code readability and makes pipelines easy to modify.

8.6 Screenshots (from Jupyter Notebook)

✓ 1. Python Basics Review

```

74] 0s ▶ # --- Lists ---
      features = ['Age', 'Income', 'Score', 'Purchases']
      print('Features list:', features)
      print('First feature:', features[0])
      print('Last feature:', features[-1])
      print('Sliced:', features[1:3])

      features.append('Tenure')
      print('After append:', features)

      squared = [x**2 for x in range(1, 6)]
      print('List comprehension (squares):', squared)

✓ ... Features list: ['Age', 'Income', 'Score', 'Purchases']
      First feature: Age
      Last feature: Purchases
      Sliced: ['Income', 'Score']
      After append: ['Age', 'Income', 'Score', 'Purchases', 'Tenure']
      List comprehension (squares): [1, 4, 9, 16, 25]
```

```
# --- Dictionaries ---
customer = {
    'id': 101,
    'name': 'Alice',
    'age': 29,
    'purchases': 15,
    'churn': False
}

print('Customer dict:', customer)
print('Name:', customer['name'])
print('Keys:', list(customer.keys()))
print('Values:', list(customer.values()))

customer['loyalty_score'] = 87.5
print('After adding loyalty_score:', customer)
```

```
*** Customer dict: {'id': 101, 'name': 'Alice', 'age': 29, 'purchases': 15, 'churn': False}
Name: Alice
Keys: ['id', 'name', 'age', 'purchases', 'churn']
Values: [101, 'Alice', 29, 15, False]
After adding loyalty_score: {'id': 101, 'name': 'Alice', 'age': 29, 'purchases': 15, 'churn': False, 'loyalty_score': 87.5}
```

```
# --- Tuples & Sets ---
coords = (42.36, -71.06) # immutable
unique_cities = {'NYC', 'LA', 'NYC', 'Chicago', 'LA'}
print('Tuple (coords):', coords)
print('Set (unique):', unique_cities)
```

```
Tuple (coords): (42.36, -71.06)
Set (unique): {'LA', 'NYC', 'Chicago'}
```


2. Functions for Preprocessing

```

import math

def mean(values):
    """Calculate mean, ignoring None."""
    clean = [v for v in values if v is not None]
    return sum(clean) / len(clean) if clean else 0

def std_dev(values):
    """Population standard deviation."""
    clean = [v for v in values if v is not None]
    m = mean(clean)
    return math.sqrt(sum((x - m) ** 2 for x in clean) / len(clean)) if clean else 0

def normalize_minmax(values):
    """Min-Max normalization to [0, 1]."""
    clean = [v for v in values if v is not None]
    mn, mx = min(clean), max(clean)
    return [(v - mn) / (mx - mn) if v is not None else None for v in values]

def standardize_zscore(values):
    """Z-score standardization."""
    m, s = mean(values), std_dev(values)
    return [(v - m) / s if v is not None and s != 0 else None for v in values]

def fill_missing(values, strategy='mean'):
    """Fill None values with mean, median, or mode."""
    clean = [v for v in values if v is not None]
    if strategy == 'mean':
        fill_val = mean(clean)
    elif strategy == 'median':
        s = sorted(clean)
        n = len(s)
        fill_val = (s[n//2] + s[~(n//2)]) / 2
    elif strategy == 'mode':
        from collections import Counter
        fill_val = Counter(clean).most_common(1)[0][0]
    else:
        fill_val = 0
    return [fill_val if v is None else v for v in values]

def fill_missing(values, strategy='mean'):
    """Fill None values with mean, median, or mode."""
    clean = [v for v in values if v is not None]
    if strategy == 'mean':
        fill_val = mean(clean)
    elif strategy == 'median':
        s = sorted(clean)
        n = len(s)
        fill_val = (s[n//2] + s[~(n//2)]) / 2
    elif strategy == 'mode':
        from collections import Counter
        fill_val = Counter(clean).most_common(1)[0][0]
    else:
        fill_val = 0
    return [fill_val if v is None else v for v in values]

# Test
sample = [10, 20, None, 40, 50, None, 70]
print('Original:', sample)
print('Filled (mean):', fill_missing(sample, 'mean'))
print('Normalized:', [round(v, 3) if v is not None else None for v in normalize_minmax(fill_missing(sample))])
print('Standardized:', [round(v, 3) if v is not None else None for v in standardize_zscore(fill_missing(sample))])

.. Original: [10, 20, None, 40, 50, None, 70]
   Filled (mean): [10, 20, 38.0, 40, 50, 38.0, 70]
   Normalized: [0.0, 0.167, 0.467, 0.5, 0.667, 0.467, 1.0]
   Standardized: [-1.551, -0.997, 0.0, 0.111, 0.665, 0.0, 1.773]

```

```

class DataProcessor:
    """
    Reusable data preprocessing module.
    Supports: loading, summary stats, missing value handling,
              normalization, standardization, encoding, and reporting.
    """

    def __init__(self, data: list[dict]):
        self.data = data
        self.original_size = len(data)
        self.log = []

    # — Utility —————
    def _col(self, col):
        return [row.get(col) for row in self.data]

    def _set_col(self, col, values):
        for row, v in zip(self.data, values):
            row[col] = v

    # — Summary —————
    def columns(self):
        return list(self.data[0].keys()) if self.data else []

    def summary(self):
        print(f'\n{"="*50}')
        print(f' DataProcessor Summary')
        print(f'{"="*50}')
        print(f' Rows: {len(self.data)} (original: {self.original_size})')
        print(f' Columns: {self.columns()}')
        for col in self.columns():
            vals = self._col(col)
            missing = sum(1 for v in vals if v is None)
            numeric = [v for v in vals if isinstance(v, (int, float))]
            if numeric:
                print(f' {col:15s}: numeric | missing={missing} | mean={mean(numeric):.2f} | std={std_dev(numeric):.2f}')
            else:
                uniq = len(set(v for v in vals if v is not None))
                print(f' {col:15s}: string | missing={missing} | unique={uniq}')
        print(f'{"="*50}\n')

    return self

    # — Remove Duplicates —————
    def drop_duplicates(self):
        before = len(self.data)
        seen = set()
        unique = []
        for row in self.data:
            key = tuple(sorted(row.items()))
            if key not in seen:
                seen.add(key)
                unique.append(row)
        self.data = unique
        removed = before - len(self.data)
        self.log.append(f'drop_duplicates: removed {removed} rows')
        return self

    # — Label Encoding —————
    def label_encode(self, col):
        vals = self._col(col)
        unique_vals = sorted(set(v for v in vals if v is not None))
        mapping = {v: i for i, v in enumerate(unique_vals)}
        encoded = [mapping.get(v, -1) for v in vals]
        self._set_col(col, encoded)
        self.log.append(f'label_encode({col}): mapping={mapping}')
        return self

    # — Report —————
    def report(self):
        print('\n=== Processing Log ===')
        for i, entry in enumerate(self.log, 1):
            print(f' {i}. {entry}')

    # — To DataFrame —————
    def to_dataframe(self):
        import pandas as pd
        return pd.DataFrame(self.data)

print('DataProcessor class defined!')
DataProcessor class defined!

```

4. Demo — Use DataProcessor on Sample Dataset

```

import random
random.seed(42)

raw_data = [
    {'CustomerID': i,
     'Age': random.choice([random.randint(18, 70), None, None]),
     'Income': random.choice([random.randint(20000, 120000), None]),
     'Purchases': random.randint(1, 50),
     'Gender': random.choice(['Male', 'Female', None]),
     'Churn': random.choice(['Yes', 'No'])}
    for i in range(1, 31)
]

# Add duplicates
raw_data += raw_data[0:3]

dp = DataProcessor(raw_data)
dp.summary()

```

```

...
=====
DataProcessor Summary
=====
Rows: 33 (original: 33)
Columns: ['CustomerID', 'Age', 'Income', 'Purchases', 'Gender', 'Churn']
CustomerID    : numeric | missing=0 | mean=14.27 | std=9.12
Age           : numeric | missing=17 | mean=49.06 | std=16.05
Income        : numeric | missing=11 | mean=80511.86 | std=28429.32
Purchases     : numeric | missing=0 | mean=24.82 | std=14.29
Gender        : string  | missing=12 | unique=2
Churn         : string  | missing=0 | unique=2
=====

```

```

# Method chaining – all preprocessing steps
dp.handle_missing('Age', 'mean') \
  .handle_missing('Income', 'median') \
  .handle_missing('Gender', 'mode') \
  .drop_duplicates() \
  .normalize('Income') \
  .standardize('Age') \
  .label_encode('Gender') \
  .label_encode('Churn')

dp.summary()
dp.report()

```

```

...
=====
DataProcessor Summary
=====
Rows: 30 (original: 33)
Columns: ['CustomerID', 'Age', 'Income', 'Purchases', 'Gender', 'Churn']
CustomerID    : numeric | missing=0 | mean=15.50 | std=8.66
Age           : numeric | missing=0 | mean=0.00 | std=1.00
Income        : numeric | missing=0 | mean=0.56 | std=0.26
Purchases     : numeric | missing=0 | mean=24.40 | std=14.63
Gender        : numeric | missing=0 | mean=0.23 | std=0.42
Churn         : numeric | missing=0 | mean=0.47 | std=0.50
=====

=== Processing Log ===
1. handle_missing(Age, mean): filled 17 values
2. handle_missing(Income, median): filled 11 values
3. handle_missing(Gender, mode): filled 12 values
4. drop_duplicates: removed 3 rows
5. normalize(Income): min-max applied
6. standardize(Age): z-score applied
7. label_encode(Gender): mapping={'Female': 0, 'Male': 1}
8. label_encode(Churn): mapping={'No': 0, 'Yes': 1}

```

```
# Convert to DataFrame for visualization
df = dp.to_dataframe()
print('Processed DataFrame:')
print(df.head(10).to_string(index=False))
```

Processed DataFrame:

CustomerID	Age	Income	Purchases	Gender	Churn
1	0.893848	0.559373	16	1	1
2	1.607737	0.874314	38	0	1
3	-3.083533	0.203375	33	0	1
4	0.383928	0.559373	15	0	0
5	2.015674	0.994580	45	0	0
6	-1.451787	0.559373	7	1	0
7	-0.017635	0.559373	3	0	0
8	0.281944	0.437582	36	0	0
9	0.485912	0.915137	3	0	1
10	-0.017635	0.000000	7	0	0

5. KPI Visualization

```
import matplotlib.pyplot as plt


fig, axes = plt.subplots(1, 3, figsize=(15, 4))

# Age distribution after standardization
axes[0].hist(df['Age'].dropna(), bins=15, color='#3498db', edgecolor='black')
axes[0].set_title('Age (Z-score Standardized)')
axes[0].set_xlabel('Z-score')

# Income distribution after normalization
axes[1].hist(df['Income'].dropna(), bins=15, color='#2ecc71', edgecolor='black')
axes[1].set_title('Income (Min-Max Normalized)')
axes[1].set_xlabel('[0, 1]')

# Gender encoded counts
gender_counts = df['Gender'].value_counts()
axes[2].bar(gender_counts.index.astype(str), gender_counts.values, color=['#e74c3c', '#9b59b6'], edgecolor='black')
axes[2].set_title('Gender (Label Encoded)')
axes[2].set_xlabel('Encoded Value')

plt.suptitle('Lab 8 – Processed Feature Distributions', fontsize=13, fontweight='bold')
```



6. KPI Summary

```

print('=== Lab 8 KPI Summary ===')
print(f'Original rows:      {dp.original_size}')
print(f'Processed rows:      {len(dp.data)}')
print(f'Methods in class:      {len([m for m in dir(dp) if not m.startswith("_")])}')
print(f'Preprocessing steps:  {len(dp.log)}')
print(f'Missing values left:   {df.isnull().sum().sum()}')
print(f'Method chaining:      Supported (returns self)')
print(f'Reusability:          High – works on any list-of-dicts dataset')

```

```

... === Lab 8 KPI Summary ===
Original rows:      33
Processed rows:      30
Methods in class:    12
Preprocessing steps:  8
Missing values left: 0
Method chaining:     Supported (returns self)
Reusability:         High – works on any list-of-dicts dataset

```

KPI / Metric	Expected Value	Achieved Value
Missing Values Handled	All columns filled	_____
Preprocessing Steps	8 methods called	_____
Method Reusability	Works on any dict dataset	_____
Class Methods Count	> 8 methods	_____
Missing Values Remaining	0	_____

8.8 Discussion

Building a preprocessing module from scratch forces a deep understanding of data transformation logic. The `DataProcessor` class demonstrates key software engineering principles: encapsulation (data hidden inside the class), single responsibility (each method does one thing), and open/closed principle (new methods can be added without modifying existing ones). The method chaining pattern mirrors scikit-learn's Pipeline API, preparing students for professional machine learning workflows.

8.9 Conclusion

In this lab, we developed a fully functional, object-oriented `DataProcessor` module in Python using only core language features. The module handles missing values, normalizes and standardizes numeric features, removes duplicates, and encodes categorical variables — all via a clean method-chaining interface. This hands-on exercise built a strong foundation in Python programming and object-oriented design, essential skills for all subsequent machine learning labs.

Lab 9: Decision Tree Classifier

Topic	Decision Tree Classifier
Real-World Problem	Predict whether a customer will make a purchase based on marketing data
Dataset	Marketing Dataset
Tool(s)	Python / Google Colab
Objectives	Create dataset, train decision tree, evaluate performance, validate accuracy, and visualize the tree structure

9.1 Theory and Concepts

A Decision Tree is a supervised learning algorithm that recursively splits data using feature values. It creates an interpretable tree-like model that is easy to understand and useful for classification problems such as predicting customer purchase decisions.

Key concepts used in this lab:

- Entropy: Measures impurity in a dataset. Lower entropy means the data is more pure.
- Information Gain: Measures how much uncertainty is reduced after splitting on a feature.
- Gini Impurity: An alternative impurity measure used by decision trees.
- Pruning: Reduces unnecessary branches and helps prevent overfitting.
- Max Depth: Controls the maximum number of levels in the tree.

9.2 Dataset Description

The marketing dataset contains customer profile attributes and a target label named Purchase. The goal is to predict whether a customer is likely to make a purchase.

Feature	Type	Description
Age	Numeric	Customer age in years
Income	Numeric	Annual income in thousands
Gender	Categorical	Male / Female
Married	Binary	Yes / No
Has_Children	Binary	Yes / No
Education	Categorical	High School, Bachelor, Master, PhD
Purchase	Target	Final class label: Yes or No

9.3 Lab Objective

The objective of this lab is to create a synthetic marketing dataset, preprocess categorical attributes, train a Decision Tree Classifier, evaluate model accuracy, perform cross-validation, identify important features, visualize the tree, and test the model with a new customer record.

9.4 Implementation Procedure

Follow the steps below in Google Colab or Jupyter Notebook. Paste the related code screenshot and output screenshot in the spaces provided under each heading.

1. Import Required Libraries

Import all required Python libraries for data handling, numerical operations, model training, evaluation metrics, and visualization.

Code Picture Space:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier, export_text, plot_tree
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    classification_report,
    confusion_matrix
)
```

2. Create Sample Marketing Dataset

Create a synthetic marketing dataset with customer attributes such as age, income, gender, marital status, children status, and education level.

Code Picture Space:

```
# Set random seed for reproducible results
np.random.seed(42)

# Number of samples
n = 200

# Create sample marketing dataset
data = {
    'Age': np.random.randint(18, 70, n),
    'Income': np.random.randint(20, 150, n),
    'Gender': np.random.choice(['Male', 'Female'], n),
    'Married': np.random.choice(['Yes', 'No'], n),
    'Has_Children': np.random.choice(['Yes', 'No'], n),
    'Education': np.random.choice(['High School', 'Bachelor', 'Master', 'PhD'], n)
}

# Convert dictionary to DataFrame
df = pd.DataFrame(data)

print("Dataset Shape:", df.shape)
df.head()
```

Output Picture Space:

Dataset Shape: (200, 6)

	Age	Income	Gender	Married	Has_Children	Education
0	56	81	Male	Yes	Yes	Bachelor
1	69	70	Male	Yes	No	Bachelor
2	46	78	Male	Yes	Yes	Master
3	32	137	Female	No	No	PhD
4	60	115	Female	Yes	No	Master

3. Generate Target Variable

Generate the target class Purchase using a logical scoring pattern. Customers meeting the threshold are labeled as Yes; otherwise, they are labeled as No.

Code Picture Space:

Updated Dataset Shape: (200, 7)

	Age	Income	Gender	Married	Has_Children	Education	Purchase
0	56	81	Male	Yes	Yes	Bachelor	Yes
1	69	70	Male	Yes	No	Bachelor	Yes
2	46	78	Male	Yes	Yes	Master	Yes
3	32	137	Female	No	No	PhD	Yes
4	60	115	Female	Yes	No	Master	Yes

Output Picture Space:

Updated Dataset Shape: (200, 7)

	Age	Income	Gender	Married	Has_Children	Education	Purchase
0	56	81	Male	Yes	Yes	Bachelor	Yes
1	69	70	Male	Yes	No	Bachelor	Yes
2	46	78	Male	Yes	Yes	Master	Yes
3	32	137	Female	No	No	PhD	Yes
4	60	115	Female	Yes	No	Master	Yes

4. Check Class Distribution

Check how many records belong to each class of the Purchase target variable before training the model.

Code Picture Space:

```
print("Class Distribution:")
print(df['Purchase'].value_counts())
```

Output Picture Space:

```
... Class Distribution:
Purchase
Yes      112
No        88
Name: count, dtype: int64
```

5. Encode Categorical Variables

Convert categorical variables into numerical dummy variables because machine learning models require numeric input.

Code Picture Space:

```
# Encode categorical variables
categorical_columns = ['Gender', 'Married', 'Has_Children', 'Education']
df_encoded = pd.get_dummies(df, columns=categorical_columns, drop_first=True)

# Separate features and target variable
X = df_encoded.drop('Purchase', axis=1)
y = df_encoded['Purchase']

print("Encoded Dataset Shape:", df_encoded.shape)
df_encoded.head()
```

Output Picture Space:

```
... Encoded Dataset Shape: (200, 9)
```

	Age	Income	Purchase	Gender_Male	Married_Yes	Has_Children_Yes	Education_High School
0	56	81	Yes	True	True	True	False
1	69	70	Yes	True	True	False	False
2	46	78	Yes	True	True	True	False
3	32	137	Yes	False	False	False	False
4	60	115	Yes	False	True	False	False

6. Split Dataset into Training and Testing Sets

Split the dataset into 80% training data and 20% testing data. Stratification helps keep the class distribution balanced in both sets.

Code Picture Space:

```
# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print(f"Train size: {len(X_train)}, Test size: {len(X_test)}")

Train size: 160, Test size: 40
```

7. Train Decision Tree Classifier

Train the Decision Tree model using the entropy criterion, a maximum depth of 4, and a minimum sample split value to control tree complexity.

Code Picture Space:

```
# Create Decision Tree model
dt = DecisionTreeClassifier(
    criterion='entropy',
    max_depth=4,
    min_samples_split=10,
    random_state=42
)

# Train model
dt.fit(X_train, y_train)

print("Decision Tree model trained successfully.")

... Decision Tree model trained successfully.
```

8. Make Predictions

Use the trained model to predict the Purchase class for the test dataset and compare actual versus predicted values.

Code Picture Space:

```
# Predict on test data
y_pred = dt.predict(X_test)

# Show first 10 predictions
prediction_results = pd.DataFrame({
    'Actual': y_test.values,
    'Predicted': y_pred
})

prediction_results.head(10)
```

Output Picture Space:

	Actual	Predicted
0	No	No
1	Yes	Yes
2	No	No
3	No	No
4	Yes	No
5	Yes	Yes
6	Yes	Yes
7	Yes	Yes
8	Yes	Yes
9	Yes	Yes

9. Evaluate Model Performance

Evaluate the model using accuracy, precision, recall, F1-score, classification report, and confusion matrix.

Code Picture Space:

```
print("=== Model Evaluation ===")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f"Precision: {precision_score(y_test, y_pred, pos_label='Yes'):.4f}")
print(f"Recall: {recall_score(y_test, y_pred, pos_label='Yes'):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred, pos_label='Yes'):.4f}")

print("\n=== Classification Report ===")
print(classification_report(y_test, y_pred))

print("\n=== Confusion Matrix ===")
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

Output Picture Space:

```
=== Model Evaluation ===
Accuracy: 0.9250
Precision: 1.0000
Recall: 0.8636
F1-Score: 0.9268

=== Classification Report ===
              precision    recall  f1-score   support

     No         0.86      1.00      0.92         18
     Yes         1.00      0.86      0.93         22

   accuracy         0.93
  macro avg         0.93
weighted avg         0.94

=== Confusion Matrix ===
[[18  0]
 [ 3 19]]
```

10. Cross-Validation

Use 5-fold cross-validation to estimate model performance across multiple train-test splits.

Code Picture Space:

```
# Perform 5-fold cross-validation
cv_scores = cross_val_score(dt, X, y, cv=5)

print("=== 5-Fold Cross Validation ===")
print(f"CV Scores: {cv_scores}")
print(f"Mean CV Accuracy: {cv_scores.mean():.4f} (+/- {cv_scores.std():.4f})")

=== 5-Fold Cross Validation ===
CV Scores: [0.925 0.825 0.8  0.85 0.8 ]
Mean CV Accuracy: 0.8400 (+/- 0.0464)
```

11. Feature Importance

Identify which input features contributed most to the Decision Tree model. This helps explain the model decision process.

Code Picture Space:

```
# Calculate feature importance
importance = pd.Series(
    dt.feature_importances_,
    index=X.columns
).sort_values(ascending=False)

print("=== Feature Importance ===")
print(importance)
```

Output Picture Space:

```
=== Feature Importance ===
Married_Yes      0.289417
Age              0.250369
Income           0.179312
Education_PhD    0.171582
Education_Master 0.109320
Gender_Male      0.000000
Education_High School 0.000000
Has_Children_Yes 0.000000
dtype: float64
```


12. Visualize Feature Importance

Create a bar chart of feature importance values to make the influential variables easier to interpret visually.

Code Picture Space:

```
plt.figure(figsize=(10, 6))
importance.plot(kind='bar')
plt.title('Feature Importance - Decision Tree')
plt.xlabel('Features')
plt.ylabel('Importance Score')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

4

Output Picture Space:**13. Decision Tree Structure in Text Form**

Print the trained decision tree rules in text format. This gives a readable view of split conditions and predicted classes.

Code Picture Space:

```
# Visualize tree as text representation
tree_rules = export_text(
    dt,
    feature_names=list(X.columns),
    max_depth=3
)

print("=== Tree Structure (Text) ===")
print(tree_rules[:1000])
```

Output Picture Space:

```
... === Tree Structure (Text) ===
|--- Income <= 51.50
|   |--- Married_Yes <= 0.50
|   |   |--- class: No
|   |   |--- Married_Yes > 0.50
|   |       |--- Education_PhD <= 0.50
|   |       |   |--- Education_Master <= 0.50
|   |       |   |   |--- class: No
|   |       |   |   |--- Education_Master > 0.50
|   |       |   |       |--- class: Yes
|   |       |   |--- Education_PhD > 0.50
|   |       |       |--- class: Yes
|   |--- Income > 51.50
|   |   |--- Age <= 25.50
|   |   |   |--- Married_Yes <= 0.50
|   |   |   |   |--- class: No
|   |   |   |   |--- Married_Yes > 0.50
|   |   |       |--- Education_Master <= 0.50
|   |   |       |   |--- class: No
|   |   |       |   |--- Education_Master > 0.50
|   |   |       |       |--- class: Yes
|   |   |--- Age > 25.50
|   |       |--- Married_Yes <= 0.50
|   |       |   |--- Education_PhD <= 0.50
|   |       |   |   |--- class: No
|   |       |   |   |--- Education_PhD > 0.50
|   |       |   |       |--- class: Yes
|   |       |--- Married_Yes > 0.50
|   |       |   |--- Age <= 30.50
|   |       |   |   |--- class: Yes
|   |       |   |--- Age > 30.50
|   |       |       |--- class: Yes
```


14. Decision Tree Visualization

Visualize the Decision Tree graphically using `plot_tree()`. The diagram shows split nodes, entropy, samples, class values, and final predictions.

Code Picture Space:

```
plt.figure(figsize=(20, 12))
plot_tree(
    dt,
    feature_names=X.columns,
    class_names=['No', 'Yes'],
    filled=True,
    rounded=True,
    fontsize=10
)
plt.title('Decision Tree for Purchase Prediction', fontsize=16)
plt.tight_layout()
plt.show()
```

Output Picture Space:



15. Test Model with a New Customer

Create a new customer record with the same encoded columns as the training data and predict the purchase decision for that customer.

Code Picture Space:

```
# Create a sample new customer using the same encoded columns
new_customer = pd.DataFrame({
    'Age': [35],
    'Income': [80],
    'Gender_Male': [1],
    'Married_Yes': [1],
    'Has_Children_Yes': [0],
    'Education_Master': [1],
    'Education_PhD': [0]
})

# Ensure columns are in the same order as training data
new_customer = new_customer.reindex(columns=X.columns, fill_value=0)

# Predict purchase decision
new_prediction = dt.predict(new_customer)
print("Predicted Purchase Decision for New Customer:", new_prediction[0])
```

Output Picture Space:

```
Predicted Purchase Decision for New Customer: Yes
```

Conclusion

In this lab, a synthetic marketing dataset was created and used to train a Decision Tree Classifier for purchase prediction. The categorical variables were encoded, the data was split into training and testing sets, and the model was evaluated using standard classification metrics. Cross-validation was also performed to check generalization performance. Feature importance and tree visualization helped explain how the model made decisions. Overall, the Decision Tree model provided an interpretable and practical method for predicting whether a customer is likely to purchase based on marketing data.

Lab 10: Naive Bayes and KNN Classifiers

Topic	Naive Bayes and KNN Classifiers
Real-World Problem	Classify emails as Ham or Spam using email-related features
Dataset	Synthetic Spam Email Dataset
Tool(s)	Python, Google Colab, NumPy, Pandas, Scikit-learn, Matplotlib
Objective	Train, evaluate, compare, and tune Naive Bayes and KNN classification models
Source	Screenshots are taken from the uploaded Colab PDF: lab-10-dm.pdf

10.1 Theory and Concepts

This lab focuses on supervised classification. The objective is to classify an email as Ham or Spam using numerical features extracted from the message content. Two classifiers are used: Gaussian Naive Bayes and K-Nearest Neighbors (KNN).

- **Naive Bayes:** A probabilistic classifier based on Bayes theorem. Gaussian Naive Bayes is suitable when numeric feature values are assumed to follow a normal distribution.
- **K-Nearest Neighbors:** A distance-based classifier that assigns a class by looking at the nearest training samples. Feature scaling is important for KNN.
- **Confusion Matrix:** A table that shows correct and incorrect predictions for each class.
- **ROC-AUC:** A performance measure showing how well a model separates the positive and negative classes.
- **Cross-Validation:** A validation method that evaluates performance across multiple train-test splits.

10.2 Dataset Description

The lab uses a synthetic spam email dataset with 500 records. Ham emails are generated with lower spam-indicator frequencies, while spam emails are generated with higher frequencies. The target variable is `is_spam`, where 0 means Ham and 1 means Spam.

Feature	Type	Description
word_freq_make	Numeric	Frequency of the word make in the email
word_freq_address	Numeric	Frequency of the word address
word_freq_all	Numeric	Frequency of the word all
char_freq_\$	Numeric	Frequency of the \$ character
capital_run_length	Numeric	Capital letter run-length feature
is_spam	Binary Target	0 = Ham, 1 = Spam

10.3 Python Implementation

The following implementation screenshots show the Colab code and corresponding outputs used to build, train, evaluate, and compare the two machine learning models.

1. Import Required Libraries

Required libraries are imported for data handling, visualization, model training, scaling, evaluation metrics, and ROC/confusion-matrix display.

Code Screenshot

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    roc_auc_score,
    confusion_matrix,
    classification_report,
    ConfusionMatrixDisplay,
    RocCurveDisplay
)
```

2. Create Synthetic Spam Dataset

A balanced dataset is created with 250 Ham and 250 Spam emails. Spam emails are generated with higher values for spam-related indicators.

Code Screenshot

```
# Set random seed for reproducible results
np.random.seed(42)

# Number of samples
n = 500
n_ham = n // 2
n_spam = n - n_ham

# Ham emails (label = 0): lower frequencies for spam indicators
ham = pd.DataFrame({
    'word_freq_make': np.random.exponential(0.1, n_ham),
    'word_freq_address': np.random.exponential(0.3, n_ham),
    'word_freq_all': np.random.exponential(0.5, n_ham),
    'char_freq_$': np.random.exponential(0.05, n_ham),
    'capital_run_length': np.random.exponential(5, n_ham),
    'is_spam': [0] * n_ham
})

# Spam emails (label = 1): higher frequencies for spam indicators
spam = pd.DataFrame({
    'word_freq_make': np.random.exponential(1.5, n_spam),
    'word_freq_address': np.random.exponential(1.0, n_spam),
    'word_freq_all': np.random.exponential(2.0, n_spam),
    'char_freq_$': np.random.exponential(2.5, n_spam),
    'capital_run_length': np.random.exponential(25, n_spam),
    'is_spam': [1] * n_spam
})

# Combine ham and spam data
df = pd.concat([ham, spam], ignore_index=True)

print("Dataset Shape:", df.shape)
df.head()
```

Output Screenshot

Dataset Shape: (500, 6)

	word_freq_make	word_freq_address	word_freq_all	char_freq_\$	capital_run_length	is_spam	
0	0.046927	0.104633	0.598932	0.006216	1.023651	0	
1	0.301012	0.145888	0.384039	0.140531	3.903349	0	
2	0.131675	0.571418	0.185190	0.049404	10.315709	0	
3	0.091294	0.114344	0.840454	0.020391	6.588039	0	
4	0.016962	0.055716	0.577165	0.007499	8.213969	0	

3. Check Class Distribution


The target class balance is checked before training. The output confirms that the dataset contains an equal number of Ham and Spam emails.

Code and Output Screenshot

```
print("Class Distribution:")
print(df['is_spam'].value_counts())

df['is_spam'].value_counts().plot(kind='bar')
plt.title("Class Distribution")
plt.xlabel("Class: 0 = Ham, 1 = Spam")
plt.ylabel("Number of Emails")
plt.xticks(rotation=0)
plt.show()
```

```
Class Distribution:
is_spam
0      250
1      250
Name: count, dtype: int64
```



4. Separate Features and Target Variable

The input feature matrix X is separated from the target vector y so that the models can learn from features and predict the class label.

Code Screenshot

```
# Separate input features and target variable
X = df.drop('is_spam', axis=1)
y = df['is_spam']

print("Feature Matrix Shape:", X.shape)
print("Target Vector Shape:", y.shape)

X.head()
```

Output Screenshot

Feature Matrix Shape: (500, 5)
Target Vector Shape: (500,)

	word_freq_make	word_freq_address	word_freq_all	char_freq_\$	capital_run_length
0	0.046927	0.104633	0.598932	0.006216	1.023651
1	0.301012	0.145888	0.384039	0.140531	3.903349
2	0.131675	0.571418	0.185190	0.049404	10.315709
3	0.091294	0.114344	0.840454	0.020391	6.588039
4	0.016962	0.055716	0.577165	0.007499	8.213969

5. Split Dataset into Training and Testing Sets

The dataset is divided into 80% training data and 20% testing data. Stratification keeps the Ham/Spam ratio balanced in both sets.

Code and Output Screenshot

```
# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print(f"Training size: {len(X_train)}")
print(f"Testing size: {len(X_test)}")
```

Training size: 400
Testing size: 100

6. Standardize Features for KNN

KNN uses distance calculations, so the features are standardized using StandardScaler. Naive Bayes is trained on the original numeric values.

Code and Output Screenshot

```
# Standardize features for KNN
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Feature scaling completed successfully.")

Feature scaling completed successfully.
```

7. Train and Evaluate Gaussian Naive Bayes

Gaussian Naive Bayes is trained and evaluated using accuracy, precision, recall, F1-score, ROC-AUC, confusion matrix, and classification report.

Code Screenshot

```
# Create and train Naive Bayes classifier
nb = GaussianNB()
nb.fit(X_train, y_train)

# Make predictions
y_pred_nb = nb.predict(X_test)
y_prob_nb = nb.predict_proba(X_test)[:, 1]

# Evaluate Naive Bayes model
print("=== NAIVE BAYES CLASSIFIER ===")
print(f"Accuracy: {accuracy_score(y_test, y_pred_nb):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_nb):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_nb):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_nb):.4f}")
print(f"ROC-AUC: {roc_auc_score(y_test, y_prob_nb):.4f}")

print("\nConfusion Matrix:")
```

Output Screenshot

```
print(confusion_matrix(y_test, y_pred_nb))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_nb, target_names=['Ham', 'Spam']))

=== NAIVE BAYES CLASSIFIER ===
Accuracy: 0.9900
Precision: 1.0000
Recall: 0.9800
F1-Score: 0.9899
ROC-AUC: 0.9988

Confusion Matrix:
[[50  0]
 [ 1 49]]

Classification Report:
              precision    recall  f1-score   support

   Ham           0.98         1.00         0.99         50
   Spam           1.00         0.98         0.99         50

   accuracy              0.99
  macro avg              0.99
 weighted avg              0.99
```

8. Train and Evaluate KNN Classifier

KNN is trained with $K = 5$ and Euclidean distance. The model predicts the test labels and reports the same evaluation metrics as Naive Bayes.

Code and Output Screenshot

```
# Create and train KNN classifier
knn = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
knn.fit(X_train_scaled, y_train)

# Make predictions
y_pred_knn = knn.predict(X_test_scaled)
y_prob_knn = knn.predict_proba(X_test_scaled)[:, 1]

# Evaluate KNN model
print("=== KNN CLASSIFIER (K=5) ===")
print(f"Accuracy: {accuracy_score(y_test, y_pred_knn):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_knn):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_knn):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_knn):.4f}")
print(f"ROC-AUC: {roc_auc_score(y_test, y_prob_knn):.4f}")

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_knn))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_knn, target_names=['Ham', 'Spam']))
```

```
=== KNN CLASSIFIER (K=5) ===
Accuracy: 0.9500
Precision: 1.0000
Recall: 0.9000
F1-Score: 0.9474
ROC-AUC: 0.9764

Confusion Matrix:
[[50  0]
 [ 5 45]]

Classification Report:

```

	precision	recall	f1-score	support
Ham	0.91	1.00	0.95	50
Spam	1.00	0.90	0.95	50
accuracy			0.95	100
macro avg	0.95	0.95	0.95	100
weighted avg	0.95	0.95	0.95	100

9. Model Comparison

Both models are compared using common metrics. A bar chart makes it easier to compare overall model performance visually.

Comparison Table Screenshot

```
# Compare model performance
results = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC-AUC'],
    'Naive Bayes': [
        accuracy_score(y_test, y_pred_nb),
        precision_score(y_test, y_pred_nb),
        recall_score(y_test, y_pred_nb),
        f1_score(y_test, y_pred_nb),
        roc_auc_score(y_test, y_prob_nb)
    ],
    'KNN (K=5)': [
        accuracy_score(y_test, y_pred_knn),
        precision_score(y_test, y_pred_knn),
        recall_score(y_test, y_pred_knn),
        f1_score(y_test, y_pred_knn),
        roc_auc_score(y_test, y_prob_knn)
    ]
})


print("=== MODEL COMPARISON ===")
print(results.round(4))
```

```
=== MODEL COMPARISON ===
Metric Naive Bayes KNN (K=5)
0 Accuracy 0.9900 0.9500
1 Precision 1.0000 1.0000
2 Recall 0.9800 0.9000
3 F1-Score 0.9899 0.9474
4 ROC-AUC 0.9988 0.9764
```

Comparison Chart Screenshot

```
# Visualize model comparison
results_plot = results.set_index('Metric')

results_plot.plot(kind='bar', figsize=(10, 6))
plt.title("Naive Bayes vs KNN Model Comparison")
plt.xlabel("Evaluation Metric")
plt.ylabel("Score")
plt.ylim(0, 1.05)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



10. KNN Parameter Tuning

Different values of K are tested using cross-validation. This helps identify which neighbor count gives better validation performance.

Tuning Code and Output Screenshot

```
# Tune KNN parameter K using cross-validation
print("=== KNN PARAMETER TUNING ===")

k_values = [3, 5, 7, 9, 11]
cv_results = []

for k in k_values:
    knn_k = KNeighborsClassifier(n_neighbors=k, metric='euclidean')
    scores = cross_val_score(knn_k, X_train_scaled, y_train, cv=5)
    cv_results.append({
        'K': k,
        'Mean CV Accuracy': scores.mean(),
        'Standard Deviation': scores.std()
    })
print(f"K={k}: CV Accuracy = {scores.mean():.4f} (+/- {scores.std():.4f})")

cv_results_df = pd.DataFrame(cv_results)
cv_results_df
```


```
=== KNN PARAMETER TUNING ===
K=3: CV Accuracy = 0.9700 (+/- 0.0061)
K=5: CV Accuracy = 0.9650 (+/- 0.0122)
K=7: CV Accuracy = 0.9475 (+/- 0.0184)
K=9: CV Accuracy = 0.9500 (+/- 0.0112)
K=11: CV Accuracy = 0.9525 (+/- 0.0094)
```

	K	Mean CV Accuracy	Standard Deviation	
0	3	0.9700	0.006124	
1	5	0.9650	0.012247	
2	7	0.9475	0.018371	
3	9	0.9500	0.011180	
4	11	0.9525	0.009354	

Tuning Chart Code Screenshot

```
# Visualize KNN tuning results
plt.figure(figsize=(8, 5))
plt.plot(cv_results_df['K'], cv_results_df['Mean CV Accuracy'], marker='o')
plt.title("KNN Parameter Tuning")
plt.xlabel("Number of Neighbors (K)")
plt.ylabel("Mean Cross-Validation Accuracy")
plt.xticks(k_values)
plt.grid(True)
plt.show()
```

Tuning Chart Output Screenshot



11. Confusion Matrix Visualization


The confusion matrices show how many Ham and Spam emails were classified correctly or incorrectly by each model.

Code Screenshot

```
# Naive Bayes confusion matrix
ConfusionMatrixDisplay.from_predictions(
    y_test,
    y_pred_nb,
    display_labels=['Ham', 'Spam']
)
plt.title("Confusion Matrix - Naive Bayes")
plt.show()

# KNN confusion matrix
ConfusionMatrixDisplay.from_predictions(
    y_test,
    y_pred_knn,
    display_labels=['Ham', 'Spam']
)
plt.title("Confusion Matrix - KNN")
plt.show()
```

Output Screenshot



12. ROC Curve Visualization

ROC curves show class-separation ability. The uploaded PDF contains the ROC code cell, while the chart output area is clipped/blank, so a clean ROC output is recreated from the same code logic.

Code Screenshot from PDF


```
# Plot ROC curves for both models
ax = plt.gca()

RocCurveDisplay.from_predictions(
    y_test,
    y_prob_nb,
    name='Naive Bayes',
    ax=ax
)

RocCurveDisplay.from_predictions(
    y_test,
    y_prob_knn,
    name='KNN (K=5)',
    ax=ax
)

plt.title("ROC Curve Comparison")
plt.show()
```

Recreated Output Chart



10.4 Result Summary

The results show that both models performed well on the synthetic spam classification task. Naive Bayes achieved slightly better performance than KNN in this run.

Metric	Naive Bayes	KNN (K=5)
Accuracy	0.9900	0.9500
Precision	1.0000	1.0000
Recall	0.9800	0.9000
F1-Score	0.9899	0.9474
ROC-AUC	0.9988	0.9764

For KNN parameter tuning, K = 3 produced the highest mean cross-validation accuracy in the displayed results, while larger K values slightly reduced validation accuracy.

Conclusion

In this lab, a synthetic email dataset was created and used to classify messages as Ham or Spam. Gaussian Naive Bayes and KNN classifiers were trained, evaluated, and compared using multiple performance metrics. Naive Bayes performed slightly better on this dataset, while KNN required feature scaling and parameter tuning to achieve good results. The lab also demonstrated practical model evaluation through classification reports, confusion matrices, ROC-AUC, and cross-validation.

Lab 11: Support Vector Machine (SVM) Classifier

Fraud Detection using Linear SVM and RBF SVM

Topic	Support Vector Machine (SVM) Classifier
Real-World Problem	Detect fraudulent transactions from an imbalanced transaction dataset
Dataset	Synthetic Fraud Detection Dataset
Tool(s)	Python, Google Colab, scikit-learn, Matplotlib
Models	Linear SVM, RBF SVM, Tuned RBF SVM
Objective	Train, evaluate, tune, and compare SVM classifiers using classification metrics and visualization

11.1 Theory and Concepts

Support Vector Machine is a supervised machine learning algorithm used for classification. It finds a decision boundary that separates classes with the maximum possible margin. In this lab, SVM is applied to an imbalanced fraud detection problem.

Linear SVM: Creates a straight hyperplane and works well when classes are approximately linearly separable.

RBF Kernel SVM: Uses a non-linear kernel to handle complex class boundaries.

StandardScaler: Scales features to zero mean and unit variance, which is important for distance-based and margin-based models like SVM.

GridSearchCV: Tests multiple hyperparameter combinations and selects the best model based on cross-validation score.

ROC-AUC: Measures how well the model separates legitimate and fraudulent classes across thresholds.

Precision-Recall: Useful for imbalanced datasets because fraud cases are rare and recall is often critical.

11.2 Dataset Description

The dataset is generated synthetically. It contains legitimate transactions and fraudulent transactions. Legitimate records are clustered near the origin, while fraud records are shifted to a different cluster. Eight extra random features are added to make the dataset more realistic.

Item	Value	Description
Total Samples	1050	1000 legitimate transactions and 50 fraud transactions
Target Variable	is_fraud	0 = Legitimate, 1 = Fraud
Features	10	Two main features and eight additional generated features
Class Balance	Imbalanced	Fraud cases represent about 4.76% of the dataset
Train/Test Split	80% / 20%	Stratification keeps the class ratio balanced in both sets

1. Import Required Libraries

Required libraries are imported for data generation, preprocessing, SVM model training, evaluation metrics, and visualization.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    roc_auc_score,
    confusion_matrix,
    classification_report,
    ConfusionMatrixDisplay,
    roc_curve,
    precision_recall_curve
)
```

2. Generate Synthetic Fraud Detection Dataset

A synthetic dataset is created with legitimate and fraudulent transactions. The target classes are intentionally imbalanced to represent a realistic fraud detection scenario.

```
n_legit = 1000
n_fraud = 50 # Imbalanced dataset

# Legitimate transactions clustered around origin
legit = np.random.multivariate_normal(
    mean=[0, 0],
    cov=[[1, 0.3], [0.3, 1]],
    size=n_legit
)
legit_labels = np.zeros(n_legit)

# Fraudulent transactions shifted to another cluster
fraud = np.random.multivariate_normal(
    mean=[1, 1],
    cov=[[1, 0.3], [0.3, 1]],
    size=n_fraud
)
fraud_labels = np.ones(n_fraud)

# Combine both classes
X = np.vstack([legit, fraud])
y = np.hstack([legit_labels, fraud_labels])

# Add 8 more features for realism
X_extra = np.random.randn(len(X), 8) * 0.5
X = np.hstack([X, X_extra])

feature_names = ["feature_1", "feature_2"] + [f"extra_feature_{i}" for i in range(1, 9)]
df = pd.DataFrame(X, columns=feature_names)
df["is_fraud"] = y.astype(int)

print(f"Dataset: {len(X)} samples, {X.shape[1]} features")
print(f"Fraud: {int(y.sum())} ({100*y.mean():.2f}%)")
print(f"Legitimate: {int((1-y).sum())} ({100*(1-y).mean():.2f}%)")

df.head()
```

Dataset: 1050 samples, 10 features
 Fraud: 50 (4.76%)
 Legitimate: 1000 (95.24%)

	feature_1	feature_2	extra_feature_1	extra_feature_2	extra_feature_3	extra_feature_4	extra_feature_5	extra_feature_6
0	-0.318665	-0.482262	0.188205	-0.451026	-0.434832	0.562717	-0.594706	0.821336
1	-1.423220	0.378853	-0.164309	0.301594	-0.272057	-0.081396	0.020460	-0.501094
2	0.327298	0.050263	-0.114300	-0.497175	-1.281167	-0.095514	1.206308	0.392302
3	-1.727223	-0.819182	0.011233	0.273560	-0.590406	0.557161	0.357691	0.359093
4	0.057520	0.699485	0.336431	0.295907	-0.177020	-0.286801	0.050928	0.774510

Next steps: [Generate code with df](#) [New interactive sheet](#)

3. Check Class Distribution

The class distribution is checked before modeling. The output confirms that the dataset has 1000 legitimate records and 50 fraud records.

```
print("Class Distribution:")
print(df["is_fraud"].value_counts().sort_index())

plt.figure(figsize=(6, 4))
df["is_fraud"].value_counts().sort_index().plot(kind="bar")
plt.title("Class Distribution")
plt.xlabel("Class: 0 = Legitimate, 1 = Fraud")
plt.ylabel("Number of Transactions")
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```

4. Split Dataset and Standardize Features

The data is split into training and testing sets using stratification. Features are then standardized because SVM performance depends strongly on feature scale.

```
# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
```

```
print(f"Training size: {len(X_train)}")
print(f"Testing size: {len(X_test)}")
```

```
Training size: 840
Testing size: 210
```

5. Train and Evaluate Linear SVM

A Linear SVM is trained and evaluated using accuracy, precision, recall, F1-score, ROC-AUC, confusion matrix, and classification report.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix, classification_report

# Train the Linear SVM model
svm = LinearSVC()
svm.fit(X_train, y_train)

# Evaluate the model
y_pred = svm.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, svm.decision_function(X_test))
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

6. Train and Evaluate RBF SVM

The RBF SVM uses a non-linear kernel and is evaluated on the same metrics to compare its performance against the Linear SVM.

```
# === LINEAR SVM ===
print("\n=== LINEAR SVM ===")

svm_linear = SVC(kernel="linear", C=1.0, probability=True, random_state=42)
svm_linear.fit(X_train_scaled, y_train)

y_prob_lin = svm_linear.predict_proba(X_test_scaled)[:, 1]

# Evaluation metrics
print(f"Accuracy: {accuracy_score(y_test, y_pred_lin):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_lin):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_lin):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_lin):.4f}")
print(f"ROC-AUC: {roc_auc_score(y_test, y_prob_lin):.4f}")

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_lin))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_lin, target_names=["Legitimate", "Fraud"]))
```

```
=== LINEAR SVM ===
Accuracy: 0.9952
Precision: 0.9091
Recall: 1.0000
F1-Score: 0.9524
ROC-AUC: 0.9990

Confusion Matrix:
[[199  1]
 [ 0 10]]

Classification Report:
              precision    recall  f1-score   support

Legitimate    1.00        0.99        1.00        200
Fraud          0.91        1.00        0.95         10

 accuracy          0.95        1.00        0.97        210
 macro avg         0.95        1.00        0.97        210
 weighted avg      1.00        1.00        1.00        210
```

7. RBF SVM Hyperparameter Tuning

GridSearchCV is used to test different C and gamma values. The best parameters are selected using ROC-AUC scoring.

```
# === RBF SVM ===
print("\n=== RBF SVM ===")

svm_rbf = SVC(kernel="rbf", C=1.0, gamma="scale", probability=True, random_state=42)
svm_rbf.fit(X_train_scaled, y_train)

# Predictions
y_pred_rbf = svm_rbf.predict(X_test_scaled)
y_prob_rbf = svm_rbf.predict_proba(X_test_scaled)[:, 1]

# Evaluation metrics
print(f"Accuracy: {accuracy_score(y_test, y_pred_rbf):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_rbf):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_rbf):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_rbf):.4f}")
print(f"ROC-AUC: {roc_auc_score(y_test, y_prob_rbf):.4f}")

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_rbf))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_rbf, target_names=["Legitimate", "Fraud"]))
```

```
=== RBF SVM ===
Accuracy: 0.9905
Precision: 1.0000
Recall: 0.8000
F1-Score: 0.8889
ROC-AUC: 0.9975

Confusion Matrix:

Classification Report:
              precision    recall  f1-score   support

Legitimate    0.99        1.00        1.00        200
Fraud          1.00        0.80        0.89         10

 accuracy          0.99        0.90        0.94        210
 macro avg         1.00        0.90        0.94        210
 weighted avg      0.99        0.99        0.99        210
```

8. Model Comparison Table

The table compares Linear SVM, RBF SVM, and Tuned RBF SVM using common classification metrics.

```
AUC because fraud detection is an imbalanced classification task.

# === HYPERPARAMETER TUNING FOR RBF SVM ===
print("\n=== RBF Hyperparameter Tuning ===")

param_grid = {
    "C": [0.1, 1, 10],
    "gamma": ["scale", 0.01, 0.1]
}

grid = GridSearchCV(
    SVC(kernel="rbf", probability=True, random_state=42),
    param_grid=param_grid,
    cv=3,
    scoring="roc_auc",
    n_jobs=-1
)

grid.fit(X_train_scaled, y_train)

print(f"Best Parameters: {grid.best_params_}")
print(f"Best CV ROC-AUC: {grid.best_score_:.4f}")

# Best model predictions
best_svm = grid.best_estimator_
y_pred_best = best_svm.predict(X_test_scaled)
y_prob_best = best_svm.predict_proba(X_test_scaled)[:, 1]

print(f"\nBest Model Test ROC-AUC: {roc_auc_score(y_test, y_prob_best):.4f}")

=== RBF Hyperparameter Tuning ===
Best Parameters: {'C': 0.1, 'gamma': 'scale'}
Best CV ROC-AUC: 0.9921
```

9. Visualize Model Comparison

A bar chart is used to compare model performance visually across accuracy, precision, recall, F1-score, and ROC-AUC.

```
# === COMPARISON TABLE ===
print("\n=== MODEL COMPARISON ===")

comparison = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score", "ROC-AUC"],
    "Linear SVM": [
        accuracy_score(y_test, y_pred_lin),
        precision_score(y_test, y_pred_lin),
        recall_score(y_test, y_pred_lin),
        f1_score(y_test, y_pred_lin),
        roc_auc_score(y_test, y_prob_lin)
    ],
    "RBF SVM": [
        accuracy_score(y_test, y_pred_rbf),
        precision_score(y_test, y_pred_rbf),
        recall_score(y_test, y_pred_rbf),
        f1_score(y_test, y_pred_rbf),
        roc_auc_score(y_test, y_prob_rbf)
    ],
    "Tuned RBF SVM": [
        accuracy_score(y_test, y_pred_best),
        precision_score(y_test, y_pred_best),
        recall_score(y_test, y_pred_best),
        f1_score(y_test, y_pred_best),
        roc_auc_score(y_test, y_prob_best)
    ]
})



print(comparison.round(4))
comparison.round(4)
```

```
=== MODEL COMPARISON ===
Metric Linear SVM RBF SVM Tuned RBF SVM
0 Accuracy 0.9952 0.9905 0.9524
1 Precision 0.9091 1.0000 0.0000
2 Recall 1.0000 0.8000 0.0000
3 F1-Score 0.9524 0.8889 0.0000
4 ROC-AUC 0.9990 0.9975 0.9975
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	Metric	Linear SVM	RBF SVM	Tuned RBF SVM
0	Accuracy	0.9952	0.9905	0.9524
1	Precision	0.9091	1.0000	0.0000
2	Recall	1.0000	0.8000	0.0000
3	F1-Score	0.9524	0.8889	0.0000
4	ROC-AUC	0.9990	0.9975	0.9975

10. Confusion Matrix Visualization

Confusion matrices show correct and incorrect predictions for each model. For fraud detection, false negatives are especially important because they represent missed fraud cases.



11. ROC Curve and Precision-Recall Curve

ROC and Precision-Recall curves help evaluate classification performance across thresholds. ROC-AUC is especially useful for separation quality, while Precision-Recall is important for imbalanced data.

```
# ROC Curve
plt.figure(figsize=(10, 5))


fpr_lin, tpr_lin, _ = roc_curve(y_test, y_prob_lin)
fpr_rbf, tpr_rbf, _ = roc_curve(y_test, y_prob_rbf)
fpr_best, tpr_best, _ = roc_curve(y_test, y_prob_best)

plt.plot(fpr_lin, tpr_lin, label=f"Linear SVM (AUC={roc_auc_score(y_test, y_prob_lin):.4f})")
plt.plot(fpr_rbf, tpr_rbf, label=f"RBF SVM (AUC={roc_auc_score(y_test, y_prob_rbf):.4f})")
plt.plot(fpr_best, tpr_best, label=f"Tuned RBF SVM (AUC={roc_auc_score(y_test, y_prob_best):.4f})")
plt.plot([0, 1], [0, 1], "k--", label="Random")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve Comparison")
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

# Precision-Recall Curve
plt.figure(figsize=(10, 5))

precision_lin, recall_lin, _ = precision_recall_curve(y_test, y_prob_lin)
precision_rbf, recall_rbf, _ = precision_recall_curve(y_test, y_prob_rbf)
precision_best, recall_best, _ = precision_recall_curve(y_test, y_prob_best)

plt.plot(recall_lin, precision_lin, label="Linear SVM")
plt.plot(recall_rbf, precision_rbf, label="RBF SVM")
plt.plot(recall_best, precision_best, label="Tuned RBF SVM")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve")
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



11.5 Result Summary

The model results show that Linear SVM achieved very strong performance and captured all fraud cases in the test set. RBF SVM also performed well, but it missed more fraud records. The tuned RBF model produced high ROC-AUC but did not improve fraud recall in this run.

Metric	Linear SVM	RBF SVM	Tuned RBF SVM
Accuracy	0.9952	0.9905	0.9524
Precision	0.9091	1.0000	0.0000
Recall	1.0000	0.8000	0.0000
F1-Score	0.9524	0.8889	0.0000
ROC-AUC	0.9990	0.9975	0.9975

11.6 Conclusion

In this lab, a synthetic imbalanced fraud detection dataset was created and Support Vector Machine classifiers were trained using Linear and RBF kernels. Feature scaling was applied before training because SVM is sensitive to feature scale. The Linear SVM achieved the best recall for fraud detection in this experiment, correctly identifying all fraud cases in the test set. The RBF SVM performed well overall but missed some fraud cases. Hyperparameter tuning with GridSearchCV demonstrated how ROC-AUC can guide model selection, especially for imbalanced classification tasks. Overall, the lab shows how SVM models can be trained, evaluated, tuned, and interpreted using standard machine learning metrics and visualizations.

Lab 12: K-Means and Hierarchical Clustering

Topic	Customer Segmentation using Clustering
Real-World Problem	Group mall customers into useful marketing segments
Dataset	Synthetic Mall Customer Dataset
Tool(s)	Python, Google Colab, scikit-learn, Matplotlib, SciPy
Algorithms	K-Means Clustering and Hierarchical Clustering
Objective	Find customer clusters, compare clustering methods, visualize results, and define marketing strategies

12.1 Theory and Concepts

Clustering is an unsupervised machine learning technique used to discover natural groups in data without using predefined class labels. In this lab, clustering is applied to mall customer data to identify customer segments based on age, annual income, and spending score.

K-Means Clustering: Partitions data into K groups by assigning each record to the nearest cluster centroid.

Hierarchical Clustering: Builds clusters step by step and represents relationships using a dendrogram.

StandardScaler: Standardizes features so that Age, Annual Income, and Spending Score contribute equally.

Elbow Method: Uses inertia values to help select the suitable number of clusters.

Silhouette Score: Measures how well each point fits within its cluster compared with other clusters.

Calinski-Harabasz Index: Measures the ratio of between-cluster separation to within-cluster compactness.

Dendrogram: Visualizes how samples or clusters merge in hierarchical clustering.

12.2 Dataset Description

The dataset is generated synthetically with 200 customer records. It contains five customer groups created from Age, Annual Income, and Spending Score. These groups are then used to test K-Means and Hierarchical Clustering.

Feature	Type	Description
Age	Numeric	Customer age in years
Annual_Income	Numeric	Annual income value used for segmentation
Spending_Score	Numeric	Customer spending behavior score
Number of Records	200	Synthetic records divided into five natural customer groups
Optimal K	5	Selected because the generated data contains five customer segments

1. Import Required Libraries

The required Python libraries are imported for numerical processing, data handling, clustering algorithms, feature scaling, evaluation metrics, hierarchical clustering, and visualization.

In this section, we import all Python libraries required for data handling, preprocessing, clustering, evaluation, and visualization.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import silhouette_score, calinski_harabasz_score
from scipy.cluster.hierarchy import dendrogram, linkage

# Display settings
pd.set_option('display.max_columns', None)
```

2. Generate Synthetic Mall Customer Dataset

A synthetic mall customer dataset is created with five natural customer segments. Each segment has different age, income, and spending behavior patterns.

```
# Generate synthetic mall customer data
np.random.seed(42)
n = 200

# Create 5 customer segments
segments = []

# Segment 1: Young, low income, high spenders
segments.append(np.column_stack([
    np.random.normal(22, 3, 40),
    np.random.normal(25, 5, 40),
    np.random.normal(80, 8, 40)
]))

# Segment 2: Middle-aged, high income, low spenders
segments.append(np.column_stack([
    np.random.normal(45, 5, 40),
    np.random.normal(85, 8, 40),
    np.random.normal(20, 6, 40)
]))

# Segment 3: Young, high income, high spenders
segments.append(np.column_stack([
    np.random.normal(28, 4, 40),
    np.random.normal(90, 7, 40),
    np.random.normal(85, 6, 40)
]))

# Segment 4: Senior, medium income, medium spenders
segments.append(np.column_stack([
    np.random.normal(55, 4, 40),
    np.random.normal(50, 6, 40),
    np.random.normal(50, 7, 40)
]))

# Segment 5: Middle-aged, medium income, medium spenders
segments.append(np.column_stack([
    np.random.normal(35, 4, 40),
    np.random.normal(55, 6, 40),
    np.random.normal(55, 7, 40)
]))

# Combine all segments
X = np.vstack(segments)

# Create DataFrame
df = pd.DataFrame(X, columns=['Age', 'Annual_Income', 'Spending_Score'])

print(f"Dataset: {len(df)} customers")
print(df.describe().round(2))
df.head()
```

```
Dataset: 200 customers
   Age  Annual_Income  Spending_Score
count  200.00         200.00         200.00
mean    36.79          61.12          58.07
std     12.97          25.31          23.93
min     15.03          11.90           7.85
25%     25.95          44.03          43.11
50%     34.66          56.79          57.31
75%     47.43          86.43          80.05
max     63.76         106.76          99.71
```

	Age	Annual_Income	Spending_Score
0	23.490142	28.692333	78.242625
1	21.585207	25.856841	82.856901
2	23.943066	24.421759	91.823152
3	26.569090	23.494482	75.853838
4	21.297540	17.807390	73.532051

Next steps: [Generate code with df](#) [New Interactive sheet](#)

3. Standardize Features

Since clustering algorithms are distance-based, the features are standardized so each variable contributes equally to distance calculations.

```
# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print("Feature scaling completed successfully.")
print("Scaled data shape:", X_scaled.shape)

Feature scaling completed successfully.
Scaled data shape: (200, 3)
```

4. Elbow Method for Optimal K

The Elbow Method tests different values of K and records inertia and silhouette scores. The selected optimal value is K = 5.

```
# Elbow Method for optimal K
print("\n=== Elbow Method ===")
```

```
inertias = []
```

```
silhouettes = []
K_range = range(2, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    labels = kmeans.fit_predict(X_scaled)

    inertias.append(kmeans.inertia_)
    silhouettes.append(silhouette_score(X_scaled, labels))

print(f"K={k}: Inertia={kmeans.inertia_:.1f}, Silhouette={silhouette_score(X_scaled, labels):.3f}")

# Optimal K based on generated segments
optimal_k = 5
print(f"\nOptimal K selected: {optimal_k}")
```

```
=== Elbow Method ===
K=2: Inertia=318.2, Silhouette=0.430
K=3: Inertia=190.4, Silhouette=0.516
K=4: Inertia=101.9, Silhouette=0.592
K=5: Inertia=45.3, Silhouette=0.660
K=6: Inertia=40.9, Silhouette=0.599
K=7: Inertia=37.7, Silhouette=0.527
K=8: Inertia=35.5, Silhouette=0.448
K=9: Inertia=33.1, Silhouette=0.353
K=10: Inertia=31.4, Silhouette=0.351
```

```
Optimal K selected: 5
```

5. K-Means Clustering

K-Means clustering groups customers by assigning each customer to the nearest centroid. The model is evaluated using Silhouette Score and Calinski-Harabasz Index.

5. K-Means Clustering

K-Means groups customers by assigning them to the nearest cluster centroid. In this lab, we use **K = 5** because the dataset was generated with five customer segments.

```
# K-Means Clustering
print(f"\n=== K-MEANS (K={optimal_k}) ===")

kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init=10)
kmeans_labels = kmeans.fit_predict(X_scaled)

# Evaluation metrics
kmeans_silhouette = silhouette_score(X_scaled, kmeans_labels)
kmeans_ch = calinski_harabasz_score(X_scaled, kmeans_labels)

print(f"Silhouette Score: {kmeans_silhouette:.4f}")
print(f"Calinski-Harabasz Index: {kmeans_ch:.2f}")

# Cluster centers converted back to original scale
centers = scaler.inverse_transform(kmeans.cluster_centers_)
centers_df = pd.DataFrame(centers, columns=['Age', 'Annual_Income', 'Spending_Score'])
centers_df['Cluster'] = range(optimal_k)

print("\nCluster Centers:")
print(centers_df.round(2))

# Cluster sizes
print("\nCluster Sizes:")
print(pd.Series(kmeans_labels).value_counts().sort_index())
```

```
=== K-MEANS (K=5) ===
Silhouette Score: 0.6598
Calinski-Harabasz Index: 596.61
```

```
Cluster Centers:
   Age  Annual_Income  Spending_Score  Cluster
0  44.84           85.53           21.14         0
1  27.56           91.54           84.51         1
2  21.34           24.85           80.08         2
3  55.84           50.32           49.44         3
4  34.37           53.37           55.19         4
```

```
Cluster Sizes:
0    40
1    40
2    40
```

```
3    40
4    40
Name: count, dtype: int64
```

6. Hierarchical Clustering

Hierarchical clustering builds clusters step by step using Ward linkage. Its results are compared with K-Means using the same number of clusters.

```
# Hierarchical Clustering
print(f"\n=== HIERARCHICAL CLUSTERING (K={optimal_k}) ===")

hier = AgglomerativeClustering(n_clusters=optimal_k, linkage='ward')
hier_labels = hier.fit_predict(X_scaled)

# Evaluation metrics
hier_silhouette = silhouette_score(X_scaled, hier_labels)
hier_ch = calinski_harabasz_score(X_scaled, hier_labels)

print(f"Silhouette Score: {hier_silhouette:.4f}")
print(f"Calinski-Harabasz Index: {hier_ch:.2f}")
```

```
=== HIERARCHICAL CLUSTERING (K=5) ===
Silhouette Score: 0.6572
Calinski-Harabasz Index: 589.63
```

7. Clustering Model Comparison

This comparison table evaluates K-Means and Hierarchical Clustering using Silhouette Score and Calinski-Harabasz Index.

```
# Model comparison
print("\n=== CLUSTERING COMPARISON ===")
print(f"{'Metric':<25} {'K-Means':<12} {'Hierarchical':<15}")
print("-" * 55)
print(f"{'Silhouette Score':<25} {kmeans_silhouette:<12.4f} {hier_silhouette:<15.4f}")
print(f"{'Calinski-Harabasz':<25} {kmeans_ch:<12.2f} {hier_ch:<15.2f}")

comparison_df = pd.DataFrame({
    'Metric': ['Silhouette Score', 'Calinski-Harabasz'],
    'K-Means': [kmeans_silhouette, kmeans_ch],
    'Hierarchical': [hier_silhouette, hier_ch]
})

comparison_df
```

```
=== CLUSTERING COMPARISON ===
```

Metric	K-Means	Hierarchical
Silhouette Score	0.6598	0.6572
Calinski-Harabasz	596.61	589.63

	Metric	K-Means	Hierarchical
0	Silhouette Score	0.659763	0.657157
1	Calinski-Harabasz	596.606477	589.631936

8. Visualization Code

The visualization code creates a 2x2 figure containing the elbow curve, silhouette score curve, K-Means scatter plot, and dendrogram.

```
# Visualization
fig, axes = plt.subplots(2, 2, figsize=(14, 12))

# Elbow curve
axes[0, 0].plot(K_range, inertias, 'bo-')

axes[0, 0].set_xlabel('Number of Clusters (K)')
axes[0, 0].set_ylabel('Inertia (WCSS)')
axes[0, 0].set_title('Elbow Method for Optimal K')
axes[0, 0].axvline(x=5, color='r', linestyle='--', alpha=0.5, label='Optimal K=5')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)

# Silhouette scores
axes[0, 1].plot(K_range, silhouettes, 'go-')
axes[0, 1].set_xlabel('Number of Clusters (K)')
axes[0, 1].set_ylabel('Silhouette Score')
axes[0, 1].set_title('Silhouette Score vs K')
axes[0, 1].grid(True, alpha=0.3)

# K-Means scatter plot
scatter = axes[1, 0].scatter(
    df['Annual_Income'],
    df['Spending_Score'],
    c=kmeans_labels,
    cmap='viridis',
    alpha=0.7,
    edgecolors='k',
    s=50
)
axes[1, 0].scatter(
    centers[:, 1],
    centers[:, 2],
    c='red',
    marker='x',
    s=200,
    label='Centroids'
)
axes[1, 0].set_xlabel('Annual Income ($K)')
axes[1, 0].set_ylabel('Spending Score')
axes[1, 0].set_title('K-Means Customer Clusters')
axes[1, 0].legend()
plt.colorbar(scatter, ax=axes[1, 0], label='Cluster')


# Dendrogram sample for visualization
sample_idx = np.random.choice(len(X_scaled), 50, replace=False)
Z = linkage(X_scaled[sample_idx], method='ward')

dendrogram(Z, ax=axes[1, 1], truncate_mode='lastp', p=5, leaf_rotation=90)
axes[1, 1].set_title('Hierarchical Clustering Dendrogram (Sample)')
axes[1, 1].set_xlabel('Cluster Size')
axes[1, 1].set_ylabel('Distance')

plt.tight_layout()
plt.savefig('clustering_results.png', dpi=150, bbox_inches='tight')
plt.show()
```


9. Visualization Output

The final visualization summarizes the clustering analysis. It shows the best K, cluster separation quality, customer groups, centroids, and hierarchical dendrogram.



10. Customer Segment Interpretation

The cluster centers are interpreted into meaningful customer segment names along with suggested marketing strategies for each group.

```
# Customer segment interpretation table
segment_table = pd.DataFrame({
    'Cluster': [0, 1, 2, 3, 4],
    'Segment Name': [
        'Carefree Spenders',
        'Conservative Earners',
        'Premium Customers',
        'Steady Seniors',
        'Average Families'
    ],
    'Age': ['Young', 'Middle', 'Young', 'Senior', 'Middle'],
    'Income': ['Low', 'High', 'High', 'Medium', 'Medium'],
    'Spending': ['High', 'Low', 'High', 'Medium', 'Medium'],
    'Marketing Strategy': [
        'Youth discounts and social media ads',
        'Premium products and trust-building',
        'VIP programs and exclusive offers',
        'Reliability and value propositions',
        'Family bundles and seasonal sales'
    ]
})
```

```
    'Premium Customers',
    'Steady Seniors',
    'Average Families'
],
'Age': ['Young', 'Middle', 'Young', 'Senior', 'Middle'],
'Income': ['Low', 'High', 'High', 'Medium', 'Medium'],
'Spending': ['High', 'Low', 'High', 'Medium', 'Medium'],
'Marketing Strategy': [
    'Youth discounts and social media ads',
    'Premium products and trust-building',
    'VIP programs and exclusive offers',
    'Reliability and value propositions',
    'Family bundles and seasonal sales'
]
})
```

Cluster	Segment Name	Age	Income	Spending	Marketing Strategy
0	Carefree Spenders	Young	Low	High	Youth discounts and social media ads
1	Conservative Earners	Middle	High	Low	Premium products and trust-building
2	Premium Customers	Young	High	High	VIP programs and exclusive offers
3	Steady Seniors	Senior	Medium	Medium	Reliability and value propositions
4	Average Families	Middle	Medium	Medium	Family bundles and seasonal sales

11. KPIs and Evaluation Metrics

The KPI table summarizes target values for clustering quality, optimal K, revenue contribution, and cluster separation.

```
# KPI / Evaluation Metrics
kpi_table = pd.DataFrame({
    'KPI / Metric': [
        'Silhouette Score',
        'Cluster Revenue Contribution',
        'Optimal K',
        'Calinski-Harabasz'
    ],
    'Description': [
        'Cluster separation quality',
        'Revenue share by segment',
        'Number of clusters from elbow',
        'Cluster density vs separation'
    ],
    'Target / Value': [
        '>= 0.35',
        'All segments > 15%',
        'K = 5',
        '>= 120'
    ]
})

kpi_table
```

	KPI / Metric	Description	Target / Value	
0	Silhouette Score	Cluster separation quality	>= 0.35	
1	Cluster Revenue Contribution	Revenue share by segment	All segments > 15%	
2	Optimal K	Number of clusters from elbow	K = 5	
3	Calinski-Harabasz	Cluster density vs separation	>= 120	

12.5

Result Summary

The experiment produced five meaningful customer clusters. K-Means achieved slightly stronger performance than Hierarchical Clustering based on both Silhouette Score and Calinski-Harabasz Index, while Hierarchical Clustering provided a useful dendrogram view of the cluster structure.

Metric	K-Means	Hierarchical
Optimal K	5	5
Silhouette Score	0.6598	0.6572
Calinski-Harabasz Index	596.61	589.63
Cluster Sizes	40 customers in each cluster	Comparable five-cluster grouping

12.6 Customer Segment Summary

Based on the cluster centers, the customers can be interpreted as practical marketing groups. This helps convert clustering output into business decisions.

Cluster	Segment Name	Age	Income	Spending	Marketing Strategy
0	Carefree Spenders	Young	Low	High	Youth discounts and social media ads
1	Conservative Earners	Middle	High	Low	Premium products and trust-building
2	Premium Customers	Young	High	High	VIP programs and exclusive offers
3	Steady Seniors	Senior	Medium	Medium	Reliability and value propositions
4	Average Families	Middle	Medium	Medium	Family bundles and seasonal sales

12.7 Conclusion

In this lab, a synthetic mall customer dataset was created and customer segmentation was performed using K-Means and Hierarchical Clustering. The data was standardized before clustering to ensure that all features contributed equally. The Elbow Method and Silhouette Score supported the selection of $K = 5$. K-Means produced a Silhouette Score of 0.6598 and a Calinski-Harabasz Index of 596.61, slightly outperforming Hierarchical Clustering. The visualization clearly showed five customer groups, and the segment interpretation table converted the cluster results into practical marketing strategies. Overall, this lab demonstrates how unsupervised learning can be used to identify customer groups and support data-driven marketing decisions.

Lab 13: Outlier Detection and Anomaly Detection

Topic	Outlier Detection and Anomaly Detection
Real-World Problem	Detect fraudulent banking transactions from unusual transaction behavior
Dataset	Synthetic Banking Transactions Dataset
Tool(s)	Python, NumPy, Pandas, Scikit-learn, Matplotlib
Methods	Z-Score, IQR Method, Isolation Forest
Objective	Compare anomaly detection methods using detection rate, false alarm rate, and confusion matrix results

13.1 Theory and Concepts

Outlier detection identifies records that behave very differently from the majority of the data. In banking fraud detection, these anomalous records can indicate suspicious transactions such as unusually high amounts, high transaction frequency, risky merchants, or very short time gaps between transactions.

Concept	Description
Z-Score Method	Flags values that are far from the mean using standard deviation. In this lab, a transaction is marked as an outlier if any feature has an absolute Z-score greater than 2.5.
IQR Method	Uses the interquartile range to detect extreme values below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$.
Isolation Forest	A machine learning anomaly detection model that isolates abnormal observations using random decision trees.
Detection Rate	Percentage of actual fraudulent transactions correctly detected.
False Alarm Rate	Percentage of normal transactions incorrectly flagged as fraud.

13.2 Dataset Description

The dataset contains 1000 synthetic banking transactions. It includes 950 normal transactions and 50 fraudulent transactions, giving an imbalanced fraud detection scenario similar to real banking data.

Feature	Type	Description
Transaction_Amount	Numeric	Amount involved in the transaction
Account_Balance	Numeric	Account balance associated with the transaction
Transaction_Frequency	Numeric	Number or frequency of transactions
Merchant_Risk_Score	Numeric	Risk score assigned to the merchant
Time_Since_Last	Numeric	Time gap since the previous transaction
Label	Target	0 = normal transaction, 1 = fraud/anomaly

1. Import Required Libraries

This step imports libraries for numerical processing, data handling, anomaly detection, feature scaling, evaluation, and visualization.

```
import numpy as np
import pandas as pd
from sklearn.ensemble import IsolationForest
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt

pd.set_option('display.max_columns', None)
```

2. Generate Synthetic Banking Transactions

A synthetic banking dataset is created with normal and fraudulent transactions. Fraud transactions have unusually high amounts, lower balances, high frequency, risky merchants, and short time since the last transaction.

```
# Generate synthetic banking transactions
np.random.seed(42)
n_normal = 950
n_fraud = 50

# Normal transactions
normal = pd.DataFrame({
    'Transaction_Amount': np.random.lognormal(5, 0.8, n_normal),
    'Account_Balance': np.random.lognormal(8, 0.5, n_normal),
    'Transaction_Frequency': np.random.poisson(10, n_normal),
    'Merchant_Risk_Score': np.random.beta(2, 5, n_normal) * 100,
    'Time_Since_Last': np.random.exponential(48, n_normal)
})
normal['Label'] = 0 # Normal

# Fraudulent transactions (anomalous patterns)
fraud = pd.DataFrame({
    'Transaction_Amount': np.random.lognormal(8, 0.3, n_fraud), # Very high amounts
    'Account_Balance': np.random.lognormal(6, 0.5, n_fraud), # Low balances
    'Transaction_Frequency': np.random.poisson(50, n_fraud), # Unusually high frequency
    'Merchant_Risk_Score': np.random.beta(7, 2, n_fraud) * 100, # High risk merchants
    'Time_Since_Last': np.random.exponential(2, n_fraud) # Very frequent transactions
})
fraud['Label'] = 1 # Fraud / Anomaly

# Combine normal and fraudulent transactions
df = pd.concat([normal, fraud], ignore_index=True)
X = df.drop('Label', axis=1)
```

2. Generate Synthetic Banking Transactions (continued)

This is the continued code/output screenshot for the same section.

```
y_true = df['Label']

print(f"Dataset: {len(df)} transactions")
print(f"Normal: {(y_true == 0).sum()} ({100 * (y_true == 0).mean():.1f}%)")
print(f"Fraud: {(y_true == 1).sum()} ({100 * (y_true == 1).mean():.1f}%)")
print("\nFeature Statistics:")
print(df.describe().round(2))

df.head()
```

```
Dataset: 1000 transactions
Normal: 950 (95.0%)
Fraud: 50 (5.0%)
```

Feature Statistics:

	Transaction_Amount	Account_Balance	Transaction_Frequency
count	1000.00	1000.00	1000.00
mean	348.88	3325.02	11.99
std	685.54	1872.93	9.51
min	11.10	167.60	2.00
25%	91.81	2097.76	8.00
50%	158.17	2955.19	10.00
75%	274.41	4183.67	12.00
max	6574.81	14713.99	73.00

	Merchant_Risk_Score	Time_Since_Last	Label
count	1000.00	1000.00	1000.00
mean	31.70	43.72	0.05
std	19.02	45.71	0.22
min	0.54	0.01	0.00
25%	17.25	11.27	0.00
50%	28.68	30.15	0.00
75%	42.74	61.63	0.00
max	99.74	365.34	1.00

	Transaction_Amount	Account_Balance	Transaction_Frequency	Merchant_Risk_Score	Time_Since_Last	Label
0	220.825174	2642.239075	8	10.232871	16.876951	0
1	132.872306	3554.706213	11	38.179449	15.777137	0
2	249.173846	1594.364157	18	7.953991	7.576506	0
3	501.911538	6135.717499	9	36.081494	16.384012	0
4	123.060647	2860.994043	10	18.806663	17.794734	0

3. Standardize Features

StandardScaler converts features to a comparable scale. This step is important because anomaly detection methods can be affected by features with large numeric ranges.

```
# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print("Feature scaling completed successfully.")
print("Scaled feature matrix shape:", X_scaled.shape)
```

```
Feature scaling completed successfully.
Scaled feature matrix shape: (1000, 5)
```

4. Method 1: Z-Score Outlier Detection

The Z-Score method identifies transactions that are more than 2.5 standard deviations away from the mean in any feature.

```
# === METHOD 1: Z-SCORE METHOD ===
print("\n=== Z-SCORE METHOD ===")

z_scores = np.abs((X - X.mean()) / X.std())
z_outliers = (z_scores > 2.5).any(axis=1).astype(int)
```

```
print(f"Detected Outliers: {z_outliers.sum()} ({100 * z_outliers.mean():.1f}%")
print(f"True Positives: {((z_outliers == 1) & (y_true == 1)).sum()}")
print(f"False Positives: {((z_outliers == 1) & (y_true == 0)).sum()}")
print(f"Detection Rate: {100 * ((z_outliers == 1) & (y_true == 1)).sum() / (y_true == 1).sum():.1f}%")
print(f"False Alarm Rate: {100 * ((z_outliers == 1) & (y_true == 0)).sum() / (y_true == 0).sum():.1f}%")
```

```
=== Z-SCORE METHOD ===
Detected Outliers: 104 (10.4%)
True Positives: 50
False Positives: 54
Detection Rate: 100.0%
False Alarm Rate: 5.7%
```

5. Method 2: IQR Outlier Detection

The IQR method detects outliers using lower and upper bounds based on Q1, Q3, and the interquartile range.

```
# === METHOD 2: IQR METHOD ===
print("\n=== IQR METHOD ===")

iqr_outliers = np.zeros(len(df), dtype=int)

for col in X.columns:
    Q1 = X[col].quantile(0.25)
    Q3 = X[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    iqr_outliers = np.logical_or(iqr_outliers, ((X[col] < lower) | (X[col] > upper))).astype(int)

print(f"Detected Outliers: {iqr_outliers.sum()} ({100 * iqr_outliers.mean():.1f}%")
print(f"True Positives: {((iqr_outliers == 1) & (y_true == 1)).sum()}")
print(f"False Positives: {((iqr_outliers == 1) & (y_true == 0)).sum()}")
print(f"Detection Rate: {100 * ((iqr_outliers == 1) & (y_true == 1)).sum() / (y_true == 1).sum():.1f}%")
print(f"False Alarm Rate: {100 * ((iqr_outliers == 1) & (y_true == 0)).sum() / (y_true == 0).sum():.1f}%")
```

```
=== IQR METHOD ===
Detected Outliers: 188 (18.8%)
True Positives: 50
False Positives: 138
Detection Rate: 100.0%
False Alarm Rate: 14.5%
```

6. Method 3: Isolation Forest

Isolation Forest is trained with 5% contamination because the dataset contains approximately 5% fraudulent transactions. It provides anomaly scores and classifies unusual observations as fraud.

lab, contamination is set to **0.05**, meaning the model expects around 5% anomalous transactions.

```
# === METHOD 3: ISOLATION FOREST ===
print("\n=== ISOLATION FOREST ===")

iso_forest = IsolationForest(
    n_estimators=100,
    contamination=0.05,
    random_state=42,
    n_jobs=-1
)

iso_labels = iso_forest.fit_predict(X_scaled)
iso_outliers = (iso_labels == -1).astype(int)  # -1 means anomaly

print(f"Detected Outliers: {iso_outliers.sum()} ({100 * iso_outliers.mean():.1f}%)")
print(f"True Positives: {(iso_outliers == 1) & (y_true == 1).sum()}")
print(f"False Positives: {(iso_outliers == 1) & (y_true == 0).sum()}")
print(f"Detection Rate: {100 * ((iso_outliers == 1) & (y_true == 1)).sum() / (y_true == 1).sum():.1f}%")
print(f"False Alarm Rate: {100 * ((iso_outliers == 1) & (y_true == 0)).sum() / (y_true == 0).sum():.1f}%")

# Isolation Forest anomaly scores
scores = iso_forest.decision_function(X_scaled)
print(f"\nAnomaly Score Range: [{scores.min():.3f}, {scores.max():.3f}]")
```

6. Method 3: Isolation Forest (continued)

This is the continued code/output screenshot for the same section.

```
print(f"Mean Score (Normal): {scores[y_true == 0].mean():.3f}")
print(f"Mean Score (Fraud): {scores[y_true == 1].mean():.3f}")
```

```
=== ISOLATION FOREST ===
Detected Outliers: 50 (5.0%)
True Positives: 46
False Positives: 4
Detection Rate: 92.0%
False Alarm Rate: 0.4%

Anomaly Score Range: [-0.086, 0.244]
Mean Score (Normal): 0.176
Mean Score (Fraud): -0.029
```


7. Method Comparison Table

The three methods are compared using detected outliers, true positives, false positives, detection rate, and false alarm rate.

```
# === COMPARISON TABLE ===
print("\n=== METHOD COMPARISON ===")

comparison = pd.DataFrame({
    'Method': ['Z-Score', 'IQR', 'Isolation Forest'],
    'Detected': [z_outliers.sum(), iqr_outliers.sum(), iso_outliers.sum()],
    'True_Pos': [
        ((z_outliers == 1) & (y_true == 1)).sum(),
        ((iqr_outliers == 1) & (y_true == 1)).sum(),
        ((iso_outliers == 1) & (y_true == 1)).sum()
    ],
    'False_Pos': [
        ((z_outliers == 1) & (y_true == 0)).sum(),
        ((iqr_outliers == 1) & (y_true == 0)).sum(),
        ((iso_outliers == 1) & (y_true == 0)).sum()
    ],
    'Detection_Rate': [
        100 * ((z_outliers == 1) & (y_true == 1)).sum() / (y_true == 1).sum(),
        100 * ((iqr_outliers == 1) & (y_true == 1)).sum() / (y_true == 1).sum(),
        100 * ((iso_outliers == 1) & (y_true == 1)).sum() / (y_true == 1).sum()
    ],
    'False_Alarm_Rate': [
        100 * ((z_outliers == 1) & (y_true == 0)).sum() / (y_true == 0).sum(),
        100 * ((iqr_outliers == 1) & (y_true == 0)).sum() / (y_true == 0).sum(),
        100 * ((iso_outliers == 1) & (y_true == 0)).sum() / (y_true == 0).sum()
    ]
})

print(comparison.round(2))
comparison.round(2)
```

```
=== METHOD COMPARISON ===
      Method  Detected  True_Pos  False_Pos  Detection_Rate \
0      Z-Score      104       50        54         100.0
1         IQR      188       50       138         100.0
2 Isolation Forest    50       46         4          92.0

      False_Alarm_Rate
0             5.68
1            14.53
2             0.42
```

	Method	Detected	True_Pos	False_Pos	Detection_Rate	False_Alarm_Rate
0	Z-Score	104	50	54	100.0	5.68
1	IQR	188	50	138	100.0	14.53
2	Isolation Forest	50	46	4	92.0	0.42

8. Visualization of Outlier Detection Results

The visualization code creates four plots: transaction amount versus balance, anomaly score distribution, method comparison, and feature-level outlier counts.

```
# Visualization Code
fig, axes = plt.subplots(2, 2, figsize=(14, 12))

# Transaction Amount vs Account Balance
axes[0, 0].scatter(
    df[df['Label'] == 0]['Transaction_Amount'],
    df[df['Label'] == 0]['Account_Balance'],
    c='blue', alpha=0.5, label='Normal', s=30
)
axes[0, 0].scatter(
    df[df['Label'] == 1]['Transaction_Amount'],
    df[df['Label'] == 1]['Account_Balance'],
    c='red', alpha=0.8, label='Fraud', s=50, marker='x'
)
axes[0, 0].set_xlabel('Transaction Amount')
axes[0, 0].set_ylabel('Account Balance')
axes[0, 0].set_title('Transactions: Amount vs Balance')
axes[0, 0].legend()
axes[0, 0].set_xscale('log')
axes[0, 0].set_yscale('log')

# Isolation Forest scores distribution
axes[0, 1].hist(scores[y_true == 0], bins=50, alpha=0.7, label='Normal', color='blue')
axes[0, 1].hist(scores[y_true == 1], bins=20, alpha=0.7, label='Fraud', color='red')
axes[0, 1].axvline(x=0, color='black', linestyle='--', label='Threshold')
axes[0, 1].set_xlabel('Anomaly Score')
axes[0, 1].set_ylabel('Frequency')
axes[0, 1].set_title('Isolation Forest Score Distribution')
axes[0, 1].legend()

# Detection Rate and False Alarm Rate Comparison
det_rates = comparison['Detection_Rate'].values
far_rates = comparison['False_Alarm_Rate'].values
methods = comparison['Method'].values
x = np.arange(len(methods))
width = 0.35
```

8. Visualization of Outlier Detection Results (continued)

This is the continued code/output screenshot for the same section.

```
axes[1, 0].bar(x - width/2, det_rates, width, label='Detection Rate (%)', color='green', alpha=0.7)
axes[1, 0].bar(x + width/2, far_rates, width, label='False Alarm Rate (%)', color='orange', alpha=0.7)
axes[1, 0].set_ylabel('Percentage (%)')
axes[1, 0].set_title('Method Comparison')
axes[1, 0].set_xticks(x)
axes[1, 0].set_xticklabels(methods)
axes[1, 0].legend()
axes[1, 0].set_ylim(0, 100)

# Feature outlier counts using Z-score
feature_outliers = {}
for col in X.columns:
    z = np.abs((X[col] - X[col].mean()) / X[col].std())
    feature_outliers[col] = (z > 2.5).sum()

axes[1, 1].barh(list(feature_outliers.keys()), list(feature_outliers.values()), color='teal', alpha=0.7)
axes[1, 1].set_xlabel('Outlier Count')
axes[1, 1].set_title('Outliers by Feature (Z-Score)')

plt.tight_layout()
plt.savefig('outlier_detection.png', dpi=150, bbox_inches='tight')
plt.show()
```

9. Confusion Matrix for Isolation Forest





The confusion matrix shows correct and incorrect classifications for the best-performing anomaly detection method.

```
# Confusion Matrix for Isolation Forest
cm = confusion_matrix(y_true, iso_outliers)
cm_df = pd.DataFrame(
    cm,
    index=['Actual Normal', 'Actual Fraud'],
    columns=['Predicted Normal', 'Predicted Fraud']
)
```

```
print("Confusion Matrix - Isolation Forest")
print(cm_df)
```

```
Confusion Matrix - Isolation Forest
              Predicted Normal  Predicted Fraud
Actual Normal                946                4
Actual Fraud                   4                 46
```

Next steps: [Generate code with cm_df](#) [New interactive sheet](#)



10. KPIs / Evaluation Metrics

The KPI table summarizes target values for anomaly detection performance, including detection rate, false alarm rate, best method, and processing speed.

```
# KPIs / Evaluation Metrics
kpi_table = pd.DataFrame({
    'KPI / Metric': [
        'Anomaly Detection Rate',
        'False Alarm Rate',
        'Best Method',
        'Processing Speed'
    ],
    'Description': [
        'Fraudulent transactions caught',
        'Normal transactions incorrectly flagged',
        'Highest detection with lowest false positive rate',
        'Time to analyze 1000 records'
    ],
    'Target / Value': [
        '>= 85%',
        '<= 5%',
        'Isolation Forest',
        '< 2 seconds'
    ]
})

kpi_table
```

	KPI / Metric	Description	Target / Value	
0	Anomaly Detection Rate	Fraudulent transactions caught	>= 85%	
1	False Alarm Rate	Normal transactions incorrectly flagged	<= 5%	
2	Best Method	Highest detection with lowest false positive rate	Isolation Forest	
3	Processing Speed	Time to analyze 1000 records	< 2 seconds	

13.5 Result Summary

The statistical methods detected many anomalies, but they also produced a higher number of false positives. Isolation Forest detected 50 outliers, caught 46 fraudulent transactions, and produced only 4 false positives. This makes it the most balanced method in this lab.

Method	Detected	True Positives	False Positives	Main Observation
Z-Score	104	50	54	High detection but more false alarms
IQR	188	50	138	Detects all fraud but flags many normal records
Isolation Forest	50	46	4	Best balance between detection and false alarms

13.6 Conclusion

In this lab, a synthetic banking transaction dataset was created and three anomaly detection methods were applied: Z-Score, IQR, and Isolation Forest. The Z-Score and IQR methods were simple and easy to interpret, but they produced more false alarms. Isolation Forest performed better overall because it detected most fraudulent transactions while keeping the false alarm rate low. Therefore, Isolation Forest is more suitable for banking fraud anomaly detection when fraudulent patterns are rare and different from normal behavior.

Lab 14: Data Scraping & Sentiment Analysis

Product Review Sentiment Classification

Topic	Web Scraping and Sentiment Analysis
Real-World Problem	Classify product review feedback into Positive, Neutral, or Negative sentiment.
Dataset	Sample product review dataset containing 10 product reviews with ratings and text.
Tool(s)	Python, Google Colab, pandas, TextBlob, VADER, BeautifulSoup, requests, matplotlib
Main Methods	Text preprocessing, TextBlob polarity scoring, VADER compound sentiment scoring, comparison and visualization
Objective	Clean review text, apply two sentiment-analysis techniques, compare outputs, and visualize sentiment distribution.

14.1 Theory and Concepts

Sentiment analysis is a natural language processing task used to identify the emotional tone of text. In this lab, product reviews are cleaned and classified using two lexicon-based approaches: TextBlob and VADER.

Web Scraping: The process of collecting content from web pages. This lab uses a controlled sample review structure instead of a live website.

Text Preprocessing: Review text is normalized by lowercasing and removing unnecessary symbols before analysis.

TextBlob Polarity: TextBlob returns a polarity score where positive values suggest positive sentiment and negative values suggest negative sentiment.

VADER Compound Score: VADER is effective for short review-style text and uses a compound score to label positive, neutral, or negative sentiment.

14.2 Dataset Description

The sample dataset contains product names, star ratings, and customer review text. It includes both positive and negative opinions so that sentiment methods can be compared.

Feature	Type	Description
product	Categorical	Name of the reviewed product.
rating	Numeric	User rating value from 1 to 5.
text	Text	Customer review statement used for sentiment analysis.
clean_text	Text	Normalized version of review text after preprocessing.
tb_label / vader_label	Categorical	Predicted sentiment labels from TextBlob and VADER.

14.3 Python Implementation

The following implementation screenshots are taken from the uploaded Google Colab PDF. Each part includes a short explanation followed by the relevant code/output image.

14.3.1 Install and Import Required Libraries

This section installs required packages and imports libraries for data handling, web scraping support, sentiment analysis, visualization, and text cleaning.

```
# Install required packages in Google Colab
!pip install textblob vaderSentiment beautifulsoup4 requests -q

import requests
from bs4 import BeautifulSoup
import pandas as pd
import numpy as np
from textblob import TextBlob
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
import matplotlib.pyplot as plt
from collections import Counter
import re
```

126.0/126.0 kB 3.7 MB/s eta 0:00:00

14.3.2 Web Scraping Demo Using Sample Review Data

A sample review dataset is created to simulate web-scraped product reviews. The dataset is converted into a pandas DataFrame and initial rows are displayed.

```
# =====
# PART 1: Web Scraping Demo using sample review structure
# =====

sample_reviews = [
    {
        "product": "Wireless Headphones",
        "rating": 5,
        "text": "Absolutely love these headphones! Great sound quality and battery life."
    },
    {
        "product": "Wireless Headphones",
        "rating": 4,
        "text": "Good value for money. Comfortable to wear for long periods."
    },
    {
        "product": "Smart Watch",
        "rating": 2,
        "text": "Disappointed with the battery. Does not last as advertised."
    },
    {
        "product": "Smart Watch",
        "rating": 3,
```

```

        "text": "Average product. Some features are useful, others feel gimmicky."
    },
    {
        "product": "Bluetooth Speaker",
        "rating": 5,
        "text": "Amazing sound for such a compact device! Highly recommend."
    },
    {
        "product": "Bluetooth Speaker",
        "rating": 1,
        "text": "Stopped working after two weeks. Terrible quality control."
    },
    {
        "product": "Laptop Stand",
        "rating": 4,
        "text": "Sturdy and well-built. Helps with posture during long work sessions."
    },
    {
        "product": "USB-C Hub",
        "rating": 2,
        "text": "Gets uncomfortably hot during use. Connectivity is unreliable."
    },
    {
        "product": "Mechanical Keyboard",
        "rating": 5,
        "text": "Best keyboard I have ever owned. The tactile feedback is perfect."
    },
    {
        "product": "Webcam",
        "rating": 3,
        "text": "Decent video quality but the microphone picks up too much background noise."
    }
]

df = pd.DataFrame(sample_reviews)

print(f"Dataset: {len(df)} reviews")
print(df[["product", "rating"]].head())

```

```

Dataset: 10 reviews
   product  rating
0  Wireless Headphones    5
1  Wireless Headphones    4
2    Smart Watch         2
3    Smart Watch         3
4  Bluetooth Speaker     5

```

14.3.3 Text Preprocessing

Before scoring sentiment, review text is normalized by converting it to lowercase and removing unnecessary characters. This prepares cleaner text for analysis.

```

# =====
# PART 2: Text Preprocessing
# =====

def preprocess_text(text):
    """Clean and normalize text for analysis."""
    text = text.lower()
    text = re.sub(r"^[a-zA-Z\s]", "", text)
    text = re.sub(r"\s+", " ", text).strip()
    return text

df["clean_text"] = df["text"].apply(preprocess_text)

print("\n=== Cleaned Text Samples ===")
print(df[["text", "clean_text"]].head())

=== Cleaned Text Samples ===
      text \
0  Absolutely love these headphones! Great sound ...
1  Good value for money. Comfortable to wear for ...
2  Disappointed with the battery. Does not last a...
3  Average product. Some features are useful, oth...

```


14.3.4 Sentiment Analysis with TextBlob

TextBlob calculates a polarity score for each review and assigns Positive, Neutral, or Negative labels based on threshold values.

```
# =====
# PART 3: Sentiment Analysis with TextBlob
# =====

def textblob_sentiment(text):
    blob = TextBlob(text)
    polarity = blob.sentiment.polarity

    if polarity > 0.1:
        return "Positive", polarity
    elif polarity < -0.1:
        return "Negative", polarity
    else:
        return "Neutral", polarity

tb_results = df["text"].apply(textblob_sentiment)

df["tb_label"] = [r[0] for r in tb_results]
df["tb_score"] = [r[1] for r in tb_results]
```

14.3.5 Sentiment Analysis with VADER

VADER calculates a compound score that is useful for short review-style text. The score is converted into a sentiment label using positive and negative thresholds.

```
# =====
# PART 4: Sentiment Analysis with VADER
# =====

analyzer = SentimentIntensityAnalyzer()

def vader_sentiment(text):
    scores = analyzer.polarity_scores(text)
    compound = scores["compound"]

    if compound >= 0.05:
        return "Positive", compound, scores
    elif compound <= -0.05:
        return "Negative", compound, scores
    else:
        return "Neutral", compound, scores

vader_results = df["text"].apply(vader_sentiment)

df["vader_label"] = [r[0] for r in vader_results]
df["vader_score"] = [r[1] for r in vader_results]
df["vader_details"] = [r[2] for r in vader_results]
```

14.3.6 Results Comparison

This section compares TextBlob and VADER predictions by displaying labels, sentiment scores, and the count of each sentiment class.

```
print("\n== Sentiment Analysis Results ==")
print(df[["product",
          "rating",
          "tb_label",
          "tb_score",
          "vader_label",
          "vader_score"]
      ].round(3))

print("\n== TextBlob Distribution ==")
print(df["tb_label"].value_counts())

print("\n== VADER Distribution ==")
print(df["vader_label"].value_counts())
```

```
== Sentiment Analysis Results ==
   product  rating  tb_label  tb_score  vader_label  vader_score
0  Wireless Headphones    5  Positive    0.608    Positive    0.872
1  Wireless Headphones    4  Positive    0.350    Positive    0.822
2    Smart Watch    2  Negative   -0.375    Negative   -0.477
3    Smart Watch    3   Neutral   -0.017    Positive    0.440
4  Bluetooth Speaker    5  Positive    0.290    Positive    0.784
5  Bluetooth Speaker    1  Negative   -1.000    Negative   -0.612
6    Laptop Stand    4   Neutral   -0.050    Positive    0.382
7    USB-C Hub    2  Positive    0.250    Negative   -0.402
8  Mechanical Keyboard    5  Positive    1.000    Positive    0.836
9      Webcam    3  Positive    0.183    Neutral    0.000
```

```
== TextBlob Distribution ==
tb_label
Positive    6
Negative    2
Neutral     2
Name: count, dtype: int64
```

```
== VADER Distribution ==
vader_label
Positive    6
Negative    3
Neutral     1
Name: count, dtype: int64
```

14.3.7 Product-wise Sentiment Summary

Reviews are grouped by product to compute average rating, average TextBlob score, average VADER score, and review count.

```
# Product-wise average sentiment score
product_summary = df.groupby("product").agg(
    Average_Rating=("rating", "mean"),
    TextBlob_Avg_Score=("tb_score", "mean"),
    VADER_Avg_Score=("vader_score", "mean"),
    Review_Count=("text", "count")
).round(3)

print("== Product-wise Sentiment Summary ==")
print(product_summary)
```

```
== Product-wise Sentiment Summary ==
   product  Average_Rating  TextBlob_Avg_Score  VADER_Avg_Score  \
product
Bluetooth Speaker         3.0             -0.355             0.086
Laptop Stand              4.0             -0.050             0.382
Mechanical Keyboard       5.0              1.000             0.836
Smart Watch               2.5             -0.196             -0.018
USB-C Hub                 2.0              0.250             -0.402
Webcam                   3.0              0.183              0.000
Wireless Headphones       4.5              0.479              0.847
```

```
   product  Review_Count
product
Bluetooth Speaker         2
Laptop Stand              1
Mechanical Keyboard       1
Smart Watch               2
```

```
   product  Review_Count
product
USB-C Hub         1
Webcam            1
Wireless Headphones  2
```

14.3.8 Sentiment Visualization


Bar charts are used to compare the distribution of Positive, Neutral, and Negative reviews for both TextBlob and VADER methods.

```
# Visualization of sentiment distributions
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

df["tb_label"].value_counts().reindex(["Positive", "Neutral", "Negative"], fill_value=0).plot(
    kind="bar",
    ax=axes[0],
    title="TextBlob Sentiment Distribution"
)
axes[0].set_xlabel("Sentiment")
axes[0].set_ylabel("Number of Reviews")
axes[0].tick_params(axis="x", rotation=0)

df["vader_label"].value_counts().reindex(["Positive", "Neutral", "Negative"], fill_value=0).plot(
    kind="bar",
    ax=axes[1],
    title="VADER Sentiment Distribution"
)
axes[1].set_xlabel("Sentiment")
axes[1].set_ylabel("Number of Reviews")
axes[1].tick_params(axis="x", rotation=0)

plt.tight_layout()
plt.show()
```



14.4 Result Summary

The analysis shows that both methods detect mostly positive reviews, but VADER marks more reviews as negative than TextBlob. This difference occurs because VADER considers sentiment intensity and review-style cues.

Method	Positive	Neutral	Negative
TextBlob	6	2	2
VADER	6	1	3

14.5 Conclusion

In this lab, a sample product review dataset was created and analyzed using TextBlob and VADER sentiment analysis. The workflow included review-data preparation, text preprocessing, sentiment scoring, method comparison, product-wise summarization, and visualization. TextBlob is simple and easy to interpret using polarity scores, while VADER is especially useful for short product-review text because it captures intensity and sentiment cues. Overall, both methods successfully classified reviews into positive, neutral, and negative categories.

Lab 15: Performance Comparison of Models

Topic	Medical Diagnosis Model Performance Comparison
Real-World Problem	Predict whether a patient is healthy or has a disease using medical features.
Dataset	Synthetic medical diagnosis dataset with 1,000 patient records.
Tool(s)	Python, Google Colab, scikit-learn, pandas, NumPy, matplotlib
Algorithms	Logistic Regression, Decision Tree, Random Forest, Gradient Boosting, SVM, Naive Bayes, KNN
Objective	Compare models using stratified cross-validation and select the best model using F1-Score and ROC-AUC.

1. Theory and Concepts

Model performance comparison is used to evaluate several machine learning algorithms on the same dataset. In this lab, stratified cross-validation is applied so that each fold keeps the healthy and disease class ratio balanced. The final selection is based mainly on F1-Score and ROC-AUC because these metrics summarize both class prediction quality and separation power.

- Classification predicts a category label, such as Healthy or Disease.
- Stratified cross-validation gives a more reliable model estimate by testing each model on multiple folds.
- F1-Score balances precision and recall, which is useful when both false positives and false negatives matter.
- ROC-AUC measures how well the model separates the two classes across decision thresholds.

2. Dataset Description

The synthetic dataset contains two groups: healthy patients and disease patients. Each patient record is represented by numerical medical features that are commonly used in diagnosis-style prediction tasks.

Feature	Type	Description
Age	Numeric	Patient age in years
Blood Pressure	Numeric	Blood pressure measurement
Cholesterol	Numeric	Cholesterol level
Glucose	Numeric	Glucose level
BMI	Numeric	Body Mass Index
Diagnosis	Binary Target	0 = Healthy, 1 = Disease

4.1 Import Required Libraries

This section imports all Python libraries required for data handling, preprocessing, model training, cross-validation, evaluation, and visualization.

```
import pandas as pd
import numpy as np

from sklearn.model_selection import cross_val_score, StratifiedKFold, train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, roc_auc_score, classification_report,
    confusion_matrix, ConfusionMatrixDisplay, RocCurveDisplay
)

import matplotlib.pyplot as plt
```

4.2 Generate Synthetic Medical Diagnosis Dataset

This section creates healthy and disease patient records, combines them into one dataset, separates the feature matrix and target vector, and displays the class distribution and sample records.

```
np.random.seed(42)

# Generate synthetic medical diagnosis dataset
n_samples = 1000
n_healthy = 600
n_disease = 400

# Healthy patient records
healthy = pd.DataFrame({
    'Age': np.random.normal(45, 12, n_healthy),
    'Blood_Pressure': np.random.normal(120, 15, n_healthy),
    'Cholesterol': np.random.normal(190, 25, n_healthy),
    'Glucose': np.random.normal(95, 12, n_healthy),
    'BMI': np.random.normal(24, 3, n_healthy),
    'Diagnosis': [0] * n_healthy
})

# Disease patient records
disease = pd.DataFrame({
    'Age': np.random.normal(58, 10, n_disease),
    'Blood_Pressure': np.random.normal(145, 18, n_disease),
    'Cholesterol': np.random.normal(240, 30, n_disease),
    'Glucose': np.random.normal(130, 22, n_disease),
    'BMI': np.random.normal(29, 4, n_disease),
    'Diagnosis': [1] * n_disease
})

# Combine both classes
df = pd.concat([healthy, disease], ignore_index=True)

# Separate features and target variable
X = df.drop('Diagnosis', axis=1)
y = df['Diagnosis']

print('Dataset Shape:', df.shape)
print('\nClass Distribution:')
print(df['Diagnosis'].value_counts().sort_index())
print('\nSample Data:')
df.head()
```

Dataset Shape: (1000, 6)

Class Distribution:
Diagnosis
0 600
1 400
Name: count, dtype: int64

Sample Data:

	Age	Blood_Pressure	Cholesterol	Glucose	BMI	Diagnosis
0	50.960570	131.354829	193.130613	117.814288	21.834787	0
1	43.340828	106.167520	179.264861	94.272070	24.530463	0
2	52.772262	133.044089	193.057438	86.499119	22.359960	0
3	63.276358	140.334568	203.582451	76.835427	23.185035	0
4	42.190160	126.201524	191.221502	73.362324	29.020356	0

4.3 Standardize Features

StandardScaler is applied to put numerical features on a similar scale. This is important for algorithms such as SVM, KNN, and Logistic Regression.

```
# Standardise features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print('Original feature shape:', X.shape)
print('Scaled feature shape:', X_scaled.shape)

Original feature shape: (1000, 5)
Scaled feature shape: (1000, 5)
```

4.4 Define Classification Models

Multiple classifiers are defined for comparison, including linear, tree-based, ensemble, probabilistic, kernel-based, and distance-based models.

```
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42),
    'Decision Tree': DecisionTreeClassifier(max_depth=5, random_state=42),
    'Random Forest': RandomForestClassifier(n_estimators=100, random_state=42),
    'Gradient Boosting': GradientBoostingClassifier(random_state=42),
    'SVM (RBF)': SVC(kernel='rbf', C=1.0, probability=True, random_state=42),
    'Naive Bayes': GaussianNB(),
```

```
    'KNN (K=5)': KNeighborsClassifier(n_neighbors=5)
}

print('Models defined successfully:')
for model_name in models.keys():
    print('-', model_name)
```

```
Models defined successfully:
- Logistic Regression
- Decision Tree
- Random Forest
- Gradient Boosting
- SVM (RBF)
- Naive Bayes
- KNN (K=5)
```

4.5 Model Evaluation using 5-Fold Stratified Cross-Validation

Each model is evaluated using 5-fold stratified cross-validation. The metrics include Accuracy, Precision, Recall, F1-Score, and ROC-AUC.

```
# 5-fold stratified cross-validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

results = []

print('=== Model Performance Comparison ===')
print(f'{{"Model":<22}} {{"Accuracy":<10}} {{"Precision":<10}} {{"Recall":<10}} {{"F1-Score":<10}} {{"ROC-AUC":<10}}')
print('-' * 75)

for name, model in models.items():
    acc = cross_val_score(model, X_scaled, y, cv=cv, scoring='accuracy').mean()
    prec = cross_val_score(model, X_scaled, y, cv=cv, scoring='precision').mean()
    rec = cross_val_score(model, X_scaled, y, cv=cv, scoring='recall').mean()
    f1 = cross_val_score(model, X_scaled, y, cv=cv, scoring='f1').mean()
    roc = cross_val_score(model, X_scaled, y, cv=cv, scoring='roc_auc').mean()

    results.append({
        'Model': name,
        'Accuracy': acc,
        'Precision': prec,
        'Recall': rec,
        'F1-Score': f1,
        'ROC-AUC': roc
    })

print(f'{{name:<22}} {{acc:<10.4f}} {{prec:<10.4f}} {{rec:<10.4f}} {{f1:<10.4f}} {{roc:<10.4f}}')
```

```
results_df = pd.DataFrame(results)

=== Model Performance Comparison ===
Model          Accuracy  Precision  Recall  F1-Score  ROC-AUC
-----
Logistic Regression  0.9630  0.9643  0.9425  0.9527  0.9909
Decision Tree      0.9220  0.9295  0.8725  0.8982  0.9262
Random Forest      0.9480  0.9472  0.9225  0.9337  0.9886
Gradient Boosting  0.9560  0.9523  0.9375  0.9441  0.9897
SVM (RBF)          0.9550  0.9587  0.9275  0.9420  0.9913
Naive Bayes        0.9620  0.9665  0.9375  0.9510  0.9936
KNN (K=5)          0.9540  0.9573  0.9275  0.9410  0.9815
```

4.6 Rank Models by F1-Score

The models are sorted by F1-Score to identify the strongest overall classifier. In the output, Logistic Regression appears as the best model with the highest F1-Score and strong ROC-AUC.

```
# Rank models by F1-Score
results_df = results_df.sort_values('F1-Score', ascending=False)

print('=== Model Ranking by F1-Score ===')
print(results_df[['Model', 'F1-Score', 'ROC-AUC']].to_string(index=False))

# Best model
best = results_df.iloc[0]
print(f"""
```

```
Best Model: {best['Model']}""")
print(f"F1-Score: {best['F1-Score']:.4f}")
print(f"ROC-AUC: {best['ROC-AUC']:.4f}")
```

```
=== Model Ranking by F1-Score ===
      Model  F1-Score  ROC-AUC
Logistic Regression  0.952674  0.990917
Naive Bayes         0.950983  0.993604
Gradient Boosting   0.944088  0.989719
SVM (RBF)           0.941951  0.991333
KNN (K=5)           0.941008  0.981510
Random Forest       0.933706  0.988604
Decision Tree       0.898153  0.926198


Best Model: Logistic Regression
F1-Score: 0.9527
ROC-AUC: 0.9909
```

4.7 Visualize Model Comparison

The bar chart compares all models across multiple evaluation metrics, making it easier to visually identify the best-performing classifier.

```
# Visualize model comparison
plot_df = results_df.set_index('Model')

plt.figure(figsize=(12, 6))
plot_df[['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC-AUC']].plot(kind='bar', figsize=(12, 6))
plt.title('Medical Diagnosis Model Performance Comparison')
plt.xlabel('Machine Learning Models')
plt.ylabel('Score')
plt.ylim(0, 1.05)
plt.xticks(rotation=45, ha='right')
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()
```



4.8 Train-Test Evaluation of the Best Model

After selecting the best model, Logistic Regression is trained and evaluated on a separate test split. The output includes a classification report and confusion matrix.

After selecting the best model using cross-validation, we train it on the training set and evaluate it on a separate test set using a classification report and confusion matrix.

```
# Split data for final best-model evaluation
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)

best_model_name = best['Model']
best_model = models[best_model_name]

best_model.fit(X_train, y_train)
y_pred = best_model.predict(X_test)

print(f'=== Best Model Test Evaluation: {best_model_name} ===')
print(classification_report(y_test, y_pred, target_names=['Healthy', 'Disease']))

cm = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(cm)


ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Healthy', 'Disease']).plot()
plt.title(f'Confusion Matrix - {best_model_name}')
plt.show()
```

```
=== Best Model Test Evaluation: Logistic Regression ===
              precision    recall  f1-score   support

   Healthy       0.97       0.97       0.97       120
   Disease       0.96       0.95       0.96        80

 accuracy              0.96       0.96       0.96       200
 macro avg              0.96       0.96       0.96       200
weighted avg              0.96       0.96       0.96       200


Confusion Matrix:
[[117  3]
 [ 4 76]]
```



4.9 ROC Curve for Best Model

The ROC curve shows the classification separation ability of the selected model. A higher ROC-AUC value indicates better discrimination between healthy and disease classes.

```
# ROC Curve for best model
RocCurveDisplay.from_estimator(best_model, X_test, y_test)
plt.title(f'ROC Curve - {best_model_name}')
plt.grid(True, alpha=0.3)
plt.show()
```



5. Result Summary

The cross-validation results show that Logistic Regression achieved the best F1-Score in this run, followed closely by Naive Bayes and Gradient Boosting. The best model was then evaluated on a separate test set for final validation.

Model	F1-Score	ROC-AUC	Rank Note
Logistic Regression	0.9527	0.9909	Best overall by F1-Score
Naive Bayes	0.9510	0.9936	Very strong ROC-AUC
Gradient Boosting	0.9441	0.9897	Strong ensemble model
SVM (RBF)	0.9420	0.9913	Good separation performance
KNN (K=5)	0.9410	0.9815	Good distance-based result
Random Forest	0.9337	0.9886	Good ensemble baseline
Decision Tree	0.8982	0.9262	Lowest score among compared models

6. Conclusion

In this lab, a synthetic medical diagnosis dataset was generated and several machine learning classification models were compared. The models were evaluated using stratified cross-validation and multiple performance metrics. Based on F1-Score and ROC-AUC, Logistic Regression was selected as the best-performing model for final evaluation. The train-test results and ROC curve show that model comparison is a useful approach for selecting a suitable algorithm for medical diagnosis prediction tasks

